

Qwen3.5-Omni Technical Report

Qwen Team



<https://huggingface.co/spaces/Qwen/Qwen3.5-Omni-Online-Demo>

<https://modelscope.cn/studios/Qwen/Qwen3.5-Omni-Online-Demo>

Abstract

In this work, we present **Qwen3.5-Omni**, the latest advancement in the Qwen-Omni model family. Representing a significant evolution over its predecessor, Qwen3.5-Omni scales to hundreds of billions of parameters and supports a 256k context length. By leveraging a massive dataset comprising heterogeneous text-vision pairs and over 100 million hours of audio-visual content, the model demonstrates robust omni-modality capabilities. Qwen3.5-Omni-Plus achieves SOTA results across 215 audio and audio-visual understanding, reasoning, and interaction subtasks and benchmarks, surpassing Gemini-3.1 Pro in key audio tasks and matching it in comprehensive audio-visual understanding. Architecturally, Qwen3.5-Omni employs a **Hybrid Attention Mixture-of-Experts (MoE) framework** for both **Thinker and Talker**, enabling efficient long-sequence inference. The model facilitates sophisticated interaction, supporting over 10 hours of audio understanding and 400 seconds of 720P video (at 1 FPS). To address the inherent instability and unnaturalness in streaming speech synthesis—often caused by encoding efficiency discrepancies between text and speech tokenizers—we introduce **ARIA (Adaptive Rate Interleave Alignment)**. ARIA dynamically aligns text and speech units, significantly enhancing the stability and prosody of conversational speech with minimal latency impact. Furthermore, Qwen3.5-Omni expands linguistic boundaries, supporting multilingual understanding and speech generation across 10 languages with human-like emotional nuance. Beyond preset voices, the model enables zero-shot voice customization via user-provided samples. Finally, Qwen3.5-Omni exhibits superior audio-visual grounding capabilities, generating script-level structured captions with precise temporal synchronization and automated scene segmentation. Remarkably, we observed the emergence of a new capability in omnimodal models: directly performing coding based on audio-visual instructions, which we call **Audio-Visual Vibe Coding**. Qwen3.5-Omni is publicly accessible via API¹.

1 Introduction

Human interaction with the world is inherently omnimodal and agentic, involving the integration of visual, auditory, and linguistic information, and the production of responses through text, speech, and goal-directed tool-mediated actions, facilitating information exchange with other organisms and demonstrating intelligence. Building on the rapid advances in the understanding and reasoning capabilities of large models across text (Brown et al., 2020; OpenAI, 2023; Gemini Team, 2024; Anthropic, 2023a;b; 2024; Bai et al., 2023a; Yang et al., 2024a; 2025a; Touvron et al., 2023; Dubey et al., 2024), vision (Li et al., 2023; Liu et al., 2023; Zhu et al., 2023; Bai et al., 2023b; 2025a), and audio (Chu et al., 2023; 2024), natively omnimodal systems that jointly process and generate across all modalities have drawn substantial attention (OpenAI, 2024; Comanici et al., 2025; Xu et al., 2025a;b). However, existing models predominantly operate within passive perception-response paradigms and exhibit limited capacity for scalable agentic behavior, real-time interaction, autonomous tool utilization, and cross-modal reasoning, which are essential prerequisites for practical deployment.

In this report, we present Qwen3.5-Omni, Qwen’s latest generation of fully omnimodal LLM, supporting the understanding of text, images, audio, and audio-visual content. Natively pretrained in an omnimodal manner on massive amounts of text, visual data, and more than 100 million hours of audio-visual data, Qwen3.5-Omni is designed as a native omni agent model: it not only perceives and reasons across all modalities, but also acts, autonomously invoking WebSearch, executing complex FunctionCall, generating speech outputs, and engaging in real-time streaming interaction. The model series includes Plus and Flash variants, all of which are instruct models with 256k-token long-context input.

Qwen3.5-Omni builds on the Thinker–Talker architecture introduced in Qwen2.5-Omni (Xu et al., 2025a) and introduces **five key technical upgrades** over Qwen3-Omni (Xu et al., 2025b): (1) both the Thinker

¹<https://www.alibabacloud.com/help/en/model-studio/qwen-omni>



Figure 1: Qwen3.5-Omni is a unified end-to-end model capable of processing multiple modalities, such as text, audio, image and video, and generating real-time text or speech response. Based on these features, Qwen3.5-Omni supports a wide range of tasks, including but not limited to voice dialogue, video dialogue, and audio-visual tool use.

and Talker adopt Hybrid-Attention Mixture-of-Experts (MoE) designs, enabling highly efficient inference; (2) supporting long-context modeling up to 256k tokens, supporting more than 10 hours of audio and over 400 seconds of 720P audio-visual content at 1 FPS; (3) on the speech generation side, a multi-codebook codec representation enables single-frame, immediate synthesis; (4) the Talker introduces **ARIA**, a technique that dynamically aligns text and speech units during streaming decoding, significantly improving naturalness and robustness; and (5) multilingual training is substantially expanded, covering 113 languages and dialects for speech recognition and 36 for speech synthesis.

Enabled by these technical advances, Qwen3.5-Omni delivers **three major new capabilities** over Qwen3-Omni: (1) controllable audio-visual captioning, capable of generating controllable, detailed, and structured captions as well as screenplay-level fine-grained descriptions, including automatic segmentation, timestamp annotation, and detailed descriptions of characters and their relationship to audio; (2) comprehensive real-time interaction, encompassing semantic interruption through native turn-taking intent recognition, end-to-end voice control over volume, speed, and emotion, and voice cloning from user-provided samples; and (3) native omnimodal agentic behavior, including autonomous WebSearch, complex FunctionCall invocation, and Audio-Visual Vibe Coding, an emergent capability wherein the model directly generates executable code from audio-visual instructions, enabling the model to respond to real-time queries without external orchestration.

Critically, Qwen3.5-Omni maintains state-of-the-art performance on text and visual modalities without degradation relative to same-size single-model Qwen counterparts. Across 215 audio and audio-visual understanding, reasoning, and interaction subtasks and benchmarks, covering audio-visual benchmarks, audio benchmarks, ASR benchmarks, language-specific speech-to-text translation tasks, and language-specific ASR tasks, Qwen3.5-Omni-Plus achieves SOTA results, surpassing Gemini-3.1 Pro across general audio understanding, reasoning, recognition, translation, and dialogue, while its overall audio-visual understanding reaches the level of Gemini-3.1 Pro.

2 Architecture

2.1 Overview

As shown in Figure 2, Qwen3.5-Omni continues to adopt the Thinker-Talker architecture (Xu et al., 2025a). Compared with Qwen3-Omni (Xu et al., 2025b), Qwen3.5-Omni introduces several key improvements in scalability, alignment, and real-time interaction:




Figure 2: The overview of Qwen3.5-Omni. Qwen3.5-Omni adopts the Thinker-Talker architecture. Thinker is tasked with text generation while Talker focuses on generating streaming speech tokens by receives high-level representations directly from Thinker. To achieve ultra-low-latency streaming, Talker autoregressively predicts a multi-codebook sequence. At each decoding step, an MTP module outputs the residual codebooks for the current frame, after which the Code2Wav renderer incrementally synthesizes the corresponding waveform, enabling frame-by-frame streaming generation.

- The overall backbone adopts a Hybrid Mixture-of-Experts (MoE) design, improving scalability while better balancing capacity and efficiency across multimodal understanding and generation.
- The Thinker receives visual and audio signals through the Vision Encoder and AuT, respectively. Audio and video inputs are interleaved for unified multimodal modeling, with explicit timestamps inserted to improve temporal perception, especially for long video or audio-video contexts. This design enables the Thinker to handle extended inputs, supporting up to 256k tokens, 10 hours of audio, or 400 seconds of 720P video at 1 FPS.
- The Talker is responsible for contextual speech generation by conditioning on multimodal inputs together with the textual outputs from the Thinker. Qwen3.5-Omni adopts the RVQ-based speech representation introduced in Qwen3-Omni (Xu et al., 2025b), which substantially improves inference efficiency.
- To support real-time interaction, Qwen3.5-Omni adopts both chunk-wise streaming input processing in the Thinker and a streaming Talker design, enabling low-latency end-to-end multimodal conversation.
- Different from the dual-track Talker input design in Qwen3-Omni (Xu et al., 2025b), the Talker in Qwen3.5-Omni adopts ARIA to dynamically align text and speech units before interleaving them. This design mitigates the instability caused by mismatched tokenization rates between text and speech, thereby reducing issues such as skipped words, incorrect pronunciations, and ambiguous rendering of numbers.

In the following sections, we first introduce with the AuT encoder, including its training methodology. Then, describe how Thinker processes various inputs. We then detail Talker’s multi-codebook streaming speech generation. Finally, we highlight a series of improvements on both the understanding and generation modules aimed at achieving ultra-low-latency, end-to-end streaming audio inference.




Figure 3: The overview of AuT. Consuming 40 million hours of supervised data especially more multilingual data, AuT encoder in Qwen3.5-Omni obtain stronger general purpose audio representation in 6.25Hz.

2.2 Audio Transformer (AuT)

We use transformer based audio encoder trained from scratch in attention-encoder-decoder model AuT, as is shown in Figure 3. The training of Qwen3.5-Omni encoder consumed 40 million hours of audio-text pair data generated by Qwen3-ASR. The filter bank features of the audio are downsampled 16 times using 4 Conv2D blocks and then fed into self-attention layers to obtain audio tokens in 6.25Hz token rate. Comparing to the training process of Qwen3-Omni encoder, the encoder of Qwen3.5-Omni adapts more multilingual data of more than 20 languages, and the proportion of Chinese, English and multilingual data comes to 3.5 : 3.5 : 3. The dynamic attention window size training mechanism is adopted for guaranting balance performance of inference under real-time prefill caching and for the offline audio understanding tasks.

2.3 Perception

Text, Audio, Image and Video (w/o Audio). The Thinker converts text, audio, image, and silent video inputs into a unified sequence of representations. For text, we use the Qwen3.5 tokenizer (Team, 2026), which adopts byte-level byte-pair encoding with a vocabulary size of 250k (up from 150k), improving encoding and decoding efficiency by 10–60% across most languages. For audio inputs, including audio extracted from video, we resample the waveform to 16 kHz and convert it into a 128-channel mel-spectrogram using a 25 ms window and a 10 ms hop size. We use AuT as the audio encoder, trained from scratch on 40 million hours of audio data, where each output frame corresponds to approximately 160 ms of the original signal. For visual inputs, we adopt the vision encoder from Qwen3.5 (Team, 2026) to process both images and videos. Trained on a mixture of image and video data, this encoder provides strong capabilities in both image understanding and video comprehension. To preserve video information as much as possible while maintaining alignment with the audio stream, we sample video frames at a dynamic frame rate.

Audio-visual Timestamp. Following Qwen3-Omni (Xu et al., 2025b), we apply TM-RoPE to endow the model with temporal awareness for audio-video synchronization. However, we find that directly encoding absolute time through temporal position IDs can lead to excessively sparse indices for visual patches from long video with audio inputs, which weakens long-range temporal modeling. In addition, such a design often requires large-scale and uniformly distributed training samples across different frame rates, increasing data construction cost. To address these issues, we prepend each video or audio-video temporal patch with an explicit timestamp represented as a formatted text string in seconds, allowing the model to learn timecode representations more naturally. For audio sequences, we further insert timestamps at random intervals to improve temporal alignment across modalities. Although this strategy

slightly increases the context length, it enables more precise and robust temporal perception, especially when extrapolating long-context multimodal inputs.

In the context of multimodal audio-visual streams, the audio component is encoded with a temporal ID for every 160 ms. The video is treated as a sequence of frames with monotonically increasing temporal IDs that are dynamically adjusted based on their actual timestamps to ensure a consistent temporal resolution of 160 ms per ID. The height and width IDs for video frames are assigned in the same manner as for still images. To prevent positional conflicts when processing multiple modalities, the position numbering is made contiguous, with each subsequent modality commencing from one plus the maximum position ID of the preceding modality. This refined approach to positional encoding enables the model to effectively integrate and jointly model information from diverse modalities. Qwen3.5-Omni aligns these representations using their temporal IDs, which are explicitly anchored to absolute time. This design choice affords the model the flexibility to support streaming inputs of arbitrary duration.

2.4 Speech Generation

Talker operates directly on the RVQ tokens produced by *Qwen3.5-Omni-Audio-Tokenizer*. To model the residual codebooks, it employs a multi-token prediction (MTP) module, which enables fine-grained modeling and control of acoustic details. Coupled with a causal ConvNet for waveform reconstruction, Talker delivers high-fidelity speech synthesis with low inference latency and modest computational overhead.

In multi-turn spoken dialogue, Talker is conditioned on the rich contextual information provided by the Thinker component, including historical text tokens, multimodal representations, and the streamed text of the current turn. Such conditioning allows Talker to dynamically modulate acoustic attributes—such as prosody, loudness, and emotion—in accordance with the evolving conversational context.

Architecturally, our approach differs from Qwen3-Omni (Xu et al., 2025b) in two key respects. First, we introduce a dedicated system prompt for Talker that specifies target voice characteristics, thereby enabling both zero-shot voice cloning and controllable speech generation. Compared with conventional speaker embeddings, this prompt can encode richer multimodal cues, including textual descriptions and codec sequences, providing substantially finer-grained control over acoustic realization. Second, we propose **ARIA** (Adaptive Rate Interleave Alignment), which unifies the conventional dual-channel generation paradigm into a single-channel formulation. Rather than relying on MFA-derived alignments or fixed interleaving rates, ARIA enforces an adaptive rate constraint: for any prefix of the generated sequence, the cumulative speech-to-text token ratio must not exceed the corresponding item-level global ratio. Despite its simplicity, this design affords flexible text-speech alignment across languages, including those with relatively low encoding efficiency, and naturally supports arbitrary text-token prefixes followed by coherent speech-token continuation.

2.5 Designs for Streaming and Concurrency

In streaming audio-visual interaction scenarios, the first-packet latency is a critical factor affecting user experience, and the model’s concurrency capability is key to reducing service costs and improving response speed. This section discusses how Qwen3.5-Omni enhances concurrency and reduces first-packet latency through algorithmic and architectural optimizations. Table 1 provides an overview of the relevant architecture of the Qwen3.5-Omni and its associated latency.

Table 1: **Architecture of Qwen3.5-Omni and end-to-end first-packet latency under audio/video settings (ms).**

Module	Architecture	Streaming	
Audio Encoder	AuT	✓	
Vision Encoder	SigLIP2	–	
Thinker	Hybrid MoE Transformer	✓	
Talker	Hybrid MoE Transformer	✓	
MTP	Dense Transformer	✓	
Code2wav	ConvNet	✓	
First-Packet Latency (Audio Input)		<i>Plus:</i> 435ms	<i>Flash:</i> 235ms
First-Packet Latency (Video Input)		<i>Plus:</i> 651ms	<i>Flash:</i> 426ms

Chunked Prefilling and Hybrid MoE Architecture. In Qwen3.5-Omni, we retain the chunked-prefilling mechanism as implemented in Qwen3-Omni and Qwen2.5-Omni, whose audio and vision encoders are

capable of outputting chunks along the temporal dimension. This approach significantly reduces the Time-To-First-Token (TTFT) for both the Thinker and the Talker. Architecturally, both the Thinker and the Talker in Qwen3.5-Omni are built upon the Hybrid MoE architecture introduced in Qwen3.5. Beyond the general efficiency advantage of Hybrid MoE, this architecture includes the Gated Delta Net (GDN) module, which is particularly effective for accelerating the modeling of long audio-video sequences. As a result, it significantly reduces KV-cache I/O overhead in long-context inference, improving generation throughput and enabling higher serving concurrency.

Streaming Generation with ARIA. For streaming speech generation and high-concurrency serving, Qwen3.5-Omni largely inherits the efficient design of Qwen3-Omni: Talker predicts RVQ codec tokens with a lightweight MTP module, and the generated multi-codebook tokens are converted to waveform by a causal and streaming ConvNet codec decoder. These components remain computationally lightweight, batch-friendly, and well-suited for low-latency deployment. Built on this shared foundation, the previously introduced ARIA further reformulates the dual-channel generation pattern in Qwen3-Omni into a unified interleaved single-stream formulation over text and speech tokens. By organizing text and speech generation under a monotonic interleaving constraint, ARIA reduces the synchronization overhead between separate generation tracks, enables more efficient token scheduling during decoding, and better matches the naturally incremental regime of streaming interaction.

In Table 2, we report the theoretical first-packet latency of Qwen3.5-Omni under different concurrency levels for audio and video input, evaluated on internal vLLM with *torch.compile* and CUDA Graph acceleration enabled for the MTP module and codec decoder. Here, Thinker TTFT (*Time-To-First-Token*) denotes the time from receiving the input stream to the first text token generated by Thinker, while Talker TTFC (*Time-To-First-Chunk*) measures the time until Talker produces the first audio chunk. TPOP (*Time-Per-Output-Token*) represents the per-output-token latency during steady-state decoding, where Talker TPOP includes the combined latency of the Talker backbone and the MTP module. TPS (*Tokens Per Second*) denotes generation throughput. Since ARIA organizes text and speech generation in a unified interleaved stream, Overall Latency cannot be obtained by simply summing several row values, but instead reflects the end-to-end critical path to the first playable audio packet. We also note that, due to the substantial scale difference between Qwen3.5-Omni-Flash and Qwen3.5-Omni-Plus, the two variants adopt different deployment-time resource allocation and parallelization strategies; therefore, their latency and throughput numbers are not intended for strict horizontal comparison. As shown in the table, Qwen3.5-Omni maintains stable latency and decoding efficiency as concurrency increases, while the low Generation RTF provides sufficient margin for smooth streaming audio generation.

Table 2: Theoretical first-packet latency of Qwen3.5-Omni under different concurrency levels. A/V denotes audio/video input.

	Qwen3.5-Omni-Flash			Qwen3.5-Omni-Plus		
	1 Conc.	4 Conc.	8 Conc.	1 Conc.	4 Conc.	8 Conc.
Thinker TTFT	80/255ms	86/446ms	103/765ms	162/377ms	183/907ms	260/1243ms
Talker TTFC	56/61ms	68/108ms	81/116ms	54/56ms	72/88ms	95/116ms
Thinker TPOP	5.6/5.9ms	8.2/9.2ms	9.6/15.8ms	17.4/18.5ms	25.6/26.9ms	33.3/40.2ms
Talker TPOP	14.2/14.2ms	16.9/17.0ms	20.5/20.6ms	14.9/14.9ms	21.0/21.3ms	25.8/27.1ms
Codec Decode			3~5ms			
Overall Latency	235/426ms	298/891ms	352/1625ms	435/651ms	619/1515ms	955/1980ms
Thinker TPS	177/171	556/457	942/598	57/54	156/149	266/240
Talker TPS	70/70	237/235	389/388	67/67	191/189	320/296
Generation RTF	0.178	0.211	0.257	0.187	0.267	0.334

3 Pretraining

Qwen3.5-Omni is pre-trained on a diverse dataset that encompasses multiple languages and dialects as shown in Table 3 and modalities, including image-text, video-text, audio-text, video-audio, video-audio-text, and pure text corpora. Following Qwen3-Omni (Xu et al., 2025b), we employ a wider range of natural language prompts to enhance both the generalization ability and instruction-following capabilities. To achieve robust performance across all modalities, our training strategy incorporates both unimodal and cross-modal data from the early pretraining stage.

In Qwen3-Omni (Xu et al., 2025b), we employ TMRoPE to endow the model with temporal awareness. However, we identify two key limitations of this approach: (1) By directly tying temporal position IDs to

Table 3: **Supported languages and dialects in Qwen3.5-Omni-Plus.**

Modality	# Varieties	Supported languages and dialects
Text	201	See Qwen3.5 for the complete list of supported languages.
Speech Input	113	74 languages: Afrikaans, Arabic, Asturian, Azerbaijani, Basque, Belarusian, Bengali, Bosnian, Bulgarian, Cantonese, Catalan, Cebuano, Chinese, Croatian, Czech, Danish, Dutch, English, Esperanto, Estonian, Filipino, Finnish, French, Galician, Georgian, German, Greek, Hebrew, Hindi, Hungarian, Icelandic, Indonesian, Interlingua, Italian, Japanese, Javanese, Kannada, Kazakh, Korean, Kyrgyz, Lingala, Latvian, Lithuanian, Macedonian, Malay, Malayalam, Maltese, Maori, Marathi, Mongolian, Norwegian Bokmål, Norwegian Nynorsk, Oriya, Persian, Polish, Portuguese, Punjabi, Romanian, Russian, Serbian, Slovak, Slovenian, Spanish, Swahili, Swedish, Tajiki, Tamil, Telugu, Thai, Turkish, Ukrainian, Urdu, Uyghur, and Vietnamese. 39 Chinese dialects: Northeastern Mandarin, Guizhou dialect, Guangdong Cantonese, Henan dialect, Hong Kong Cantonese, Shanghaiese, Shaanxi dialect, Tianjin dialect, Taiwanese Mandarin, Yunnan dialect, Anhui dialect, Fujian dialect, Gansu dialect, Guangdong Mandarin, Hubei dialect, Hunan dialect, Jiangxi dialect, Shandong dialect, Shanxi dialect, Sichuanese, Guangxi dialect, Hainan dialect, Chongqing dialect, Changsha dialect, Hangzhou dialect, Hefei dialect, Yinchuan dialect, Zhengzhou dialect, Shenyang dialect, Wenzhou dialect, Wuhan dialect, Kunming dialect, Taiyuan dialect, Nanchang dialect, Jinan dialect, Lanzhou dialect, Nanjing dialect, Hakka, and Southern Min.
Speech Output	36	29 languages: Chinese, English, German, Italian, Portuguese, Spanish, Japanese, Korean, French, Russian, Thai, Indonesian, Arabic, Vietnamese, Turkish, Finnish, Polish, Hindi, Dutch, Czech, Urdu, Tagalog, Swedish, Danish, Hebrew, Icelandic, Malay, Norwegian, and Persian. 7 Chinese dialects: Sichuanese, Beijing dialect, Tianjin dialect, Nanjing dialect, Shaanxi dialect, Cantonese, and Southern Min.

absolute time, it produces excessively large and sparse temporal position IDs for long audio-video or video inputs, which undermines the model’s ability to capture long-range temporal contexts. (2) Effective learning under this scheme typically requires large-scale and uniformly distributed sampling across different frame rates (fps), significantly increasing the cost of training data construction. To address these issues, we prepend each video or audio-video temporal patch with a timestamp represented as a formatted text string in seconds, enabling the model to better learn and interpret timecode representations. In addition, for audio sequences, we insert timestamps at random intervals to better align training across different modalities. Although this approach introduces a modest increase in context length, it allows the model to perceive temporal information more effectively and precisely.

The pre-training of Qwen3.5-Omni is structured into three distinct stages. In the first stage, we lock the LLM parameters and focus on training the vision and audio encoders, utilizing a vast corpus of audio-text and image-text pairs to enhance semantic understanding within the LLM. In the second stage, we unfreeze all parameters and train with a wider range of multimodal data for more comprehensive learning with a sequence length of 32,768. In the final stage, we use data with a sequence length of 262,144 to enhance the model’s ability to understand complex long-sequence data:

- (1) **Encoder Alignment Stage (S1):** During the initial pretraining phase, the LLM component of Qwen3.5-Omni is initialized with parameters from Qwen3.5, while the vision encoder is adopted from Qwen3.5, and the audio encoder is initialized with AuT. The two encoders are trained separately on the fixed LLM, with both initially focusing on training their respective adapters before training the encoders.
- (2) **General Stage (S2):** The second phase of pretraining utilizes a large-scale dataset containing approximately 4 trillion tokens, with the following distribution across modalities: text (0.92 trillion), audio (1.99 trillion), image (0.95 trillion), video (0.14 trillion), and video-audio (0.29 trillion). During this stage, the introduction of more diverse multimodal data and tasks enhances the model’s understanding and interaction capabilities in auditory, visual, textual, and audio-visual information.
- (3) **Long Context Stage (S3):** In the final pre-training phase, we increased the maximum token length from 32,768 to 262,144 and also raised the proportion of long audio and long video in the training

data. Experimental results indicate that these adjustments lead to significant improvements in the model’s ability to understand long sequence data.

4 Post-training

4.1 Thinker

The post-training phase employs a three-stage strategy for the Thinker, designed to preserve the model’s capabilities across all modalities without degradation, ensure high response quality under audio queries, and optimize the overall interaction experience. The training corpus, structured in the ChatML (OpenAI, 2022) format, encompasses pure text, visual, audio, and mixed-modality conversational data. Specifically, the process consists of the following stages:

- **Stage 1: Specialist Distillation** To establish a strong foundation for omnimodal capabilities, we first train a suite of domain-specialized teacher models via independent Supervised Fine-Tuning (SFT) and reinforcement learning (RL). All teacher models are fine-tuned from the pre-trained Qwen-3.5 base checkpoint. Beyond text-related tasks, including agentic, coding, and foundational reasoning tasks, we also train specialized teacher models for vision and audio. These teacher models are used to generate domain-specific data, enabling the specialized capabilities learned in each domain to be distilled into a single unified model.
- **Stage 2: On-Policy Distillation** Through the specialist distillation described above, the model already achieves strong performance in domains such as multimodal understanding and reasoning, as well as text-based dialogue, reasoning, coding, and agentic tasks. Nevertheless, a substantial gap remains between the quality of responses conditioned on audio queries and that of responses conditioned on text queries, particularly in speech dialogue. To reduce this gap, we introduce a second-stage training procedure based on on-policy distillation (OPD), with the goal of distilling the model’s stronger response capabilities under text inputs into the audio-input setting. Concretely, for each audio-text paired query, we first obtain a response generated under the text condition, which typically exhibits higher quality in terms of fluency, reasoning, and task completion. We then use this response as the distillation target for the corresponding audio-conditioned query. By training on such on-policy targets, the model gradually aligns its audio-conditioned outputs with its text-conditioned behavior, thereby improving response quality under audio inputs and promoting modality-consistent generation.
- **Stage 3: Interaction-Aligned Reinforcement Learning** Although the previous two stages substantially improve the model’s domain capabilities and cross-modal response quality, they are not sufficient to fully optimize the model for real-world interactive use. In multi-turn conversations, we observe several interaction-specific issues, including unintended language code-switching, persona inconsistency, and degraded instruction-following over extended contexts. To mitigate these issues, we introduce Interaction-Aligned RL, a third-stage reinforcement learning procedure aimed at optimizing the model for interaction quality. We construct multi-turn interaction trajectories and design reward signals around these user experience objectives, enabling the model to learn behaviors that are more stable, consistent, and aligned in prolonged interactions. By explicitly optimizing for interaction quality, this stage improves the model’s overall usability in practical conversational scenarios.

4.2 Talker

We employ a four-stage training pipeline for Talker, enabling Qwen3.5-Omni to generate natural and contextually appropriate spoken responses jointly with text. All training data is organized in the ChatML format to maintain consistency with Thinker and to facilitate voice cloning.

- (1) **General Stage:** In the initial pre-training stage, we train Qwen3.5-Omni on more than 20 million hours of multilingual speech data paired with multimodal context. In particular, the introduction of more diverse tasks, such as instruction-following speech generation, substantially enhance contextual reasoning and paralinguistic alignment, going beyond a simple monotonic mapping from multimodal representations to speech.
- (2) **Long-Context Stage:** We perform data quality stratification through a dedicated curation pipeline and conduct continual pre-training (CPT) on high-quality subsets. Augmented by Qwen3-Omni-Captioner, this stage mitigates hallucinations introduced by noisy data in the initial pre-training phase and substantially improves the naturalness and quality of generated speech. Meanwhile, we extend the maximum context length to 64k tokens, allowing the model to better handle long and complex user inputs and to produce more contextually grounded speech responses.

- (3) **Reinforcement Learning Stage:** We further align model behavior with human preferences through Direct Preference Optimization (DPO) (Rafailov et al., 2023). Concretely, we construct multilingual preference pairs based on human annotations and optimize the model with DPO. In addition, we incorporate rule-based rewards and adopt GSPO (Zheng et al., 2025) to further improve overall capability and training stability across diverse tasks.
- (4) **Speaker Fine-tuning Stage:** Finally, we perform lightweight speaker fine-tuning on top of the base model, enabling Qwen3.5-Omni to faithfully capture target speaker characteristics while further improving the naturalness, expressiveness, and controllability of its speech outputs.

5 Evaluation

A comprehensive evaluation was performed on two variants of models, including Qwen3.5-Omni-Flash and Qwen3.5-Omni-Plus. The evaluation results are divided into two main categories: understanding ($X \rightarrow \text{Text}$) and speech generation ($X \rightarrow \text{Speech}$).

5.1 Evaluation of $X \rightarrow \text{Text}$

In this section, we evaluate Qwen3.5-Omni’s ability to comprehend various multimodal inputs (text, audio, vision, and audio-visual video) and generate textual responses.

Text $\rightarrow$ Text Our evaluation of Qwen3.5-Omni on text $\rightarrow$ text primarily focuses on general knowledge tasks, instruction following, long context tasks, STEM tasks, reasoning tasks and general agent ability. Specifically, we utilize MMLU-Pro (Wang et al., 2024d), MMLU-Redux (Gema et al., 2024), SuperGPQA (Team et al., 2025) and C-Eval (Huang et al., 2023) for general knowledge tasks, IFEval (Zhou et al., 2023) and IFBench (Pyatkin et al., 2025) for instruction following, AA-LCR (Team, 2025a) and LongBench v2 (Bai et al., 2025b) for long context tasks, GPQA (Rein et al., 2023) for STEM tasks, LiveCodeBench v6 (Jain et al., 2024), HMMT Nov 25 (Balunović et al., 2025) and IMOAnswerBench (Luong et al., 2025) for reasoning tasks, BFCL-V4 (Yan et al., 2024) and TAU2Bench (Barres et al., 2025) for general agent ability.

Audio $\rightarrow$ Text To evaluate audio-to-text capabilities, we employ benchmarks across four domains: audio understanding, end-to-end speech dialogue, speech-to-text translation (S2TT), and automatic speech recognition (ASR). For audio understanding, we utilize MMAU (Sakshi et al., 2024), MMAR (Ma et al., 2025a), MMSU (Wang et al., 2025a), RUL-MuchoMusic (Zang et al., 2025), and SongFormBench (Hao et al., 2025) to assess comprehension of sound effects, speech, and music. Dialogue performance is evaluated via VoiceBench (Chen et al., 2024b), URO-Bench-pro (Yan et al., 2025), SpeechRole (Jiang et al., 2025), and WildSpeech-Bench (Zhang et al., 2025b). For S2TT, we focus on the translation of the top 59 languages in Fleurs (Conneau et al., 2022) into English and Chinese. Finally, ASR performance is measured using Fleurs (Conneau et al., 2022), Common Voice (Ardila et al., 2020), LibriSpeech (Panayotov et al., 2015), WenetSpeech (Zhang et al., 2022), KeSpeech (Tang et al., 2021), Opencpop-test (Wang et al., 2022), and MIR-1K (vocal) (Hsu & Jang, 2010), covering multilingual speech, Chinese dialects, and singing voice transcription.

Vision $\rightarrow$ Text The evaluation of the model’s vision-to-text capabilities encompasses a suite of benchmarks targeting diverse and challenging tasks. To assess performance in specialized domain of mathematical and STEM reasoning, we utilize MMMU (Yue et al., 2023), MMMU-Pro (Yue et al., 2024), MathVista (Lu et al., 2024), MathVision (Wang et al., 2024a), DynaMath (Zou et al., 2025), ZEROBench (Roberts et al., 2025). For the general visual question answering, the model is evaluated on RealWorldQA (Zhang et al., 2024), MMStar (Chen et al., 2024a), HallusionBench (Guan et al., 2024), and SimpleVQA (Cheng et al., 2025). The model’s proficiency in document understanding is measured using the CharXiv (Wang et al., 2024e), CC-OCR (Yang et al., 2024b), AI2D (Kembhavi et al., 2016), MMLongBench-Doc (Ma et al., 2024), and OCRBench (Liu et al., 2024). Furthermore, the model’s spatial intelligence is specifically tested on ERQA (Team, 2025b), CountBench (Paiss et al., 2023), RefCOCO (Kazemzadeh et al., 2014), ODInW13 (Li et al., 2022), and EmbSpatialBench (Du et al., 2024a). To evaluate performance on dynamic visual data, we report results on six video understanding benchmarks: Video-MME (Fu et al., 2024), MLVU (Zhou et al., 2025a), MVBench (Li et al., 2024), LVBench (Wang et al., 2024b), MMVU (Zhao et al., 2025) and MME-VideoOCR (Shi et al., 2025). Specifically, we evaluate the model’s performance on medical VQA across three established benchmarks: SLAKE (Liu et al., 2021), PMC-VQA (Zhang et al., 2023), and MedXpertQA-MM (Zuo et al., 2025). This assessment is designed to demonstrate the model’s comprehensive clinical reasoning capabilities and its potential utility as a reliable healthcare AI assistant.

Audio-Visual Video $\rightarrow$ Text We evaluate our model’s audio-visual understanding capabilities from multiple perspectives. For text-query evaluation, we use DailyOmni (Zhou et al., 2025b), WorldSense (Hong

et al., 2025), AVUT (Yang et al., 2025b), AV-SpeakerBench (Nguyen et al., 2025), and VideoMME (Fu et al., 2025). To assess the model’s ability in real-world audio-visual interactive scenarios, we use Qualcomm IVD (Pourreza et al., 2025) as the benchmark for audio-query-based evaluation. Beyond understanding, we also evaluate the model’s captioning capability on OmniCloze (Ma et al., 2025b) and its tool-use ability on OmniGAIA (Li et al., 2026).

5.1.1 Performance of Text→Text

We compare Qwen3.5-Omni-Plus and Qwen3.5-Omni-Flash with Qwen3.5-Plus-Instruct. As shown in Table 4, Qwen3.5-Omni-Plus demonstrates text capabilities that are on par with its text-only counterpart across multiple dimensions, including knowledge, instruction following, long-context understanding, STEM, reasoning, and general agent tasks, highlighting its strong language ability. In particular, Qwen3.5-Omni’s instruction-following performance is slightly better than the baseline. We believe that OPD and interaction-aligned RL have a positive effect on improving the instruction-following capabilities of an omni-model LLM.

Table 4: Text → Text performance of Qwen3.5-Omni and Qwen3.5-Plus-Instruct. The highest scores are shown in bold.

Datasets	Qwen3.5-Plus-Instruct	Qwen3.5-Omni-Flash	Qwen3.5-Omni-Plus
<i>Knowledge</i>			
MMLU-Pro	86.8	79.9	85.9
MMLU-Redux	94.3	90.0	94.2
SuperGPQA	67.4	54.9	66.4
C-Eval	92.3	86.0	92.0
<i>Instruction Following</i>			
IFEval	89.7	85.2	89.7
IFBench	51.1	38.4	52.6
<i>Long Context</i>			
AA-LCR	62.0	46.0	57.0
LongBench v2	60.2	46.4	59.6
<i>STEM</i>			
GPQA	85.9	76.4	83.9
<i>Reasoning</i>			
LiveCodeBench v6	67.1	56.6	65.6
HMMT Nov 25	86.2	59.0	84.4
IMOAnswerBench	68.3	51.5	65.5
<i>General Agent</i>			
BFCL-V4	66.1	55.3	63.3
TAU2Bench	82.7	78.0	81.0

5.1.2 Performance of Audio→Text

In Table 5, we compare Qwen3.5-Omni with Gemini-3.1 Pro in terms of audio-to-text performance. Compared to Gemini-3.1 Pro, Qwen3.5-Omni exhibits superior performance on MMAU, MMSU, RUL-MuchoMusic, and SongFormBench, while achieving comparable results on MMAR, demonstrating its strong comprehension capabilities across multiple audio domains. Regarding end-to-end speech dialogue, Qwen3.5-Omni significantly outperforms Gemini-3.1 Pro on VoiceBench and matches its performance on other benchmarks, further validating Qwen3.5-Omni’s robust capabilities in end-to-end voice interaction. For S2TT and ASR, Qwen3.5-Omni consistently outperforms Gemini-3.1 Pro, underscoring its superior translation and speech recognition performance across diverse languages, dialects, and domains.

5.1.3 Performance of Vision → Text

To comprehensively evaluate vision-to-text capabilities, we compare Qwen3.5-Omni-Flash and Qwen3.5-Omni-Plus with Qwen3.5-Plus-Instruct. As shown in Table 6, Qwen3.5-Omni-Plus achieves performance comparable to that of Qwen3.5-Plus-Instruct, while demonstrating stronger results on video understanding tasks involving both short and long videos. These findings highlight the strong dynamic visual

Table 5: Audio benchmark comparison across Gemini-3.1 Pro, Qwen3.5-Omni-Flash, and Qwen3.5-Omni-Plus. For most benchmarks, higher is better. For ASR benchmarks, lower WER is better. Best results are shown in bold.

Datasets	Gemini-3.1 Pro	Qwen3.5-Omni-Flash	Qwen3.5-Omni-Plus
<i>Audio Understanding (↑)</i>			
MMAU	81.1	80.4	82.2
MMAR	83.7	74.0	80.0
MMSU	81.3	72.2	82.8
RUL-MuchoMusic	59.6	60.5	72.4
SongFormBench-HarmonixSet _(acc hr.5f hr3f) ^a	75.6 46.8 77.9	80.6 67.8 83.4	81.1 72.9 85.3
SongFormBench-CN _(acc hr.5f hr3f) ^a	78.1 43.2 71.9	86.7 66.4 84.6	87.1 65.7 84.2
<i>Dialogue (↑)</i>			
VoiceBench	88.9	87.8	93.1
URO-Bench-pro _(U R O) ^b	69.1 84.0 99.2	64.1 83.8 98.7	66.3 86.3 99.8
SpeechRole	124.2	119.8	123.5
WildSpeech-Bench	76.3	72.2	75.4
<i>S2TT (↑)</i>			
Fleurs _{xx↔zh} (top59) ^c	29.5	26.9	30.2
Fleurs _{xx↔en} (top59) ^c	34.6	32.0	35.4
Fleurs _{xx↔zh/en} (top59) ^c	32.1	29.4	32.8
<i>ASR (WER↓)</i>			
Fleurs _(top60)	7.32	10.75	6.55
CV15 _(zh yue zh-tw)	8.59 13.40 6.78	4.25 3.45 2.68	3.46 1.95 2.27
CV15 _(en)	8.73	5.90	4.83
Librispeech _(clean other)	3.36 4.41	1.30 2.43	1.11 2.23
Weneetspeech _(net meeting)	11.53 14.21	4.41 5.51	4.30 5.84
Kespeech	23.67	4.47	3.46
MIR-1K _(vocal-only) ^d	8.76	4.94	4.56
Opencpop	6.83	1.11	1.49

^a SongFormBench: We use a unified prompt defining an SRT-like output timestamp format and a closed vocabulary for evaluation. The vocabulary follows the *SongForm-HX-8Class* specified in the official codebase.

^b URO-Bench-Pro: We use the pro track of URO-Bench and denote the three evaluation dimensions as follows: U for Understanding, R for Reasoning, and O for Oral Conversation. We use GenStyle-en, GenStyle-zh, Multilingual tasks for oral dimension.

^c Fleurs: The top59 languages are English, Chinese, Cantonese, Korean, Japanese, Vietnamese, Thai, Malay, German, Russian, Italian, French, Spanish, Portuguese, Dutch, Indonesian, Turkish, Arabic, Polish, Hindi, Urdu, Filipino, Persian, Czech, Greek, Swedish, Hebrew, Danish, Finnish, Norwegian, Icelandic, Bengali, Punjabi, Javanese, Marathi, Swahili, Ukrainian, Gujarati, Kannada, Azerbaijani, Malayalam, Cebuano, Romanian, Hungarian, Bulgarian, Belarusian, Catalan, Tamil, Croatian, Bosnian, Slovak, Galician, Kyrgyz, Macedonian, Slovenian, Latvian, Estonian, and Asturian; compared with the top60 list, Afrikaans is excluded because the Fleurs S2TT test set does not cover this language.

^d MIR-1K: Transcription is converted into Simplified Chinese.

perception ability of our model in real-world scenarios and suggest the effectiveness of joint video-audio training paradigms. Furthermore, we posit that audio-visual streams constitute the most naturalistic representation of real-world phenomena, wherein visual and auditory modalities are intrinsically coupled rather than independently processed.

5.1.4 Performance of Audio-Visual Video→Text

We compare Qwen3.5-Omni and Gemini-3.1 Pro across a diverse range of audio-visual tasks, as shown in Table 7. For general understanding, Qwen3.5-Omni achieves state-of-the-art performance on DailyOmni and obtains comparable results on AVUT. Our model also surpasses Gemini-3.1-Pro by a substantial margin on Qualcomm IVD, demonstrating its effectiveness in real-world audio-visual interactive scenarios. Moreover, our model shows strong performance on captioning tasks. It can provide detailed audio, visual, and audio-visual captions. In this version, we also enhance the model’s tool-use capability, achieving 57.2% on OmniGAIA.

Table 6: Vision $\rightarrow$ Text performance of Qwen3.5-Omni and Qwen3.5-Plus-Instruct. The highest scores are shown in bold.

Datasets	Qwen3.5-Plus-Instruct	Qwen3.5-Omni-Flash	Qwen3.5-Omni-Plus
<i>STEM and Puzzle</i>			
MMMU	81.0	76.9	80.1
MMMU-Pro	73.8	68.2	73.9
MathVision	73.6	65.4	73.0
Mathvista (mini)	86.9	82.9	86.1
DynaMath	84.2	79.3	83.8
ZEROBench	6	1	5
ZEROBench_sub	31.1	26.0	34.4
<i>General VQA</i>			
RealWorldQA	79.1	77.5	84.1
MMStar	80.3	75.7	79.4
MMBench _{EN-DEV-v1.1}	93.8	88.8	92.8
SimpleVQA	66.1	54.4	65.3
<i>Text Recognition and Document Understanding</i>			
CharXiv (RQ)	74.2	64.4	72.5
CC-OCR	83.0	80.8	83.4
AI2D_TEST	92.1	89.0	91.2
MMLongBench-Doc	59.7	53.6	57.5
OCRBench	91.4	89.1	91.3
<i>Spatial Intelligence</i>			
ERQA	53.8	50.0	54.8
CountBench	95.1	88.2	95.1
RefCOCO(avg)	95.2	92.6	95.0
ODInW13	50.3	46.8	49.5
EmbSpatialBench	83.4	82.7	85.4
<i>Video Understanding</i>			
VideoMME _(w/o sub.)	81.0	77.0	81.9
MLVU _(M-Avg)	85.1	81.9	86.8
MVBench	76.7	70.8	79.0
LVBench	68.6	65.7	71.2
MMVU	67.1	62.7	67.5
MME-VideoOCR	74.2	70.5	77.0
<i>Medical VQA</i>			
SLAKE	82.8	73.1	84.7
PMC-VQA	62.4	58.7	62.7
MedXpertQA-MM	55.3	44.8	54.7

5.2 Evaluation of X $\rightarrow$ Speech

In this section, we evaluate the speech generation capability of Qwen3.5-Omni. Our evaluation mainly focuses on speech generation conditioned on text and prompt speech, following a zero-shot text-to-speech (TTS) setting. We study the model from four perspectives:

- **Zero-Shot Speech Generation:** We evaluate content consistency, measured by WER, on SEED (Anastassiou et al., 2024).
- **Multilingual Speech Generation:** We evaluate both content consistency and speaker similarity in zero-shot multilingual speech generation on the TTS multilingual test set (Zhang et al., 2025a) and an internal multilingual test set built on FLEURS (Conneau et al., 2022).
- **Cross-Lingual Speech Generation:** We evaluate content consistency in zero-shot cross-lingual speech generation on CV3-Eval (Du et al., 2025).
- **Custom-Voice Speech Generation:** We evaluate the stability of our speaker fine-tuned model on the TTS multilingual test set (Zhang et al., 2025a) and our internal multilingual test set.

Table 7: **Audio-Visual** → **Text** performance of Qwen3.5-Omni and Gemini-3.1-Pro. The highest scores are shown in bold.

Datasets	Gemini-3.1 Pro	Qwen3.5-Omni-Flash	Qwen3.5-Omni-Plus
<i>Text Query QA</i>			
DailyOmni	82.7	81.8	84.6
WorldSense	65.5	57.9	62.8
AVUT	85.6	81.4	85.0
AV-SpeakerBench	75.1	65.2	71.3
VideoMME _{w/ audio} ^a	89.0	79.3	83.7
<i>Audio Query QA</i>			
Qualcomm IVD	66.2	66.3	68.5
<i>Caption</i>			
Omni-Cloze	57.2	63.0	64.8
<i>Agent (Tool Use)</i>			
OmniGAIA ^b	68.9	33.9	57.2

^a VideoMME is evaluated with use_audio_in_video=True.

^b OmniGAIA is evaluated without a thinking prompt and without <answer> formatting. All results are evaluated using DeepSeek-V3.2-Thinking as the judge.

5.2.1 Evaluation of Zero-Shot Speech Generation

We compare Qwen3.5-Omni with state-of-the-art zero-shot TTS systems. As shown in Table 8, Qwen3.5-Omni achieves highly competitive performance on the SEED-TTS benchmark, demonstrating strong content fidelity in zero-shot speech generation. These results reflect the effectiveness of our pretraining and continual pretraining pipeline in building robust speech generation and context modeling capabilities. Moreover, after RLHF optimization, Qwen3.5-Omni further improves generation stability and naturalness, achieving the best performance on the *test-en* split with a WER of 1.26.

Table 8: **Zero-shot speech generation on the SEED-TTS test set.** Performance is measured by Word Error Rate (WER, ↓), where lower is better. The best results are highlighted in bold.

Datasets	Model	Performance
<i>Content Consistency</i>		
SEED <i>test-zh</i> <i>test-en</i>	Seed-TTS _{ICL} (Anastassiou et al., 2024)	1.11 2.24
	Seed-TTS _{RL} (Anastassiou et al., 2024)	1.00 1.94
	MaskGCT (Wang et al., 2024c)	2.27 2.62
	E2 TTS (Eskimez et al., 2024)	1.97 2.19
	F5-TTS (Chen et al., 2024c)	1.56 1.83
	Spark TTS (Wang et al., 2025b)	1.20 1.98
	CosyVoice 2 (Du et al., 2024b)	1.45 2.57
	CosyVoice 3 (Du et al., 2025)	0.71 1.45
	MiniMax-Speech (Zhang et al., 2025a)	0.83 1.65
	MiMo-Audio-7B-Instruct (Zhang et al., 2025c)	1.96 5.37
	Qwen2.5-Omni-7B (Xu et al., 2025a)	1.42 2.33
	Qwen3-Omni-30B-A3B (Xu et al., 2025b)	1.07 1.39
	Qwen3.5-Omni-Plus	0.99 1.26

5.2.2 Evaluation of Multilingual Speech Generation

Qwen3.5-Omni supports speech generation in 29 languages. We compare its multilingual speech generation performance with two strong commercial systems, MiniMax-Speech and ElevenLabs. For the internal multilingual test set, we use GPT-4o-transcribe-2025-03-20 for automatic speech recognition.

As shown in Table 9 and Table 10, Qwen3.5-Omni achieves the lowest WER in 22 out of 29 evaluated languages on the multilingual test sets, outperforming the comparison systems by a clear margin in most cases. On the remaining languages, Qwen3.5-Omni remains competitive with state-of-the-art systems. In addition to content consistency, Qwen3.5-Omni also shows strong voice cloning fidelity. It obtains the highest speaker similarity scores in the majority of evaluated languages and consistently

Table 9: **Multilingual speech generation on the TTS multilingual test set.** Performance is measured by Word Error Rate (WER, ↓) for content consistency and cosine similarity (SIM, ↑) for speaker similarity. The best results are highlighted in bold.

Language	Content Consistency			Speaker Similarity		
	Qwen3.5-Omni-Plus	MiniMax	ElevenLabs	Qwen3.5-Omni-Plus	MiniMax	ElevenLabs
Chinese	0.695	2.252	16.026	0.800	0.780	0.677
English	0.631	2.164	0.756	0.833	0.756	0.613
German	0.447	1.906	0.572	0.757	0.733	0.614
Italian	0.503	1.543	1.743	0.785	0.699	0.679
Portuguese	1.221	1.877	1.331	0.792	0.805	0.711
Spanish	0.862	1.029	1.084	0.797	0.762	0.615
Japanese	3.479	3.519	10.046	0.788	0.776	0.738
Korean	1.458	1.747	1.865	0.747	0.776	0.700
French	2.430	4.099	5.216	0.730	0.628	0.535
Russian	3.182	4.281	3.878	0.790	0.761	0.676
Thai	2.170	2.701	73.936	0.788	0.800	0.588
Indonesian	0.823	1.237	1.059	0.780	0.729	0.660
Arabic	2.602	1.665	1.666	0.745	0.736	0.706
Vietnamese	1.143	0.880	73.415	0.767	0.743	0.369
Turkish	0.938	1.520	0.699	0.747	0.779	0.596
Finnish	2.784	4.666	2.964	0.859	0.835	0.759
Polish	1.427	1.415	0.766	0.839	0.802	0.729
Hindi	6.444	6.962	5.827	0.797	0.818	0.730
Dutch	1.238	1.143	0.803	0.762	0.738	0.680
Czech	2.929	3.875	2.108	0.802	0.796	0.685

outperforms both MiniMax-Speech and ElevenLabs overall. These results suggest that Qwen3.5-Omni effectively preserves speaker characteristics, such as timbre and prosodic style, while maintaining robust multilingual speech generation quality.

Furthermore, in Table 10, we report results on our internal multilingual test set, covering an additional 9 languages. Qwen3.5-Omni continues to achieve strong performance across all evaluated languages, indicating that its multilingual speech generation ability generalizes well beyond the public benchmark languages.

Table 10: **Multilingual speech generation on the internal multilingual test set.** Performance is measured by Word Error Rate (WER, ↓) for content consistency and cosine similarity (SIM, ↑) for speaker similarity. The best results are highlighted in bold.

Language	Content Consistency		Speaker Similarity	
	Qwen3.5-Omni-Plus	Ground Truth	Qwen3.5-Omni-Plus	Ground Truth
Urdu	14.819	17.822	0.775	-
Tagalog	5.193	6.885	0.870	-
Swedish	3.760	4.813	0.822	-
Danish	3.636	6.403	0.775	-
Hebrew	7.860	16.178	0.760	-
Icelandic	10.244	11.451	0.764	-
Malay	3.142	4.628	0.794	-
Norwegian	3.613	4.442	0.825	-
Persian	11.113	14.469	0.800	-

5.2.3 Evaluation of Cross-Lingual Speech Generation

Beyond multilingual voice cloning, Qwen3.5-Omni also supports cross-lingual voice cloning, where the model is required to preserve speaker identity while generating speech in a different target language. We evaluate this capability on the Cross-Lingual benchmark and compare against the CosyVoice series as well as Qwen3-Omni-30B-A3B.

In Table 11, we report the mixed error rate (WER for English and CER for the other languages) across different source–target language pairs. Overall, Qwen3.5-Omni achieves the best performance in 10 out of 12 evaluated directions and sets a new state of the art on most English-, Japanese-, and Korean-targeted pairs. In particular, for *zh-to-ko*, Qwen3.5-Omni reduces the error rate from 14.4 to 4.03 compared with

Table 11: **Cross-lingual speech generation on the Cross-Lingual benchmark.** Performance is measured by mixed error rate (WER for English and CER for the other languages, ↓). The best results are highlighted in bold.

Language	Qwen3.5-Omni-Plus	Qwen3-Omni-30B-A3B	CosyVoice3	CosyVoice2
English-to-Chinese	4.86	5.37	5.09	13.5
Japanese-to-Chinese	3.55	3.32	3.05	48.1
Korean-to-Chinese	0.84	0.99	1.06	7.70
Chinese-to-English	2.18	2.76	2.98	6.47
Japanese-to-English	2.18	3.31	4.20	17.1
Korean-to-English	2.51	3.34	4.19	11.2
Chinese-to-Japanese	5.92	8.29	7.08	13.1
English-to-Japanese	5.12	7.53	6.80	14.9
Korean-to-Japanese	2.16	4.24	3.93	5.86
Chinese-to-Korean	4.03	5.13	14.4	24.8
English-to-Korean	3.72	4.96	5.87	21.9
Japanese-to-Korean	5.12	6.23	7.92	21.5

CosyVoice3, corresponding to an approximately 72% relative reduction. Qwen3.5-Omni also performs strongly on commonly used language pairs such as *zh-to-en* and *en-to-zh*, indicating better content consistency under cross-lingual generation. These results demonstrate that Qwen3.5-Omni generalizes effectively across language boundaries while preserving target linguistic accuracy.

5.2.4 Evaluation of Custom-Voice Speech Generation

We evaluate the custom-voice speech generation capability of Qwen3.5-Omni in multilingual settings. We compare Qwen3.5-Omni with several strong commercial systems accessed through their official APIs in March 2026, including ElevenLabs Multilingual v2 (9YHcvj6GT2YYXdxww), Gemini-2.5 Pro-Preview-TTS (Achernar), GPT-Audio-2025-08-28 (Alloy), and MiniMax-Speech-2.8-HD (English_expressive_narrator).

Table 12: **Custom-voice multilingual speech generation on the multilingual test set.** Performance is measured by Word Error Rate (WER, ↓). The best results are highlighted in bold.

Language	Qwen3.5-Omni-Plus	ElevenLabs	Gemini-2.5 Pro	GPT-Audio	MiniMax
Chinese	0.785	3.801	1.890	0.829	0.786
English	0.839	1.126	0.953	1.050	1.429
German	0.182	0.500	0.509	0.558	1.581
Italian	0.458	0.513	0.991	0.769	1.063
Portuguese	1.581	1.109	2.050	1.506	1.240
Spanish	0.768	0.520	0.891	0.936	0.691
Japanese	3.306	11.685	4.420	4.317	4.254
Korean	1.309	3.981	4.110	3.999	3.635
French	2.724	2.574	3.284	2.809	3.439
Russian	4.723	4.324	3.858	4.346	3.529
Thai	1.653	114.813	2.539	4.430	1.811
Indonesian	1.596	6.094	1.498	2.362	1.585
Arabic	3.183	5.400	5.525	5.326	3.309
Vietnamese	1.320	82.849	1.699	1.854	1.058
Turkish	1.309	0.551	2.237	1.389	0.652
Finnish	4.039	2.522	5.331	3.270	2.939
Polish	1.462	0.733	1.622	1.737	0.833
Hindi	6.776	6.388	6.596	7.191	6.146
Dutch	1.135	1.005	0.973	1.561	1.406
Czech	3.769	1.916	3.380	2.859	1.766
Urdu	14.916	12.970	14.141	13.362	24.151
Tagalog	5.090	5.473	6.784	5.352	5.674
Swedish	3.588	3.132	3.196	2.898	2.833
Danish	7.183	2.604	3.876	3.846	4.951
Hebrew	7.680	102.018	4.459	5.328	8.161
Icelandic	10.322	25.110	6.348	9.648	33.431
Malay	3.738	6.448	3.731	3.406	3.955
Norwegian	5.576	7.351	4.304	3.400	9.492
Persian	12.140	20.564	12.620	13.202	12.722

As shown in Table 12, although Qwen3.5-Omni is fine-tuned only on monolingual data, it demonstrates strong cross-lingual generalization in custom-voice speech generation. The model is able to transfer the target speaker characteristics to all 29 evaluated languages while maintaining stable generation quality. Overall, Qwen3.5-Omni achieves the best WER in 10 languages and remains competitive in many others. In particular, it shows clear advantages in several challenging languages, including Japanese (3.306) and Korean (1.309), indicating strong intelligibility under cross-lingual voice transfer. These results suggest that Qwen3.5-Omni can generate custom-voice speech with robust linguistic fidelity across a wide range of languages.

6 Conclusion

In this work, we present Qwen3.5-Omni, a fully omnimodal large language model that unifies understanding, reasoning, generation, and action across text, images, audio, and audio-visual inputs. Built on the Thinker-Talker framework, Qwen3.5-Omni introduces efficient Hybrid-Attention MoE architectures, 256k long-context modeling, improved streaming speech generation with multi-codebook codec prediction and ARIA, and substantially expanded multilingual speech support. These advances enable three key capabilities: controllable audio-visual captioning, comprehensive real-time interaction, and native omnimodal agentic behavior through autonomous tool use and audio-visual code generation. Empirically, Qwen3.5-Omni achieves state-of-the-art or highly competitive performance across a broad range of audio and audio-visual benchmarks, while maintaining the strong text and vision capabilities of same-scale Qwen models. These results suggest that scaling native omnimodal training can produce unified systems that not only perceive and reason across modalities, but also interact and act in real time. We hope Qwen3.5-Omni provides a strong foundation for future research on general-purpose omnimodal agents.

References

- Philip Anastassiou, Jiawei Chen, Jitong Chen, Yuanzhe Chen, Zhuo Chen, Ziyi Chen, Jian Cong, Lelai Deng, Chuang Ding, Lu Gao, et al. Seed-tts: A family of high-quality versatile speech generation models. *arXiv preprint arXiv:2406.02430*, 2024.
- Anthropic. Introducing Claude, 2023a. URL <https://www.anthropic.com/index/introducing-claude>.
- Anthropic. Claude 2. Technical report, Anthropic, 2023b. URL <https://www-files.anthropic.com/production/images/Model-Card-Claude-2.pdf>.
- Anthropic. The Claude 3 model family: Opus, Sonnet, Haiku. Technical report, Anthropic, AI, 2024. URL https://www-cdn.anthropic.com/de8ba9b01c9ab7cbabf5c33b80b7bbc618857627/Model_Card_Claude_3.pdf.
- Rosana Ardila, Megan Branson, Kelly Davis, Michael Kohler, Josh Meyer, Michael Henretty, Reuben Morais, Lindsay Saunders, Francis M. Tyers, and Gregor Weber. Common voice: A massively-multilingual speech corpus. In Nicoletta Calzolari, Frédéric Béchet, Philippe Blache, Khalid Choukri, Christopher Cieri, Thierry Declerck, Sara Goggi, Hitoshi Isahara, Bente Maegaard, Joseph Mariani, Hélène Mazo, Asunción Moreno, Jan Odijk, and Stelios Piperidis (eds.), *Proceedings of The 12th Language Resources and Evaluation Conference, LREC 2020, Marseille, France, May 11-16, 2020*, pp. 4218–4222. European Language Resources Association, 2020. URL <https://aclanthology.org/2020.lrec-1.520/>.
- Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. Qwen technical report. *CoRR*, abs/2309.16609, 2023a.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-VL: A frontier large vision-language model with versatile abilities. *CoRR*, abs/2308.12966, 2023b.
- Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2. 5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025a.
- Yushi Bai, Shangqing Tu, Jiajie Zhang, Hao Peng, Xiaozhi Wang, Xin Lv, Shulin Cao, Jiazheng Xu, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. Longbench v2: Towards deeper understanding and reasoning on realistic long-context multitasks. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2025, Vienna, Austria, July 27 - August 1, 2025*, pp. 3639–3664. Association for Computational Linguistics, 2025b. URL <https://aclanthology.org/2025.acl-long.183/>.
- Mislav Balunović, Jasper Dekoninck, Ivo Petrov, Nikola Jovanović, and Martin Vechev. Matharena: Evaluating llms on uncontaminated math competitions, February 2025. URL <https://matharena.ai/>.
- Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. τ^2 -bench: Evaluating conversational agents in a dual-control environment, 2025. URL <https://arxiv.org/abs/2506.07982>.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. In *NeurIPS*, 2020.
- Lin Chen, Jinsong Li, Xiaoyi Dong, Pan Zhang, Yuhang Zang, Zehui Chen, Haodong Duan, Jiaqi Wang, Yu Qiao, Dahua Lin, et al. Are we on the right way for evaluating large vision-language models? *arXiv:2403.20330*, 2024a.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T Tan, and Haizhou Li. Voicebench: Benchmarking llm-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024b.
- Yushen Chen, Zhikang Niu, Ziyang Ma, Keqi Deng, Chunhui Wang, Jian Zhao, Kai Yu, and Xie Chen. F5-tts: A fairytaler that fakes fluent and faithful speech with flow matching. *arXiv preprint arXiv:2410.06885*, 2024c.

-
- Xianfu Cheng, Wei Zhang, Shiwei Zhang, Jian Yang, Xiangyuan Guan, Xianjie Wu, Xiang Li, Ge Zhang, Jiaheng Liu, Yuying Mai, Yutao Zeng, Zhoufutu Wen, Ke Jin, Baorui Wang, Weixiao Zhou, Yunhong Lu, Tongliang Li, Wenhao Huang, and Zhoujun Li. Simplevqa: Multimodal factuality evaluation for multimodal large language models. *CoRR*, abs/2502.13059, 2025.
- Yunfei Chu, Jin Xu, Xiaohuan Zhou, Qian Yang, Shiliang Zhang, Zhijie Yan, Chang Zhou, and Jingren Zhou. Qwen-Audio: Advancing universal audio understanding via unified large-scale audio-language models. *CoRR*, abs/2311.07919, 2023.
- Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, Yichong Leng, Yuanjun Lv, Jinzheng He, Junyang Lin, et al. Qwen2-audio technical report. *arXiv preprint arXiv:2407.10759*, 2024.
- Gheorghe Comanici, Eric Bieber, Mike Schaekermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.
- Alexis Conneau, Min Ma, Simran Khanuja, Yu Zhang, Vera Axelrod, Siddharth Dalmia, Jason Riesa, Clara Rivera, and Ankur Bapna. Fleurs: Few-shot learning evaluation of universal representations of speech. *2022 IEEE Spoken Language Technology Workshop (SLT)*, pp. 798–805, 2022. URL <https://api.semanticscholar.org/CorpusID:249062909>.
- Mengfei Du, Binhao Wu, Zejun Li, Xuanjing Huang, and Zhongyu Wei. Embspatial-bench: Benchmarking spatial understanding for embodied tasks with large vision-language models. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, ACL 2024, Bangkok, Thailand, August 11-16, 2024, pp. 346–355. Association for Computational Linguistics, 2024a.
- Zhihao Du, Yuxuan Wang, Qian Chen, Xian Shi, Xiang Lv, Tianyu Zhao, Zhifu Gao, Yexin Yang, Changfeng Gao, Hui Wang, et al. Cosyvoice 2: Scalable streaming speech synthesis with large language models. *arXiv preprint arXiv:2412.10117*, 2024b.
- Zhihao Du, Changfeng Gao, Yuxuan Wang, Fan Yu, Tianyu Zhao, Hao Wang, Xiang Lv, Hui Wang, Chongjia Ni, Xian Shi, Keyu An, Guanrou Yang, Yabin Li, Yanni Chen, Zhifu Gao, Qian Chen, Yue Gu, Mengzhe Chen, Yafeng Chen, Shiliang Zhang, Wen Wang, and Jieping Ye. Cosyvoice 3: Towards in-the-wild speech generation via scaling-up and post-training. *CoRR*, abs/2505.17589, 2025.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurélien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Rozière, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Graeme Nail, Grégoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel M. Kloumann, Ishan Misra, Ivan Evtimov, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Kartikeya Upasani, Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, and et al. The Llama 3 herd of models. *CoRR*, abs/2407.21783, 2024.
- Sefik Emre Eskimez, Xiaofei Wang, Manthan Thakker, Canrun Li, Chung-Hsien Tsai, Zhen Xiao, Hemin Yang, Zirun Zhu, Min Tang, Xu Tan, et al. E2 tts: Embarrassingly easy fully non-autoregressive zero-shot tts. In *2024 IEEE Spoken Language Technology Workshop (SLT)*, pp. 682–689. IEEE, 2024.
- Chaoyou Fu, Yuhan Dai, Yondong Luo, Lei Li, Shuhuai Ren, Renrui Zhang, Zihan Wang, Chenyu Zhou, Yunhang Shen, Mengdan Zhang, et al. Video-mme: The first-ever comprehensive evaluation benchmark of multi-modal llms in video analysis. *arXiv:2405.21075*, 2024.
- Chaoyou Fu, Yuhan Dai, Yongdong Luo, Lei Li, Shuhuai Ren, Renrui Zhang, Zihan Wang, Chenyu Zhou, Yunhang Shen, Mengdan Zhang, et al. Video-mme: The first-ever comprehensive evaluation benchmark of multi-modal llms in video analysis. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 24108–24118, 2025.

-
- Aryo Pradipta Gema, Joshua Ong Jun Leang, Giwon Hong, Alessio Devoto, Alberto Carlo Maria Mancino, Rohit Saxena, Xuanli He, Yu Zhao, Xiaotang Du, Mohammad Reza Ghasemi Madani, et al. Are we done with mmlu? *CoRR*, abs/2406.04127, 2024.
- Gemini Team. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. Technical report, Google, 2024. URL https://storage.googleapis.com/deepmind-media/gemini/gemini_v1_5_report.pdf.
- Tianrui Guan, Fuxiao Liu, Xiyang Wu, Ruiqi Xian, Zongxia Li, Xiaoyu Liu, Xijun Wang, Lichang Chen, Furong Huang, Yaser Yacoob, Dinesh Manocha, and Tianyi Zhou. Hallusionbench: An advanced diagnostic suite for entangled language hallucination and visual illusion in large vision-language models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2024, Seattle, WA, USA, June 16-22, 2024*, pp. 14375–14385. IEEE, 2024.
- Chunbo Hao, Ruibin Yuan, Jixun Yao, Qixin Deng, Xinyi Bai, Wei Xue, and Lei Xie. Songformer: Scaling music structure analysis with heterogeneous supervision, 2025. URL <https://arxiv.org/abs/2510.02797>.
- Jack Hong, Shilin Yan, Jiayin Cai, Xiaolong Jiang, Yao Hu, and Weidi Xie. Worldsense: Evaluating real-world omnimodal understanding for multimodal llms. *CoRR*, abs/2502.04326, 2025.
- Chao-Ling Hsu and Jyh-Shing Roger Jang. On the improvement of singing voice separation for monaural recordings using the MIR-1K dataset. *IEEE Trans. Speech Audio Process.*, 18(2):310–319, 2010. doi: 10.1109/TASL.2009.2026503. URL <https://doi.org/10.1109/TASL.2009.2026503>.
- Yuzhen Huang, Yuzhuo Bai, Zhihao Zhu, Junlei Zhang, Jinghan Zhang, Tangjun Su, Junteng Liu, Chuancheng Lv, Yikai Zhang, Jiayi Lei, Yao Fu, Maosong Sun, and Junxian He. C-Eval: A multi-level multi-discipline chinese evaluation suite for foundation models. In *NeurIPS*, 2023.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. LiveCodeBench: Holistic and contamination free evaluation of large language models for code. *CoRR*, abs/2403.07974, 2024.
- Changhao Jiang, Jiajun Sun, Yifei Cao, Jiabao Zhuang, Hui Li, Xiaoran Fan, Ming Zhang, Junjie Ye, Shihan Dou, Zhiheng Xi, et al. Speechrole: A large-scale dataset and benchmark for evaluating speech role-playing agents. *arXiv preprint arXiv:2508.02013*, 2025.
- Sahar Kazemzadeh, Vicente Ordonez, Mark Matten, and Tamara Berg. Referitgame: Referring to objects in photographs of natural scenes. In *EMNLP*, 2014.
- Aniruddha Kembhavi, Mike Salvato, Eric Kolve, Minjoon Seo, Hannaneh Hajishirzi, and Ali Farhadi. A diagram is worth a dozen images. In *ECCV*, 2016.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. *arXiv:2301.12597*, 2023.
- Kunchang Li, Yali Wang, Yinan He, Yizhuo Li, Yi Wang, Yi Liu, Zun Wang, Jilan Xu, Guo Chen, Ping Luo, et al. Mvbench: A comprehensive multi-modal video understanding benchmark. In *CVPR*, 2024.
- Liunian Harold Li, Pengchuan Zhang, Haotian Zhang, Jianwei Yang, Chunyuan Li, Yiwu Zhong, Lijuan Wang, Lu Yuan, Lei Zhang, Jenq-Neng Hwang, Kai-Wei Chang, and Jianfeng Gao. Grounded language-image pre-training. In *CVPR*, pp. 10955–10965. IEEE, 2022.
- Xiaoxi Li, Wenxiang Jiao, Jiarui Jin, Shijian Wang, Guanting Dong, Jiajie Jin, Hao Wang, Yinuo Wang, Ji-Rong Wen, Yuan Lu, et al. Omnigaia: Towards native omni-modal ai agents. *arXiv preprint arXiv:2602.22897*, 2026.
- Bo Liu, Li-Ming Zhan, Li Xu, Lin Ma, Yan Yang, and Xiao-Ming Wu. Slake: A semantically-labeled knowledge-enhanced dataset for medical visual question answering. In *18th IEEE International Symposium on Biomedical Imaging, ISBI 2021, Nice, France, April 13-16, 2021*, pp. 1650–1654. IEEE, 2021.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. *arXiv:2304.08485*, 2023.
- Yuliang Liu, Zhang Li, Mingxin Huang, Biao Yang, Wenwen Yu, Chunyuan Li, Xu-Cheng Yin, Cheng-Lin Liu, Lianwen Jin, and Xiang Bai. Ocrbench: on the hidden mystery of ocr in large multimodal models. *Science China Information Sciences*, 67(12), December 2024. ISSN 1869-1919. doi: 10.1007/s11432-024-4235-6. URL <http://dx.doi.org/10.1007/s11432-024-4235-6>.

-
- Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. Mathvista: Evaluating mathematical reasoning of foundation models in visual contexts. In *ICLR*, 2024.
- Thang Luong, Dawsen Hwang, Hoang H. Nguyen, Golnaz Ghiasi, Yuri Chervonyi, Insuk Seo, Junsu Kim, Garrett Bingham, Jonathan Lee, Swaroop Mishra, Alex Zhai, Clara Huiyi Hu, Henryk Michalewski, Jimin Kim, Jeonghyun Ahn, Junhwi Bae, Xingyou Song, Trieu H. Trinh, Quoc V. Le, and Junehyuk Jung. Towards robust mathematical reasoning. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. URL <https://aclanthology.org/2025.emnlp-main.1794/>.
- Yubo Ma, Yuhang Zang, Liangyu Chen, Meiqi Chen, Yizhu Jiao, Xinze Li, Xinyuan Lu, Ziyu Liu, Yan Ma, Xiaoyi Dong, Pan Zhang, Liangming Pan, Yu-Gang Jiang, Jiaqi Wang, Yixin Cao, and Aixin Sun. MMLONGBENCH-DOC: benchmarking long-context document understanding with visualizations. In Amir Globersons, Lester Mackey, Danielle Belgrave, Angela Fan, Ulrich Paquet, Jakub M. Tomczak, and Cheng Zhang (eds.), *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*, 2024.
- Ziyang Ma, Yinghao Ma, Yanqiao Zhu, Chen Yang, Yi-Wen Chao, Ruiyang Xu, Wenxi Chen, Yuanzhe Chen, Zhuo Chen, Jian Cong, Kai Li, Keliang Li, Siyou Li, Xinfeng Li, Xiquan Li, Zheng Lian, Yuzhe Liang, Minghao Liu, Zhikang Niu, Tianrui Wang, Yuping Wang, Yuxuan Wang, Yihao Wu, Guanrou Yang, Jianwei Yu, Ruibin Yuan, Zhisheng Zheng, Ziya Zhou, Haina Zhu, Wei Xue, Emmanouil Benetos, Kai Yu, Chng Eng Siong, and Xie Chen. MMAR: A challenging benchmark for deep reasoning in speech, audio, music, and their mix. *CoRR*, abs/2505.13032, 2025a. doi: 10.48550/ARXIV.2505.13032. URL <https://doi.org/10.48550/arXiv.2505.13032>.
- Ziyang Ma, Ruiyang Xu, Zhenghao Xing, Yunfei Chu, Yuxuan Wang, Jinzheng He, Jin Xu, Pheng-Ann Heng, Kai Yu, Junyang Lin, et al. Omni-captioner: Data pipeline, models, and benchmark for omni detailed perception. *arXiv preprint arXiv:2510.12720*, 2025b.
- Le Thien Phuc Nguyen, Zhuoran Yu, Samuel Low Yu Hang, Subin An, Jeongik Lee, Yohan Ban, SeungEun Chung, Thanh-Huy Nguyen, JuWan Maeng, Soochahn Lee, et al. See, hear, and understand: Benchmarking audiovisual human speech understanding in multimodal large language models. *arXiv preprint arXiv:2512.02231*, 2025.
- OpenAI. ChatML, 2022. URL <https://github.com/openai/openai-python/blob/e389823ba013a24b4c32ce38fa0bd87e6bccae94/chatml.md>.
- OpenAI. GPT4 technical report. *CoRR*, abs/2303.08774, 2023.
- OpenAI. Hello GPT-4o, 2024. URL <https://openai.com/index/hello-gpt-4o/>.
- Roni Paiss, Ariel Ephrat, Omer Tov, Shiran Zada, Inbar Mosseri, Michal Irani, and Tali Dekel. Teaching CLIP to count to ten. In *IEEE/CVF International Conference on Computer Vision, ICCV 2023, Paris, France, October 1-6, 2023*, pp. 3147–3157. IEEE, 2023.
- Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur. Librispeech: An ASR corpus based on public domain audio books. In *2015 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP 2015, South Brisbane, Queensland, Australia, April 19-24, 2015*. IEEE, 2015.
- Reza Pourreza, Rishit Dagli, Apratim Bhattacharyya, Sunny Panchal, Guillaume Berger, and Roland Memisevic. Can vision-language models answer face to face questions in the real-world? *arXiv preprint arXiv:2503.19356*, 2025.
- Valentina Pyatkin, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hannaneh Hajishirzi. Generalizing verifiable instruction following. *CoRR*, abs/2507.02833, 2025. doi: 10.48550/ARXIV.2507.02833. URL <https://doi.org/10.48550/arXiv.2507.02833>.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *NeurIPS*, 2023.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A graduate-level Google-proof Q&A benchmark. *CoRR*, abs/2311.12022, 2023.

-
- Jonathan Roberts, Mohammad Reza Taesiri, Ansh Sharma, Akash Gupta, Samuel Roberts, Ioana Croitoru, Simion-Vlad Bogolin, Jialu Tang, Florian Langer, Vyas Raina, Vatsal Raina, Hanyi Xiong, Vishaal Udandara, Jingyi Lu, Shiyang Chen, Sam Purkis, Tianshuo Yan, Wenye Lin, Gyungin Shin, Qiaochu Yang, Anh Totti Nguyen, Kai Han, and Samuel Albanie. Zerobench: An impossible visual benchmark for contemporary large multimodal models. *CoRR*, abs/2502.09696, 2025.
- S Sakshi, Utkarsh Tyagi, Sonal Kumar, Ashish Seth, Ramaneswaran Selvakumar, Oriol Nieto, Ramani Duraiswami, Sreyan Ghosh, and Dinesh Manocha. Mmau: A massive multi-task audio understanding and reasoning benchmark, 2024. URL <https://arxiv.org/abs/2410.19168>.
- Yang Shi, Huanqian Wang, Wulin Xie, Huanyao Zhang, Lijie Zhao, Yifan Zhang, Xinfeng Li, Chaoyou Fu, Zhuoer Wen, Wenting Liu, Zhuoran Zhang, Xinlong Chen, Bohan Zeng, Sihang Yang, Yuanxing Zhang, Pengfei Wan, Haotian Wang, and Wenjing Yang. Mme-videoocr: Evaluating ocr-based capabilities of multimodal llms in video scenarios. *CoRR*, abs/2505.21333, 2025.
- Zhiyuan Tang, Dong Wang, Yanguang Xu, Jianwei Sun, Xiaoning Lei, Shuaijiang Zhao, Cheng Wen, Xingjun Tan, Chuandong Xie, Shuran Zhou, Rui Yan, Chenjia Lv, Yang Han, Wei Zou, and Xiangang Li. Kespeech: An open source speech dataset of mandarin and its eight subdialects. In Joaquin Vanschoren and Sai-Kit Yeung (eds.), *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks 1, NeurIPS Datasets and Benchmarks 2021, December 2021, virtual*, 2021. URL <https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/0336dcbab05b9d5ad24f4333c7658a0e-Abstract-round2.html>.
- Artificial Analysis Team. Artificial analysis long context reasoning benchmark (lcr). Artificial Analysis, Inc., 2025a. Dataset.
- Gemini Robotics Team. Gemini robotics: Bringing AI into the physical world. *CoRR*, abs/2503.20020, 2025b.
- M.-A.-P. Team, Xinrun Du, Yifan Yao, Kaijing Ma, Bingli Wang, Tianyu Zheng, Kang Zhu, Minghao Liu, Yiming Liang, Xiaolong Jin, Zhenlin Wei, Chujie Zheng, Kaixin Deng, Shian Jia, Sichao Jiang, Yiyao Liao, Rui Li, Qinru Li, Sirun Li, Yizhi Li, Yunwen Li, Dehua Ma, Yuansheng Ni, Haoran Que, Qiyao Wang, Zhoufutu Wen, Siwei Wu, Tianshun Xing, Ming Xu, Zhenzhu Yang, Zekun Moore Wang, Jun Zhou, Yuelin Bai, Xingyuan Bu, Chenglin Cai, Liang Chen, Yifan Chen, Chengtuo Cheng, Tianhao Cheng, Keyi Ding, Siming Huang, Yun Huang, Yaoru Li, Yizhe Li, Zhaoqun Li, Tianhao Liang, Chengdong Lin, Hongquan Lin, Yinghao Ma, Tianyang Pang, Zhongyuan Peng, Zifan Peng, Qige Qi, Shi Qiu, Xingwei Qu, Shanghaoran Quan, Yizhou Tan, Zili Wang, Chenqing Wang, Hao Wang, Yiya Wang, Yubo Wang, Jiajun Xu, Kexin Yang, Ruibin Yuan, Yuanhao Yue, Tianyang Zhan, Chun Zhang, Jinyang Zhang, Xiyue Zhang, Xingjian Zhang, Yue Zhang, Yongchi Zhao, Xiangyu Zheng, Chenghua Zhong, Yang Gao, Zhoujun Li, Dayiheng Liu, Qian Liu, Tianyu Liu, Shiwen Ni, Junran Peng, Yujia Qin, Wenbo Su, Guoyin Wang, Shi Wang, Jian Yang, Min Yang, Meng Cao, Xiang Yue, Zhaoxiang Zhang, Wangchunshu Zhou, Jiaheng Liu, Qunshu Lin, Wenhao Huang, and Ge Zhang. Supergpqa: Scaling LLM evaluation across 285 graduate disciplines. *CoRR*, abs/2502.14739, 2025. doi: 10.48550/ARXIV.2502.14739. URL <https://doi.org/10.48550/arXiv.2502.14739>.
- Qwen Team. Qwen3.5: Accelerating productivity with native multimodal agents, February 2026. URL <https://qwen.ai/blog?id=qwen3.5>.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv:2307.09288*, 2023.
- Dingdong Wang, Jincenzi Wu, Junan Li, Dongchao Yang, Xueyuan Chen, Tianhua Zhang, and Helen Meng. MMSU: A massive multi-task spoken language understanding and reasoning benchmark. *CoRR*, abs/2506.04779, 2025a. doi: 10.48550/ARXIV.2506.04779. URL <https://doi.org/10.48550/arXiv.2506.04779>.
- Ke Wang, Juntong Pan, Weikang Shi, Zimu Lu, Mingjie Zhan, and Hongsheng Li. Measuring multimodal mathematical reasoning with math-vision dataset. *arXiv:2402.14804*, 2024a.
- Wei Han Wang, Zehai He, Wenyi Hong, Yean Cheng, Xiaohan Zhang, Ji Qi, Shiyu Huang, Bin Xu, Yuxiao Dong, Ming Ding, and Jie Tang. Lvbench: An extreme long video understanding benchmark. *CoRR*, abs/2406.08035, 2024b.
- Xinsheng Wang, Mingqi Jiang, Ziyang Ma, Ziyu Zhang, Songxiang Liu, Linqin Li, Zheng Liang, Qixi Zheng, Rui Wang, Xiaoqin Feng, Weizhen Bian, Zhen Ye, Sitong Cheng, Ruibin Yuan, Zhixian Zhao, Xinfu Zhu, Jiahao Pan, Liumeng Xue, Pengcheng Zhu, Yunlin Chen, Zhifei Li, Xie Chen, Lei Xie, Yike Guo, and Wei Xue. Spark-tts: An efficient llm-based text-to-speech model with single-stream decoupled speech tokens. *CoRR*, abs/2503.01710, 2025b.

-
- Yu Wang, Xincheng Wang, Pengcheng Zhu, Jie Wu, Hanzhao Li, Heyang Xue, Yongmao Zhang, Lei Xie, and Mengxiao Bi. Opencpop: A high-quality open source chinese popular song corpus for singing voice synthesis. In Hanseok Ko and John H. L. Hansen (eds.), *23rd Annual Conference of the International Speech Communication Association, Interspeech 2022, Incheon, Korea, September 18-22, 2022*, pp. 4242–4246. ISCA, 2022. doi: 10.21437/INTERSPEECH.2022-48. URL <https://doi.org/10.21437/Interspeech.2022-48>.
- Yuancheng Wang, Haoyue Zhan, Liwei Liu, Ruihong Zeng, Haotian Guo, Jiachen Zheng, Qiang Zhang, Xueyao Zhang, Shunsi Zhang, and Zhizheng Wu. Maskgct: Zero-shot text-to-speech with masked generative codec transformer. *arXiv preprint arXiv:2409.00750*, 2024c.
- Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhui Chen. MMLU-Pro: A more robust and challenging multi-task language understanding benchmark. *CoRR*, abs/2406.01574, 2024d.
- Zirui Wang, Mengzhou Xia, Luxi He, Howard Chen, Yitao Liu, Richard Zhu, Kaiqu Liang, Xindi Wu, Haotian Liu, Sadhika Malladi, Alexis Chevalier, Sanjeev Arora, and Danqi Chen. Charxiv: Charting gaps in realistic chart understanding in multimodal llms. *arXiv preprint arXiv:2406.18521*, 2024e.
- Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, Jialin Wang, Yang Fan, Kai Dang, et al. Qwen2. 5-omni technical report. *arXiv preprint arXiv:2503.20215*, 2025a.
- Jin Xu, Zhifang Guo, Hangrui Hu, Yunfei Chu, Xiong Wang, Jinzheng He, Yuxuan Wang, Xianzhong Shi, Ting He, Xinfu Zhu, Yuanjun Lv, Yongqi Wang, Dake Guo, He Wang, Linhan Ma, Pei Zhang, Xinyu Zhang, Hongkun Hao, Zishan Guo, Baosong Yang, Bin Zhang, Ziyang Ma, Xipin Wei, Shuai Bai, Ke Chen, Xue Lian Liu, Peng Wang, Ming Yang, Dayiheng Liu, Xingzhang Ren, Bo Zheng, Rui Men, Fan Zhou, Bowen Yu, Jianxin Yang, Le Yu, Jing-Jun Zhou, and Junyang Lin. Qwen3-omni technical report. *ArXiv*, abs/2509.17765, 2025b.
- Fanjia Yan, Huanzhi Mao, Charlie Cheng-Jie Ji, Tianjun Zhang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Berkeley function calling leaderboard. https://gorilla.cs.berkeley.edu/blogs/8_berkeley_function_calling_leaderboard.html, 2024.
- Ruiqi Yan, Xiquan Li, Wenxi Chen, Zhikang Niu, Chen Yang, Ziyang Ma, Kai Yu, and Xie Chen. Uro-bench: A comprehensive benchmark for end-to-end spoken dialogue models. *arXiv preprint arXiv:2502.17810*, 2025.
- An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. Qwen2 technical report. *arXiv:2407.10671*, 2024a.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025a.
- Yudong Yang, Jimin Zhuang, Guangzhi Sun, Changli Tang, Yixuan Li, Peihan Li, Yifan Jiang, Wei Li, Zejun Ma, and Chao Zhang. Audio-centric video understanding benchmark without text shortcut. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 6580–6598, 2025b.
- Zhibo Yang, Jun Tang, Zhaohai Li, Pengfei Wang, Jianqiang Wan, Humen Zhong, Xuejing Liu, Mingkun Yang, Peng Wang, Shuai Bai, Lianwen Jin, and Junyang Lin. CC-OCR: A comprehensive and challenging OCR benchmark for evaluating large multimodal models in literacy. *CoRR*, abs/2412.02210, 2024b.
- Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruoyi Liu, Ge Zhang, Samuel Stevens, Dongfu Jiang, Weiming Ren, Yuxuan Sun, et al. Mmmu: A massive multi-discipline multimodal understanding and reasoning benchmark for expert agi. *arXiv:2311.16502*, 2023.
- Xiang Yue, Tianyu Zheng, Yuansheng Ni, Yubo Wang, Kai Zhang, Shengbang Tong, Yuxuan Sun, Ming Yin, Botao Yu, Ge Zhang, et al. Mmmu-pro: A more robust multi-discipline multimodal understanding benchmark. *arXiv preprint arXiv:2409.02813*, 2024.
- Yongyi Zang, Sean O’Brien, Taylor Berg-Kirkpatrick, Julian McAuley, and Zachary Novack. Are you really listening? boosting perceptual awareness in music-qa benchmarks. *arXiv preprint arXiv:2504.00369*, 2025.
- Binbin Zhang, Hang Lv, Pengcheng Guo, Qijie Shao, Chao Yang, Lei Xie, Xin Xu, Hui Bu, Xiaoyu Chen, Chenchen Zeng, Di Wu, and Zhendong Peng. WENETSPEECH: A 10000+ hours multi-domain mandarin corpus for speech recognition. In *IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP 2022, Virtual and Singapore, 23-27 May 2022*, pp. 6182–6186. IEEE, 2022. doi: 10.1109/ICASSP43922.2022.9746682. URL <https://doi.org/10.1109/ICASSP43922.2022.9746682>.

- Bowen Zhang, Congchao Guo, Geng Yang, Hang Yu, Haozhe Zhang, Heidi Lei, Jialong Mai, Junjie Yan, Kaiyue Yang, Mingqi Yang, Peikai Huang, Ruiyang Jin, Sitan Jiang, Weihua Cheng, Yawei Li, Yichen Xiao, Yiyang Zhou, Yongmao Zhang, Yuan Lu, and Yucen He. Minimax-speech: Intrinsic zero-shot text-to-speech with a learnable speaker encoder. *CoRR*, abs/2505.07916, 2025a.
- Linhao Zhang, Jian Zhang, Bokai Lei, Chuhan Wu, Aiwei Liu, Wei Jia, and Xiao Zhou. Wildspeech-bench: Benchmarking end-to-end speechllms in the wild, 2025b.
- Xiaoman Zhang, Chaoyi Wu, Ziheng Zhao, Weixiong Lin, Ya Zhang, Yanfeng Wang, and Weidi Xie. PMC-VQA: visual instruction tuning for medical visual question answering. *CoRR*, abs/2305.10415, 2023.
- Xiaomi LLM-Core Team Dong Zhang, Gang Wang, Jinlong Xue, Kai Fang, Liang Zhao, Rui Ma, Shu-Qin Ren, Shuo Liu, Tao Guo, Weiwei Zhuang, Xin Zhang, Xi-Na Song, Yihan Yan, Yongzhe He, Cici, Bowen Shen, Chengxuan Zhu, Chong Ma, Chun Chen, Heyu Chen, Jiawei Li, Lei Li, Menghang Zhu, Peidian Li, Qiying Wang, Sirui Deng, Weimin Xiong, Wen Huang, Wenyu Yang, Yilin Jiang, Yixin Yang, Yu-Shi Tian, Yue Ma, Yue Yu, Zihan Zhang, Zihao Yue, Bangjun Xiao, Bin Xia, Bofei Gao, Bowen Ye, Can Cai, Chang Liu, Chenhong He, Chunan Li, Dawei Zhu, Duo Zhang, Fengyuan Shi, Guoan Wang, Hailin Zhang, Hanglong Lv, Hanyu Li, Hao Tian, Hengxu Qu, Hong-Mei Xu, Houbin Zhang, Huaqiu Liu, Jiangshan Duo, Jia Zuo, Jianyu Wei, Jiebao Xiao, Jinhao Dong, Jun Shi, Junhao Hu, Kainan Bao, Kang Zhou, Linghao Zhang, Meng Chen, Nuo Chen, Peng Zhang, Qian Chen, Qiantong Wang, Rang Li, Shao yang Liu, Shengfan Wang, Shicheng Li, Shi liang Yu, Shijie Cao, Shimao Chen, Shuhao Gu, Weikun Wang, Wen-Juan Ma, Xia Deng, Xing Yong, Xing Zhang, Xu Wang, Yi-Hao Song, Yihao Zhao, Yingbo Zhao, Yizhao Gao, Yu Cheng, Yuanfang Tu, Yudong Wang, Zhaojun Huang, Zheng-Yu Tang, Zhenrui Lin, Zhichao Song, Zhi-Yue Xu, Zhixian Zheng, and Zi-Cheng Jiang. Mimo-audio: Audio language models are few-shot learners. *ArXiv*, abs/2512.23808, 2025c. URL <https://api.semanticscholar.org/CorpusID:284351195>.
- Yi-Fan Zhang, Huanyu Zhang, Haochen Tian, Chaoyou Fu, Shuangqing Zhang, Junfei Wu, Feng Li, Kun Wang, Qingsong Wen, Zhang Zhang, et al. Mme-realworld: Could your multimodal llm challenge high-resolution real-world scenarios that are difficult for humans? *arXiv preprint arXiv:2408.13257*, 2024.
- Yilun Zhao, Haowei Zhang, Lujing Xie, Tongyan Hu, Guo Gan, Yitao Long, Zhiyuan Hu, Weiyuan Chen, Chuhan Li, Zhijian Xu, Chengye Wang, Ziyao Shangguan, Zhenwen Liang, Yixin Liu, Chen Zhao, and Arman Cohan. MMVU: measuring expert-level multi-discipline video understanding. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2025, Nashville, TN, USA, June 11-15, 2025*, pp. 8475–8489. Computer Vision Foundation / IEEE, 2025.
- Chujie Zheng, Shixuan Liu, Mingze Li, Xiong-Hui Chen, Bowen Yu, Chang Gao, Kai Dang, Yuqiong Liu, Rui Men, An Yang, et al. Group sequence policy optimization. *arXiv preprint arXiv:2507.18071*, 2025.
- Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. Instruction-following evaluation for large language models. *CoRR*, abs/2311.07911, 2023.
- Junjie Zhou, Yan Shu, Bo Zhao, Boya Wu, Zhengyang Liang, Shitao Xiao, Minghao Qin, Xi Yang, Yongping Xiong, Bo Zhang, Tiejun Huang, and Zheng Liu. MLVU: benchmarking multi-task long video understanding. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2025, Nashville, TN, USA, June 11-15, 2025*, pp. 13691–13701. Computer Vision Foundation / IEEE, 2025a.
- Ziwei Zhou, Rui Wang, and Zuxuan Wu. Daily-omni: Towards audio-visual reasoning with temporal alignment across modalities. *CoRR*, abs/2505.17862, 2025b.
- Deyao Zhu, Jun Chen, Xiaoqian Shen, Xiang Li, and Mohamed Elhoseiny. Minigpt-4: Enhancing vision-language understanding with advanced large language models. *arXiv:2304.10592*, 2023.
- Chengke Zou, Xingang Guo, Rui Yang, Junyu Zhang, Bin Hu, and Huan Zhang. Dynamath: A dynamic visual benchmark for evaluating mathematical reasoning robustness of vision language models. In *ICLR*. OpenReview.net, 2025.
- Yuxin Zuo, Shang Qu, Yifei Li, Zhang-Ren Chen, Xuekai Zhu, Ermo Hua, Kaiyan Zhang, Ning Ding, and Bowen Zhou. Medxpertqa: Benchmarking expert-level medical reasoning and understanding. In Aarti Singh, Maryam Fazel, Daniel Hsu, Simon Lacoste-Julien, Felix Berkenkamp, Tegan Maharaj, Kiri Wagstaff, and Jerry Zhu (eds.), *ICML, Proceedings of Machine Learning Research*. PMLR / OpenReview.net, 2025.

7 Authors

Core Contributors²

Bing Han	Hangrui Hu	Peng Wang	Ting He	Xinfa Zhu	Zhifang Guo
Baosong Yang	Jin Xu*	Pei Zhang	Xize Cheng	Yunfei Chu	Zishan Guo
Bin Zhang	Jianxin Yang	Qize Yang	Xuejing Liu	Yuanjun Lv	Ziyang Ma
Bo Zheng	Jingren Zhou	Rui Men	Xingzhang Ren	Yuchong Sun	
Dayiheng Liu	Keqin Chen	Ruiyang Xu	Xian Shi	Yongqi Wang	
Fan Zhou	Le Yu	Shuai Bai	Xiong Wang	Yuxuan Wang	
Hongkun Hao	Mingkun Yang	Sibo Song	Xinyu Zhang	Yang Zhang	

Contributors²

Andong Chen	Gaoji Liu	Kangxiang Xia	Qin Zhu	Xinyao Niu	Yunbao Wu
Anfeng Li	Guangdong	Kun Yan	Ruisheng Cao	Xuancheng Ren	Yu Xi
An Yang	Zhou	Kexin Yang	Rongyao Fang	Xuechun Wang	Yi Zhang
Bei Chen	Hao Ge	Lianghao Deng	Rui Hu	Xuwu Wang	Yichang Zhang
Bin Lin	Huiqiang Jiang	Lulu Hu	Ruibin Yuan	Xingzhe Wu	Yinger Zhang
Bingshen Mu	Haoran Lian	Linhan Ma	Song Chen	Xipin Wei	Yuxiang Zheng
Bohan Wang	Hongjian Tu	Lingchen Meng	Su Hao	Xiao Xu	Zeyu Cui
Buxiao Wu	Hao Yu	Lei Xie	Shen Li	Xian Yang	Ziwei Ji
Bowen Xu	Hang Zhang	Laiwen Zheng	Shixuan Liu	Yuxuan Cai	Ziyue Jiang
Beichen Zhang	Hao Zhou	Miao Hong	Shurui Li	Yizhong Cao	Zhaohai Li
Cheng Chen	Haiquan Zhao	Mei Li	Siqi Zhang	Yilei Chen	Zheng Li
Chang Gao	Humen Zhong	Mingcheng Li	Tianyi Tang	Yuxiang Chen	Zhi Li
Chengen Huang	Jiawei Chen	Mingze Li	Tingyu Xia	Yiming Dong	Zihan Qiu
Chenyang Le	Jian Guan	Minsheng Li	Wei Ding	Yang Fan	Zekun Wang
Chenhao Li	Jiayi Leng	Minghao Wu	Wenbin Ge	Yanpeng Li	Zhihai Wang
Chenglong Liu	Jiahao Li	Mingfeng Xue	Weizhou Shen	Yucheng Li	Zhenghao Xing
Chenxu Lv	Junrong Lin	Na Ni	Wei Wang	Yang Liu	Zhibo Yang
Chen Qiang	Jiawei Liu	Peng Liu	Wentao Yao	Yantao Liu	Zhuorui Ye
Chenfei Wu	Jialong Tang	Peng Wang	Xi Chen	Yuqiong Liu	Zhenru Zhang
Chenhan Yuan	Jun Tang	Pengfei Wang	Xiaotong Chen	Yuxuan Liu	Zhipeng Zhou
Chengruidong	Jianhong Tu	Peiyang Zhang	Xionghui Chen	Yuyan Luo	Zhengyang
Zhang	Jianqiang Wan	Qidong Huang	Xiaodong Deng	Yubo Ma	Zhuge
Chujie Zheng	Jinxi Wei	Qingfeng Lan	Xudong Guo	Yang Su	
Daren Chen	Jianwei Zhang	Qintong Li	Xin Le	Yuezhang Wang	
Dake Guo	Jing Zhou	Que Shen	Xiao Li	Yuhao Wang	
Fei Huang	Kai Dang	Qiuyue Wang	Xie Chen	Yi Wu	

²Alphabetical order. * denotes the corresponding author.

8 Appendix

8.1 Detailed Multilingual Evaluation Results

Multilingual ASR. As presented in Table 13, Qwen3.5-Omni demonstrates superior speech recognition capabilities compared to state-of-the-art competitors on the FLEURS test set. Qwen3.5-Omni-Plus achieves the lowest average WER of 6.6%, outperforming both Gemini-3.1-Pro (7.3%) and GPT-4o-Transcribe (10.4%). It secures the best performance in the majority of languages, with particularly significant margins in complex tonal and low-resource languages such as Cantonese (2.2% vs. 6.3% for Gemini-3.1-Pro), Thai, and Vietnamese. Meanwhile, Qwen3.5-Omni-Flash offers a highly efficient alternative, achieving an average WER of 10.8% that remains competitive against Gemini-3-Flash (10.5%). Notably, Qwen3.5-Omni-Flash exhibits exceptional robustness in challenging scenarios, drastically reducing errors in Cantonese (3.1% vs. 10.8% for Gemini-3-Flash) and maintaining strong performance in Japanese and Korean, thereby highlighting its advantage for high-value Asian language pairs.

Multilingual Translation. As shown in Tables 14 and 15, the Qwen3.5-Omni series demonstrates distinct advantages over state-of-the-art competitors on the FLEURS test set, particularly in Asian languages and specific high-resource pairs. Qwen3.5-Omni-Plus exhibits comprehensive superiority over Gemini-3.1-Pro in the many-to-many directions (en2xx/zh2xx), achieving higher average BLEU scores in both English-to-XX (33.8 vs. 31.8) and Chinese-to-XX (21.4 vs. 19.6). It also leads in key xx2en pairs such as Portuguese (49.4 vs. 47.7) and Indonesian (45.7 vs. 45.1). Although Gemini-3.1-Pro holds a slight edge in overall xx2zh averages, Qwen3.5-Omni-Plus significantly outperforms it in critical Asian languages, including Cantonese (+15.6 BLEU), Korean, and Japanese. Similarly, Qwen3.5-Omni-Flash shows targeted strengths against Gemini-3-Flash. While maintaining competitive general performance, it vastly surpasses Gemini in Cantonese translation across all directions (e.g., 37.5 vs. 22.4 in xx2zh and 37.3 vs. 26.7 in en2xx) and delivers better results in Japanese and Korean xx2zh tasks. These results underscore Qwen3.5-Omni’s robust optimization for complex Asian linguistic structures and key regional languages.

Table 13: Multilingual ASR performance on the FLEURS test set. Results are reported using Word Error Rate (WER, ↓), where lower values indicate better performance. For italicized languages, Character Error Rate (CER, ↓) is reported instead. Compared with competing models, Qwen3.5-Omni-Plus achieves the best results on the majority of languages. The best results are highlighted in bold.

Language	Qwen3.5-Omni-Plus	Qwen3.5-Omni-Flash	Gemin-3.1-Pro	GPT-4o-Transcribe	Gemini-3-Flash
<i>Chinese</i>	2.9	2.9	3.6	2.6	4.6
English	3.2	3.7	2.7	3.2	2.9
<i>Cantonese</i>	2.2	3.1	6.3	5.2	10.8
Arabic	11.7	13.6	9.2	13.0	10.1
German	2.0	2.5	2.7	2.3	3.3
French	2.6	3.3	3.7	3.7	3.9
Spanish	2.2	2.4	2.5	2.3	2.7
Portuguese	2.1	2.2	2.6	2.3	2.9
Indonesian	1.6	2.4	2.5	3.5	2.8
Italian	0.8	1.0	1.1	1.4	1.8
<i>Korean</i>	1.7	2.1	2.0	2.1	2.4
Russian	3.1	3.6	3.4	3.7	3.9
<i>Thai</i>	2.8	3.2	4.3	4.9	4.5
<i>Vietnamese</i>	1.9	2.5	2.5	3.5	3.5
<i>Japanese</i>	1.9	2.5	2.3	3.0	3.4
Turkish	3.1	4.4	3.8	4.2	4.4
Hindi	9.7	9.9	4.5	12.0	5.6
Malay	2.7	4.2	3.7	4.1	6.2
Dutch	2.8	3.5	3.5	3.7	4.7
Urdu	20.8	31.9	25.2	19.7	23.0
Norwegian	3.9	5.2	5.0	5.5	6.7
Swedish	3.1	5.0	4.6	5.2	7.7
Danish	3.5	5.3	5.7	6.5	7.9
Hebrew	12.5	16.6	15.6	19.4	20.2
Finnish	2.4	4.5	3.4	3.8	5.2
Polish	1.9	3.1	2.7	2.8	4.9
Icelandic	3.6	8.9	4.7	10.8	6.8
Czech	2.6	4.5	3.8	4.7	8.0
Filipino	5.1	7.1	7.6	7.3	8.5
Persian	12.0	12.1	8.9	9.9	10.0
Greek	4.7	8.1	5.4	6.5	7.6
Afrikaans	10.6	13.7	12.7	17.9	18.6
Asturian	15.8	25.9	23.7	23.8	48.3
Belarusian	6.7	12.2	6.7	10.2	10.9
Bulgarian	6.2	10.7	5.3	7.0	7.9
Bengali	16.2	19.8	21.9	24.1	21.9
Bosnian	5.4	9.5	6.0	13.9	11.2
Catalan	2.8	6.3	2.7	2.7	6.4
Cebuano	10.5	16.6	13.0	15.1	12.8
Estonian	6.7	16.6	4.9	7.6	7.3
Galician	5.0	8.6	4.9	6.8	14.6
Gujarati	13.9	18.4	14.8	26.9	16.3
Croatian	5.4	9.0	5.1	16.4	9.3
Hungarian	4.9	10.6	5.5	7.4	11.3
Javanese	11.8	18.3	14.1	24.7	15.7
Kazakh	6.3	16.6	6.2	11.5	12.7
Kannada	16.0	23.8	16.3	28.1	16.3
Kyrgyz	10.0	19.7	8.3	20.7	16.3
Latvian	6.7	17.8	3.7	6.3	6.8
Macedonian	4.1	7.9	4.0	6.2	7.8
Malayalam	18.8	27.0	18.3	33.5	20.3
Marathi	16.3	23.6	15.3	26.3	16.0
Punjabi	13.7	24.6	14.4	36.4	17.4
Romanian	3.2	6.1	3.4	4.5	5.9
Slovak	3.3	5.5	2.8	3.6	6.5
Slovenian	6.1	14.3	6.3	8.8	10.0
Swahili	9.4	17.5	9.9	16.3	10.7
Tajik	10.0	41.1	20.4	20.2	53.1
Azerbaijani	7.2	13.0	5.8	10.6	13.2
Ukrainian	3.2	5.4	3.4	4.3	5.1
Average	6.6	10.8	7.3	10.4	10.5

Table 14: Multilingual translation performance on the FLEURS en2xx and zh2xx test sets. Results are reported using BLEU ($\uparrow$). Compared with competing models, Qwen3.5-Omni-Plus outperforms them on the majority of language pairs. The best results are highlighted in bold.

Language	en2xx (English $\rightarrow$ Other Languages)				zh2xx (Chinese $\rightarrow$ Other Languages)			
	Qwen3.5-Omni-Plus	Qwen3.5-Omni-Flash	Gemin-3.1-Pro	Gemini-3-Flash	Qwen3.5-Omni-Plus	Qwen3.5-Omni-Flash	Gemin-3.1-Pro	Gemini-3-Flash
Chinese	47.8	46.6	47.4	46.3	–	–	–	–
English	–	–	–	–	32.2	31.2	30.1	29.5
Cantonese	40.1	37.3	25.5	26.7	36.7	35.9	23.7	24.0
Arabic	31.1	28.2	27.0	28.5	16.1	13.9	14.2	14.4
German	43.2	39.6	41.8	40.9	23.2	20.8	22.0	21.4
French	50.9	48.8	48.0	47.8	30.7	28.8	29.4	29.2
Spanish	29.1	28.9	28.3	28.8	22.2	20.4	20.8	20.6
Portuguese	51.2	48.6	47.3	47.2	28.5	26.8	25.7	25.6
Indonesian	45.3	43.7	42.1	41.7	28.8	26.9	25.4	25.2
Italian	32.7	30.7	31.9	30.9	23.1	21.1	21.9	21.3
Korean	33.9	31.8	30.8	31.7	25.1	23.4	21.9	22.8
Russian	33.8	31.8	33.1	33.2	21.5	18.9	20.0	19.9
Thai	65.4	62.9	64.8	64.4	58.0	55.5	57.2	56.4
Vietnamese	43.0	41.8	41.1	40.2	31.6	30.5	28.1	28.4
Japanese	53.2	50.6	51.3	50.8	45.6	41.6	43.0	42.0
Turkish	30.4	27.6	29.3	29.0	16.8	14.6	15.9	16.0
Hindi	33.1	29.1	28.2	29.2	19.1	14.3	17.6	17.5
Malay	39.6	37.2	35.9	36.0	24.1	21.7	21.0	20.5
Dutch	30.1	28.2	28.1	28.8	21.0	18.8	19.3	19.1
Urdu	25.0	22.1	23.0	23.0	15.5	8.6	14.9	14.7
Norwegian	35.3	32.8	33.1	33.7	20.3	17.8	18.4	18.9
Swedish	47.5	44.1	45.7	45.8	25.4	23.0	23.7	24.1
Danish	48.4	45.2	45.7	45.2	25.7	22.8	23.5	23.2
Hebrew	36.4	29.9	36.5	35.4	18.2	14.5	17.7	17.4
Finnish	30.1	26.0	32.1	32.3	18.1	15.3	18.6	17.7
Polish	25.2	22.5	24.9	23.4	17.5	15.0	15.6	15.5
Icelandic	28.5	27.2	29.6	28.2	16.2	13.5	16.0	15.6
Czech	35.9	32.5	33.3	33.7	20.4	18.1	19.4	19.0
Filipino	35.0	32.0	32.1	33.1	22.3	19.0	20.7	20.7
Persian	30.7	27.3	25.1	25.9	19.5	16.4	15.9	15.9
Greek	30.0	27.8	30.0	29.4	18.4	15.9	17.6	17.5
Asturian	32.4	27.9	31.5	30.4	20.4	16.2	18.6	18.1
Belarusian	16.4	14.7	16.4	16.5	12.6	10.8	12.1	12.2
Bulgarian	45.0	40.7	41.7	42.5	25.6	23.0	24.3	24.0
Bengali	18.6	15.7	14.3	15.0	10.6	9.2	9.0	9.3
Bosnian	37.5	34.0	36.3	35.2	21.4	18.6	19.9	19.4
Catalan	43.9	41.5	42.7	42.9	26.6	17.2	25.0	25.1
Cebuano	28.5	12.7	28.5	29.2	19.0	5.6	17.5	17.7
Estonian	30.8	26.3	31.4	30.6	18.9	13.6	17.5	17.2
Galician	37.4	35.4	36.6	35.9	23.9	22.0	22.7	22.3
Gujarati	23.8	20.9	21.0	21.5	14.3	10.8	12.3	12.3
Croatian	33.3	30.7	33.4	32.6	21.3	18.5	19.2	18.6
Hungarian	29.5	24.9	28.1	27.6	18.8	15.8	17.3	16.7
Javanese	26.8	24.4	16.4	22.1	16.5	14.8	9.0	13.0
Kazakh	24.9	21.1	20.6	22.5	15.0	12.4	12.9	13.1
Kannada	20.0	17.0	16.1	17.2	11.7	6.9	9.5	10.0
Kyrgyz	15.3	12.6	15.1	14.7	10.4	7.8	10.0	9.5
Latvian	36.1	31.0	35.4	35.3	21.9	17.8	20.1	19.5
Macedonian	38.1	34.0	38.5	38.2	22.3	20.1	22.0	21.6
Malayalam	19.3	11.1	16.1	16.1	10.4	5.2	9.9	9.4
Marathi	17.7	11.6	15.9	16.2	11.3	8.1	9.4	10.1
Punjabi	26.1	23.1	24.1	24.6	15.7	8.9	14.3	14.1
Romanian	42.0	39.9	41.4	42.0	25.6	22.8	23.6	23.4
Slovak	35.3	31.4	34.8	34.6	19.6	16.4	19.2	18.5
Slovenian	32.8	28.5	33.5	32.9	20.1	17.7	20.9	20.2
Swahili	36.3	30.9	32.1	32.2	20.4	9.3	18.5	18.3
Tajik	23.8	18.3	22.1	22.9	14.6	10.9	14.3	14.2
Azerbaijani	13.7	9.8	15.2	14.7	11.5	9.7	11.3	10.8
Ukrainian	31.7	29.2	29.8	30.1	19.5	15.0	17.0	17.6
Average	33.8	30.4	31.8	31.8	21.4	18.1	19.6	19.5

Table 15: Multilingual translation performance on the FLEURS xx2en and xx2zh test sets. Results are reported using BLEU ($\uparrow$). The best results are highlighted in bold.

Language	xx2en (Other Languages $\rightarrow$ English)				xx2zh (Other Languages $\rightarrow$ Chinese)			
	Qwen3.5-Omni-Plus	Qwen3.5-Omni-Flash	Gemin-3.1-Pro	Gemini-3-Flash	Qwen3.5-Omni-Plus	Qwen3.5-Omni-Flash	Gemin-3.1-Pro	Gemini-3-Flash
Chinese	32.2	31.2	30.1	29.5	–	–	–	–
English	–	–	–	–	47.8	46.6	47.4	46.3
Cantonese	30.3	29.9	27.7	26.3	36.8	37.5	21.2	22.4
Arabic	42.9	40.0	42.1	42.3	40.2	37.3	40.5	40.5
German	44.6	44.1	43.9	43.9	43.3	42.8	42.1	41.8
French	43.5	42.0	41.6	41.3	41.6	40.3	41.6	41.4
Spanish	32.3	31.3	30.3	30.4	38.8	38.5	38.4	38.2
Portuguese	49.4	48.2	47.7	47.5	43.6	41.8	42.5	42.4
Indonesian	45.7	43.1	45.1	44.9	43.5	41.4	42.8	42.9
Italian	34.4	31.9	30.8	31.0	40.8	39.5	39.7	39.4
Korean	34.1	32.4	32.0	32.1	39.9	37.5	37.0	37.2
Russian	38.6	37.2	36.7	36.5	41.7	39.5	41.1	40.4
Thai	34.1	32.4	34.2	33.0	40.2	37.9	40.0	39.5
Vietnamese	36.4	34.9	36.1	35.4	38.7	36.3	39.2	38.9
Japanese	30.4	29.2	29.5	29.4	38.0	35.7	35.6	34.8
Turkish	40.3	39.1	39.5	38.9	41.8	40.3	40.6	41.0
Hindi	38.8	36.2	39.3	39.2	38.9	36.9	38.0	38.4
Malay	42.9	41.1	44.8	42.4	41.0	39.5	42.4	41.5
Dutch	33.3	32.2	31.2	30.4	40.1	38.5	39.4	39.4
Urdu	35.5	31.5	32.9	32.6	37.2	33.8	36.9	36.5
Norwegian	43.5	42.1	42.6	41.4	42.2	39.6	40.9	40.8
Swedish	47.2	45.2	46.6	44.7	42.9	40.7	41.5	41.6
Danish	45.4	44.0	44.7	42.7	43.4	41.1	41.7	40.6
Hebrew	39.7	36.4	42.5	39.9	36.7	34.1	40.3	37.8
Finnish	36.9	35.0	36.7	35.5	40.8	38.7	41.0	40.6
Polish	32.1	30.5	30.4	29.7	38.4	36.0	37.9	37.4
Icelandic	31.5	27.5	35.1	35.8	38.2	31.8	37.9	37.5
Czech	42.1	39.3	40.1	39.5	40.6	39.4	40.6	40.0
Filipino	42.7	40.9	45.0	44.3	41.0	38.0	42.6	41.5
Persian	40.2	36.8	38.0	38.1	40.2	37.0	41.1	41.1
Greek	36.0	32.5	35.3	35.4	38.0	32.8	39.0	38.4
Asturian	37.0	35.1	37.2	35.5	37.7	34.0	38.1	36.7
Belarusian	23.1	19.9	20.6	19.9	33.3	31.2	33.7	33.7
Bulgarian	39.6	36.0	41.3	40.0	40.9	36.6	41.0	40.4
Bengali	32.0	27.6	34.9	34.3	35.8	32.4	38.1	37.5
Bosnian	43.1	40.8	42.7	42.4	41.5	39.1	42.3	41.4
Catalan	46.6	42.3	46.2	45.5	42.2	38.9	42.6	41.5
Cebuano	37.3	26.3	38.9	38.2	34.0	26.5	36.7	36.6
Estonian	35.7	28.3	40.1	38.3	38.0	32.2	41.7	41.2
Galician	40.8	38.6	39.6	38.6	40.9	39.1	41.6	40.9
Gujarati	33.4	28.3	40.3	39.7	35.8	31.4	39.7	39.0
Croatian	39.5	36.3	38.0	36.9	40.0	37.7	40.2	39.6
Hungarian	35.5	29.6	35.3	33.3	39.1	33.9	40.0	37.4
Javanese	35.9	28.4	38.4	36.5	34.9	28.3	36.7	36.3
Kazakh	34.4	26.9	35.7	35.5	37.1	31.4	39.3	38.8
Kannada	26.4	19.8	33.1	33.6	32.3	26.1	38.0	37.8
Kyrgyz	22.2	17.1	24.8	23.7	29.9	24.8	33.7	32.8
Latvian	33.7	25.2	38.0	37.9	37.1	30.0	41.4	40.6
Macedonian	43.3	39.6	43.1	41.9	41.6	38.0	42.1	41.8
Malayalam	31.2	25.7	34.9	34.2	36.1	31.5	38.4	38.0
Marathi	33.5	25.7	36.7	35.5	34.6	29.4	38.8	37.9
Punjabi	33.0	26.9	38.3	36.5	35.1	29.9	37.4	36.4
Romanian	43.5	39.7	42.3	41.5	42.0	38.7	42.4	41.3
Slovak	39.7	38.4	39.9	38.8	39.2	38.1	40.2	39.3
Slovenian	31.7	26.5	34.7	33.2	34.5	30.1	38.4	37.3
Swahili	35.0	27.4	42.6	40.5	33.9	26.8	39.2	37.9
Tajik	33.9	29.0	34.5	33.3	36.7	32.7	38.9	38.1
Azerbaijani	25.0	22.0	23.7	23.3	33.4	30.5	33.5	33.3
Ukrainian	42.0	40.1	41.7	41.4	41.7	39.7	41.6	40.7
Average	37.0	33.5	37.4	36.6	38.9	35.7	39.4	38.9