



新一代物理引擎 Newton

解锁高精度仿真与机器人训练

Juana Du, NVIDIA Robotics Solution Architect | Oct 2025

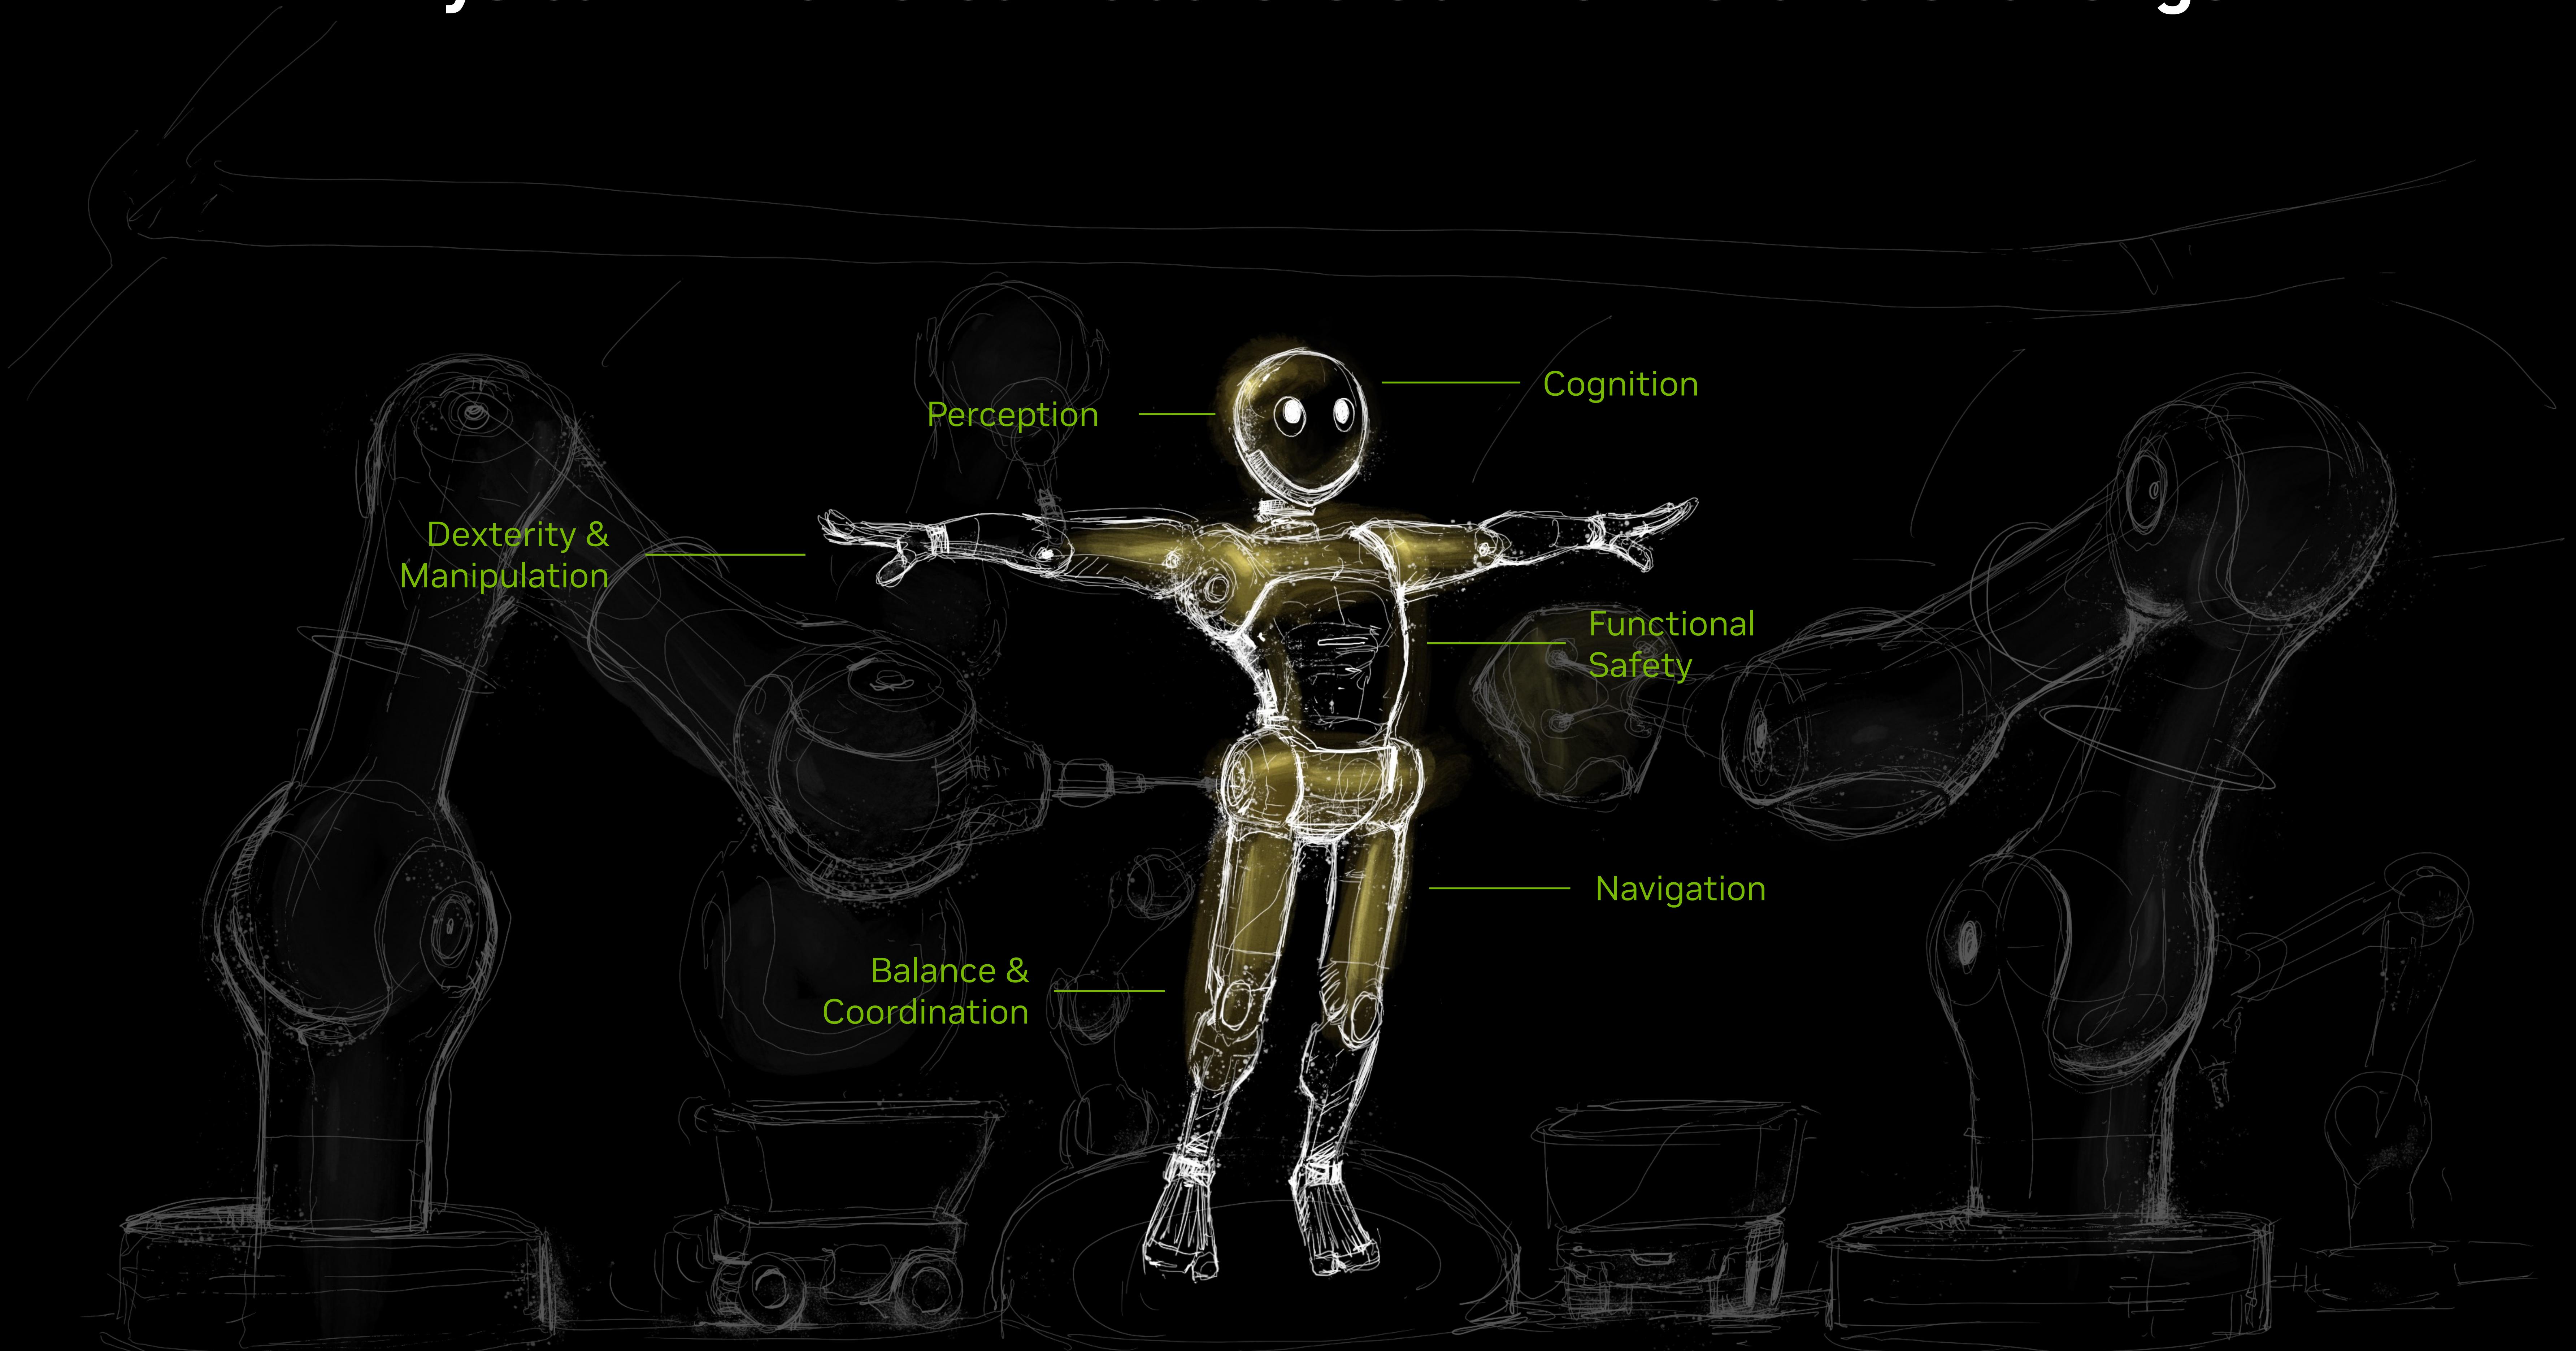


Agenda

- Newton overview
- Architecture design and GPU acceleration
- Advanced solvers and Multi-physics coupling support
- Isaac Lab integration

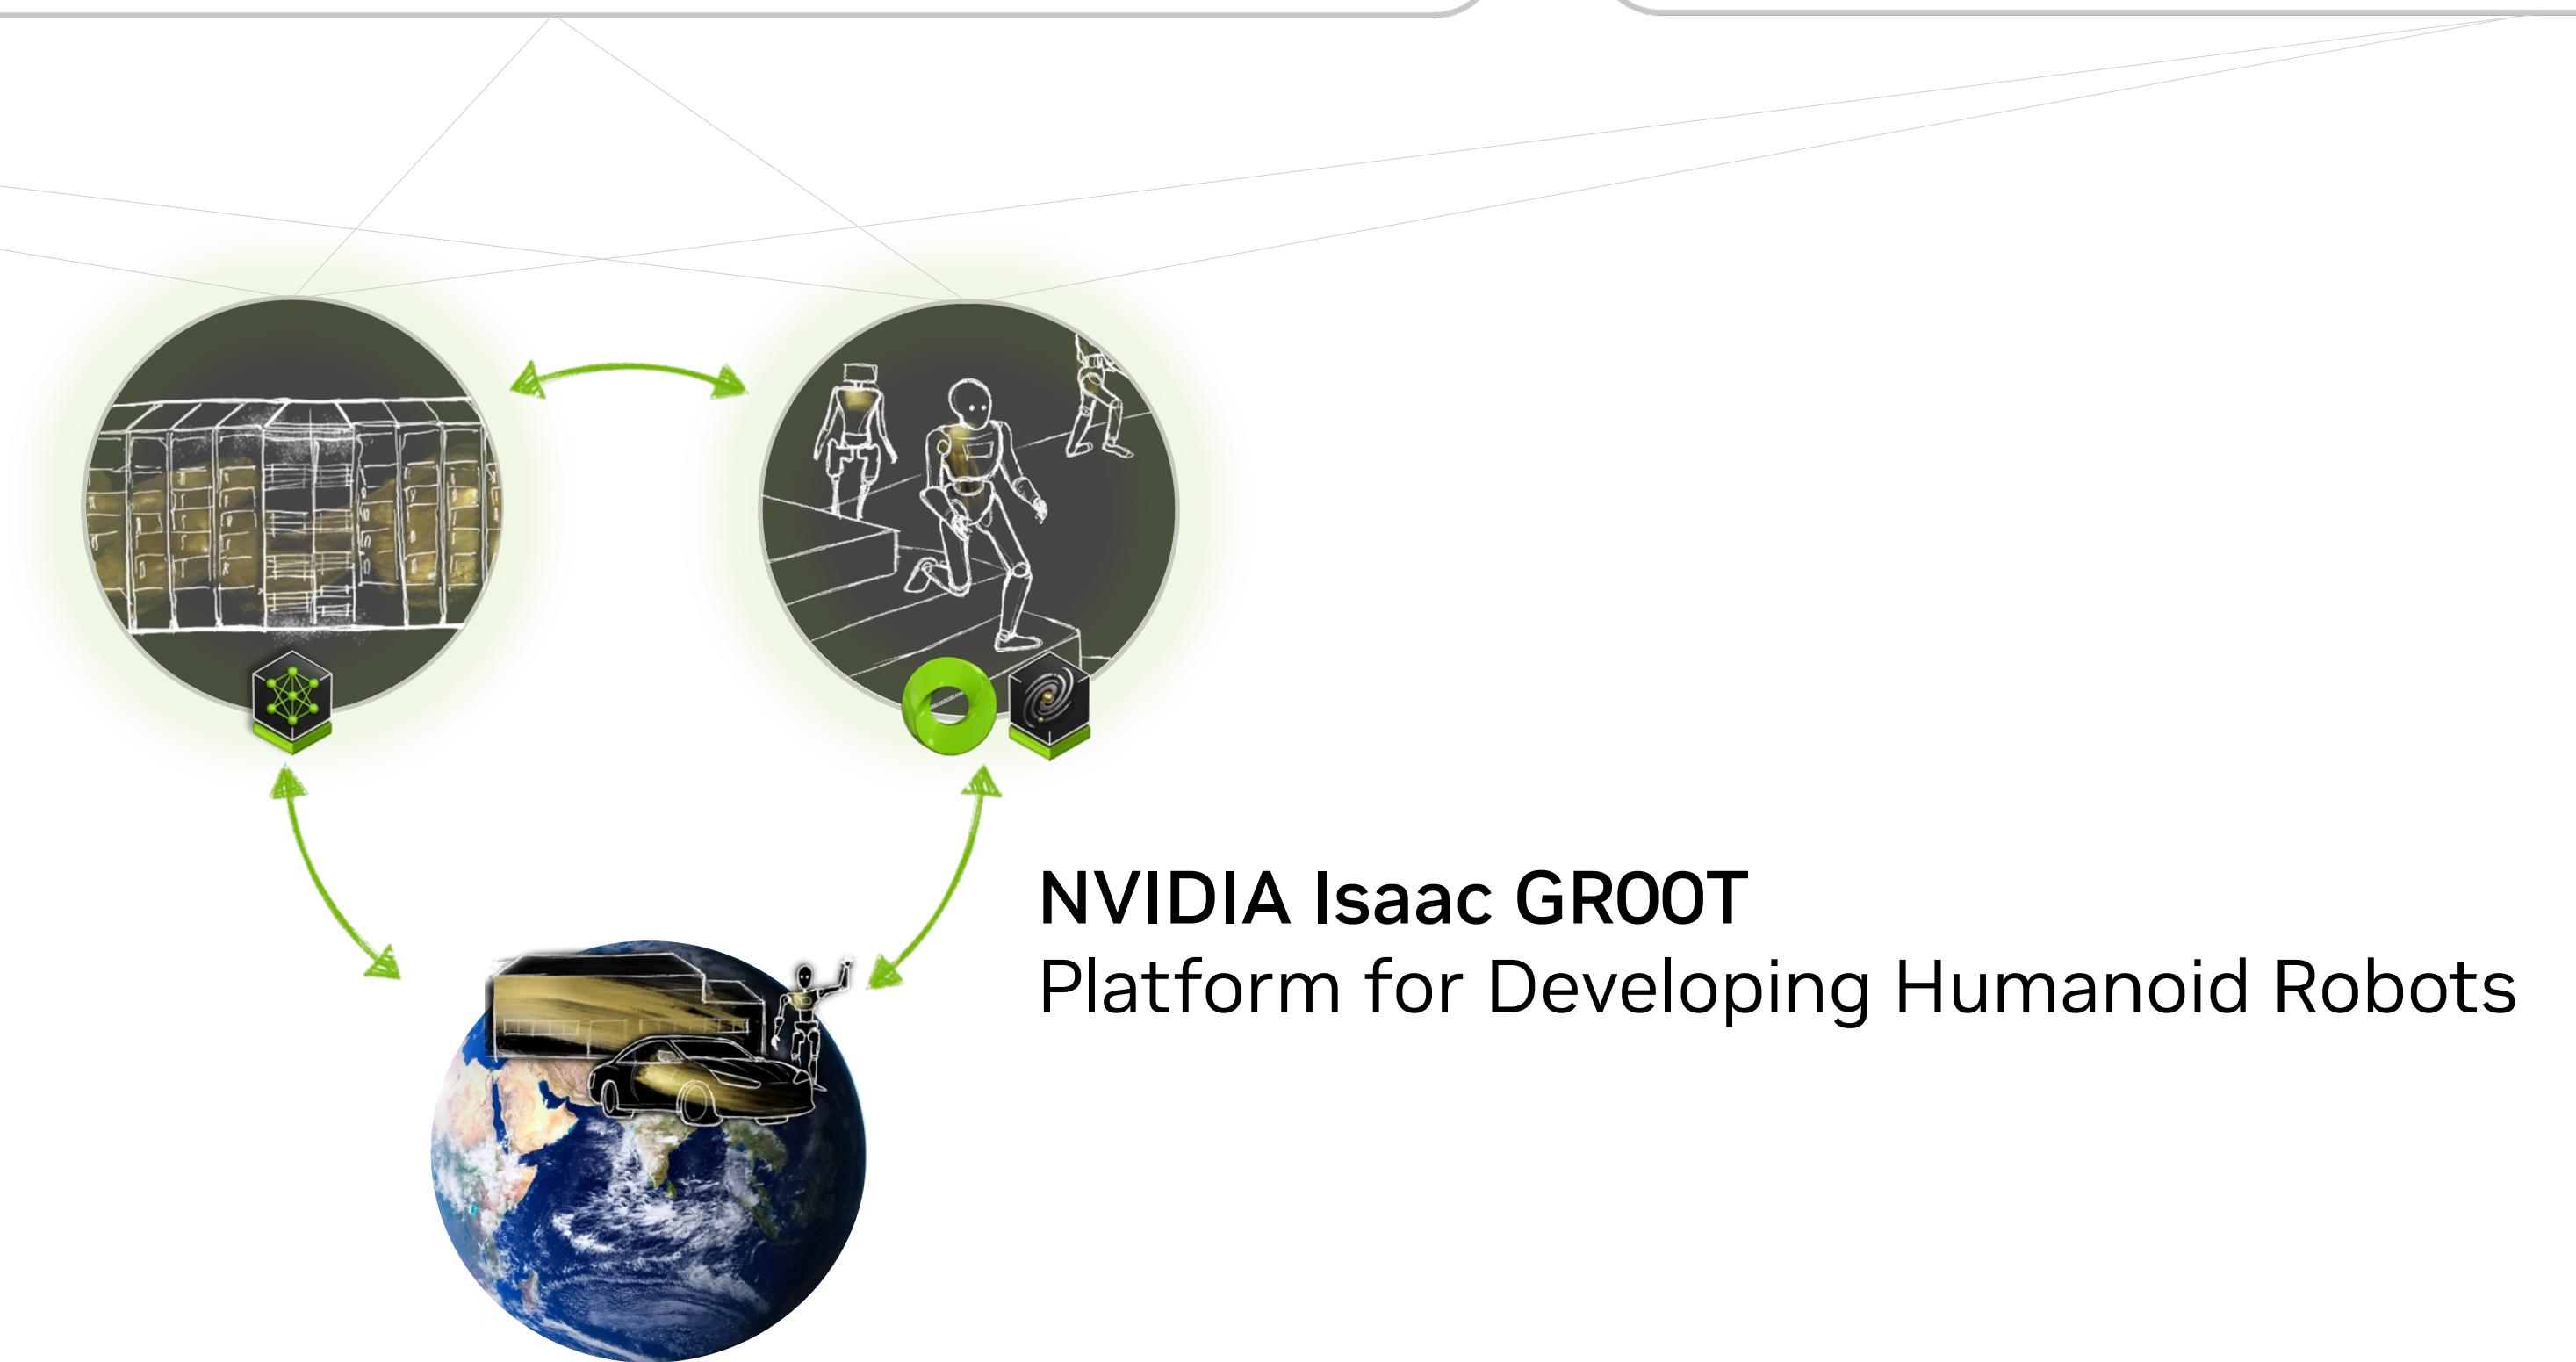
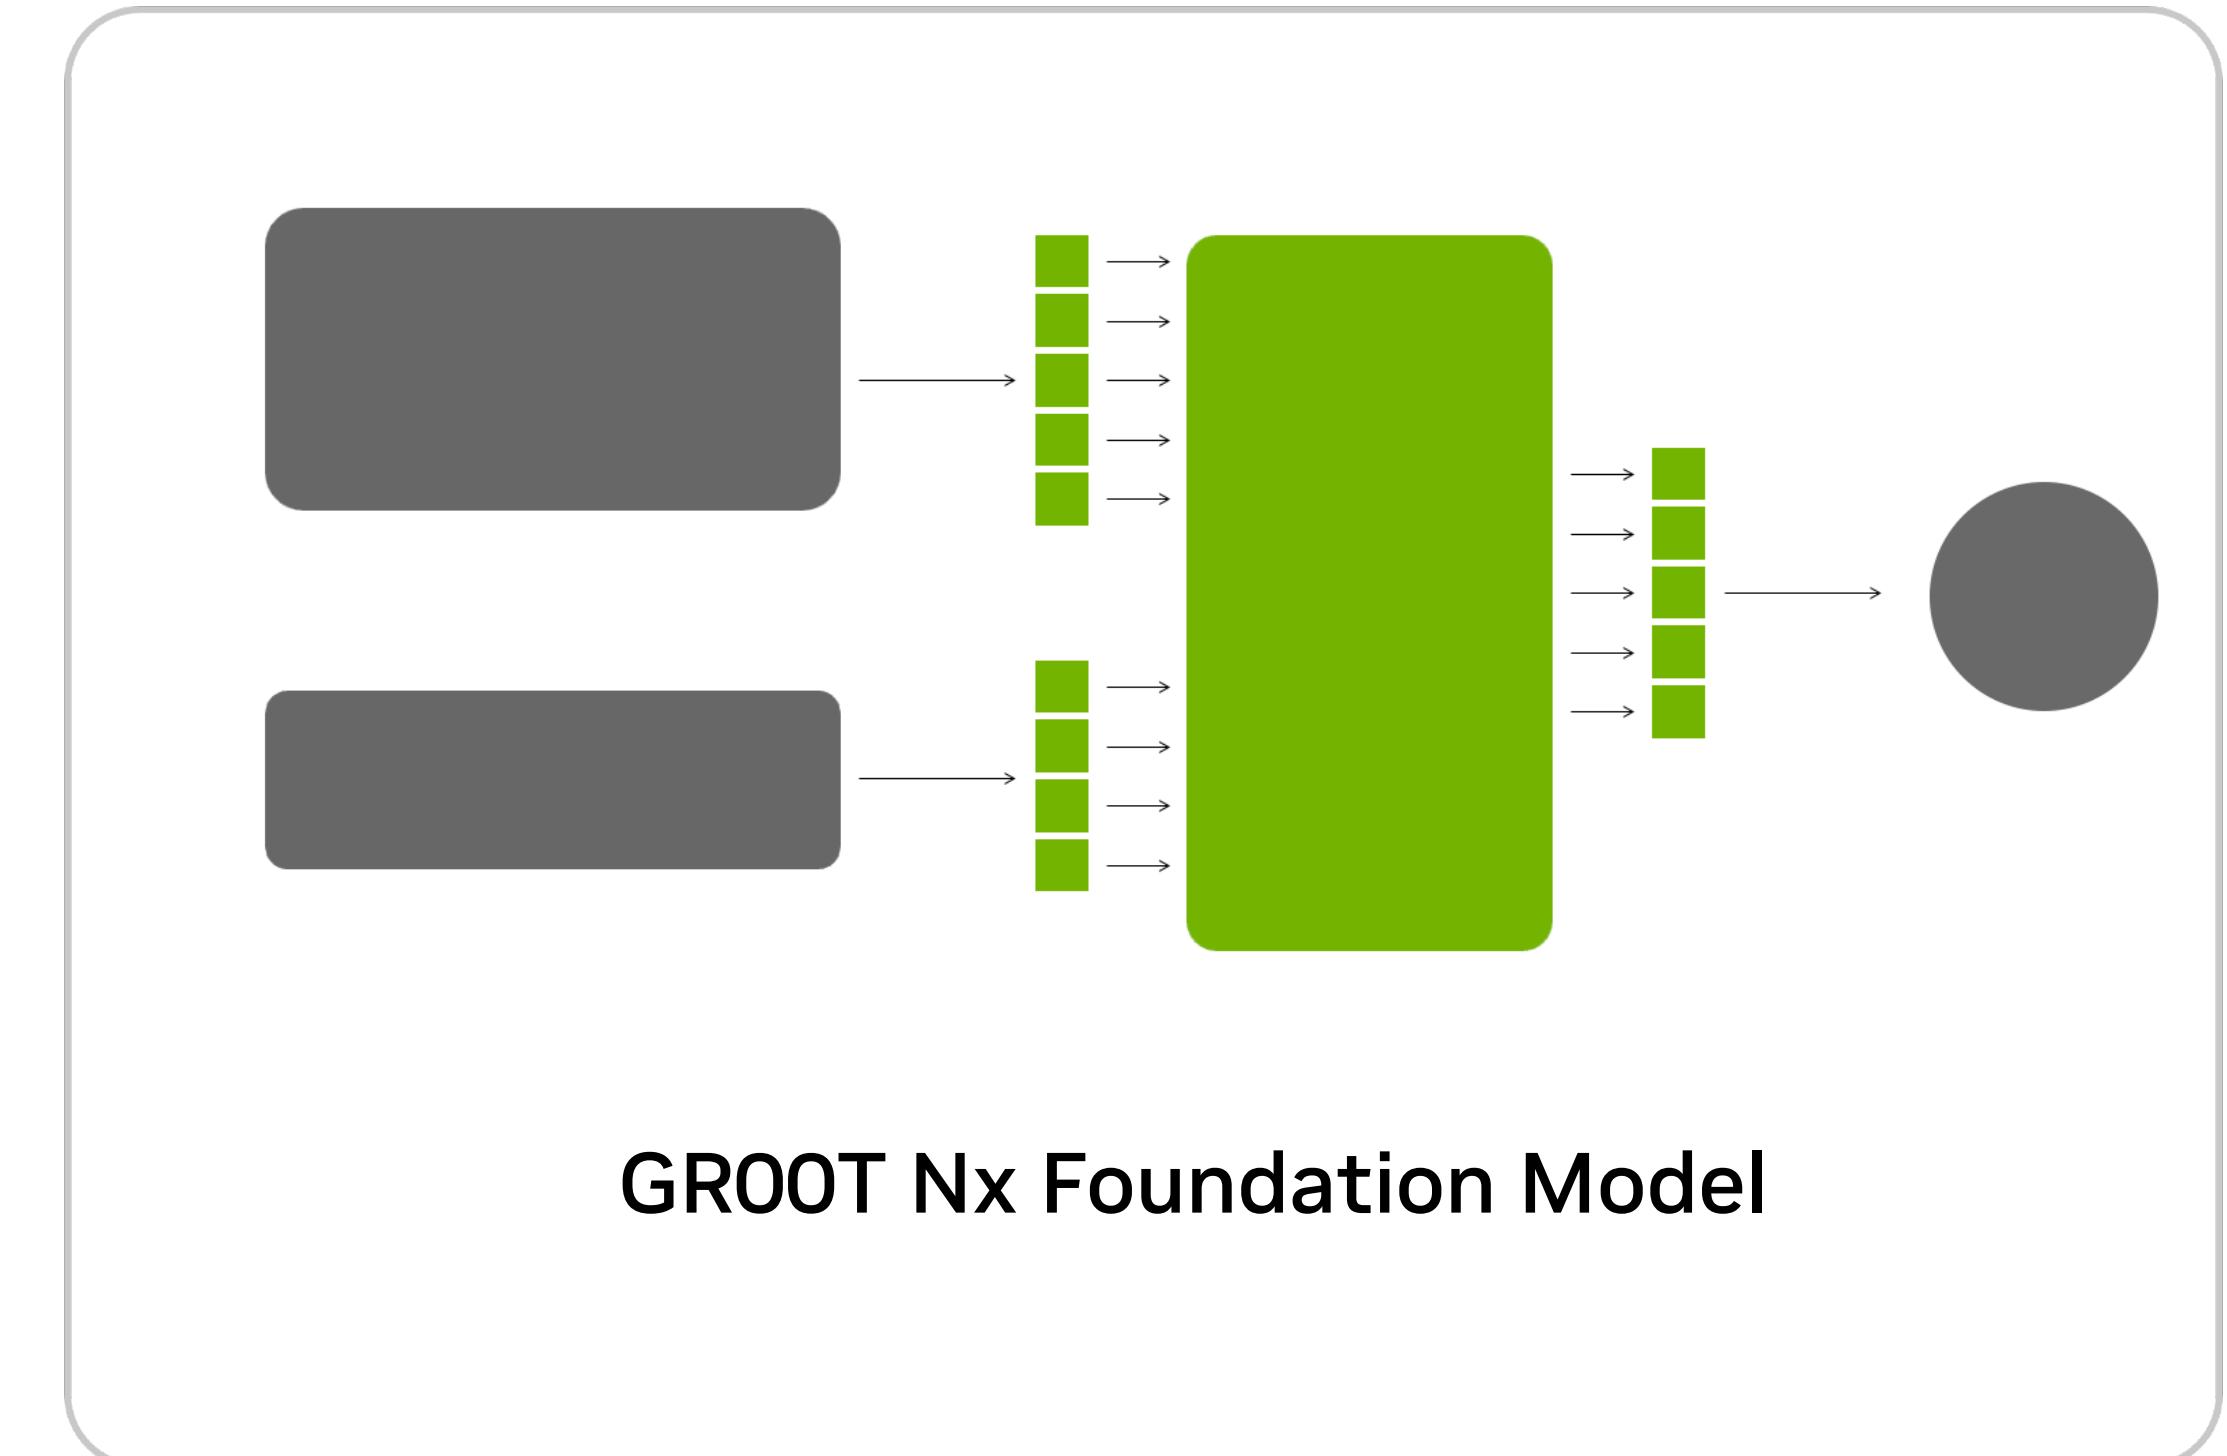
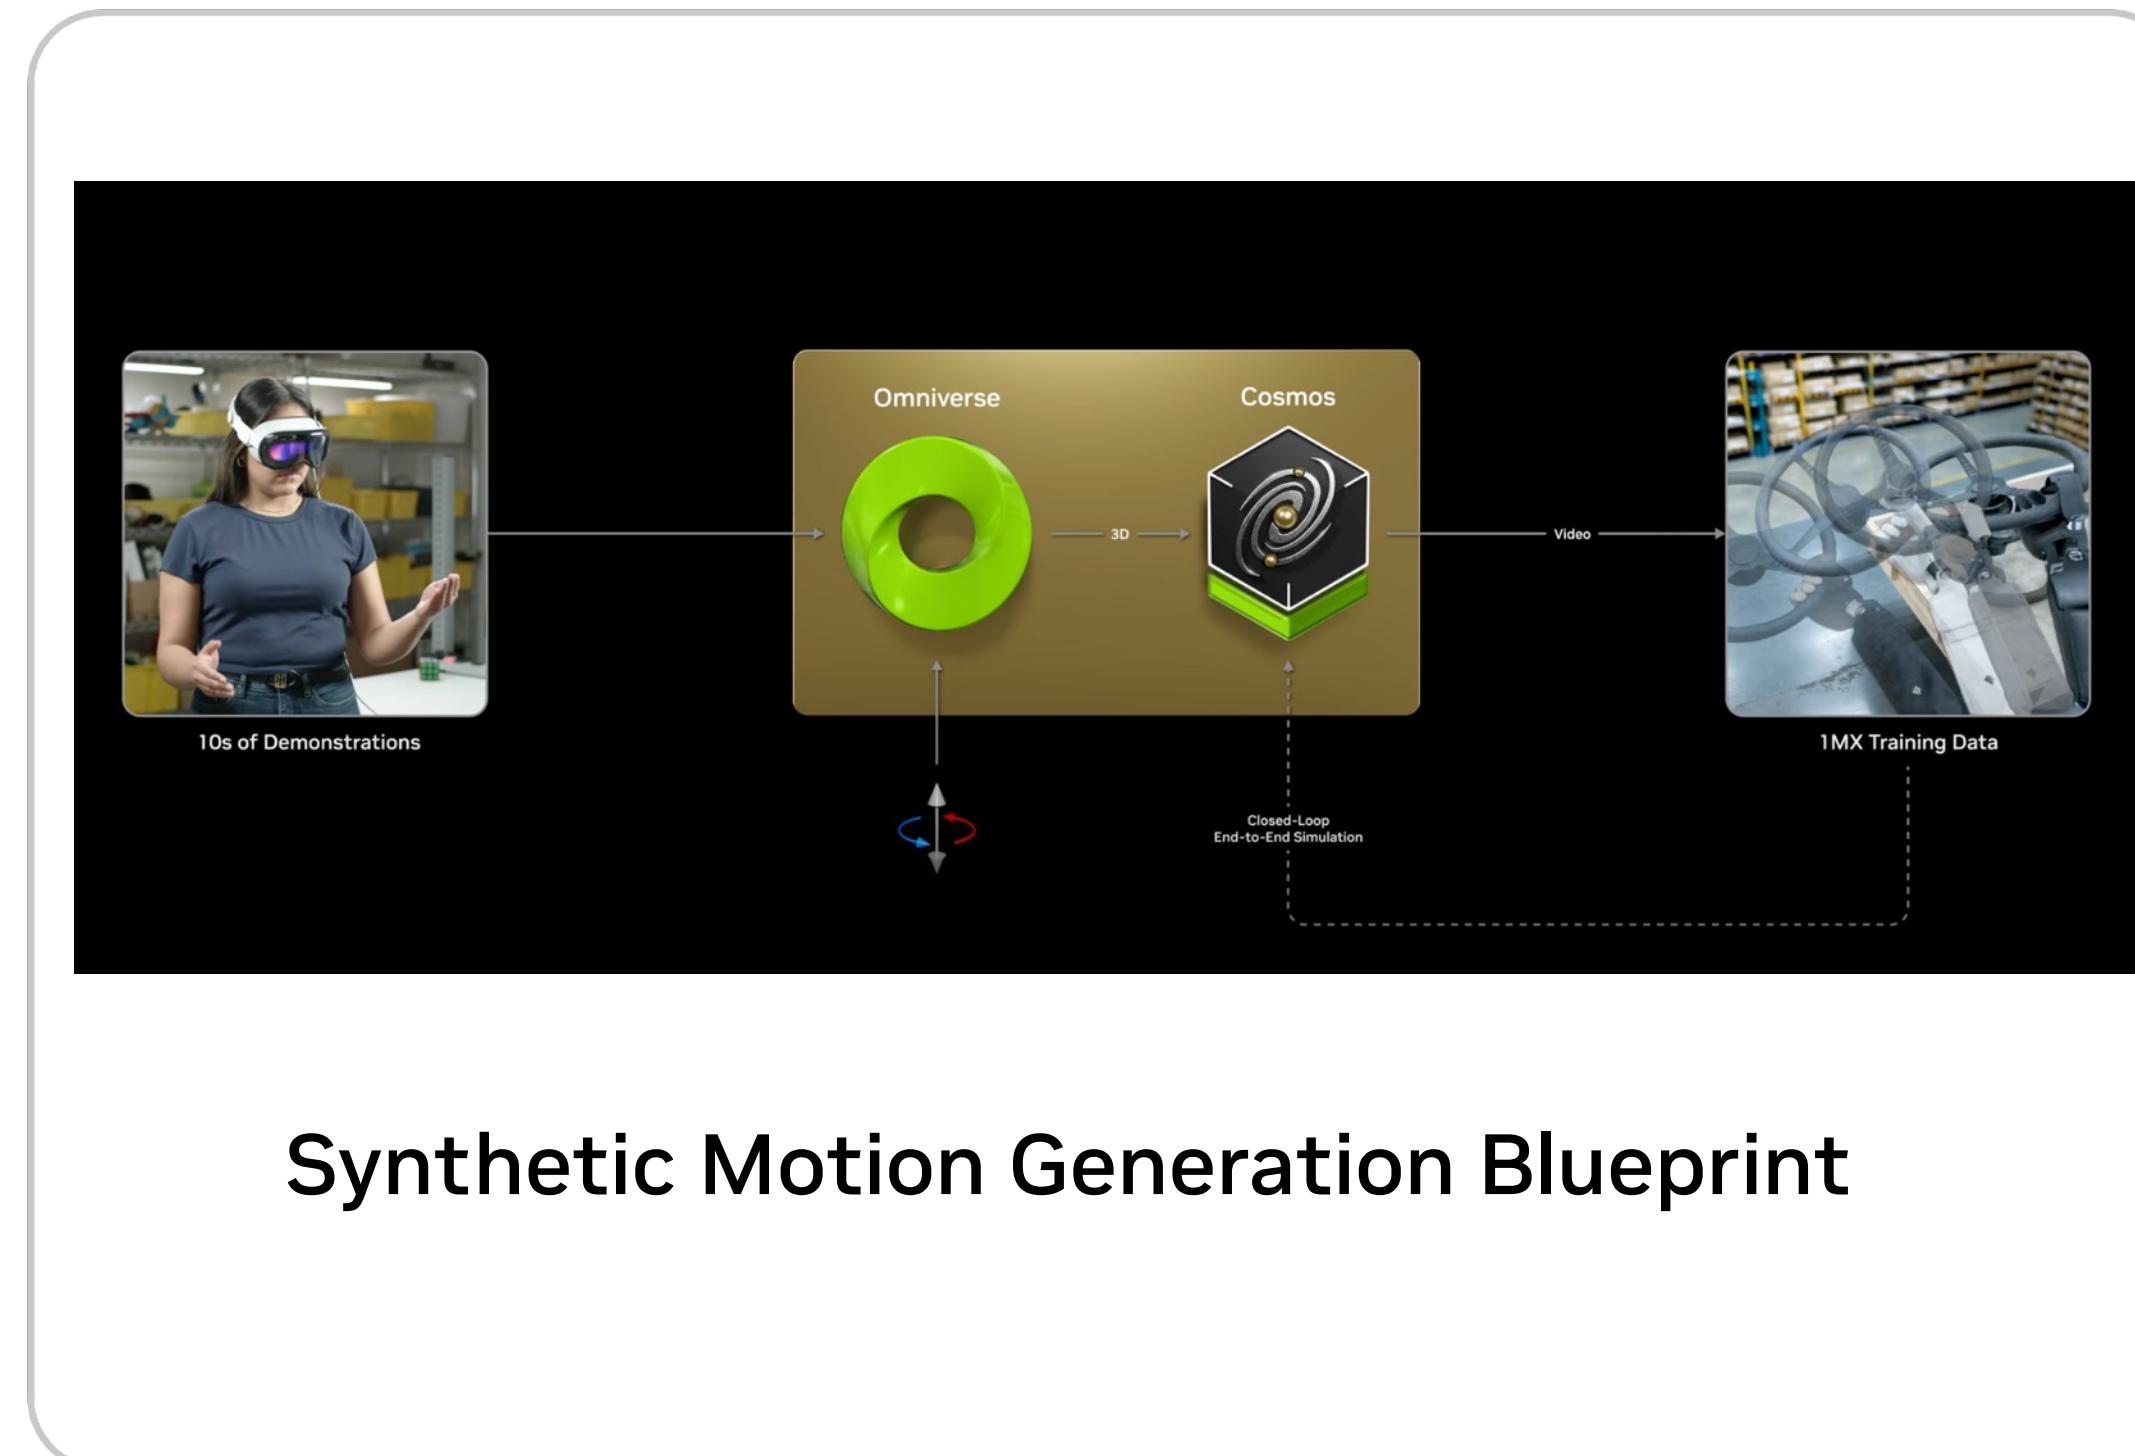
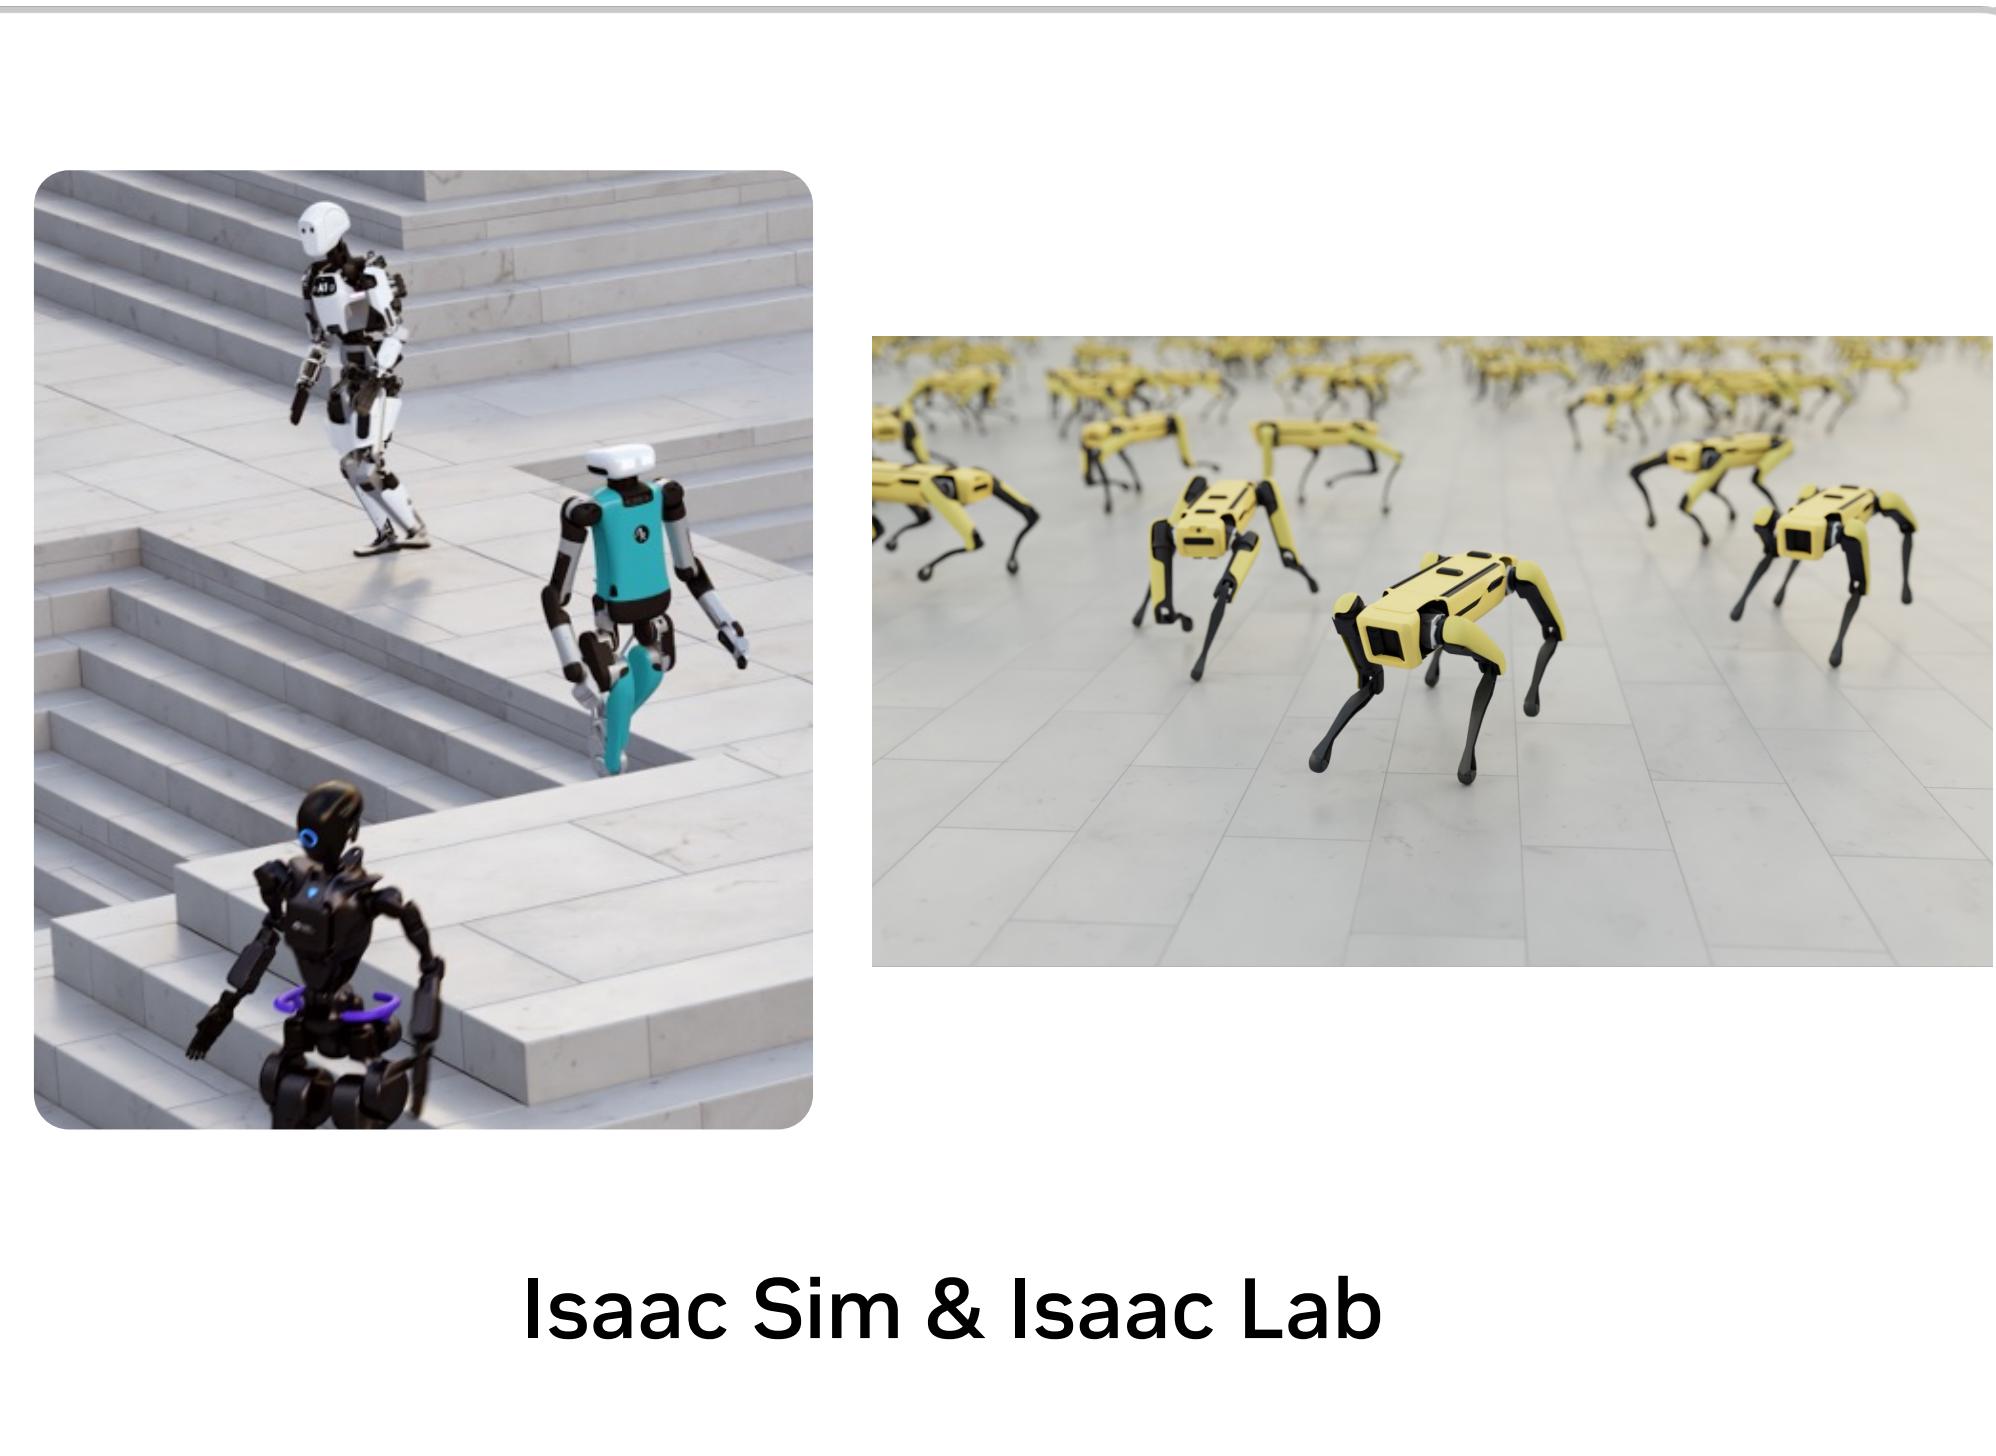
Newton overview

Physical AI-Powered Robots is Our Next Grand Challenge



Simulation is Important for Robot Learning

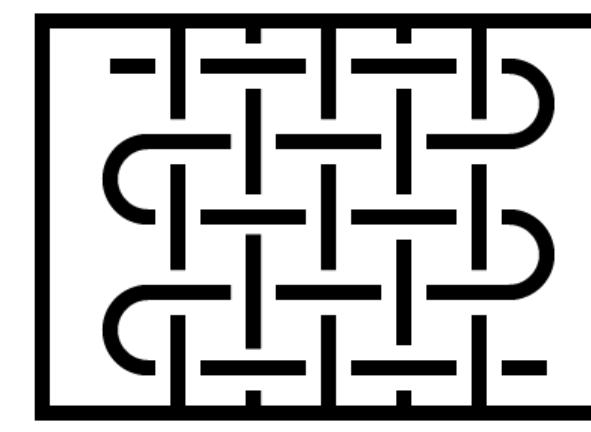
GROOT: Generalist Robot OO Technology



Next Generation Simulation for Robot Learning

Requirements

- Modular & Extensible
- Community-built
- GPU Accelerated
- Industry-ready
- Common Data Interchange
- Unified API
- Differentiable

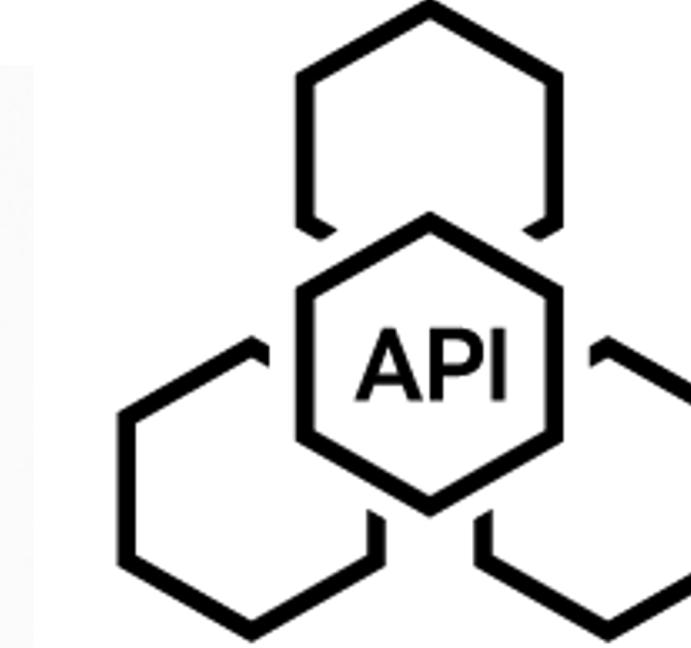
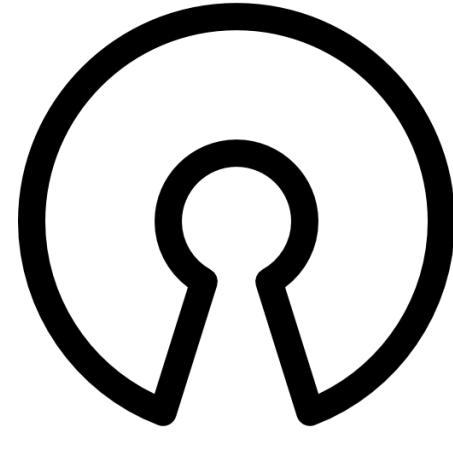


NVIDIA Warp

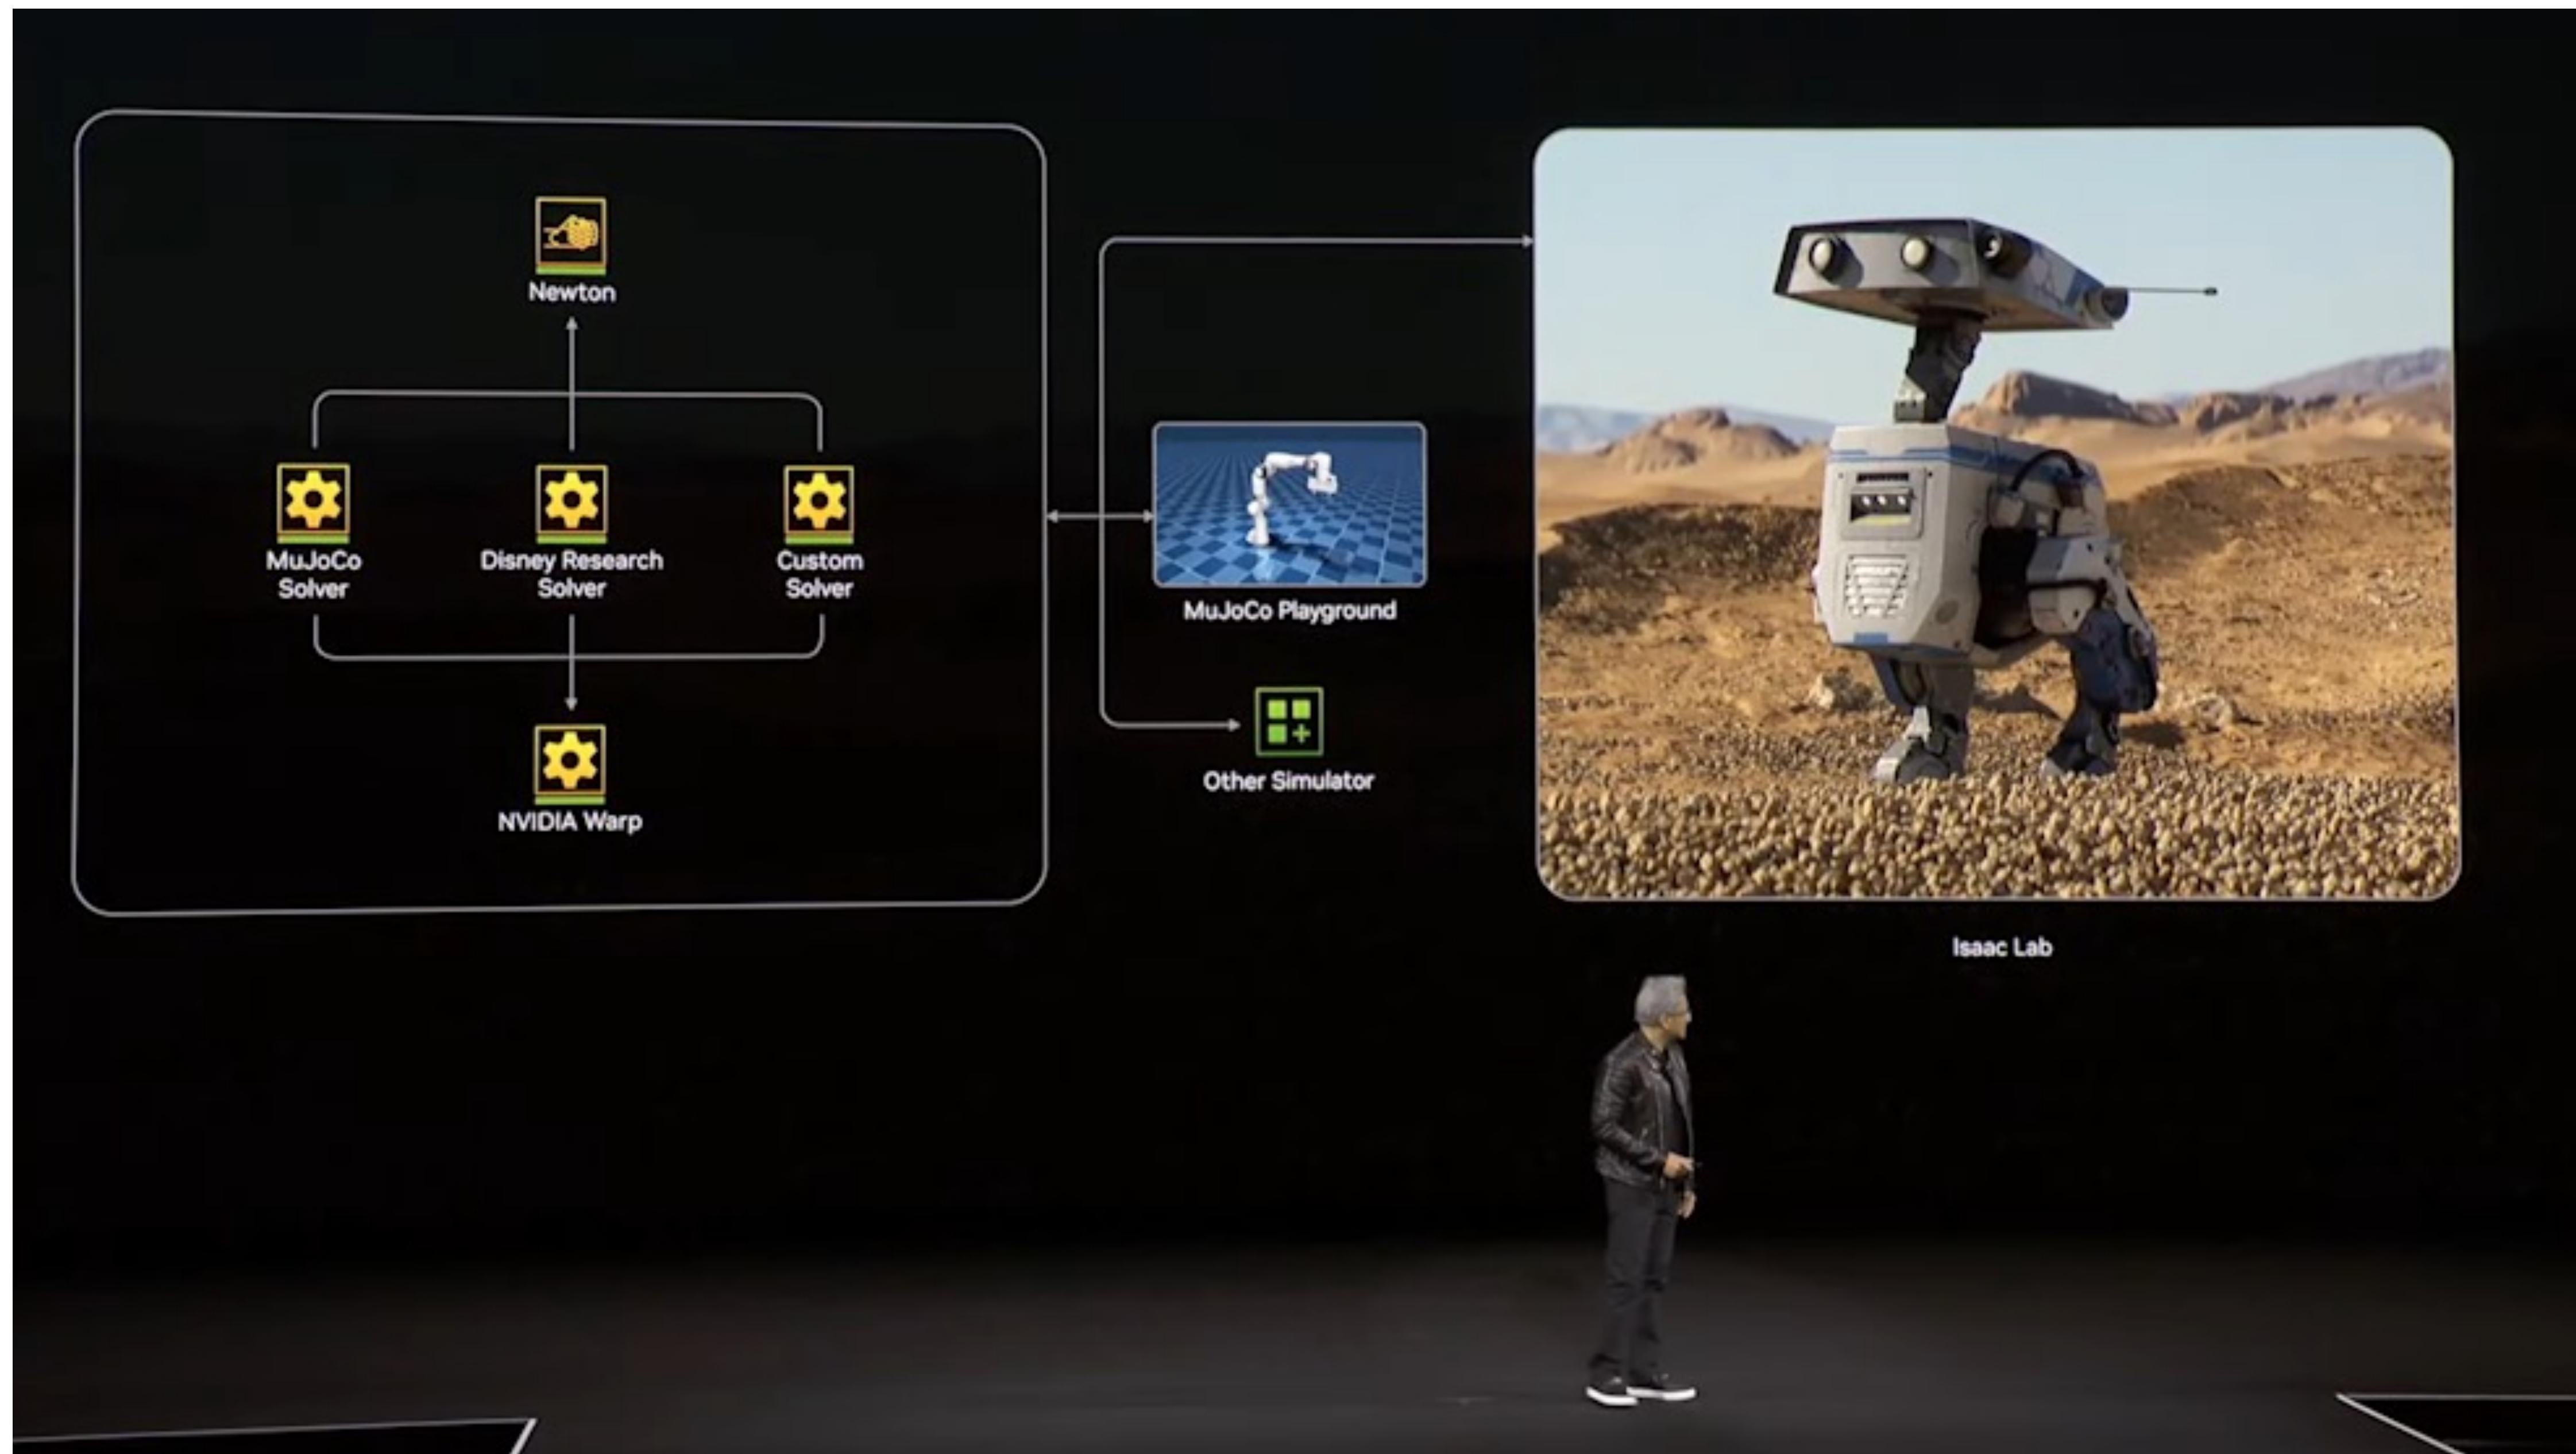


MuJoCo

$$\frac{\partial f}{\partial x}$$



Introducing Newton

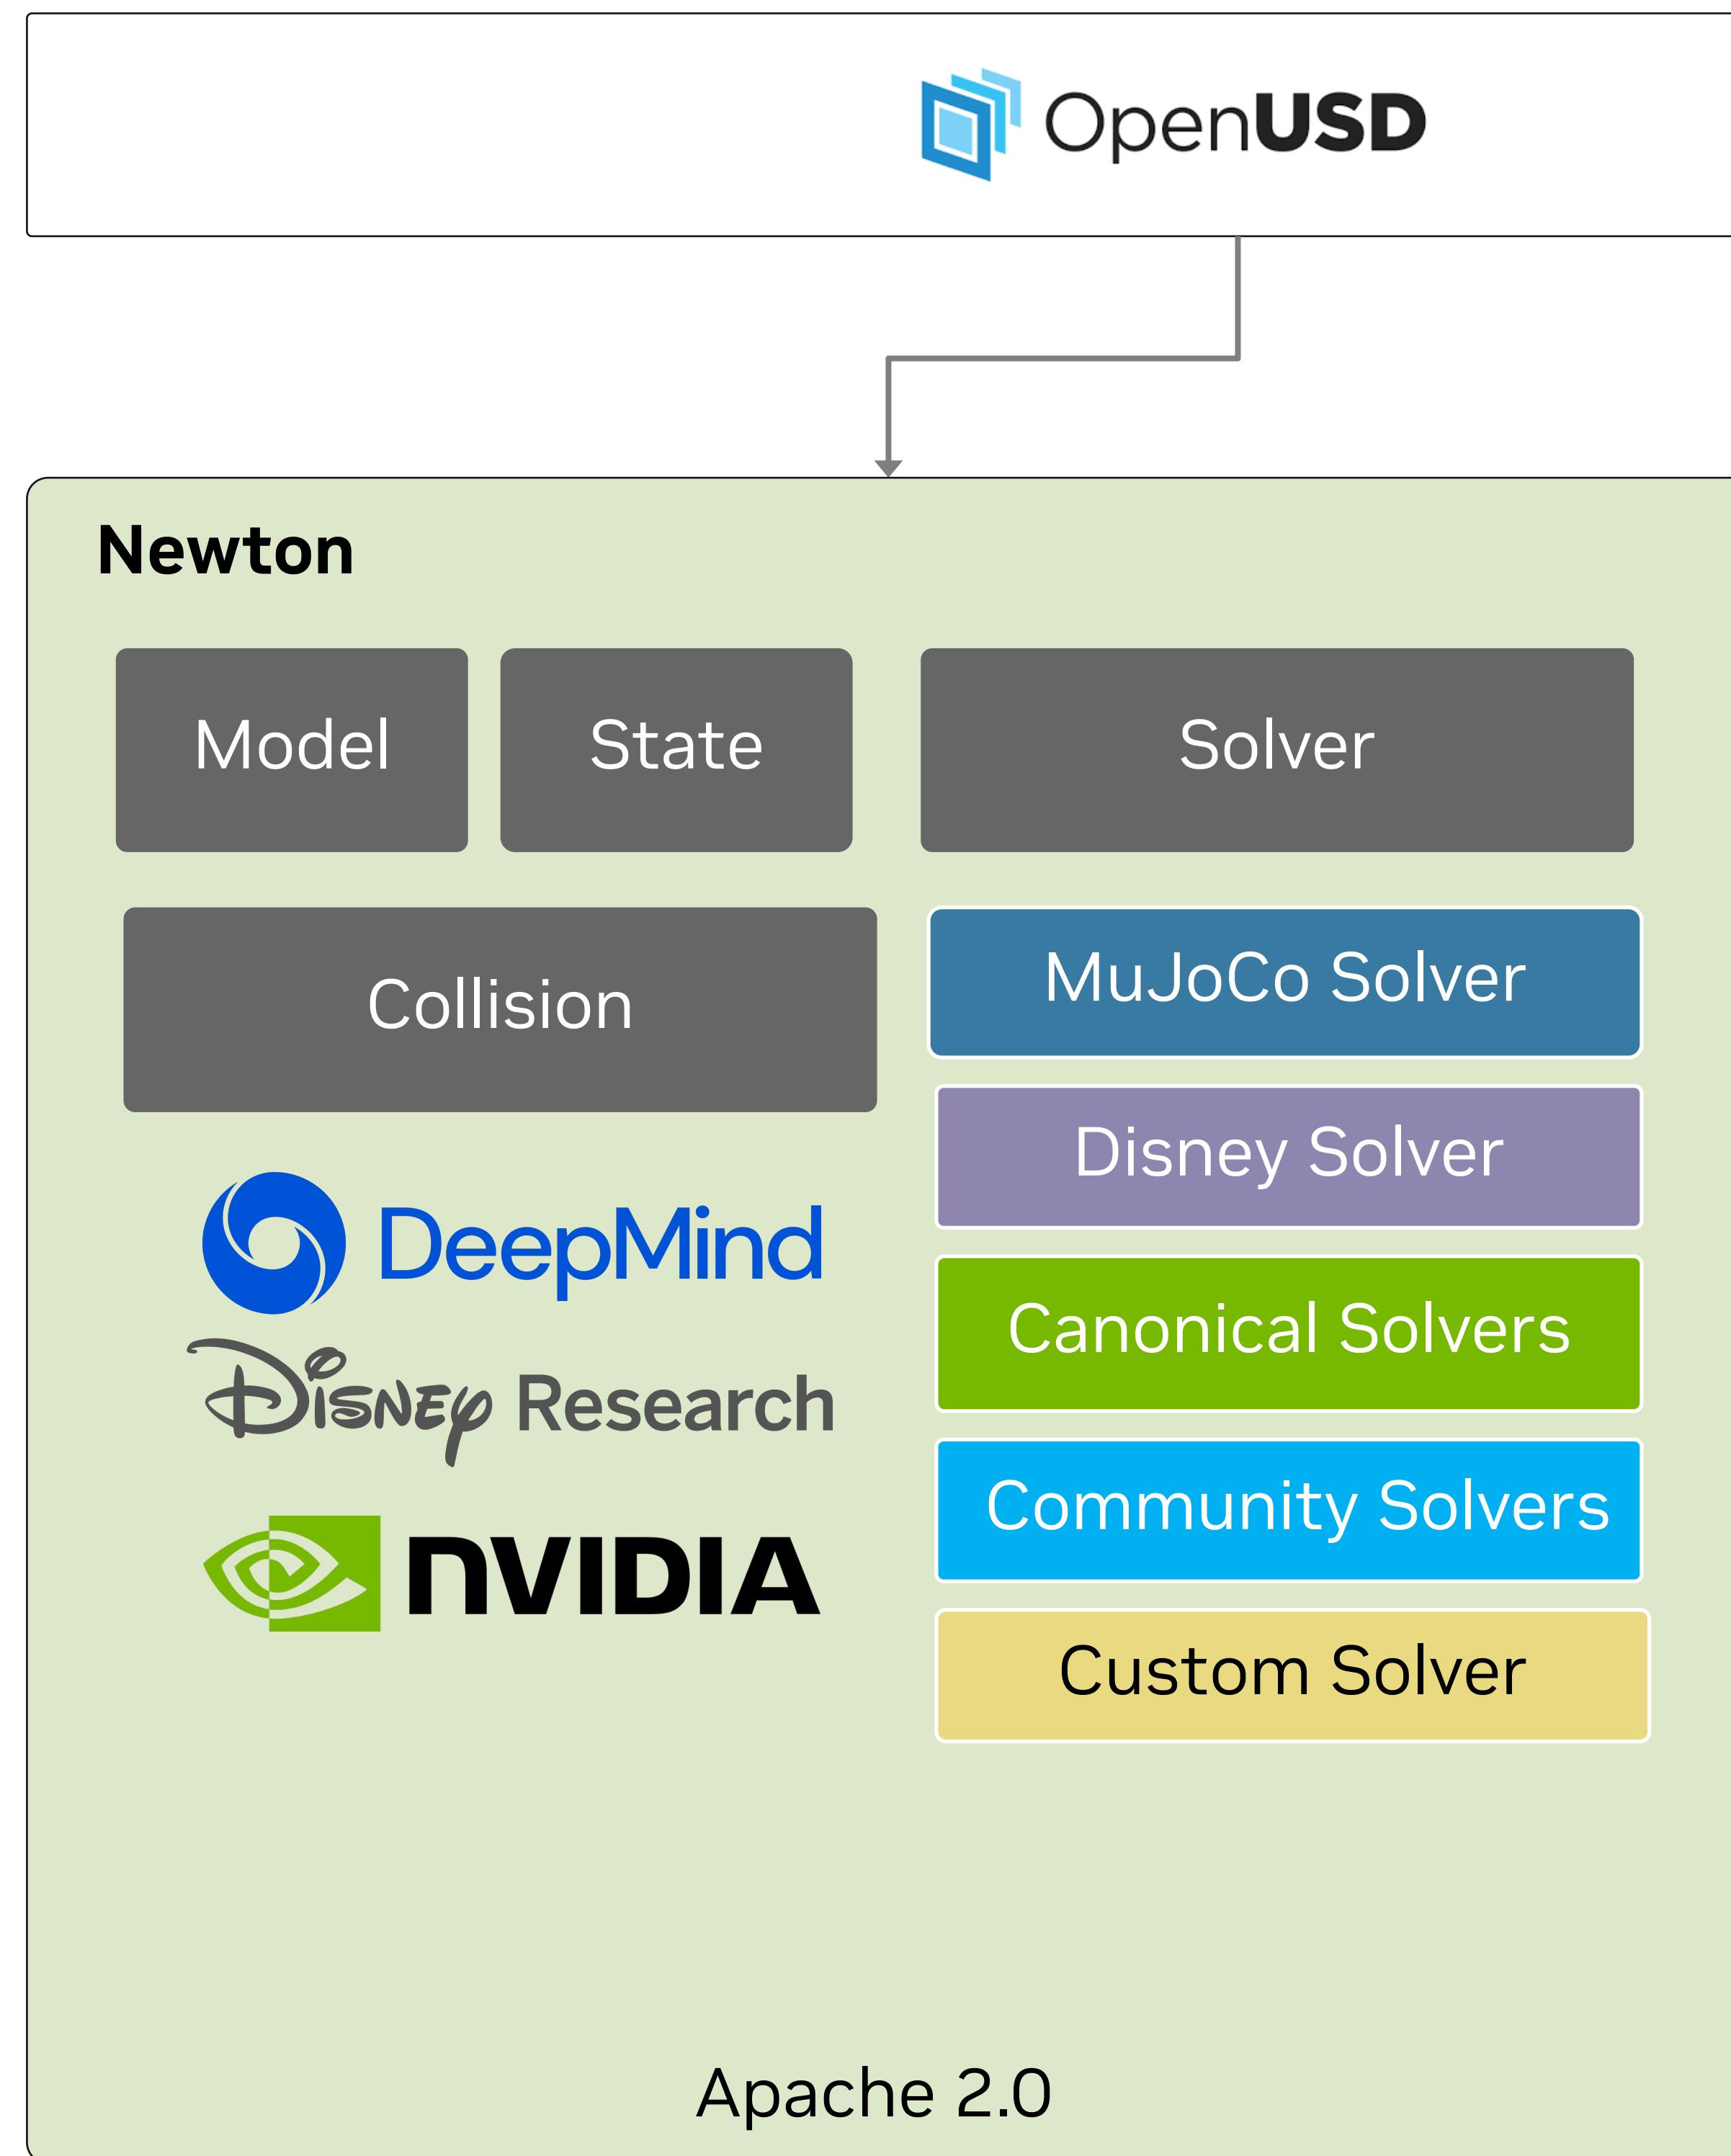


Disney Research

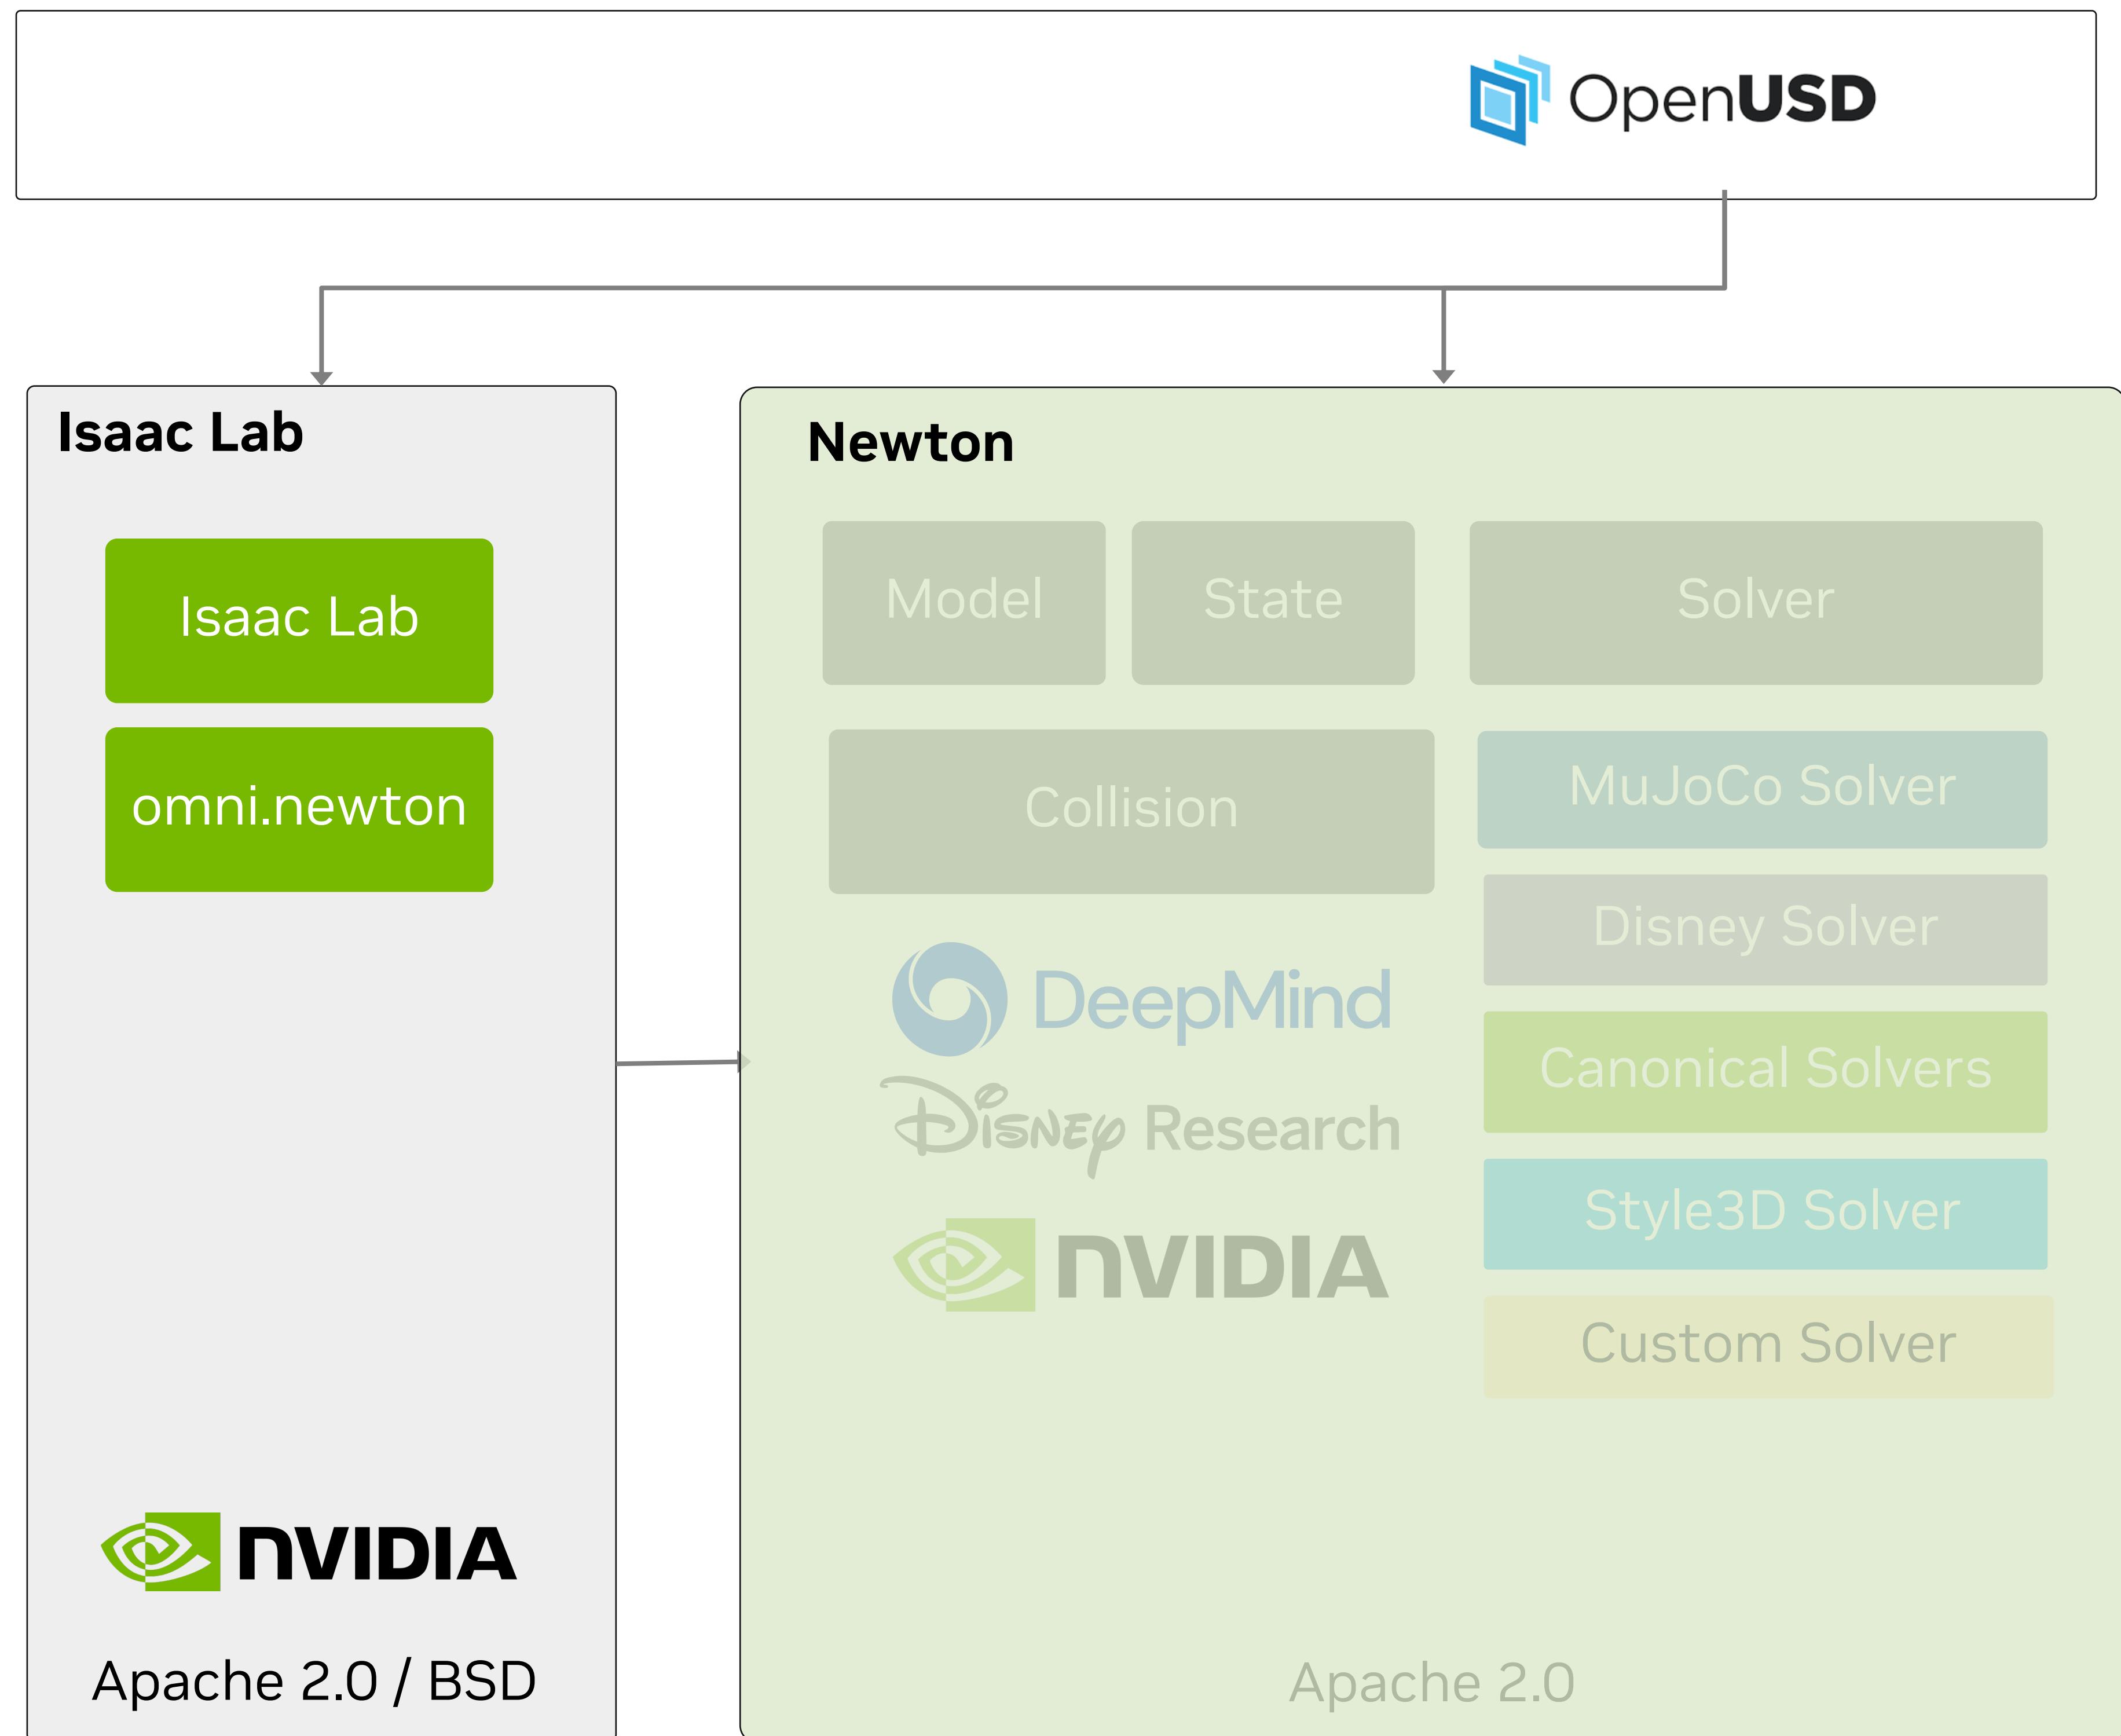
DeepMind

NVIDIA®

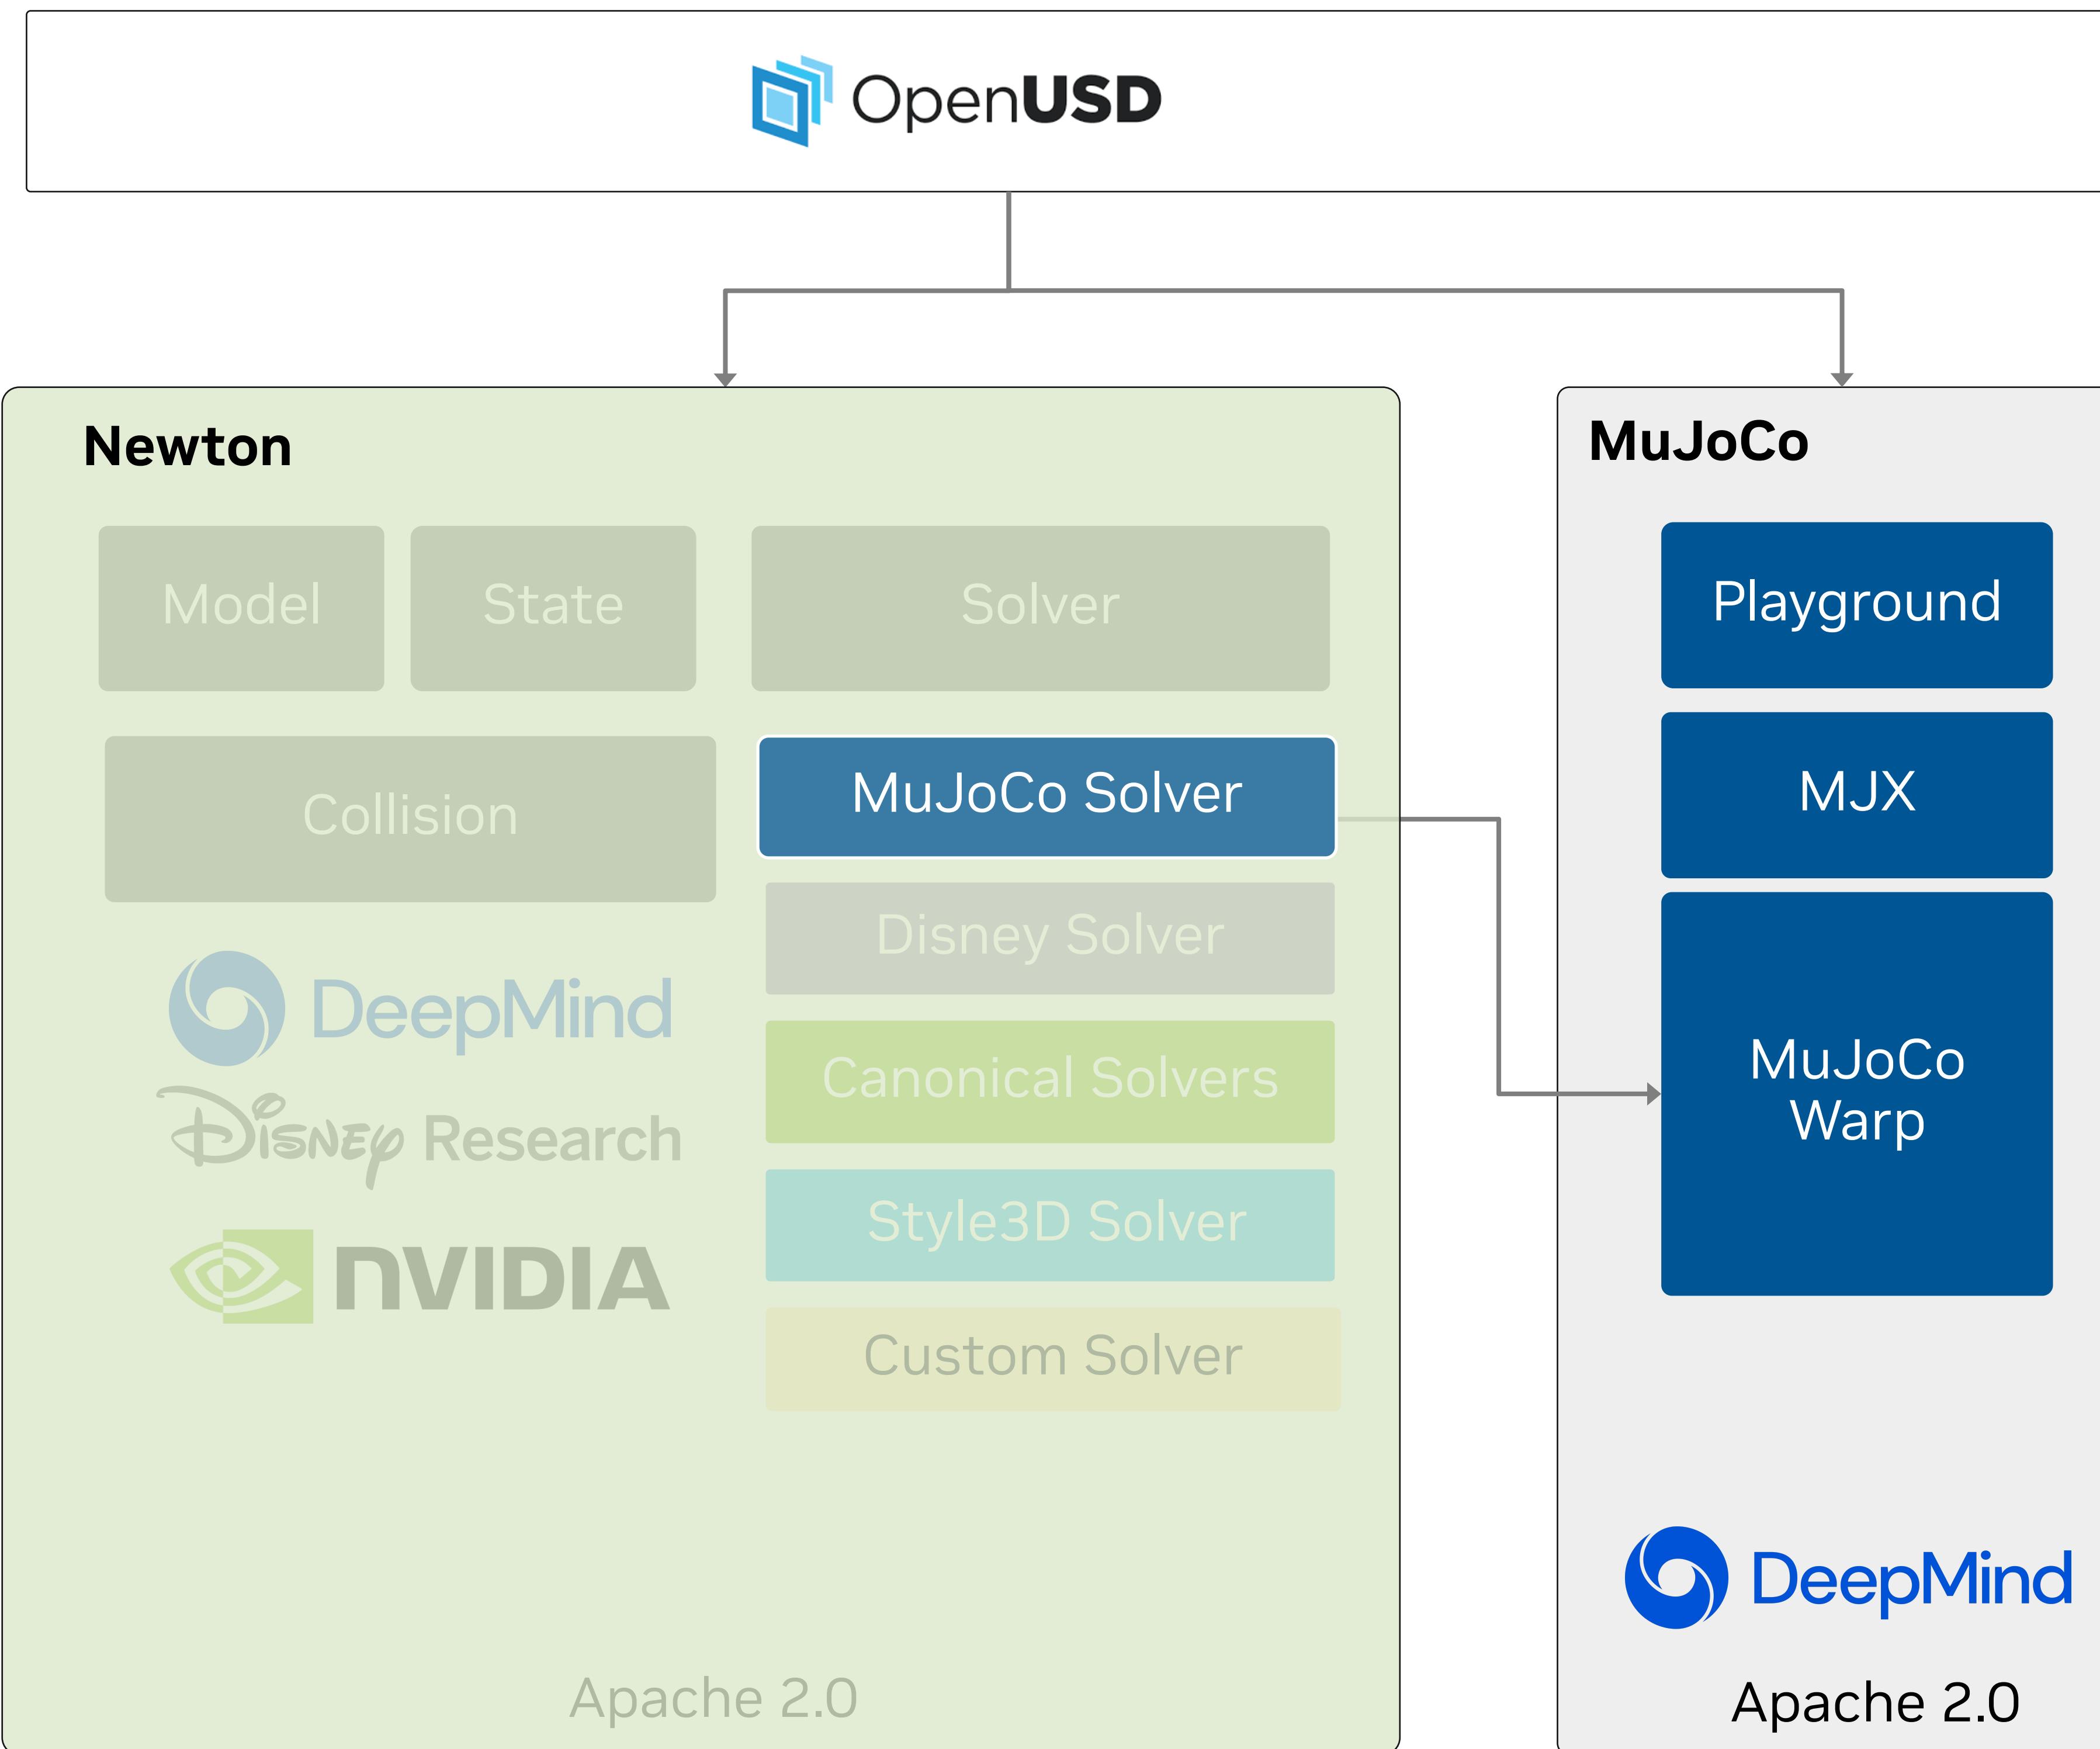
Architecture Overview



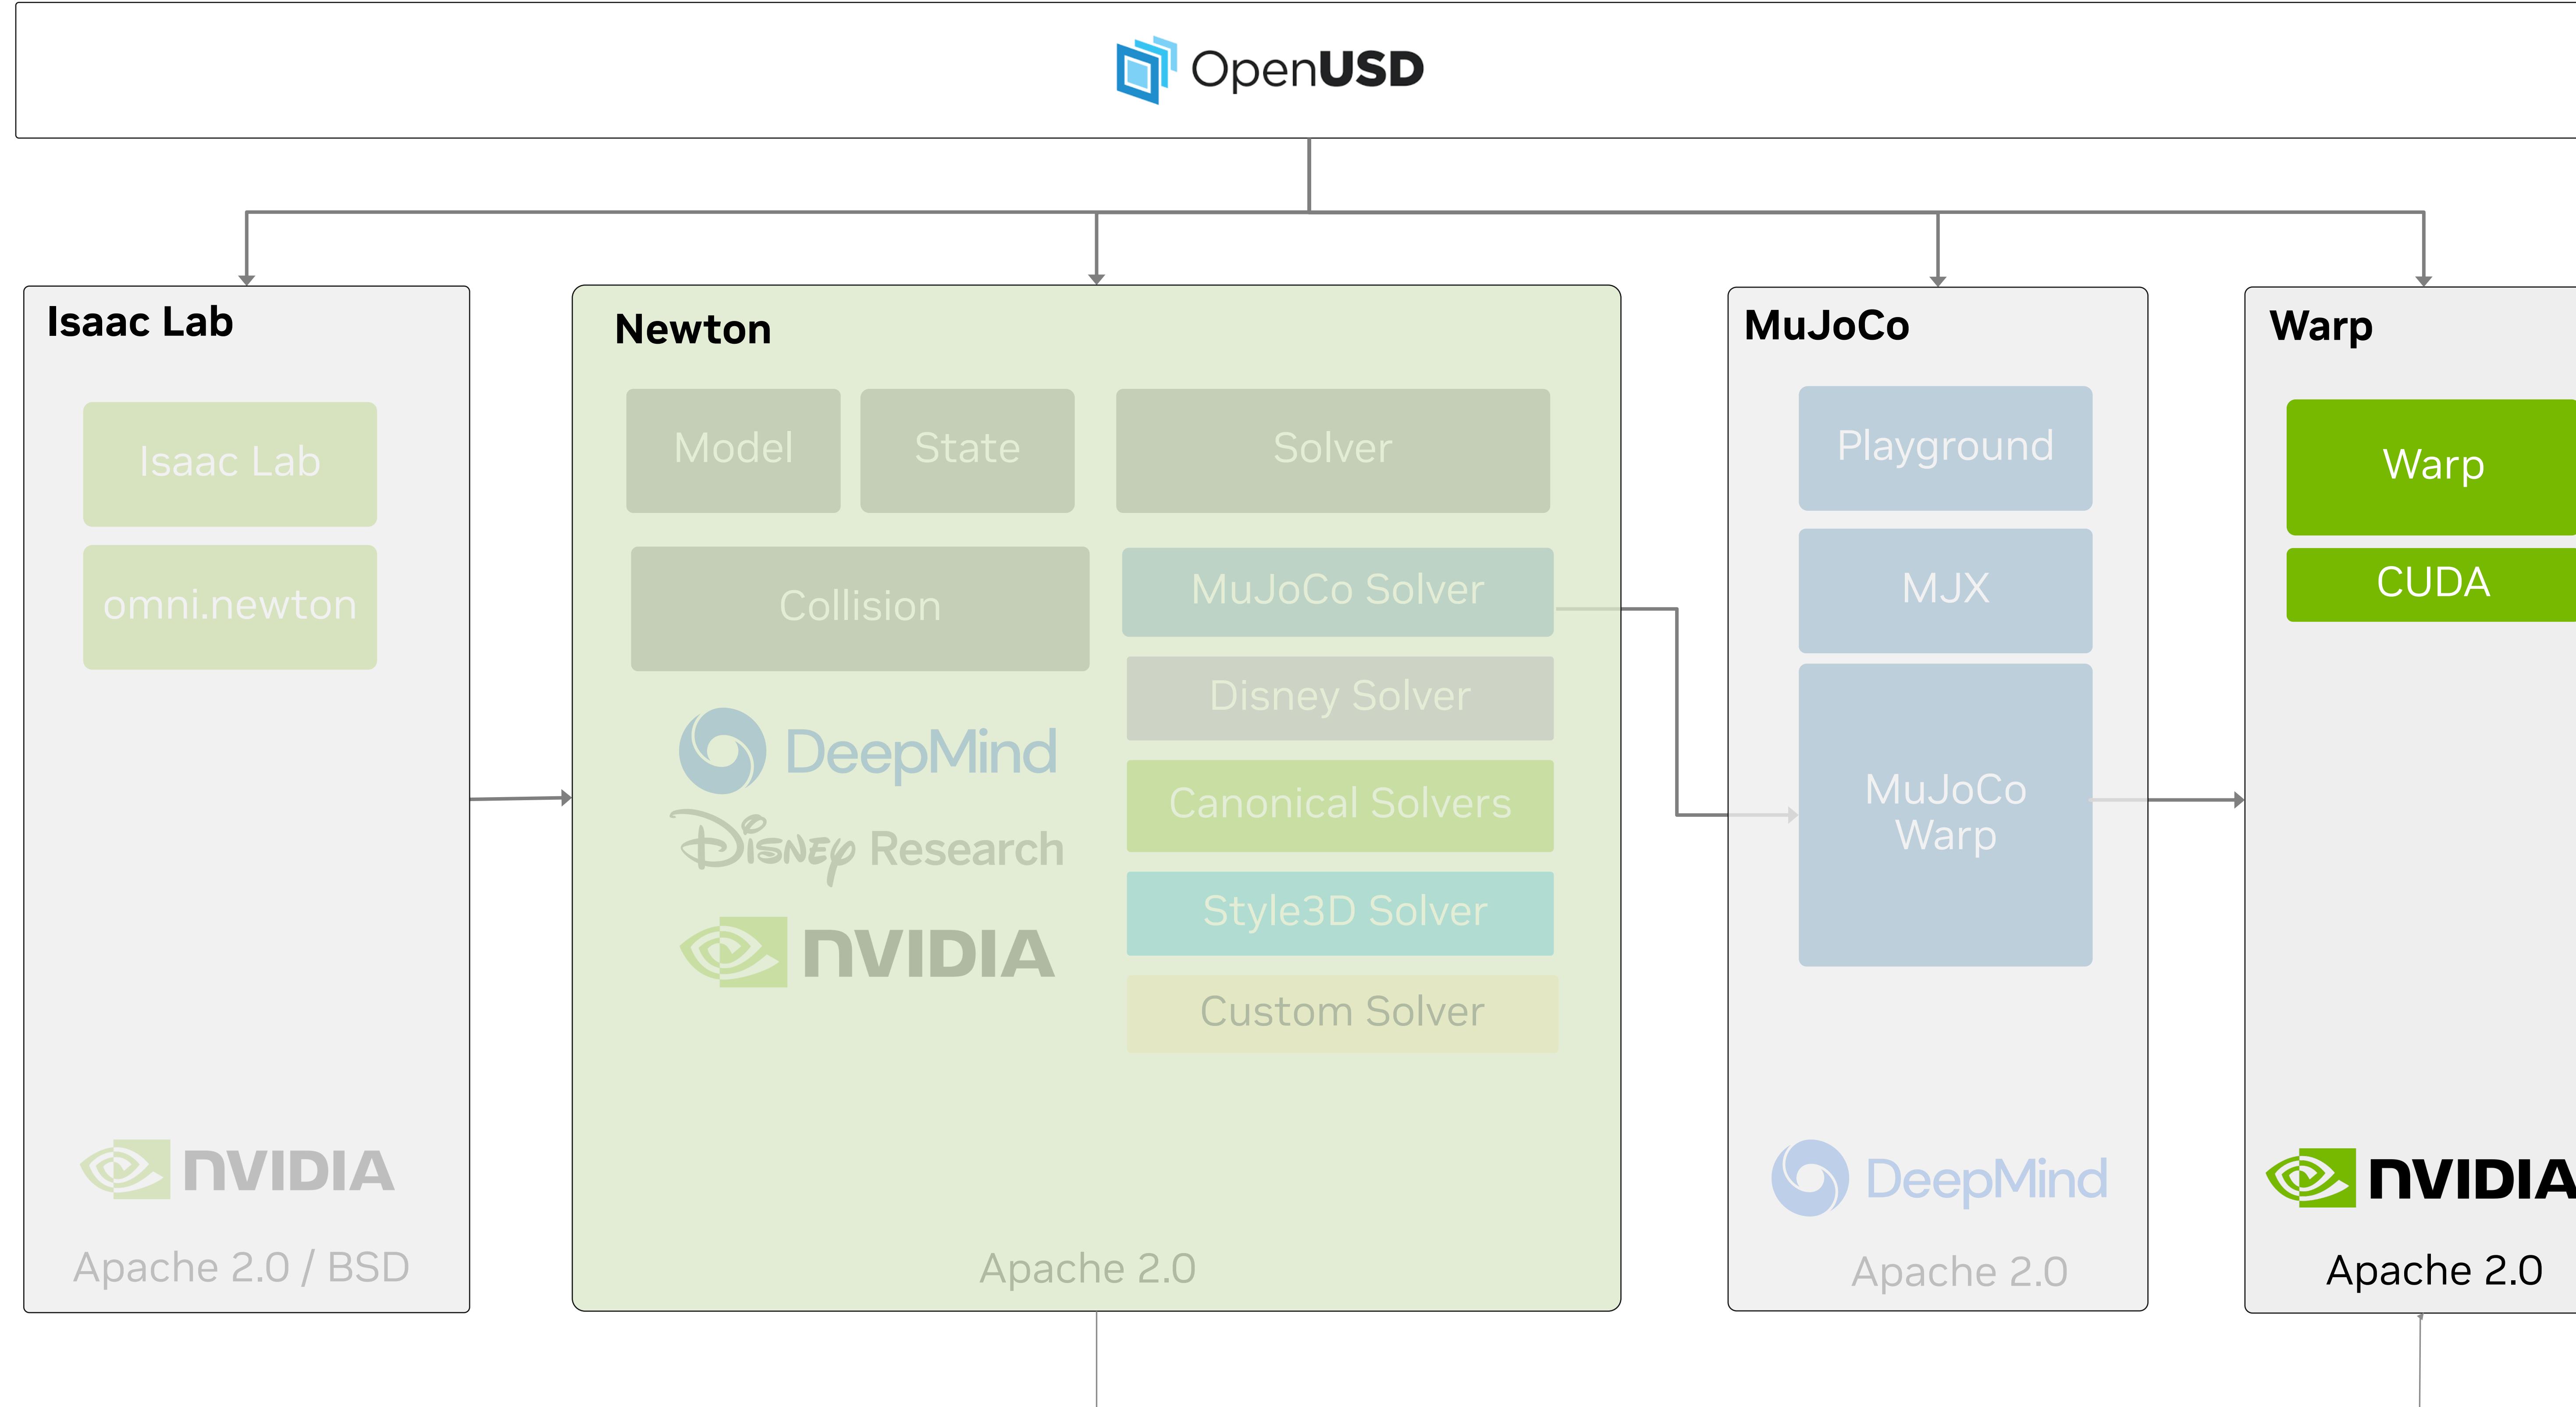
Architecture Overview



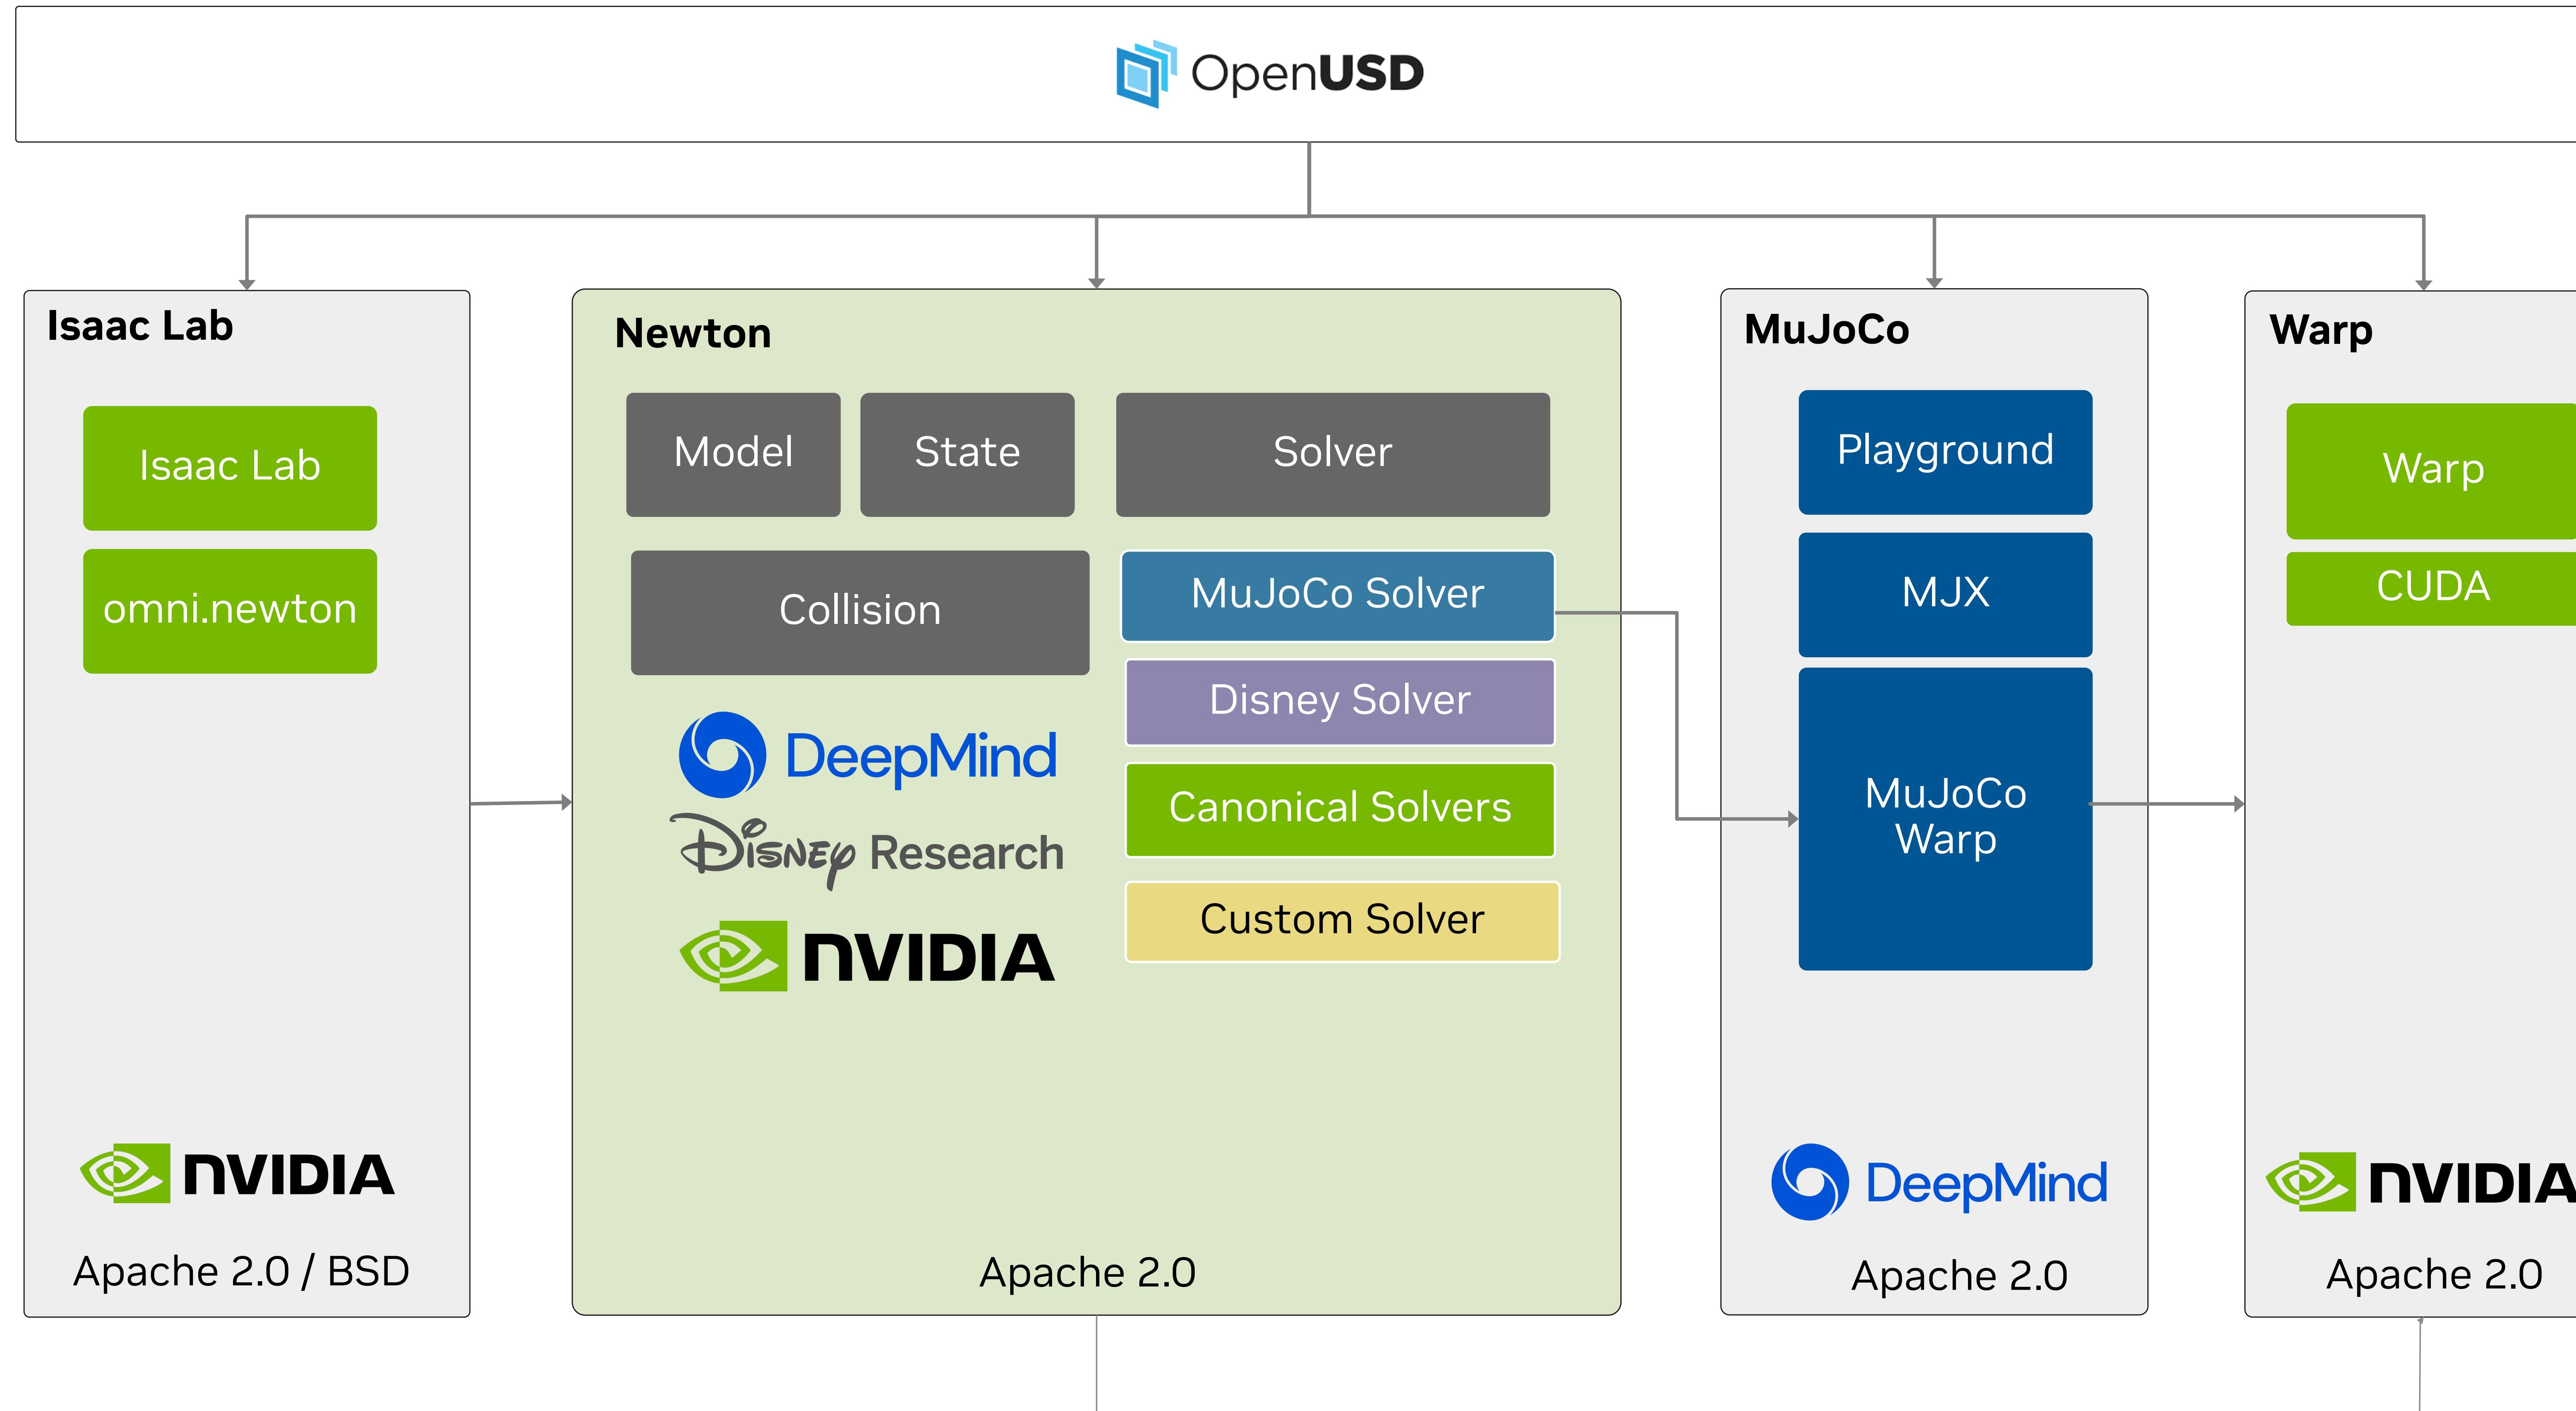
Architecture Overview



Architecture Overview



Architecture Overview



Introducing NVIDIA Warp

Performance

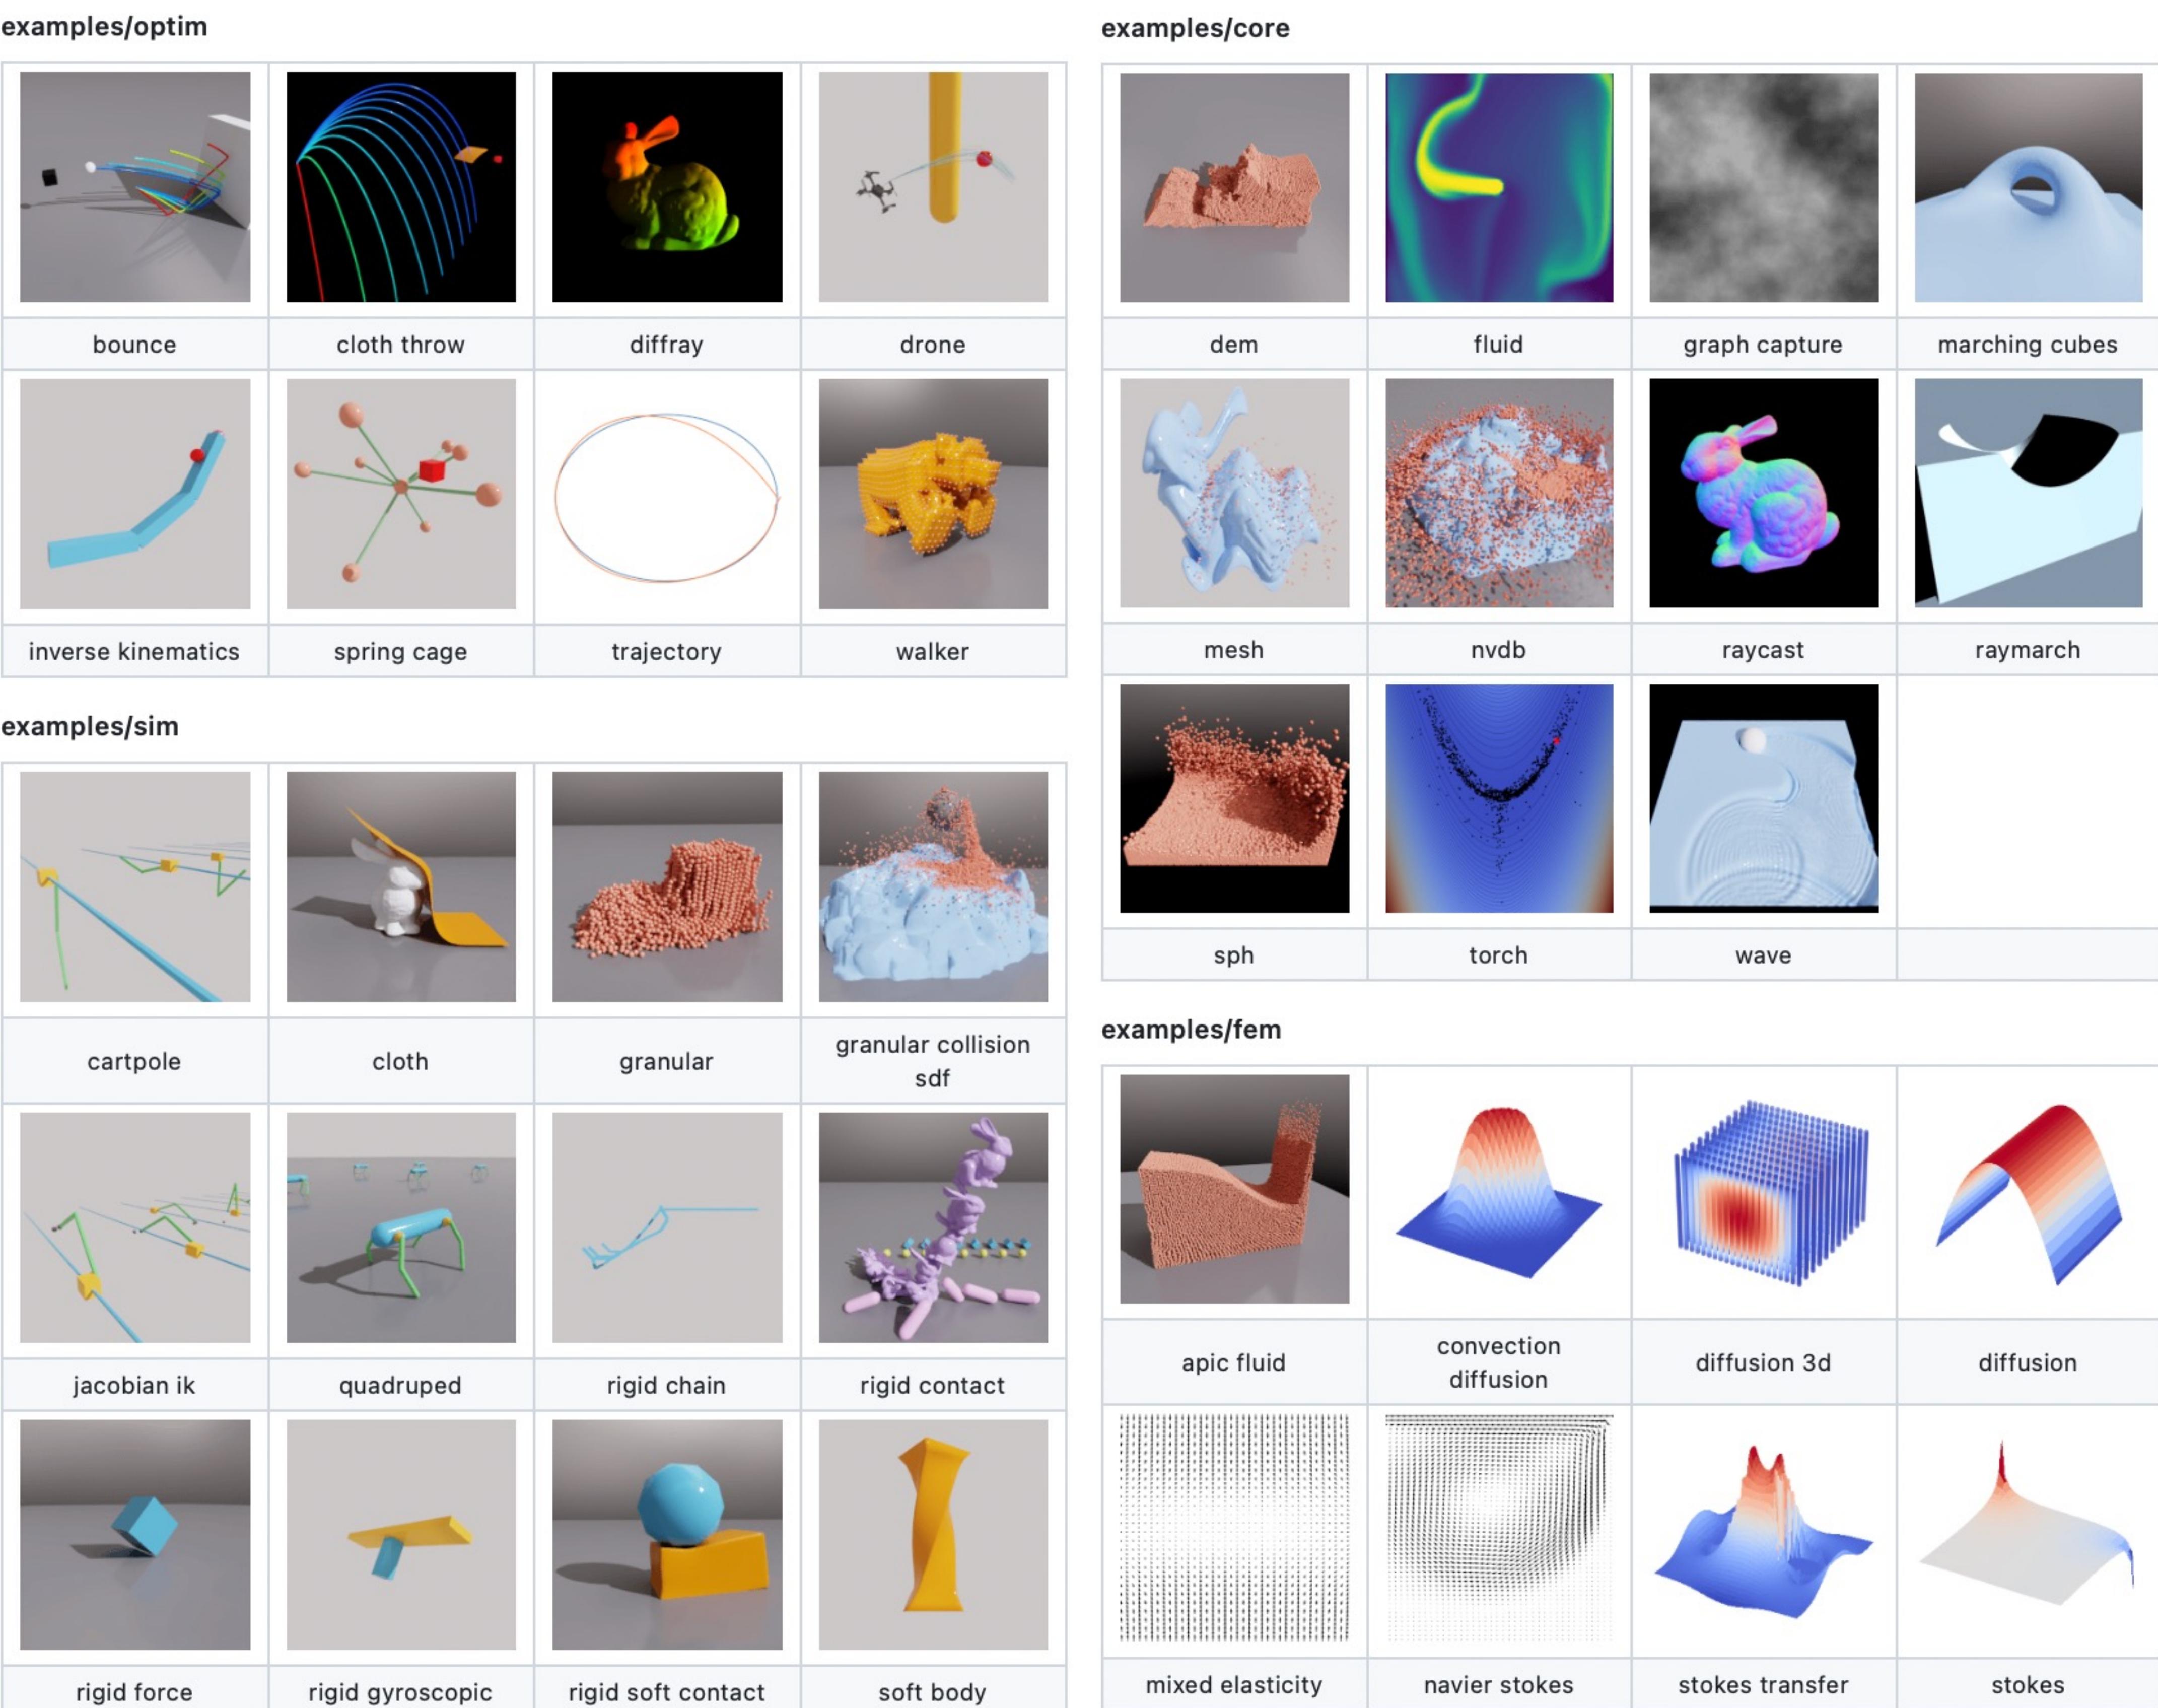
Native CUDA, JIT-compiled
CUDA-X libraries exposure

Ease of Use

Pure Python
Tile-based programming
Native simulation primitives

Capability

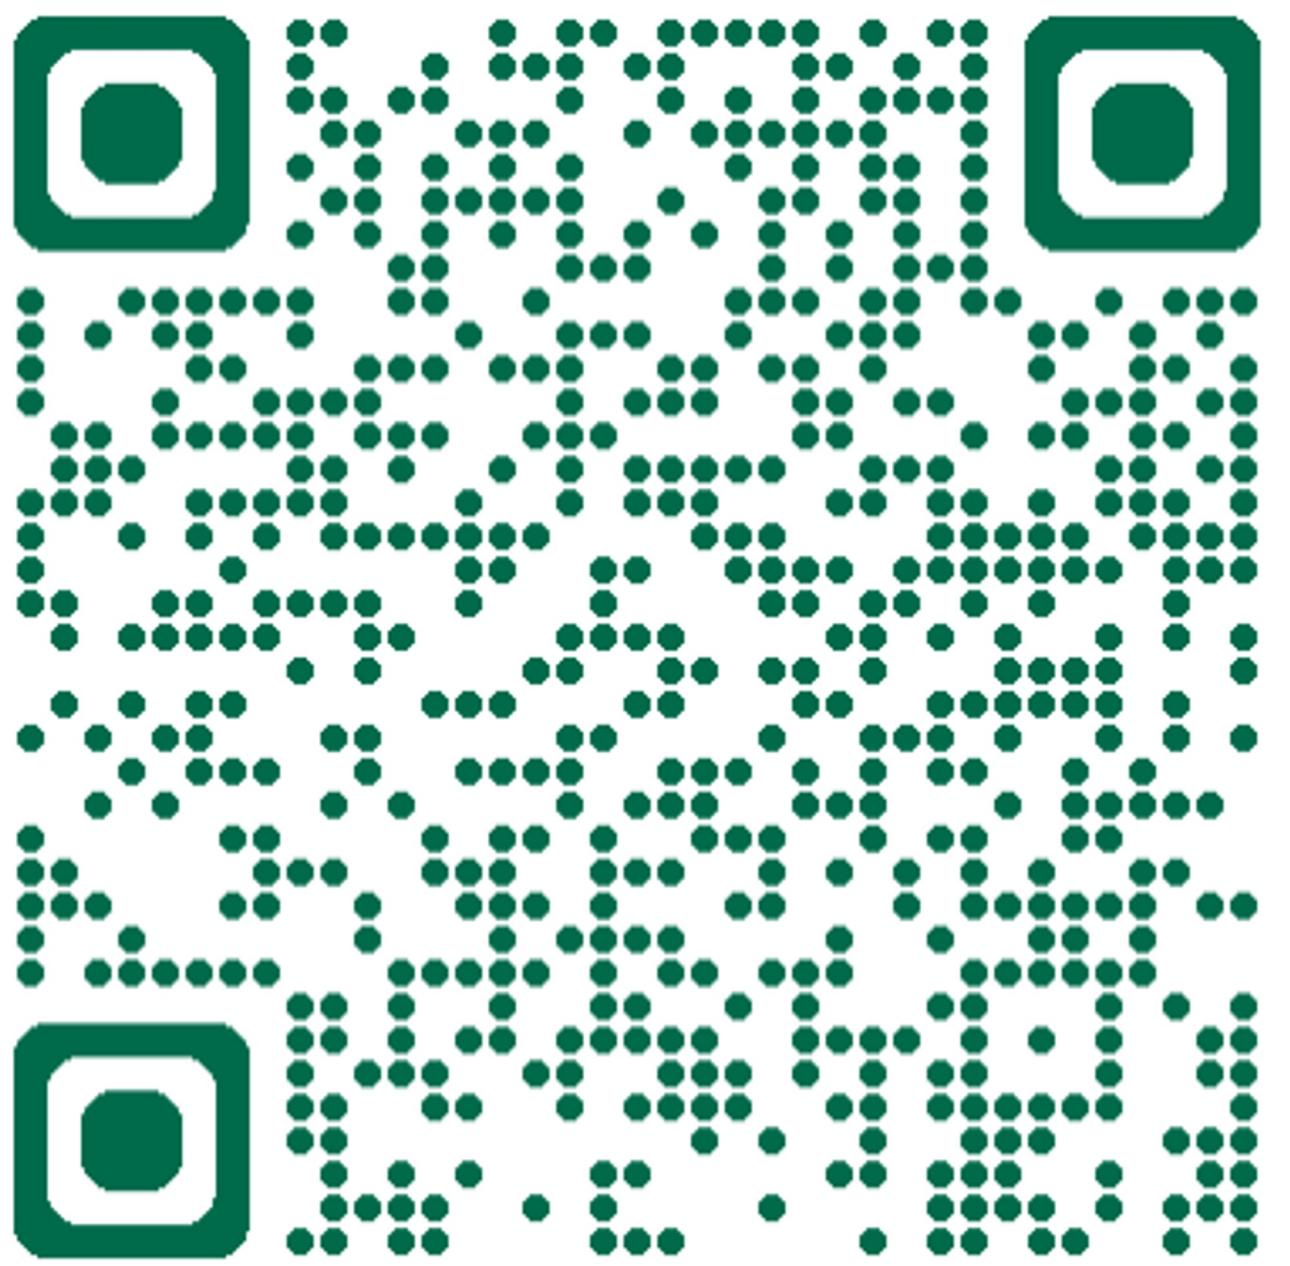
Automatic Differentiation
DL framework interop



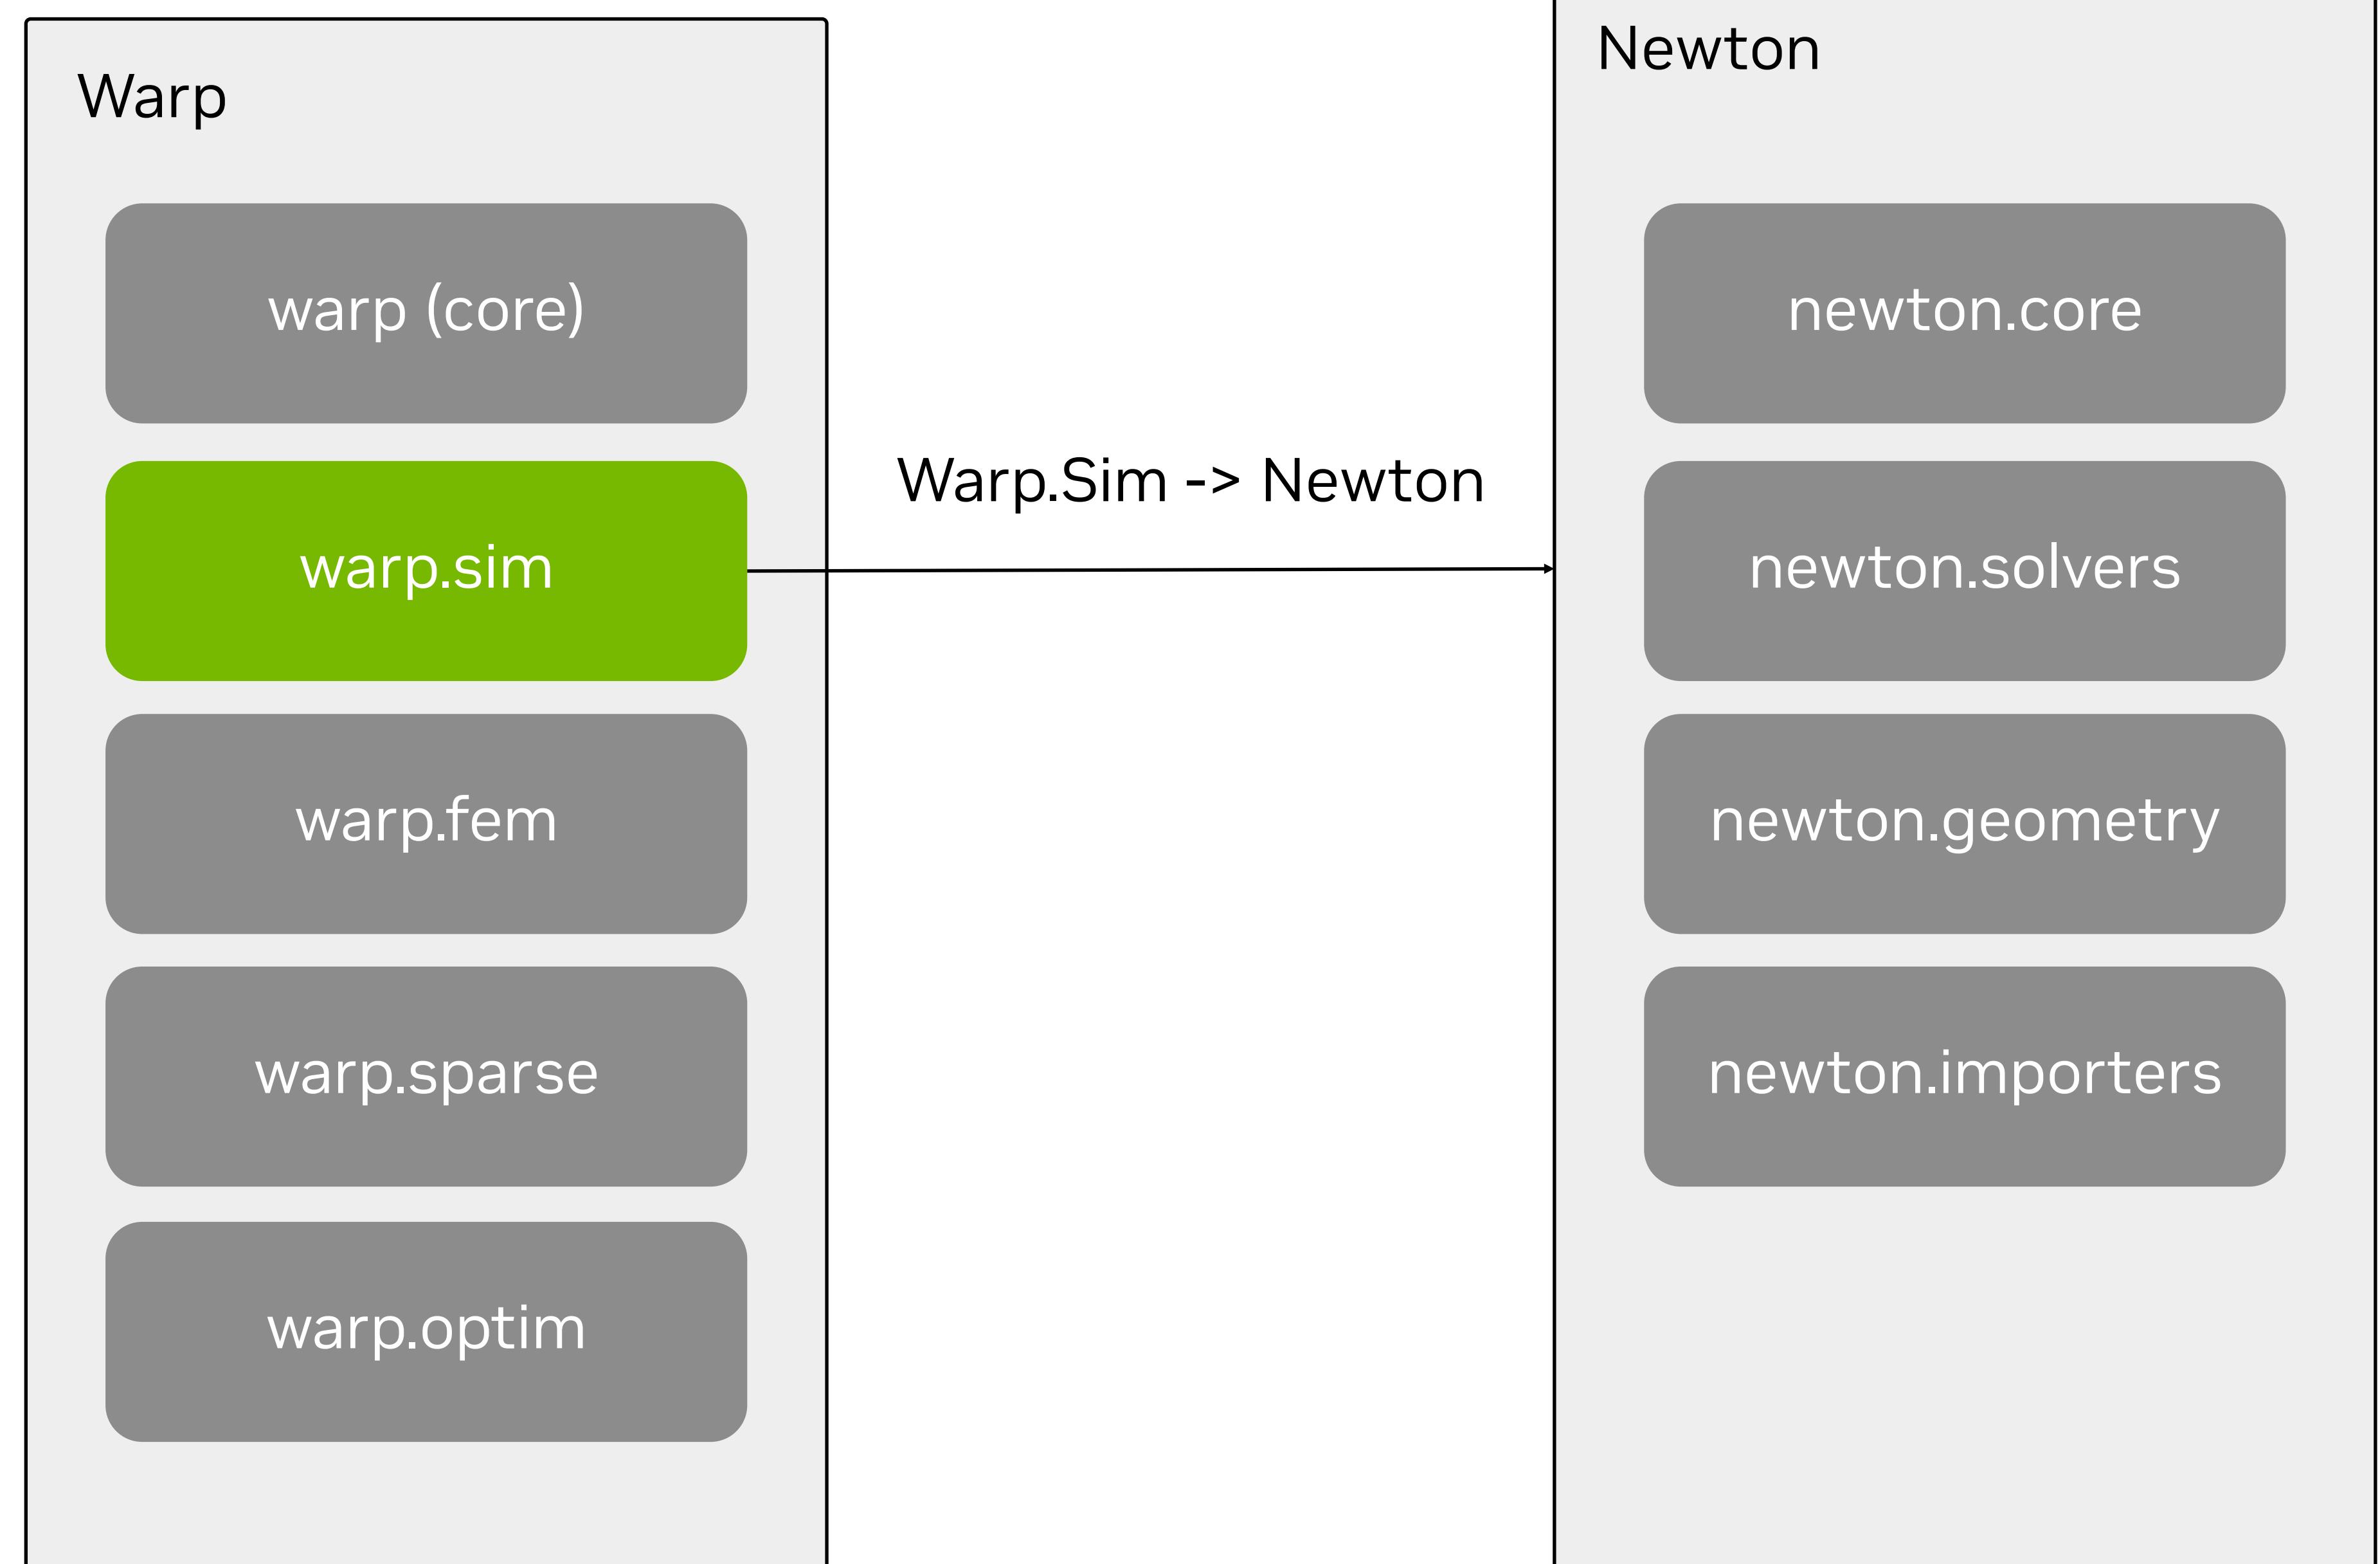
warp.sim Migration

- Newton architecture is derived from `warp.sim`
- Rebranding and deprecating over time
- Warp is a standalone framework

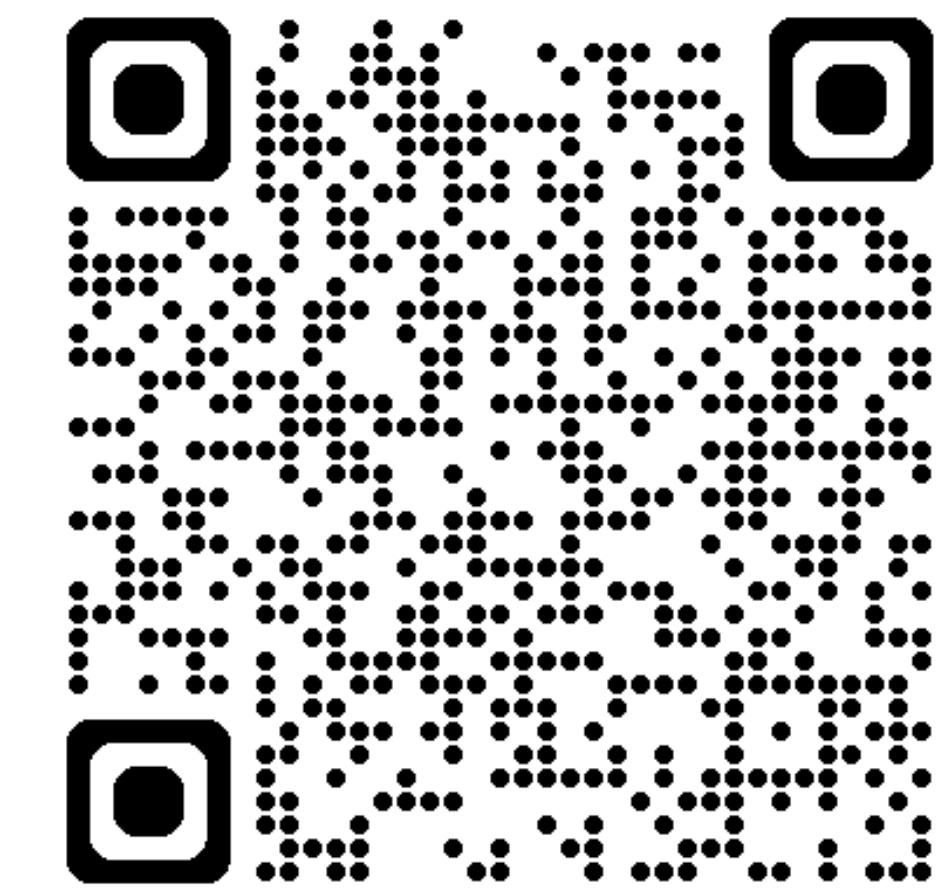
Open Source, Apache 2.0



`pip install warp-lang`



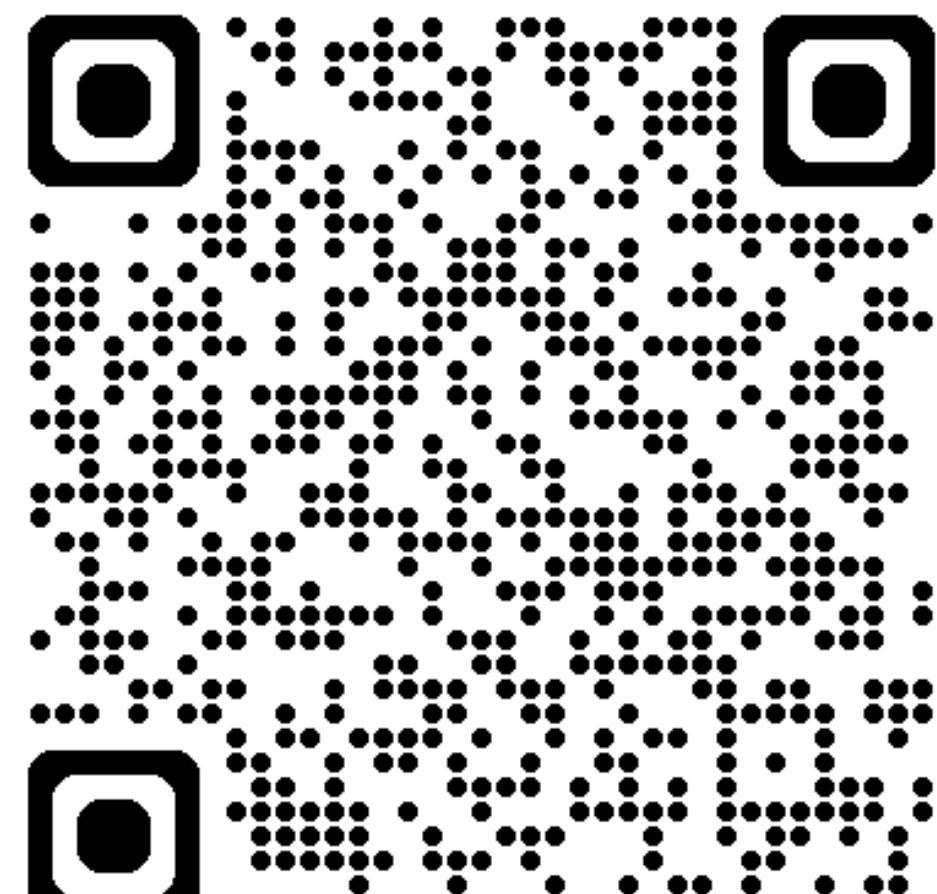
Warp Talks



Warp: A High-performance Python Framework for GPU Simulation and Graphics

Miles Macklin, NVIDIA

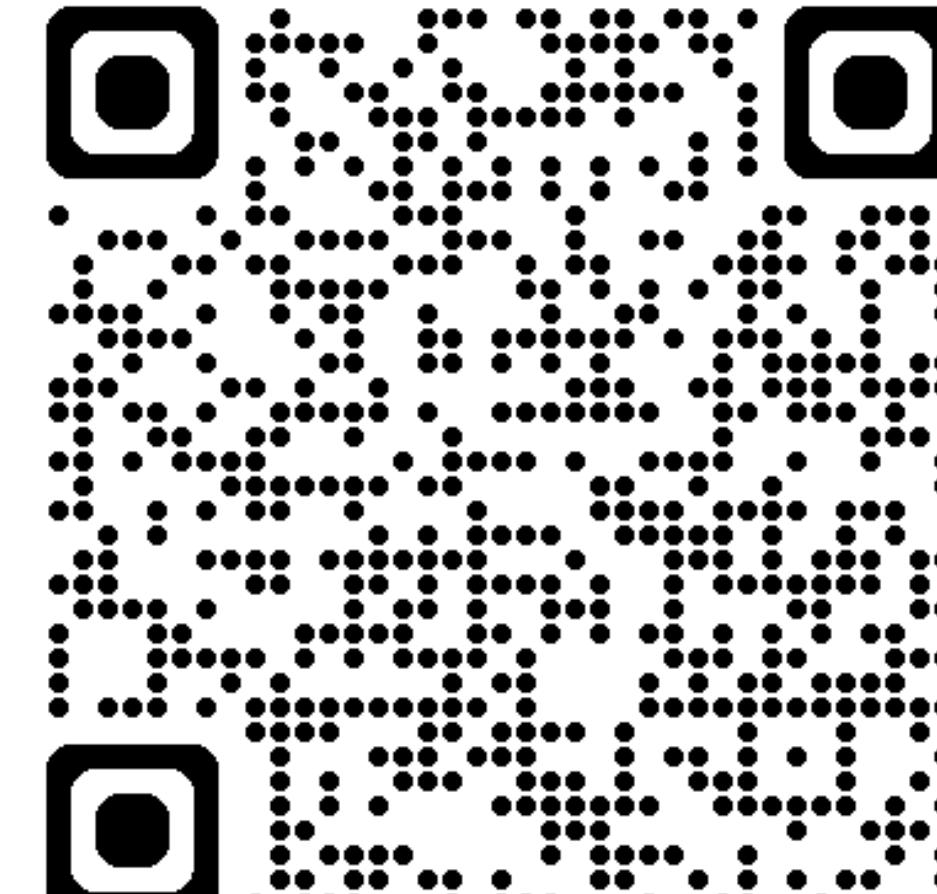
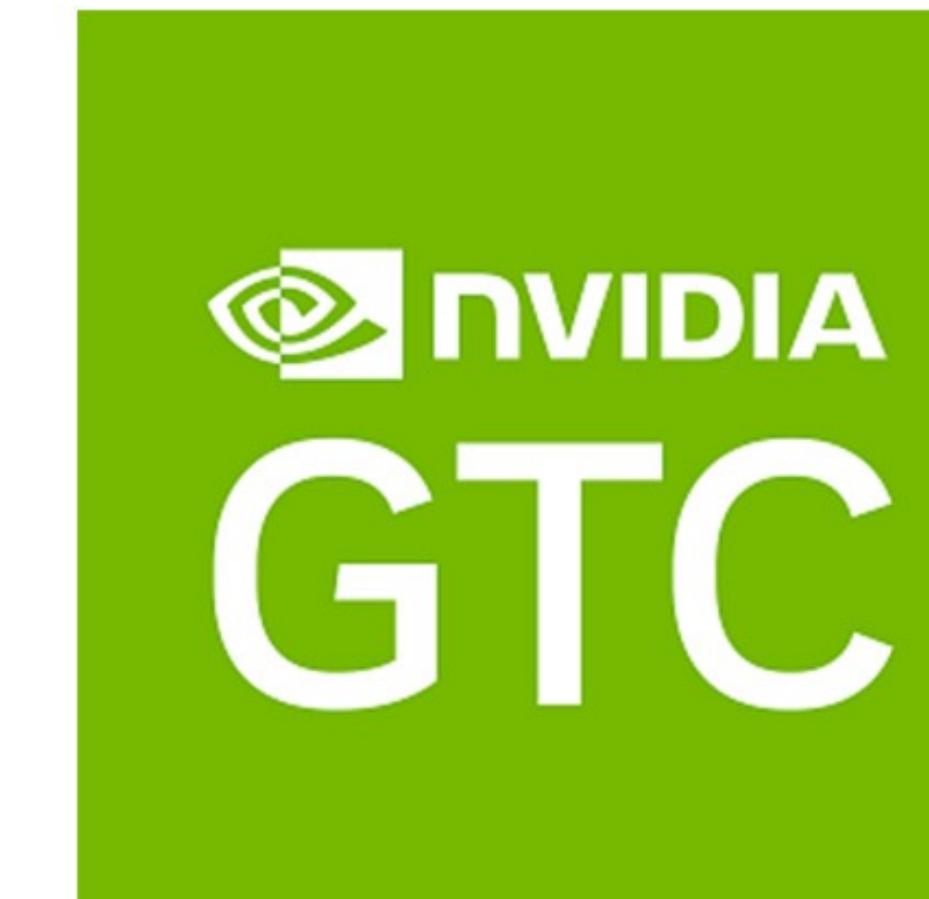
GTC, San Jose, March 2022



Warp: Advancing Simulation AI with Differentiable GPU Computing in Python

Miles Macklin, NVIDIA

GTC, San Jose, March 2024



Distributed and Differentiable Fluid Dynamics with Accelerated Python

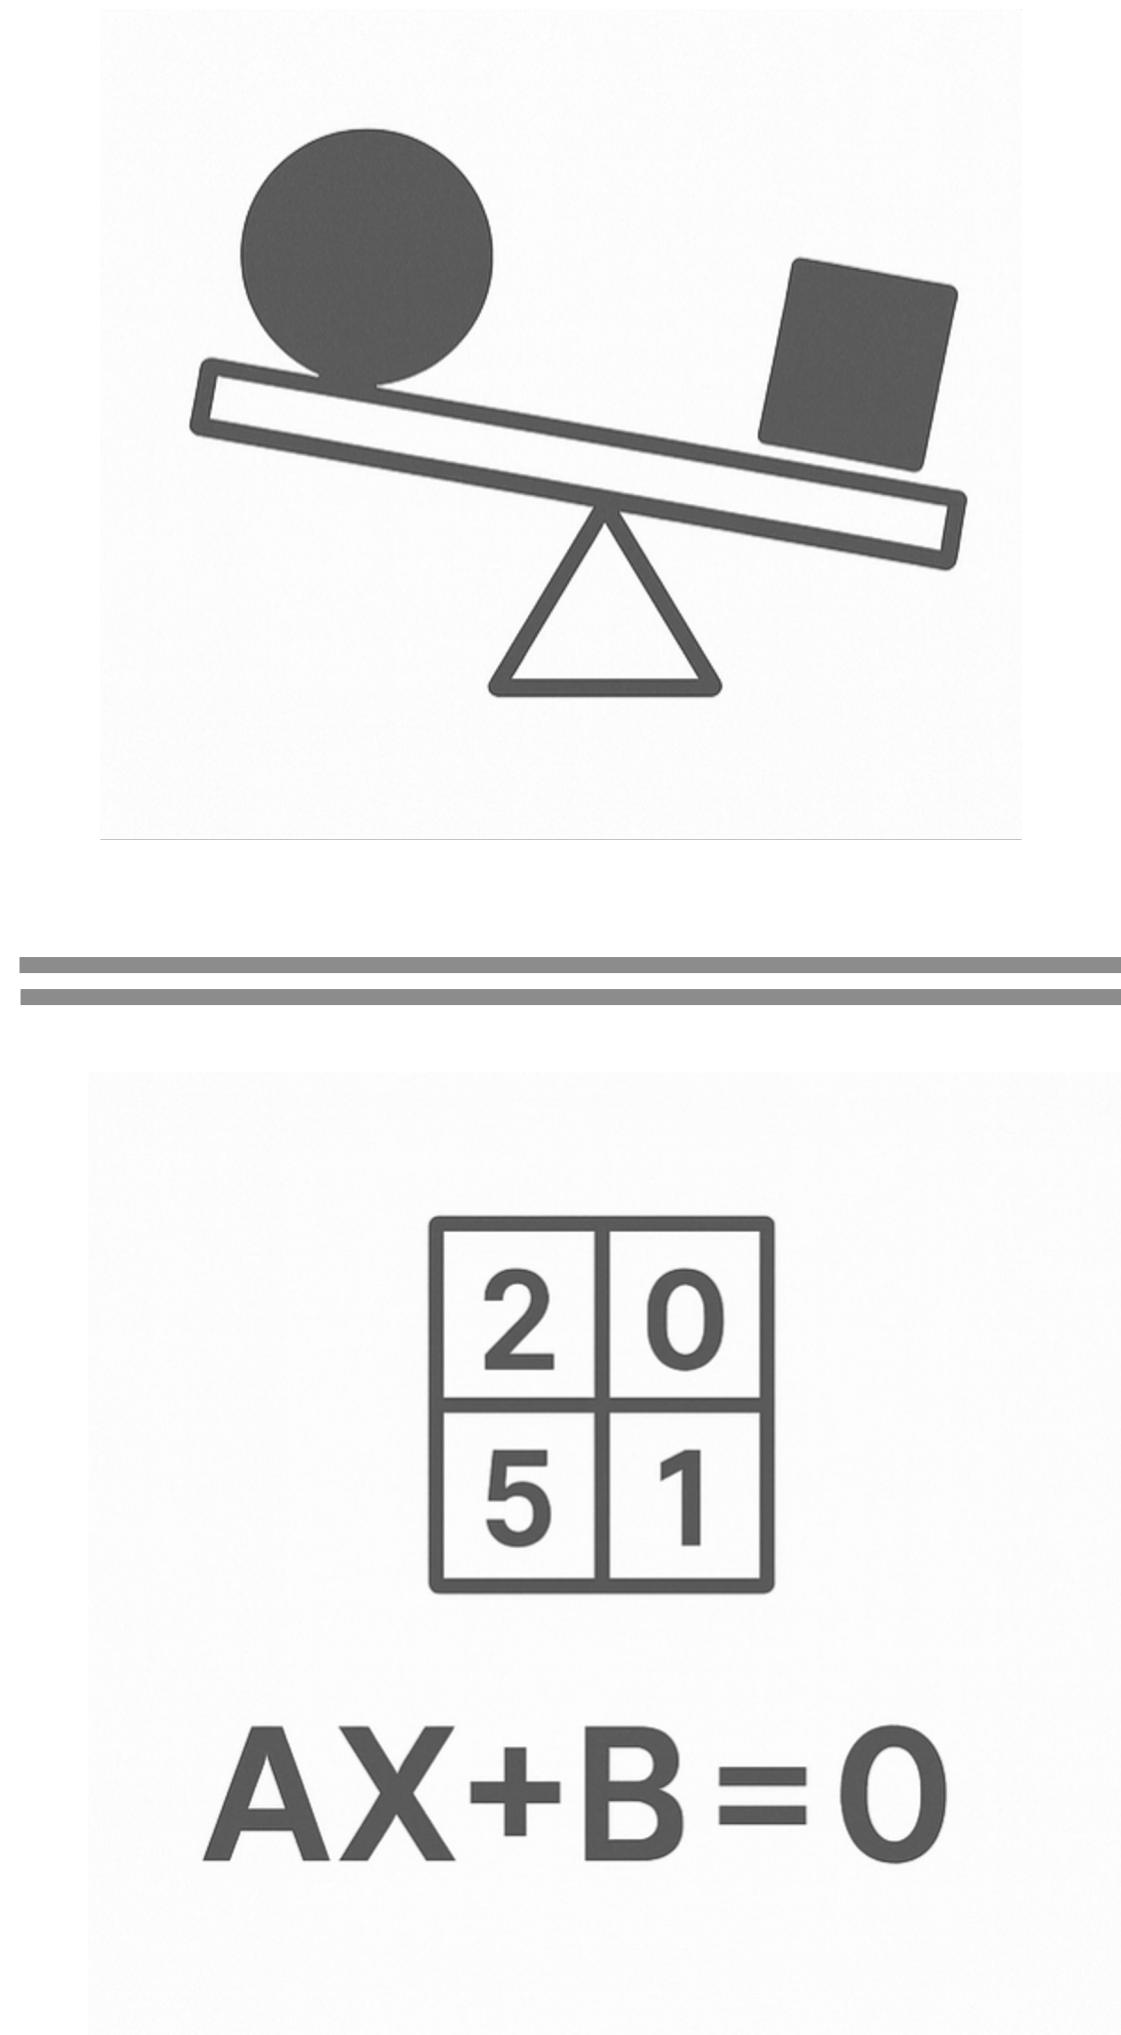
Massimiliano Meneghin, Autodesk
Lukasz Wawrzyniak, NVIDIA

GTC, San Jose, March 2025

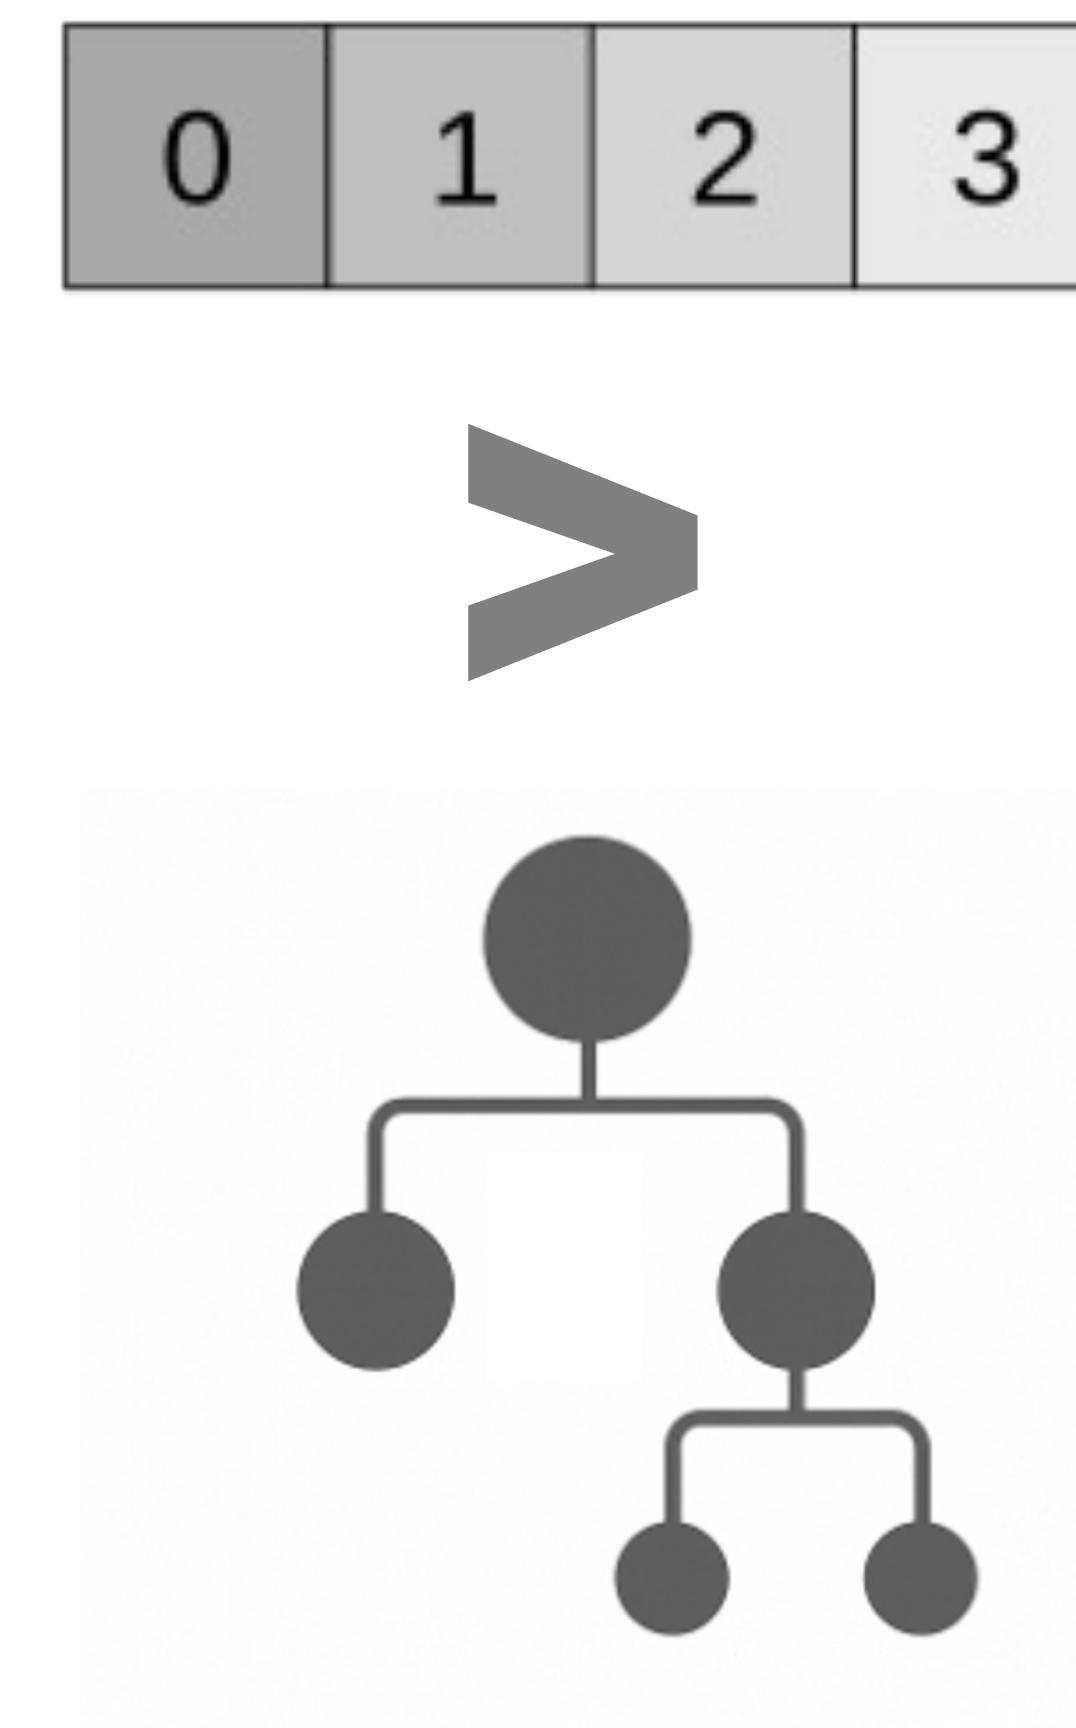
Architecture Design and GPU Acceleration

Newton Design Principles

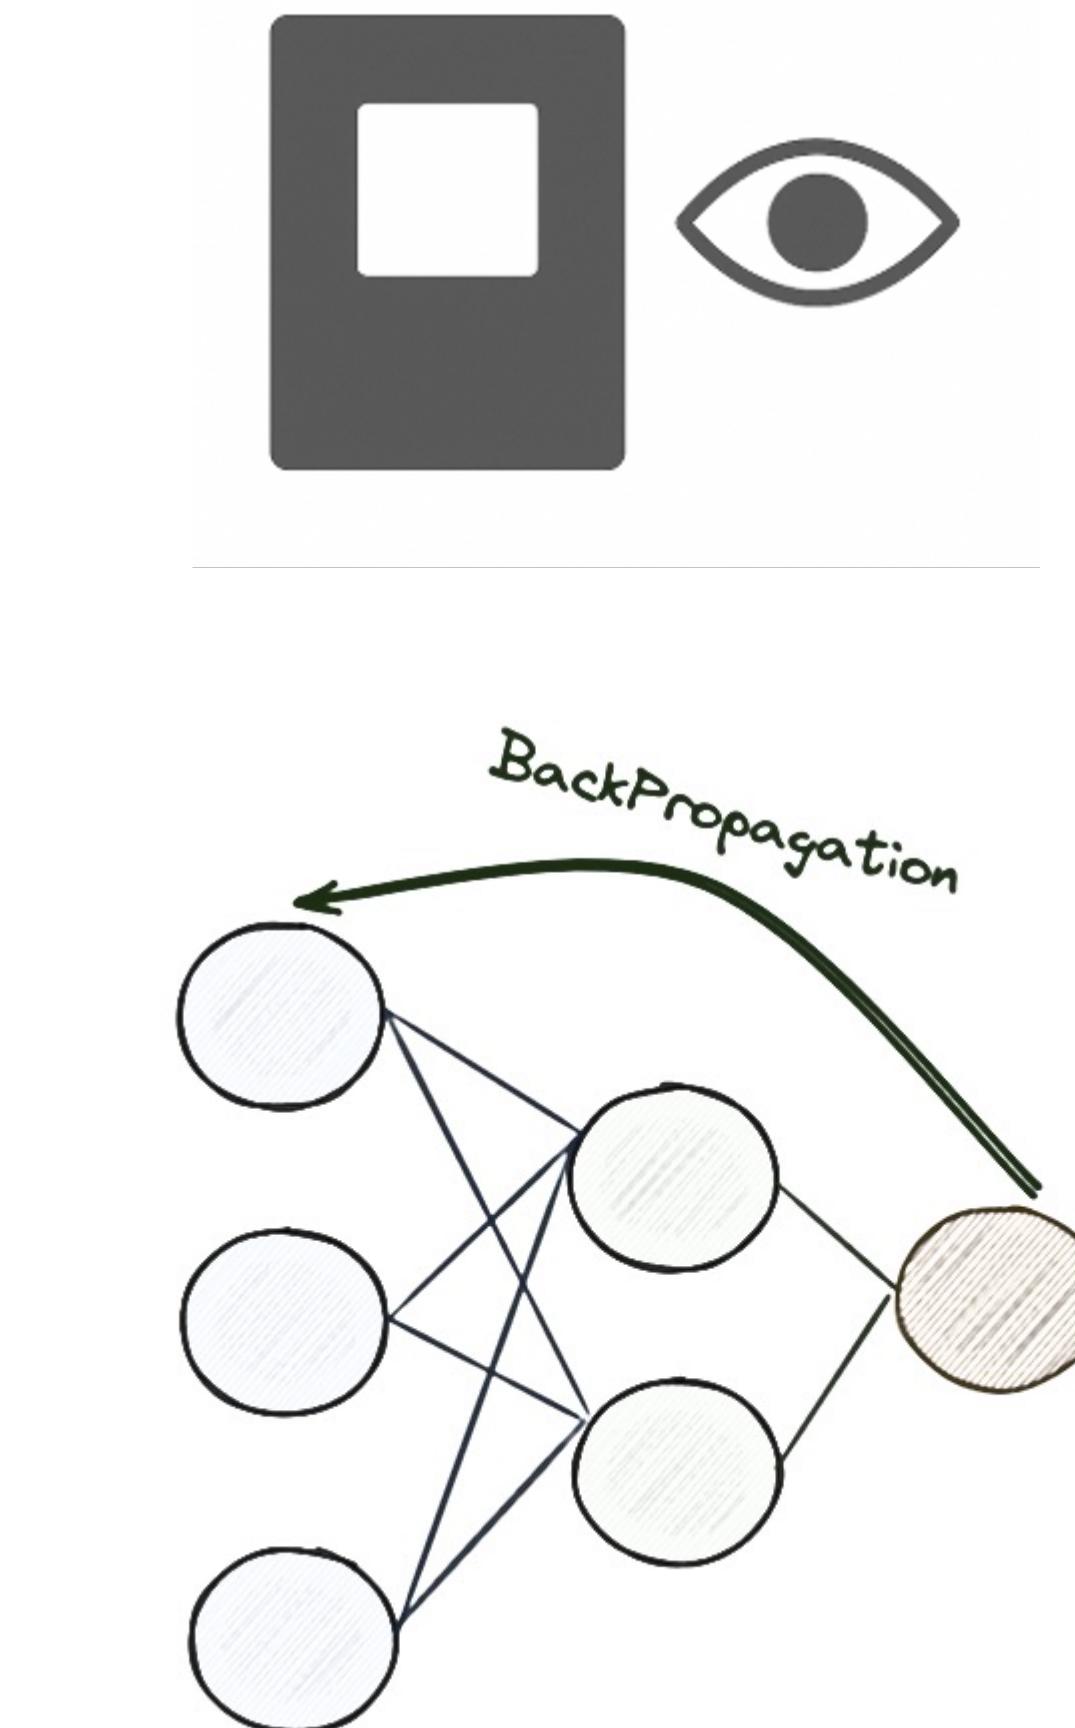
Guidelines & Lessons Learned



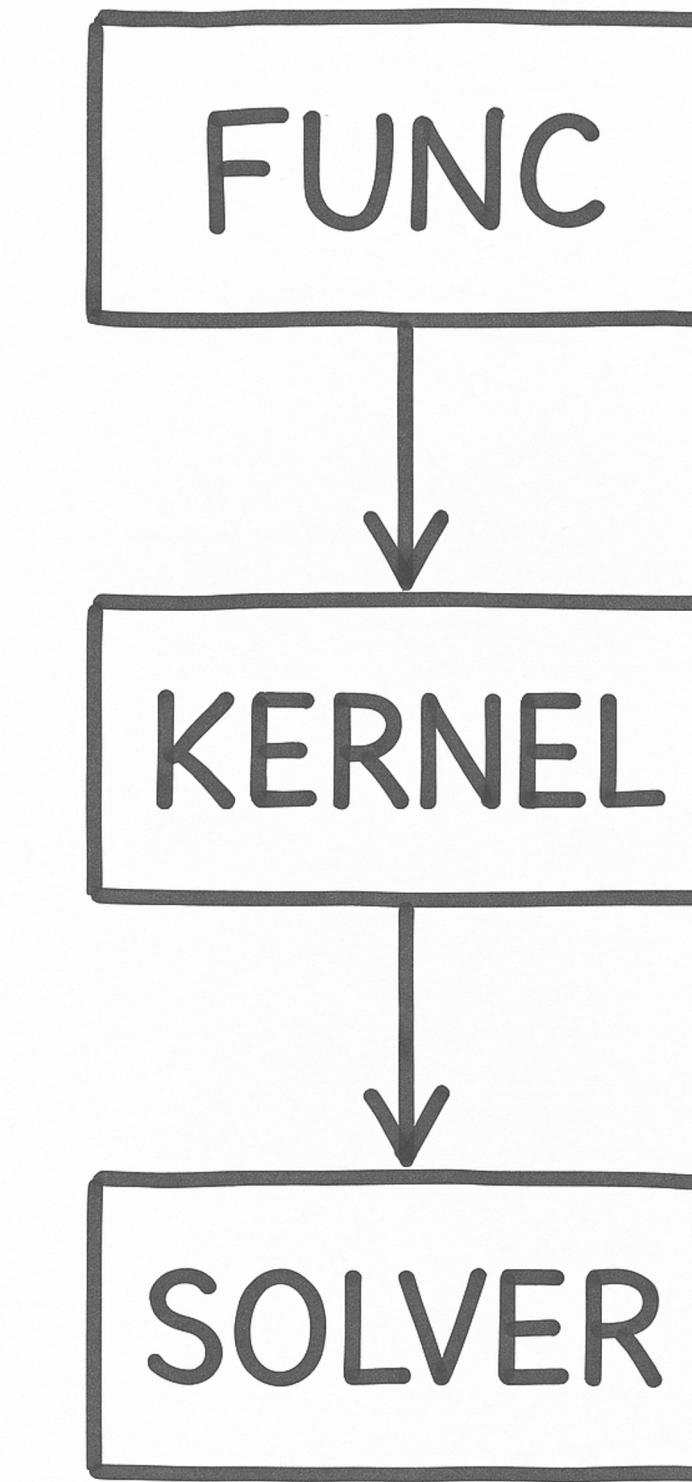
Separate physical model from numerical method



Flat data preferred over object-oriented programming



Avoid hidden state, more control over memory



Take what you need, at each level

Flat Data Design

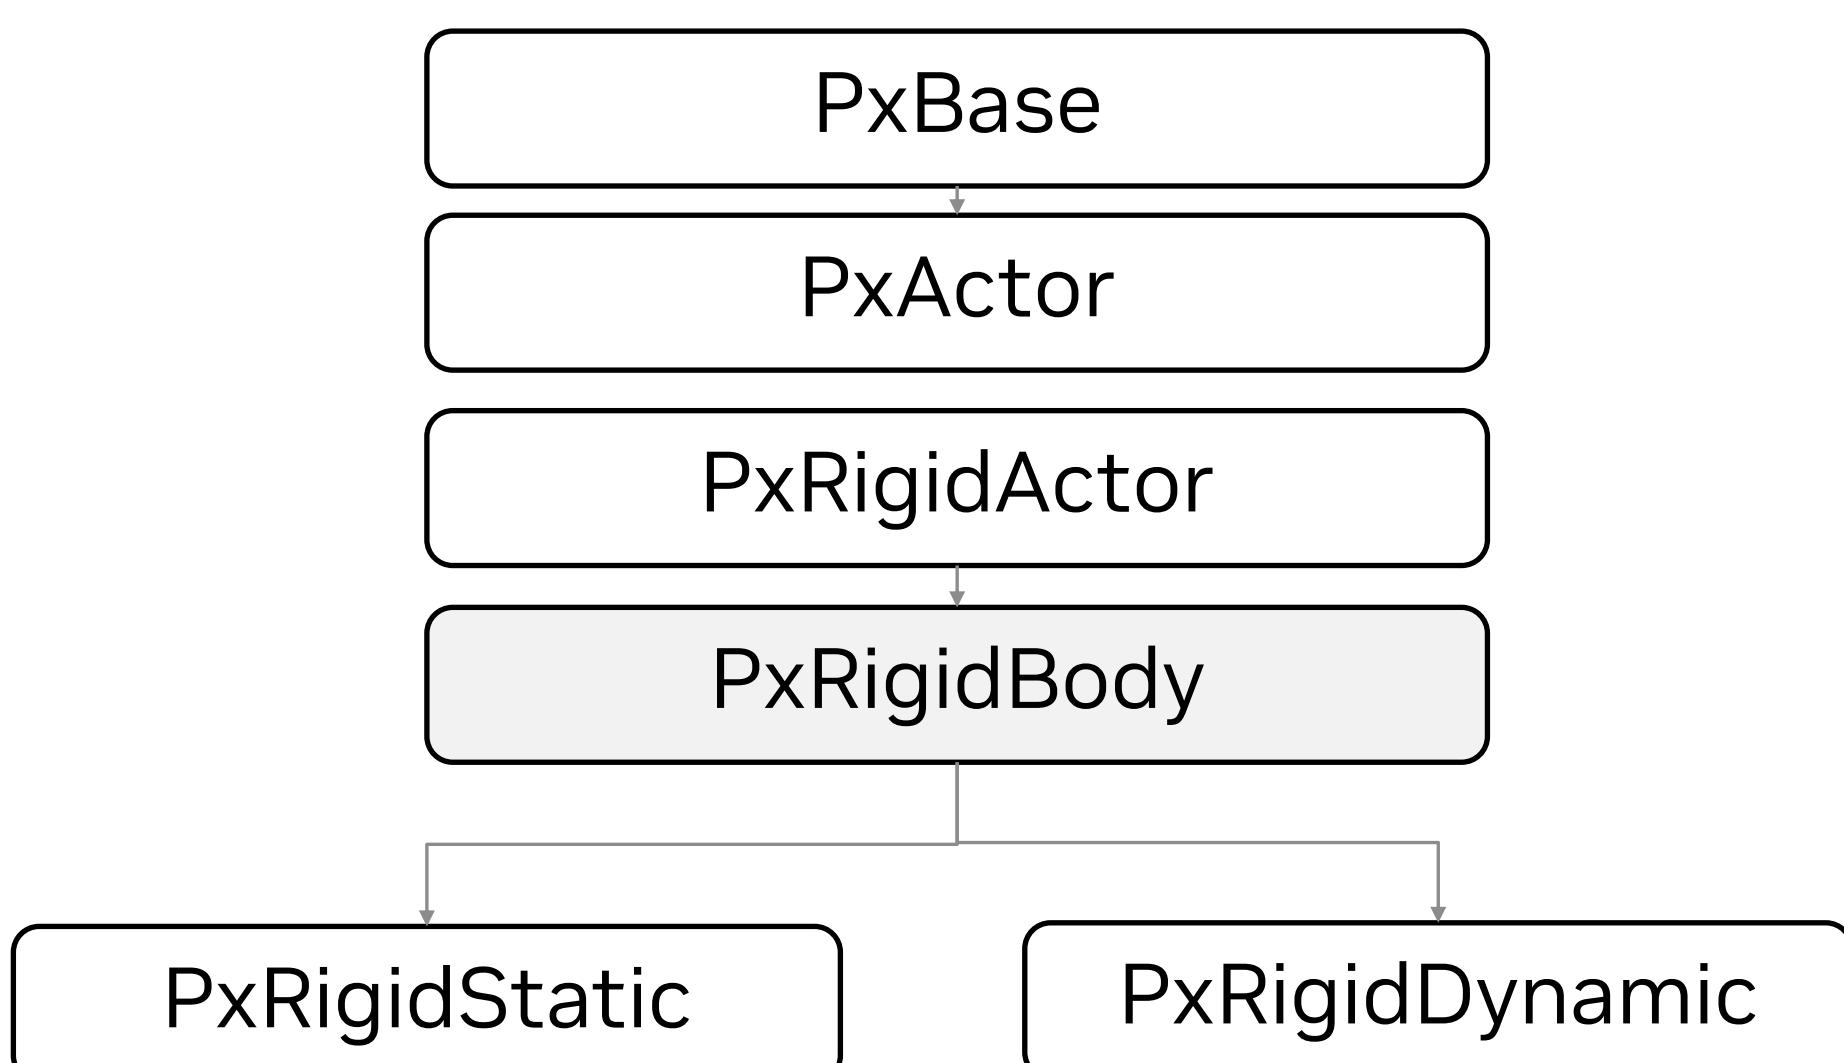
Remove layers of abstraction

PhysX 5

```
class PxRigidBody : public PxRigidActor
{
public:
    virtual void setMass(PxReal mass) = 0;
    virtual PxReal getMass() const = 0;
    virtual PxReal getInvMass() const = 0;

    virtual void setMassSpaceInertiaTensor(const PxVec3& m) = 0;
    virtual PxVec3 getMassSpaceInertiaTensor() const = 0;
    virtual PxVec3 getMassSpaceInvInertiaTensor() const = 0;

    virtual void setLinearDamping(PxReal linDamp) = 0;
    virtual PxReal getLinearDamping() const = 0;
    virtual void setAngularDamping(PxReal angDamp) = 0;
    virtual PxReal getAngularDamping() const = 0;
```



Newton

```
class Model:
    """Holds the definition of the simulation model
    This class holds the non-time varying description of
    the system, i.e.: all geometry, constraints, and
    parameters used to describe the simulation."""

    body_com : wp.array(dtype=wp.vec3)
    body_inertia : wp.array(dtype=wp.mat33)
    body_inv_inertia : wp.array(dtype=wp.mat33)
    body_mass : wp.array(dtype=float)

    joint_type : wp.array(dtype=int)
    joint_parent : wp.array(dtype=int)
    joint_child : wp.array(dtype=int)
    joint_X_p : wp.array(dtype=wp.transform)
    joint_X_c : wp.array(dtype=wp.transform)
    joint_axis : wp.array(dtype=wp.vec3)

    spring_indices : wp.array(dtype=int)
    spring_rest_length : wp.array(dtype=float)
    spring_stiffness : wp.array(dtype=float)
    spring_damping : wp.array(dtype=float)
    spring_control : wp.array(dtype=float)
```

newton.sim.Model

System description

- Flat Python struct of 1D, 2D Warp arrays
- Stores non-time varying data
- Fundamentally, we store data grouped by frequencies:
 - per-Joint (e.g.: parent body)
 - per-DOF (e.g.: armature)
 - per-Articulation (e.g.: joint dof offset)
 - per-Body (e.g.: mass, inertia)
 - per-Shape (e.g.: friction coefficients)
- Flexible thanks to Python's dynamic typing
- E.g.: add shape lubrication for custom friction
- `model.shape_lubrication = wp.ones(dtype=float, length=n)`

`class Model:`

 """Holds the definition of the simulation model

This class holds the non-time varying description of the system, i.e.: all geometry, constraints, and parameters used to describe the simulation.

Attributes:

`body_com : array(dtype=wp.vec3)`

`body_inertia : array(dtype=wp.mat33)`

`body_inv_inertia : array(dtype=wp.mat33)`

`body_mass : array(dtype=float)`

`joint_type : array(dtype=int)`

`joint_parent : array(dtype=int)`

`joint_child : array(dtype=int)`

`joint_X_p : array(dtype=wp.transform)`

`joint_X_c : array(dtype=wp.transform)`

`joint_axis : array(dtype=wp.vec3)`

`spring_indices : array(dtype=int)`

`spring_rest_length : array(dtype=float)`

`spring_stiffness : array(dtype=float)`

`spring_damping : array(dtype=float)`

`spring_control : array(dtype=float)`

...

newton.sim.State

Physical configuration

- State objects hold time-varying data
 - Primarily `joint_q`, `body_q`, `particle_q`
 - Flat struct of `wp.arrays`
- Create state objects by asking the model for a new state object using `model.state()`
- Can create as many state objects as you require, easy to save / restore state
- User can add fields
 - e.g.: `state.particle_temperature`

`class State:`

"""The State object holds all *time-varying* data for a model.

Attributes:

`joint_q` (array): Array of joint positions

`joint_qd` (array): Array of joint velocities

`body_q` (array): Array of body coordinates

`body_qd` (array): Array of body coordinates

`particle_q` (array): Array of 3D particle

`particle_qd` (array): Array of 3D particle

newton.sim.Control

External forces and actuation

- Represents actuation inputs over a timestep
- Includes joint actuation, PD targets, activations
- Supports both implicit & explicit methods
- Create Control object by calling `model.control()`

`class Control:`

"""Time-varying control data for a :class:`Model`.

Time-varying control data includes joint control inputs, and activation forces for triangle and tetrahedral elements.

Attributes:

joint actuation

`joint_f (array): Generalized joint forces`

`joint_target (array): Joint PD targets`

triangle activations

`tri_activations (array):`

tetrahedra activations

`tet_activations (array):`

newton.sim.Contacts

Contact Geometry

- Generated by collision pipeline
- User controls when/how to compute contacts
- Passed to `solver.step()`

class Contacts:

"""Provides contact information to be consumed by a solver. Stores the contact distance, position, frame, and geometry for each contact point.

.. note::

Attributes:

rigid_contact_count : array(dtype=int)
rigid_contact_point_id : array(dtype=int)
rigid_contact_shape0 : array(dtype=int)
rigid_contact_shape1 : array(dtype=int)
rigid_contact_point0 : array(dtype=wp.vec3)
rigid_contact_point1 : array(dtype=wp.vec3)
rigid_contact_normal : array(dtype=wp.vec3)
rigid_contact_thickness : array(dtype=float)
rigid_contact_tids : array(dtype=int)

soft_contact_count : array(dtype=int)
soft_contact_particle : array(dtype=int)
soft_contact_shape : array(dtype=int)
soft_contact_body_pos : array(dtype=wp.vec3)
soft_contact_body_vel : array(dtype=wp.vec3)
soft_contact_normal : array(dtype=wp.vec3)
soft_contact_tids : array(dtype=int)

Geometry Library

newton.geometry

- Goal: reusable collision detection and contact generation
- Split into standalone module for flexibility
- High-level entry points (broad phase, narrow phase)
- Expose low-level routines to users (SDF funcs, queries, GJK)

```
# geometry types
GEO_SPHERE = wp.constant(0)
GEO_BOX = wp.constant(1)
GEO_CAPSULE = wp.constant(2)
GEO_CYLINDER = wp.constant(3)
GEO_CONE = wp.constant(4)
GEO_MESH = wp.constant(5)
GEO_SDF = wp.constant(6)
GEO_PLANE = wp.constant(7)
GEO_NONE = wp.constant(8)

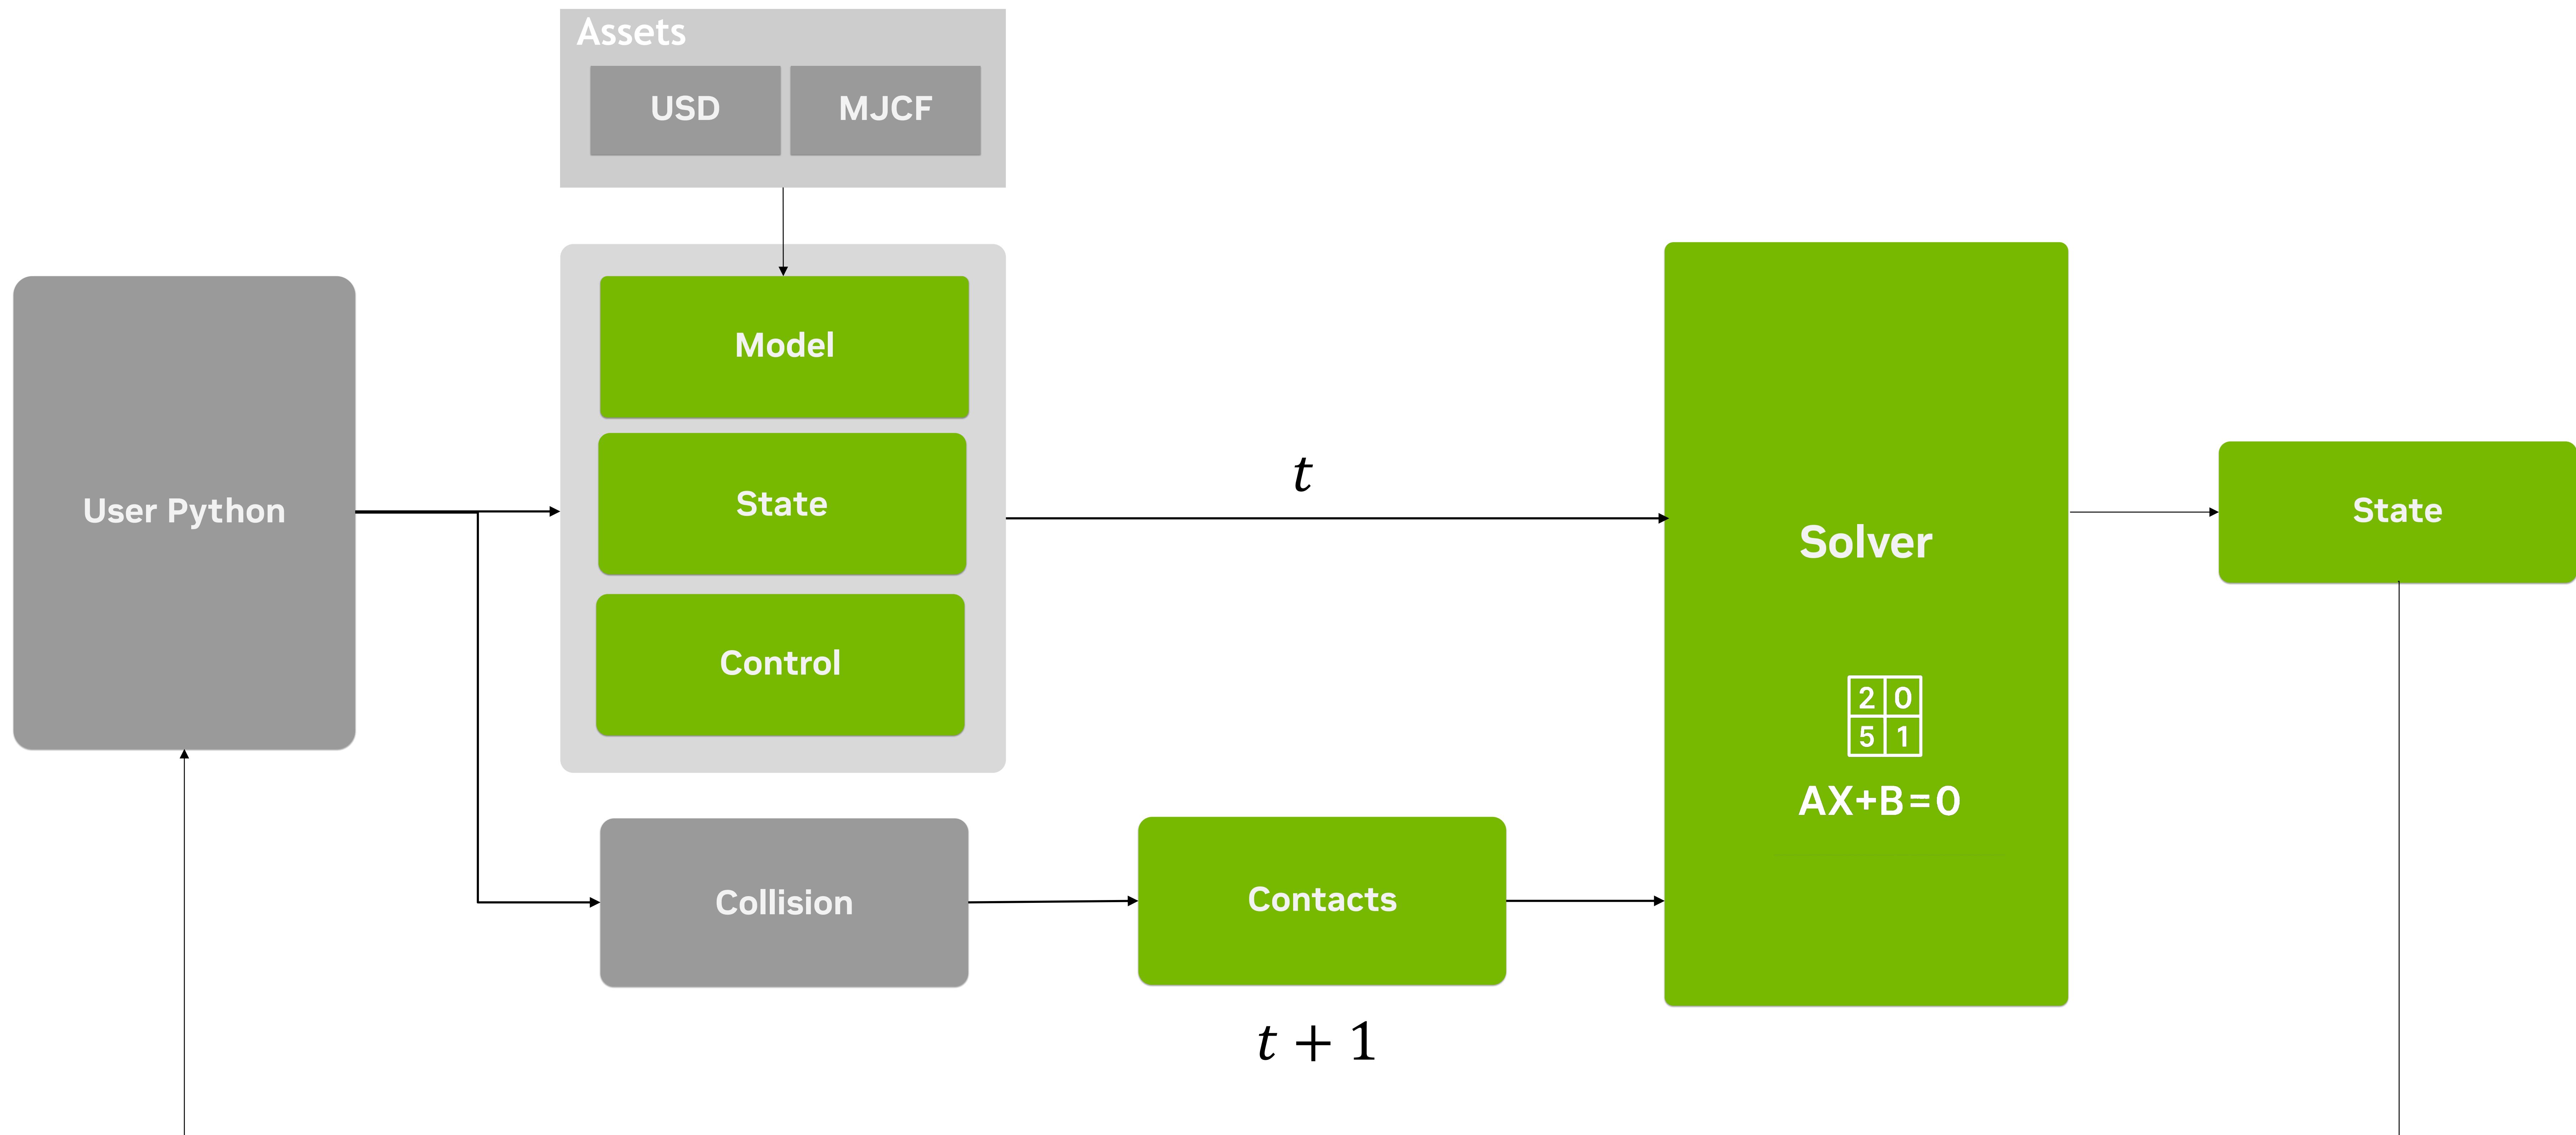
# collision pipeline
def broad_phase(
    # input
    bounds_origin: wp.array(dtype=wp.vec3),
    bounds_size: wp.array(dtype=wp.vec3),
    bounds_type: wp.array(dtype=int),
    # output
    pairs: wp.array(dtype=wp.vec2i))

def narrow_phase(
    # input
    shape_transform: wp.array(dtype=wp.transform),
    shape_type: wp.array(dtype=int),
    shape_scale: wp.array(dtype=wp.vec3),
    pairs: wp.array(dtype=wp.vec2i) ,
    # output
    contact_position: wp.array(dtype=wp.vec2i),
    contact_frame: wp.array(dtype=wp.quat),
    contact_depth: wp.array(dtype=float))

# SDF functions
@wp.func
def sphere_sdf(center: wp.vec3, radius: float, p: wp.vec3):

@wp.func
def box_sdf(upper: wp.vec3, p: wp.vec3):
...
```

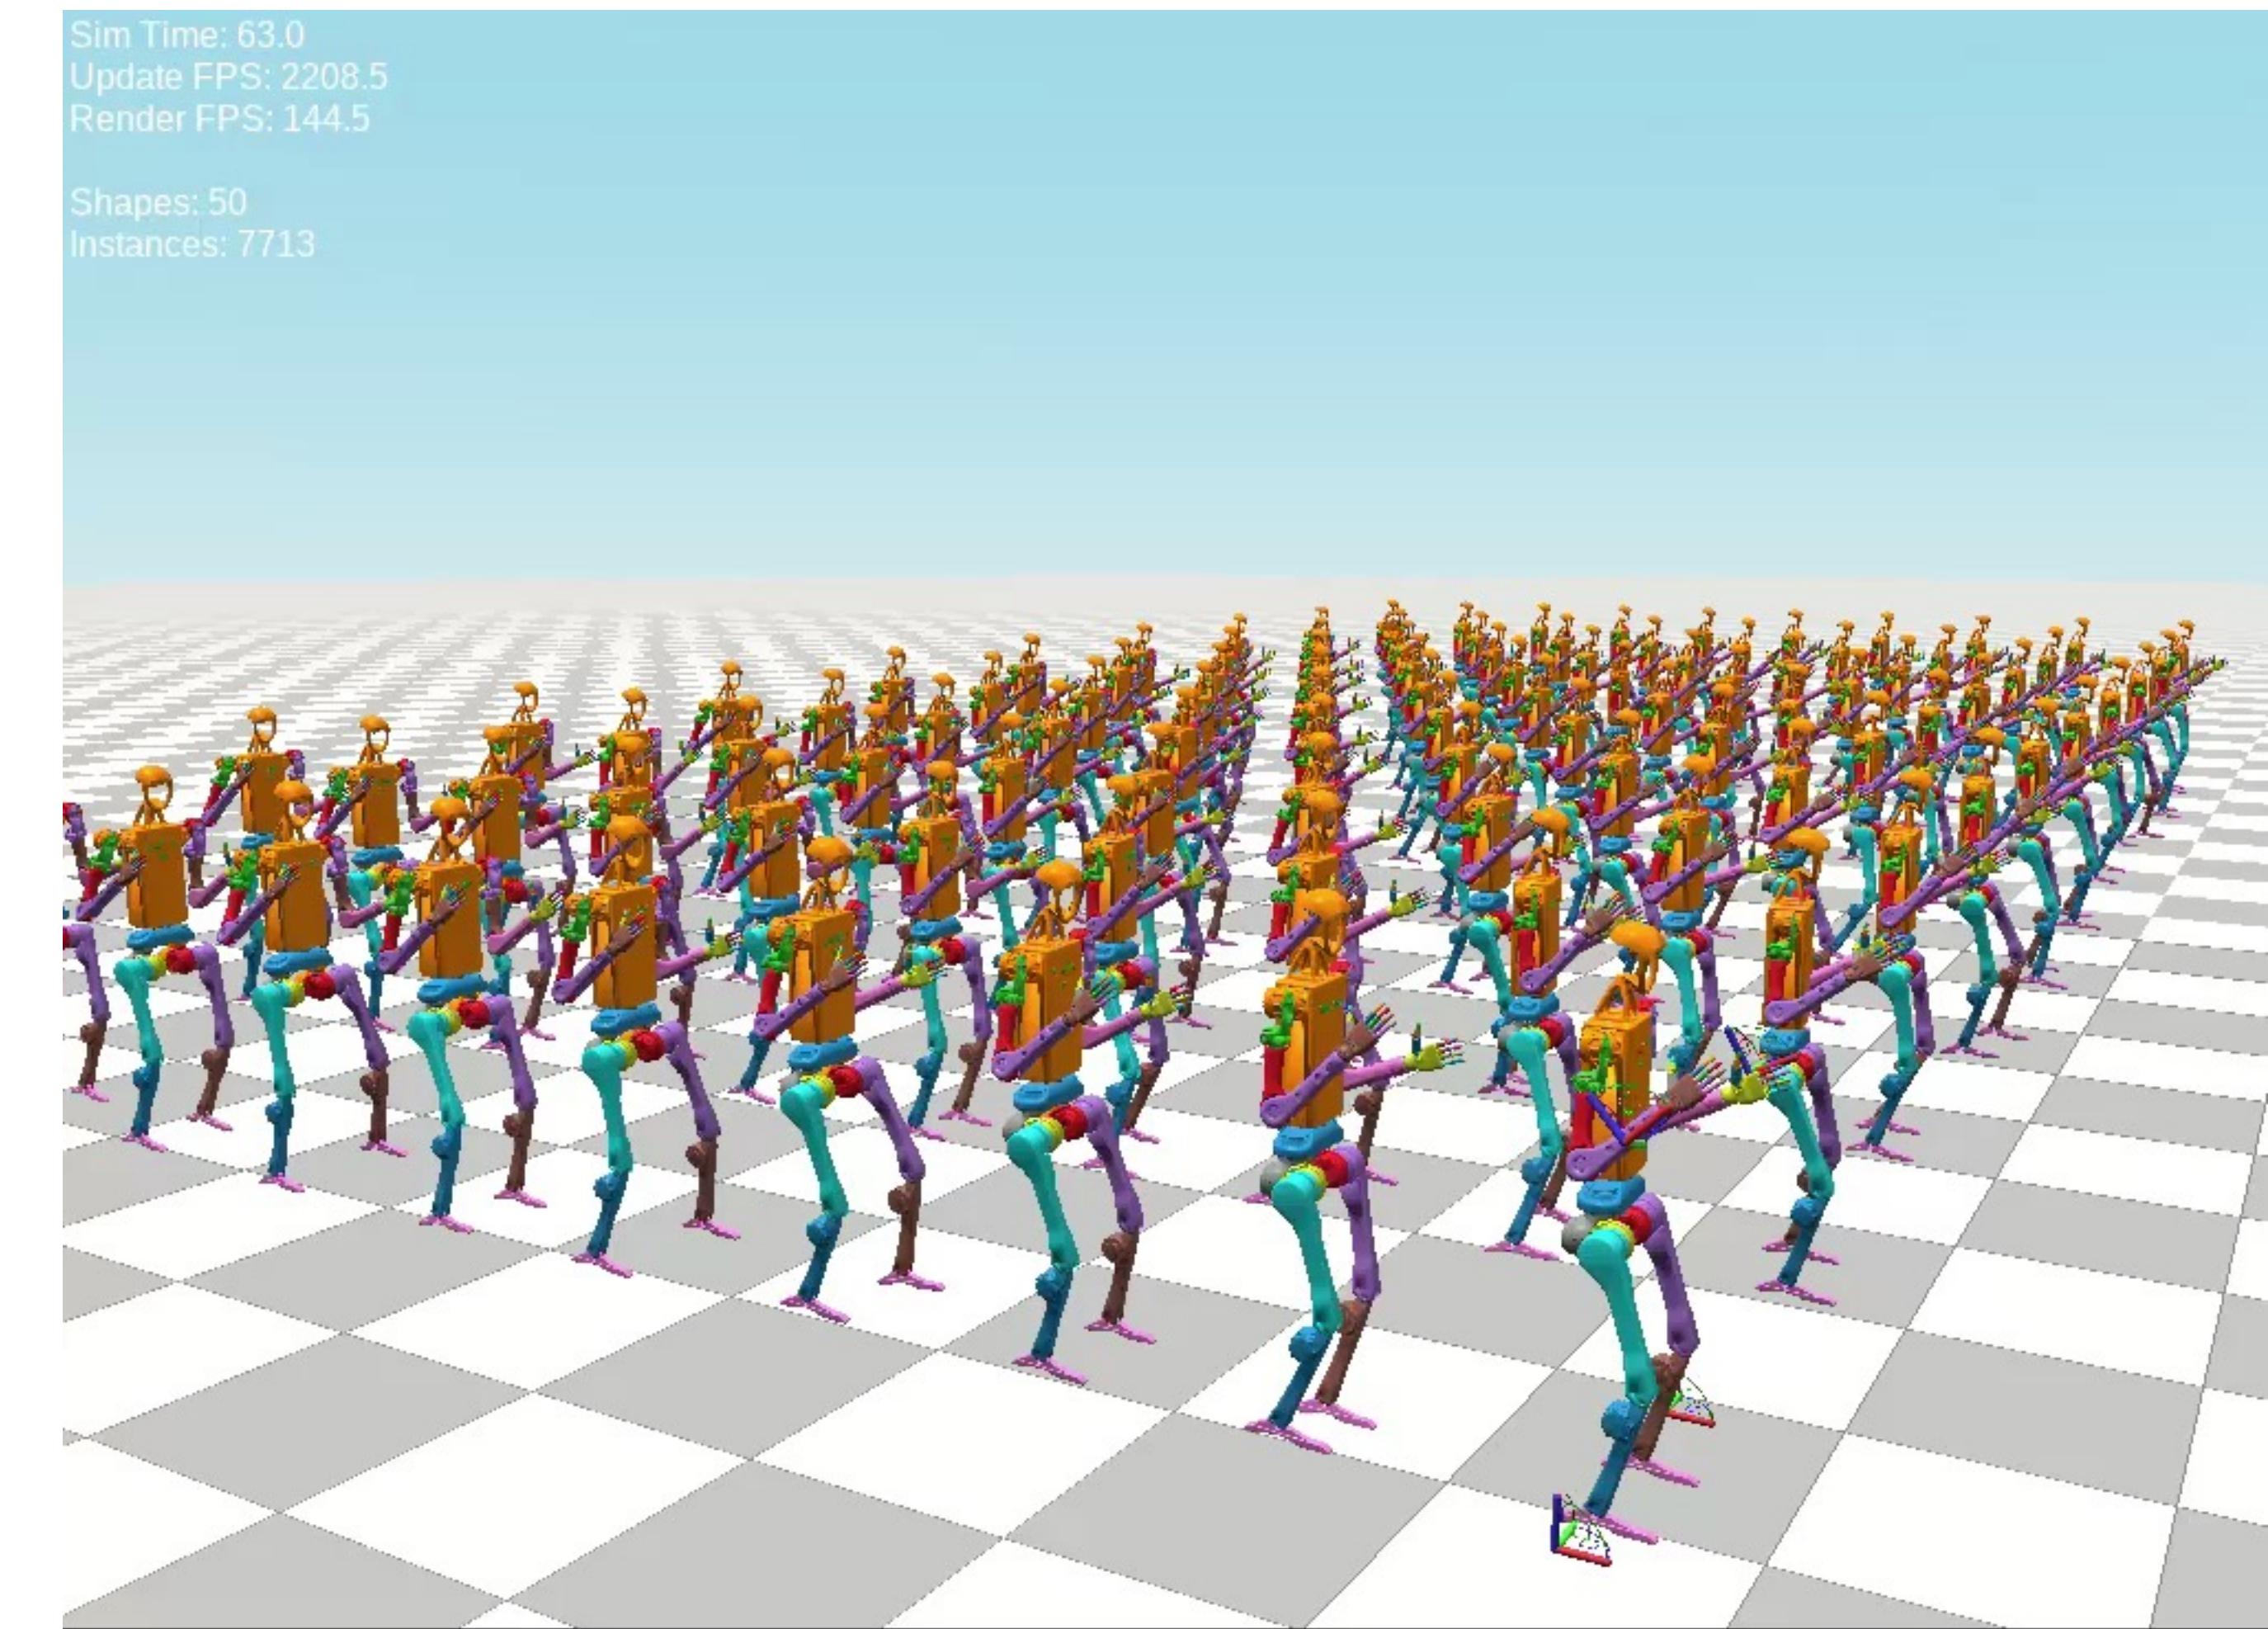
A Reference Simulation Pipeline



Inverse Kinematics (IK) Solver

- Forward/Inverse kinematics map between representations:

- `newton.ik.IKSolver()`
- `newton.ik.IKPositionObjective()`
- `newton.ik.IKRotationObjective()`
- `newton.ik. IKJointLimitObjective()`
- `Newton.ik. IKJacobianMode`

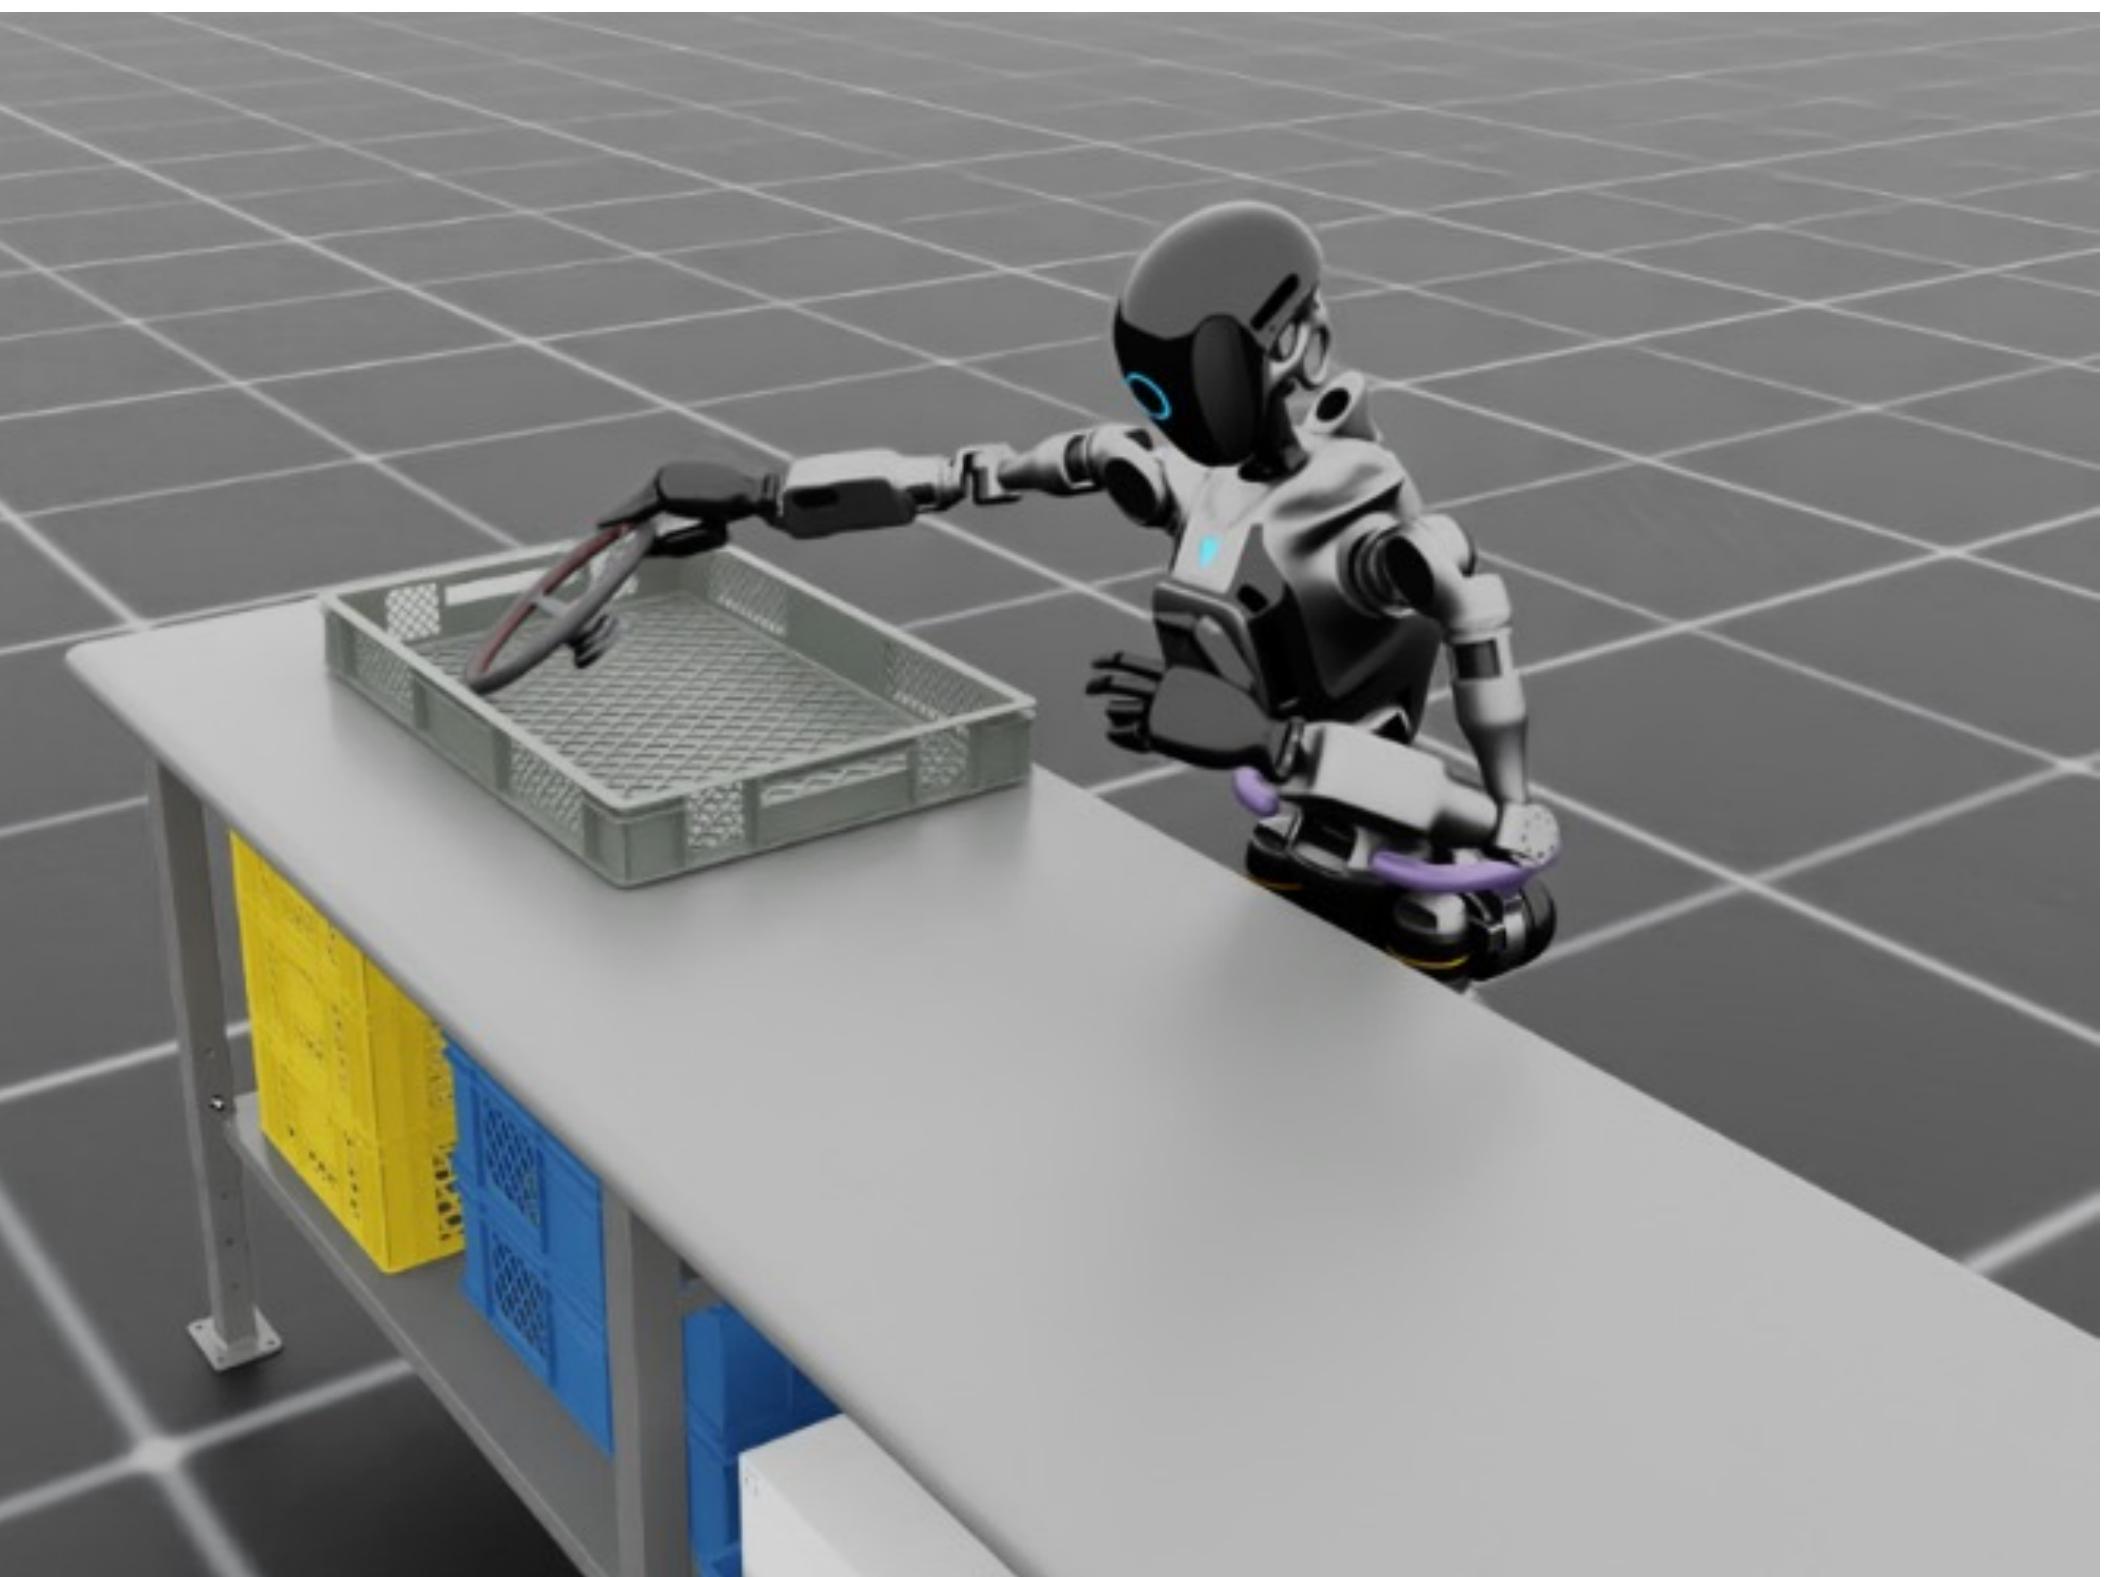


H1 Whole Body IK @ 0.03ms in Newton

GPU Acceleration

Kernal based

- Key challenges for physics simulation:
 - Possible Contacts scales $O(n^2)$
 - Actual contacts are only known at runtime
- Challenges from Tensor based frameworks
 - Branching (conditional execution)
 - Sparsity (fine-grained memory allocation)



```
@wp.kernel
def divergence(u: wp.array2d(dtype=wp.vec2),
               div: wp.array2d(dtype=float)):

    # 2D thread indices
    i, j = wp.tid()

    # boundary conditions
    if i == grid_width - 1:
        return
    if j == grid_height - 1:
        return

    # compute divergence
    dx = (u[i + 1, j][0] - u[i, j][0]) * 0.5
    dy = (u[i, j + 1][1] - u[i, j][1]) * 0.5
    div[i, j] = dx + dy

    # 2d kernel launch
    wp.launch(divergence, dim=[512, 512], inputs=[u, div])
```

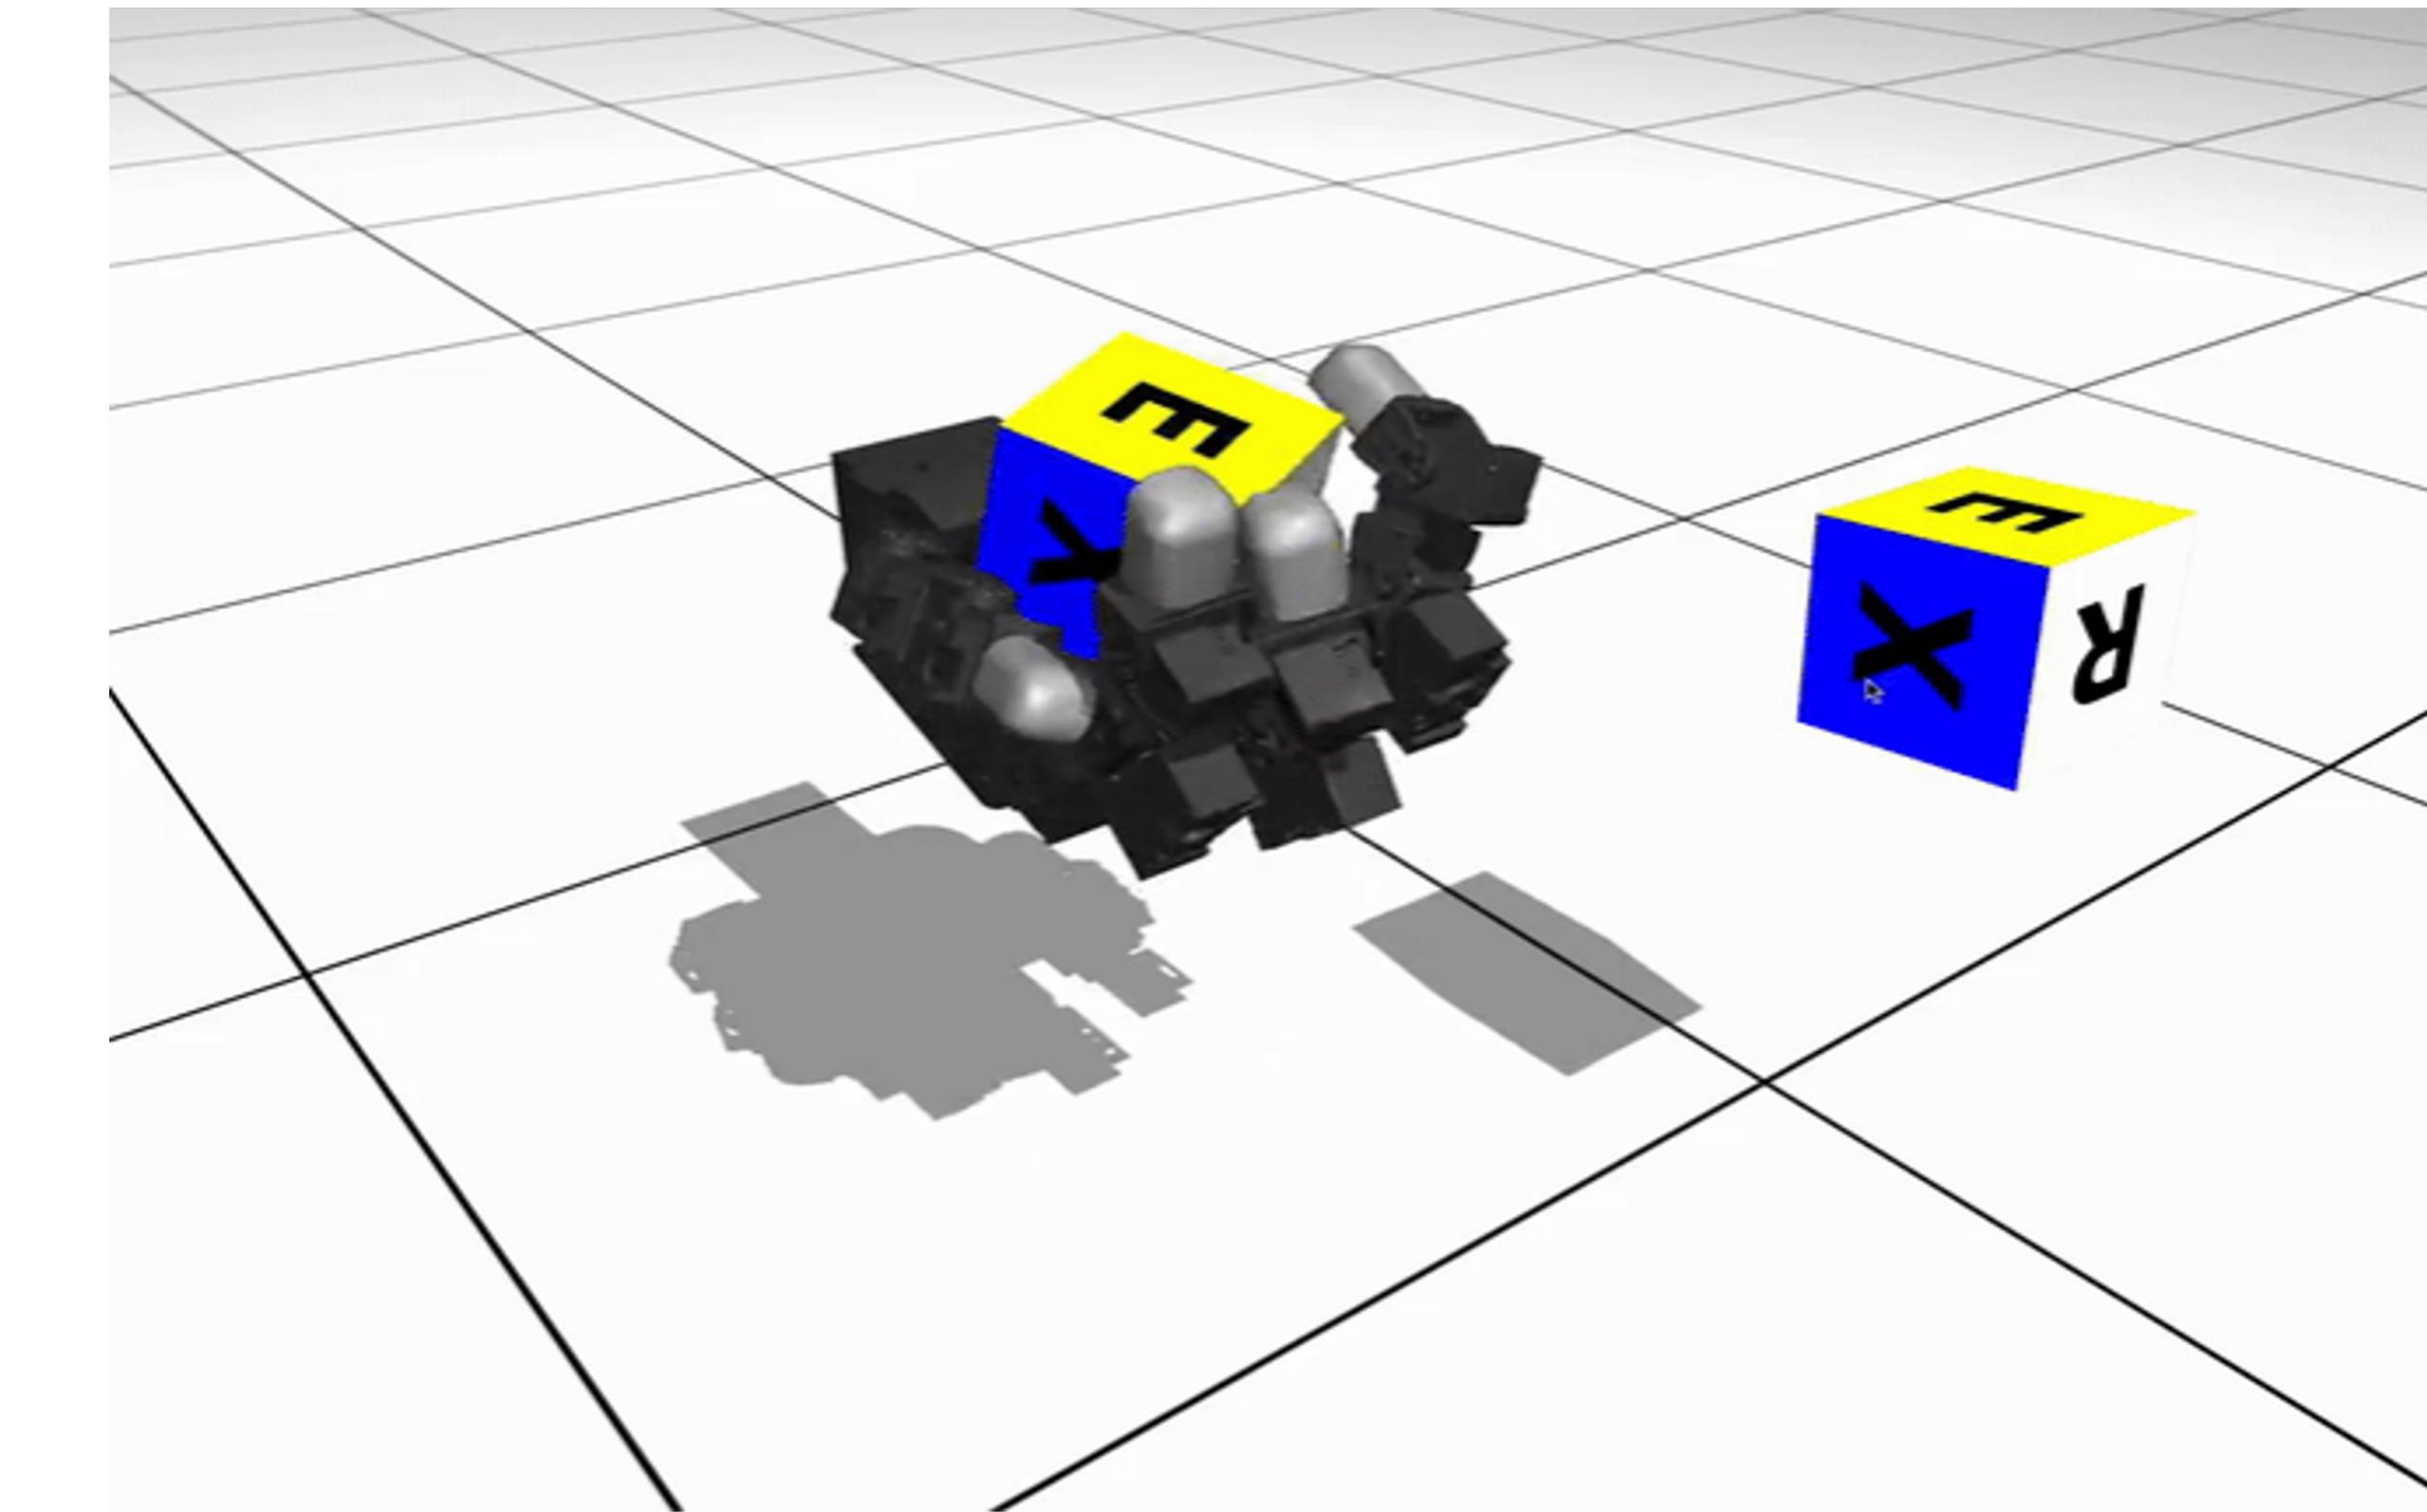
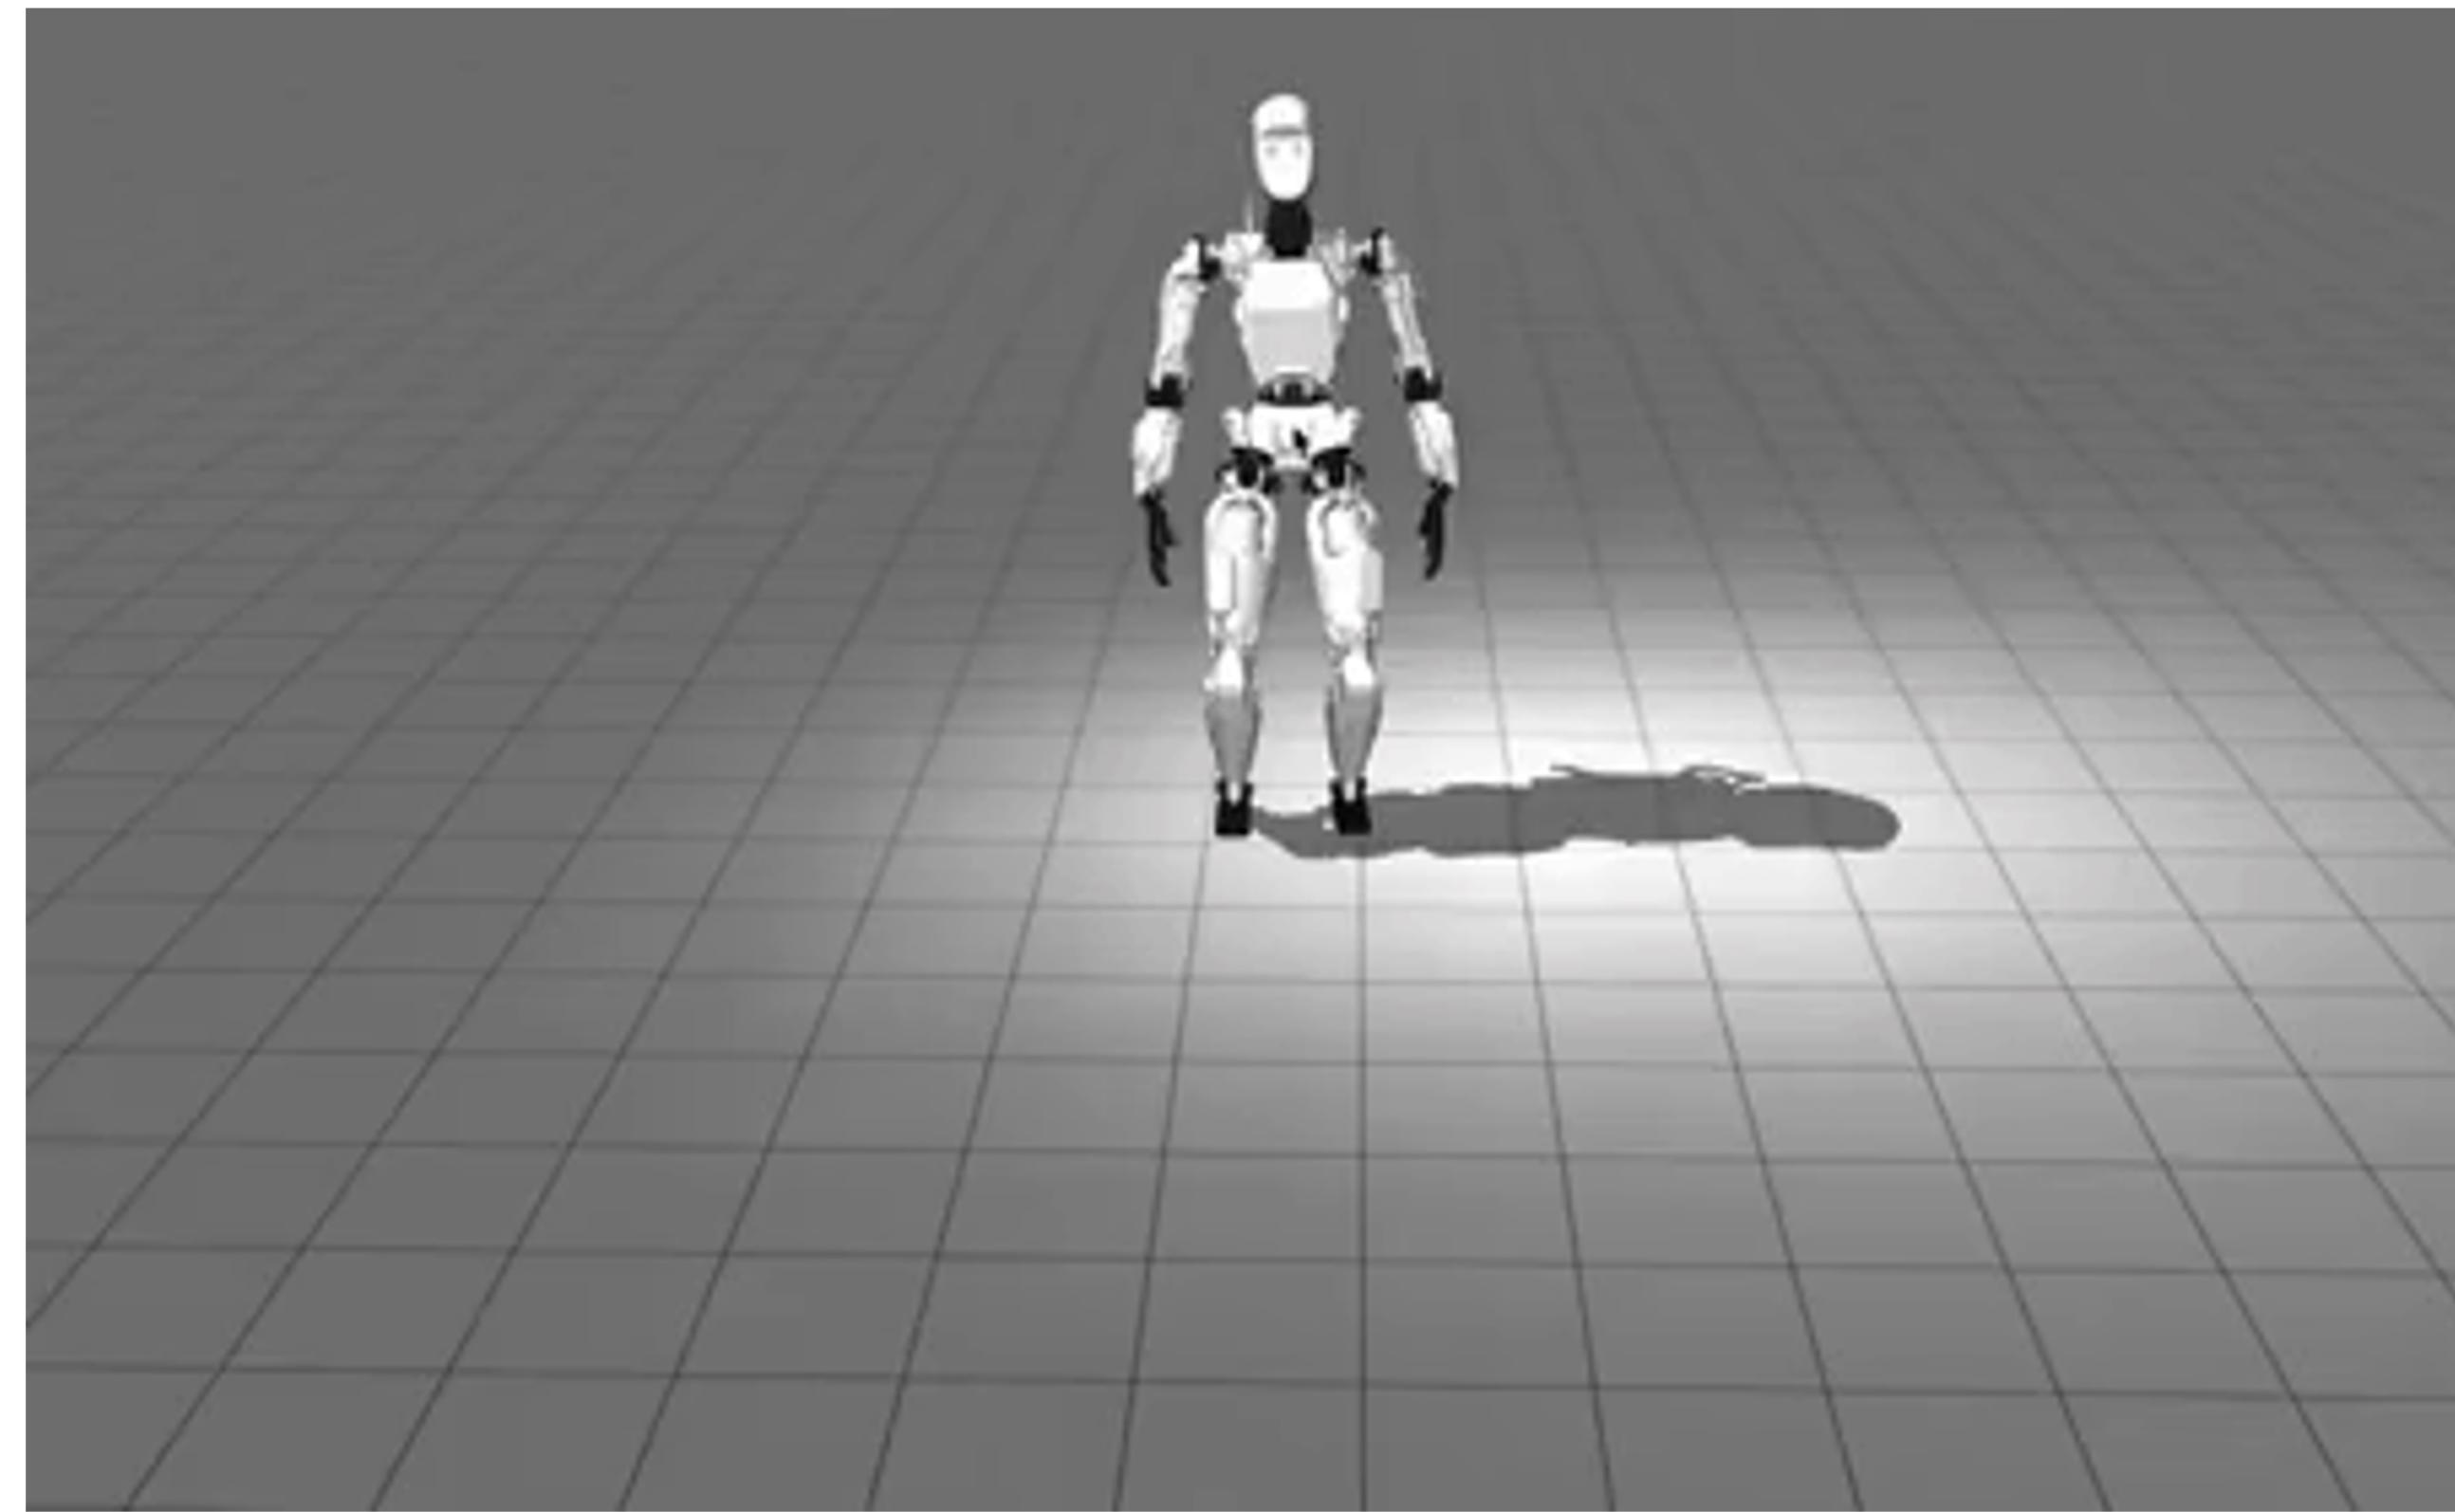
Example: 2D Divergence Calculation in Warp

GPU Acceleration

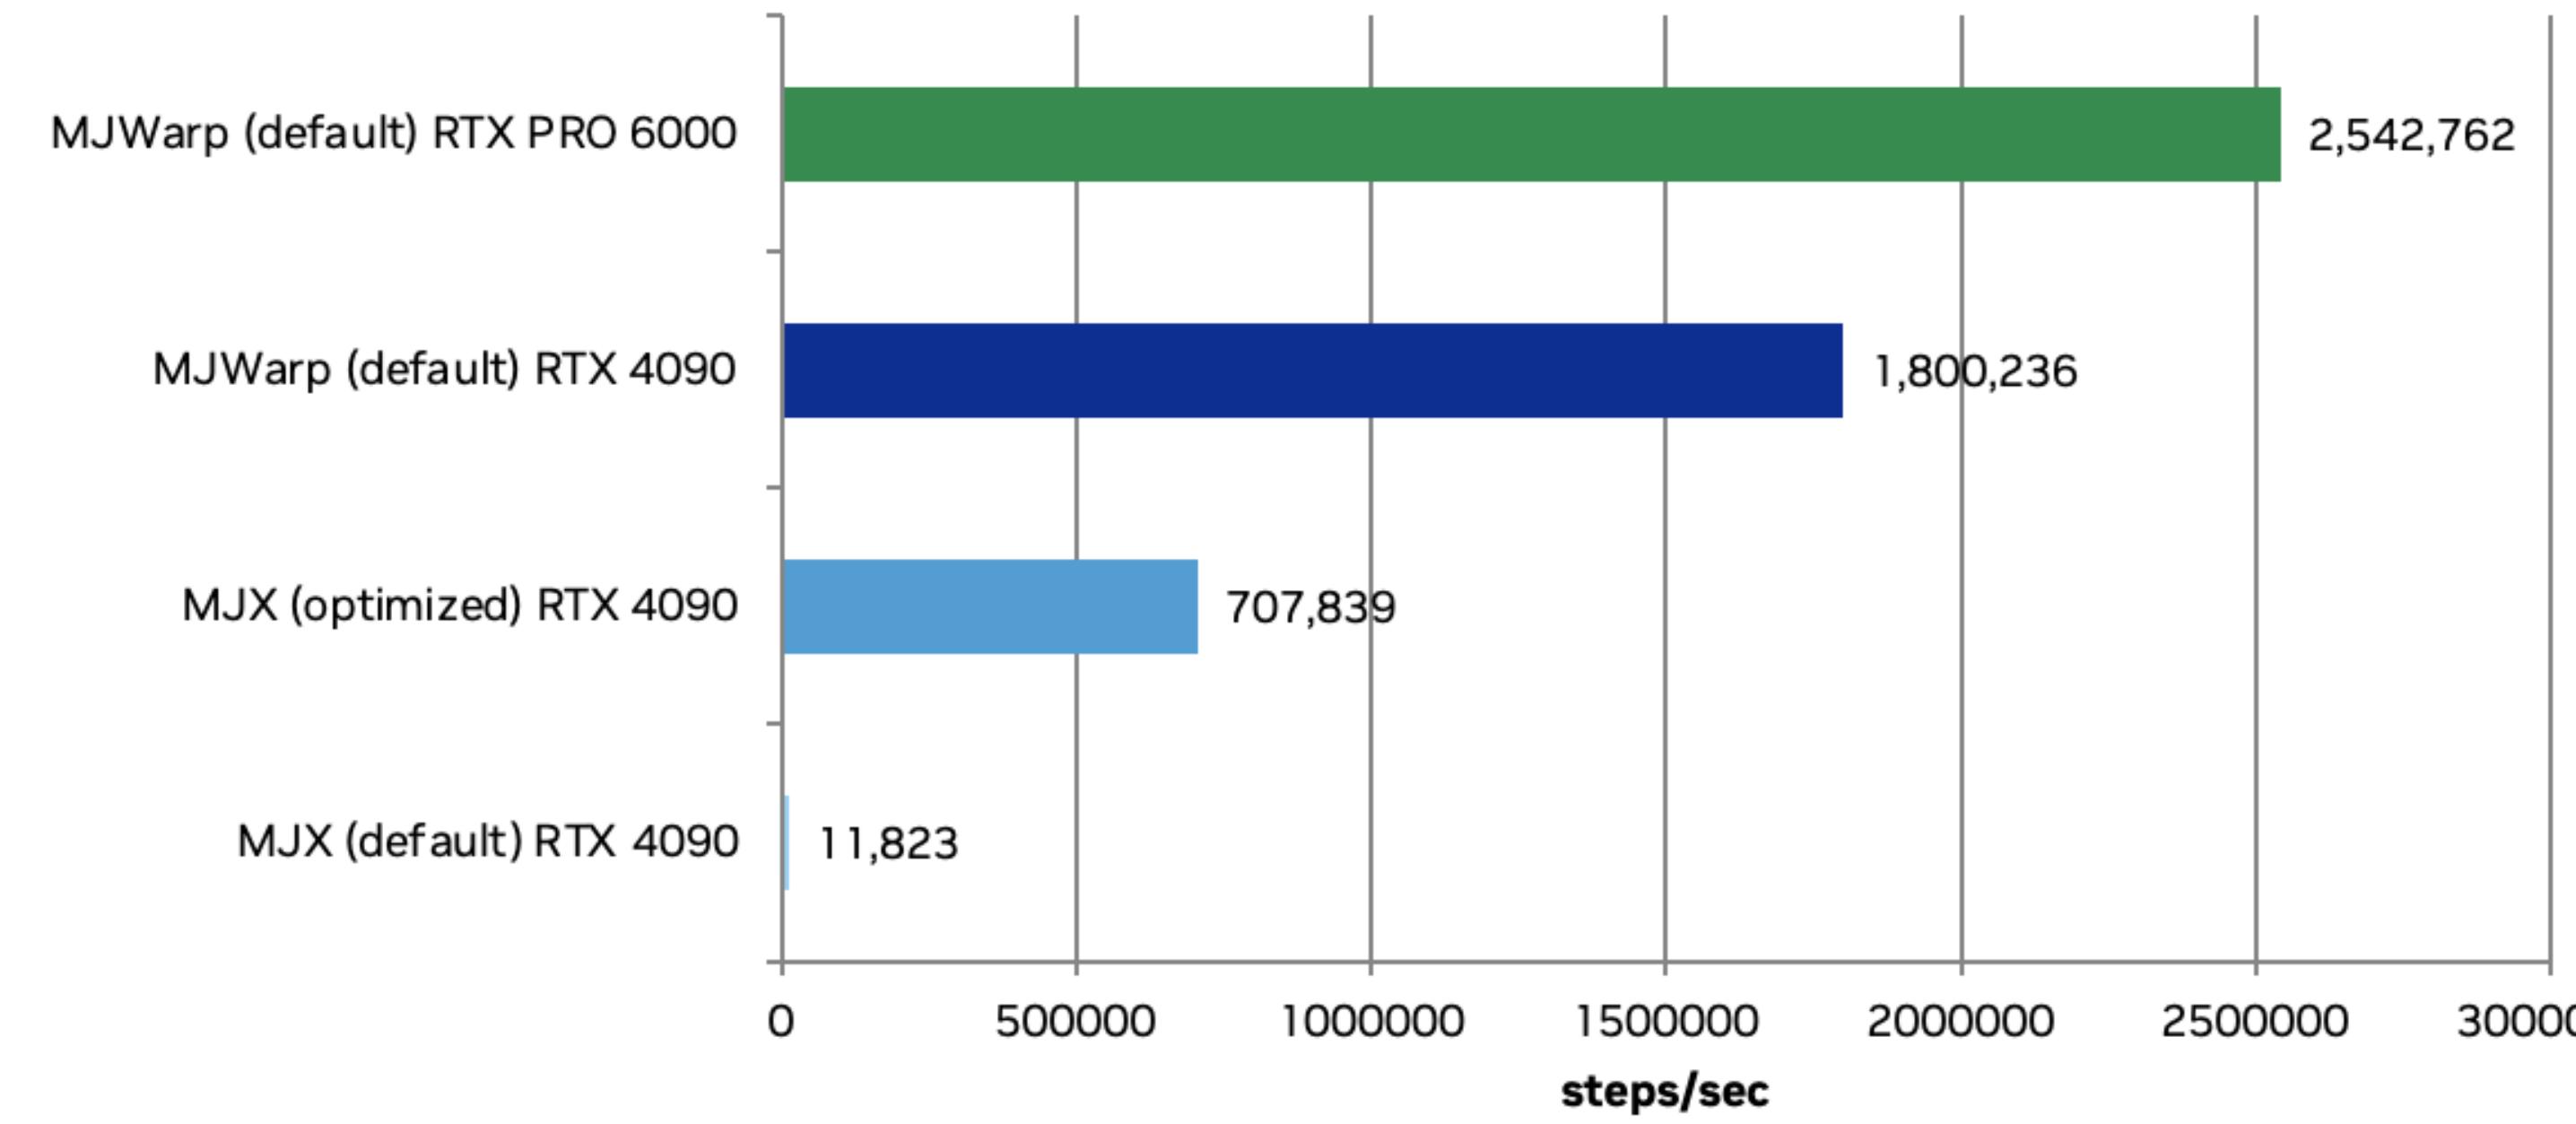
Advanced Features

- CUDA Graph Capture
 - Reduce kernel launch overhead
 - Speedups for iterative workloads like simulation loops
 - `wp.ScopedCapture()` and `wp.capture_launch(graph)`
- Tile-based Programming
 - Make use of GPU shared memory
 - Speedups for GEMM and processing adjacent elements
 - `wp.tile_matmul()` and `wp.tile_reduce()`
- NVIDIA MathDx Integration
 - Leveraging cuBLASDx and cuFFTDx
 - Build-Time Integration

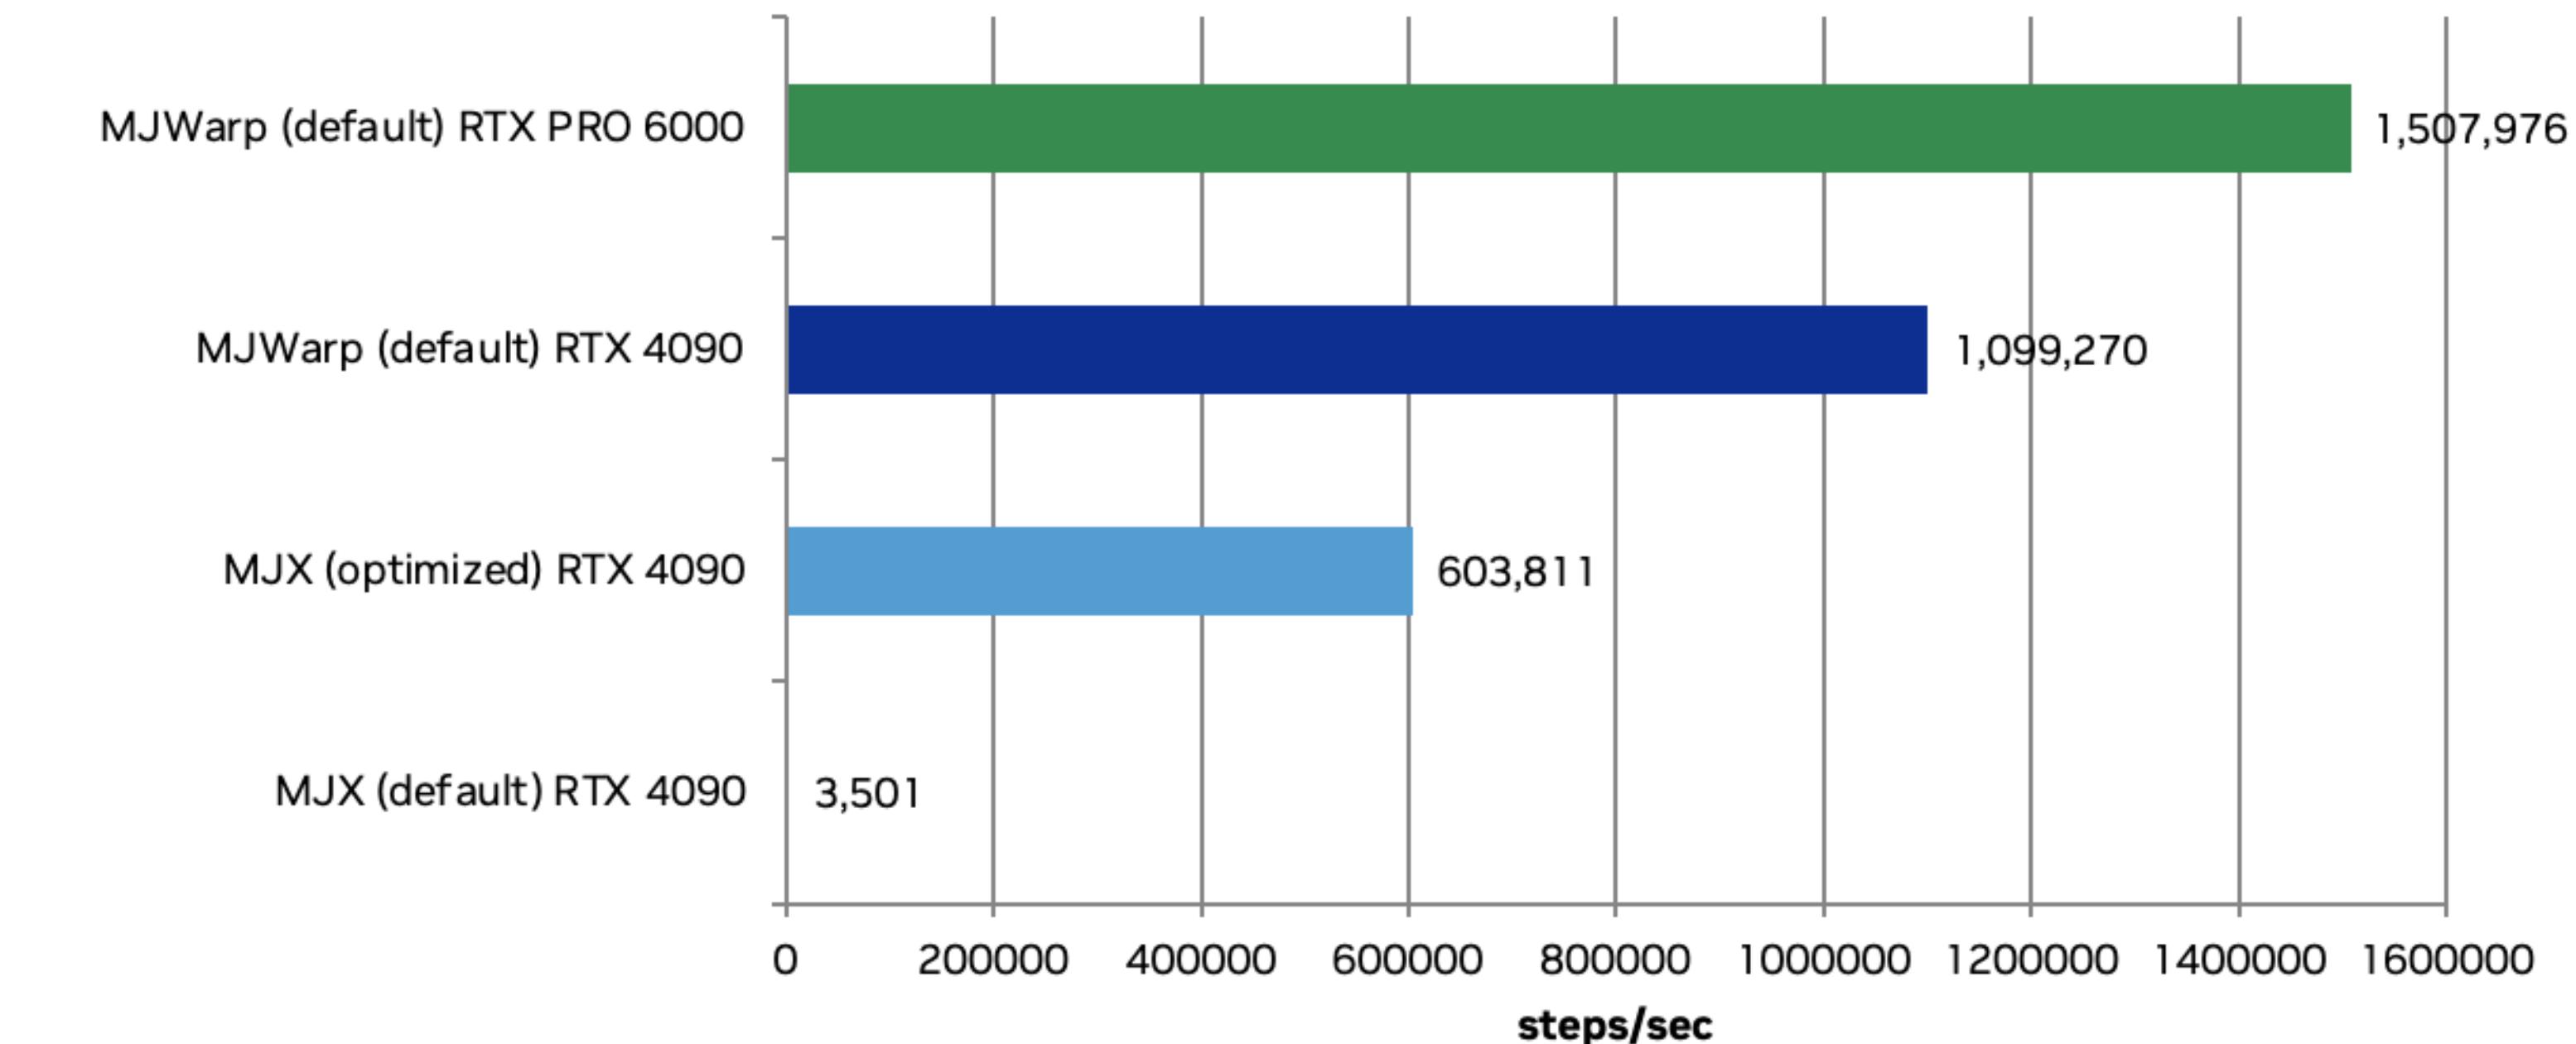
MuJoCo Solver (MJWarp)



Apptronik locomotion, physics step/sec



Leap hand manipulation, physics step/sec



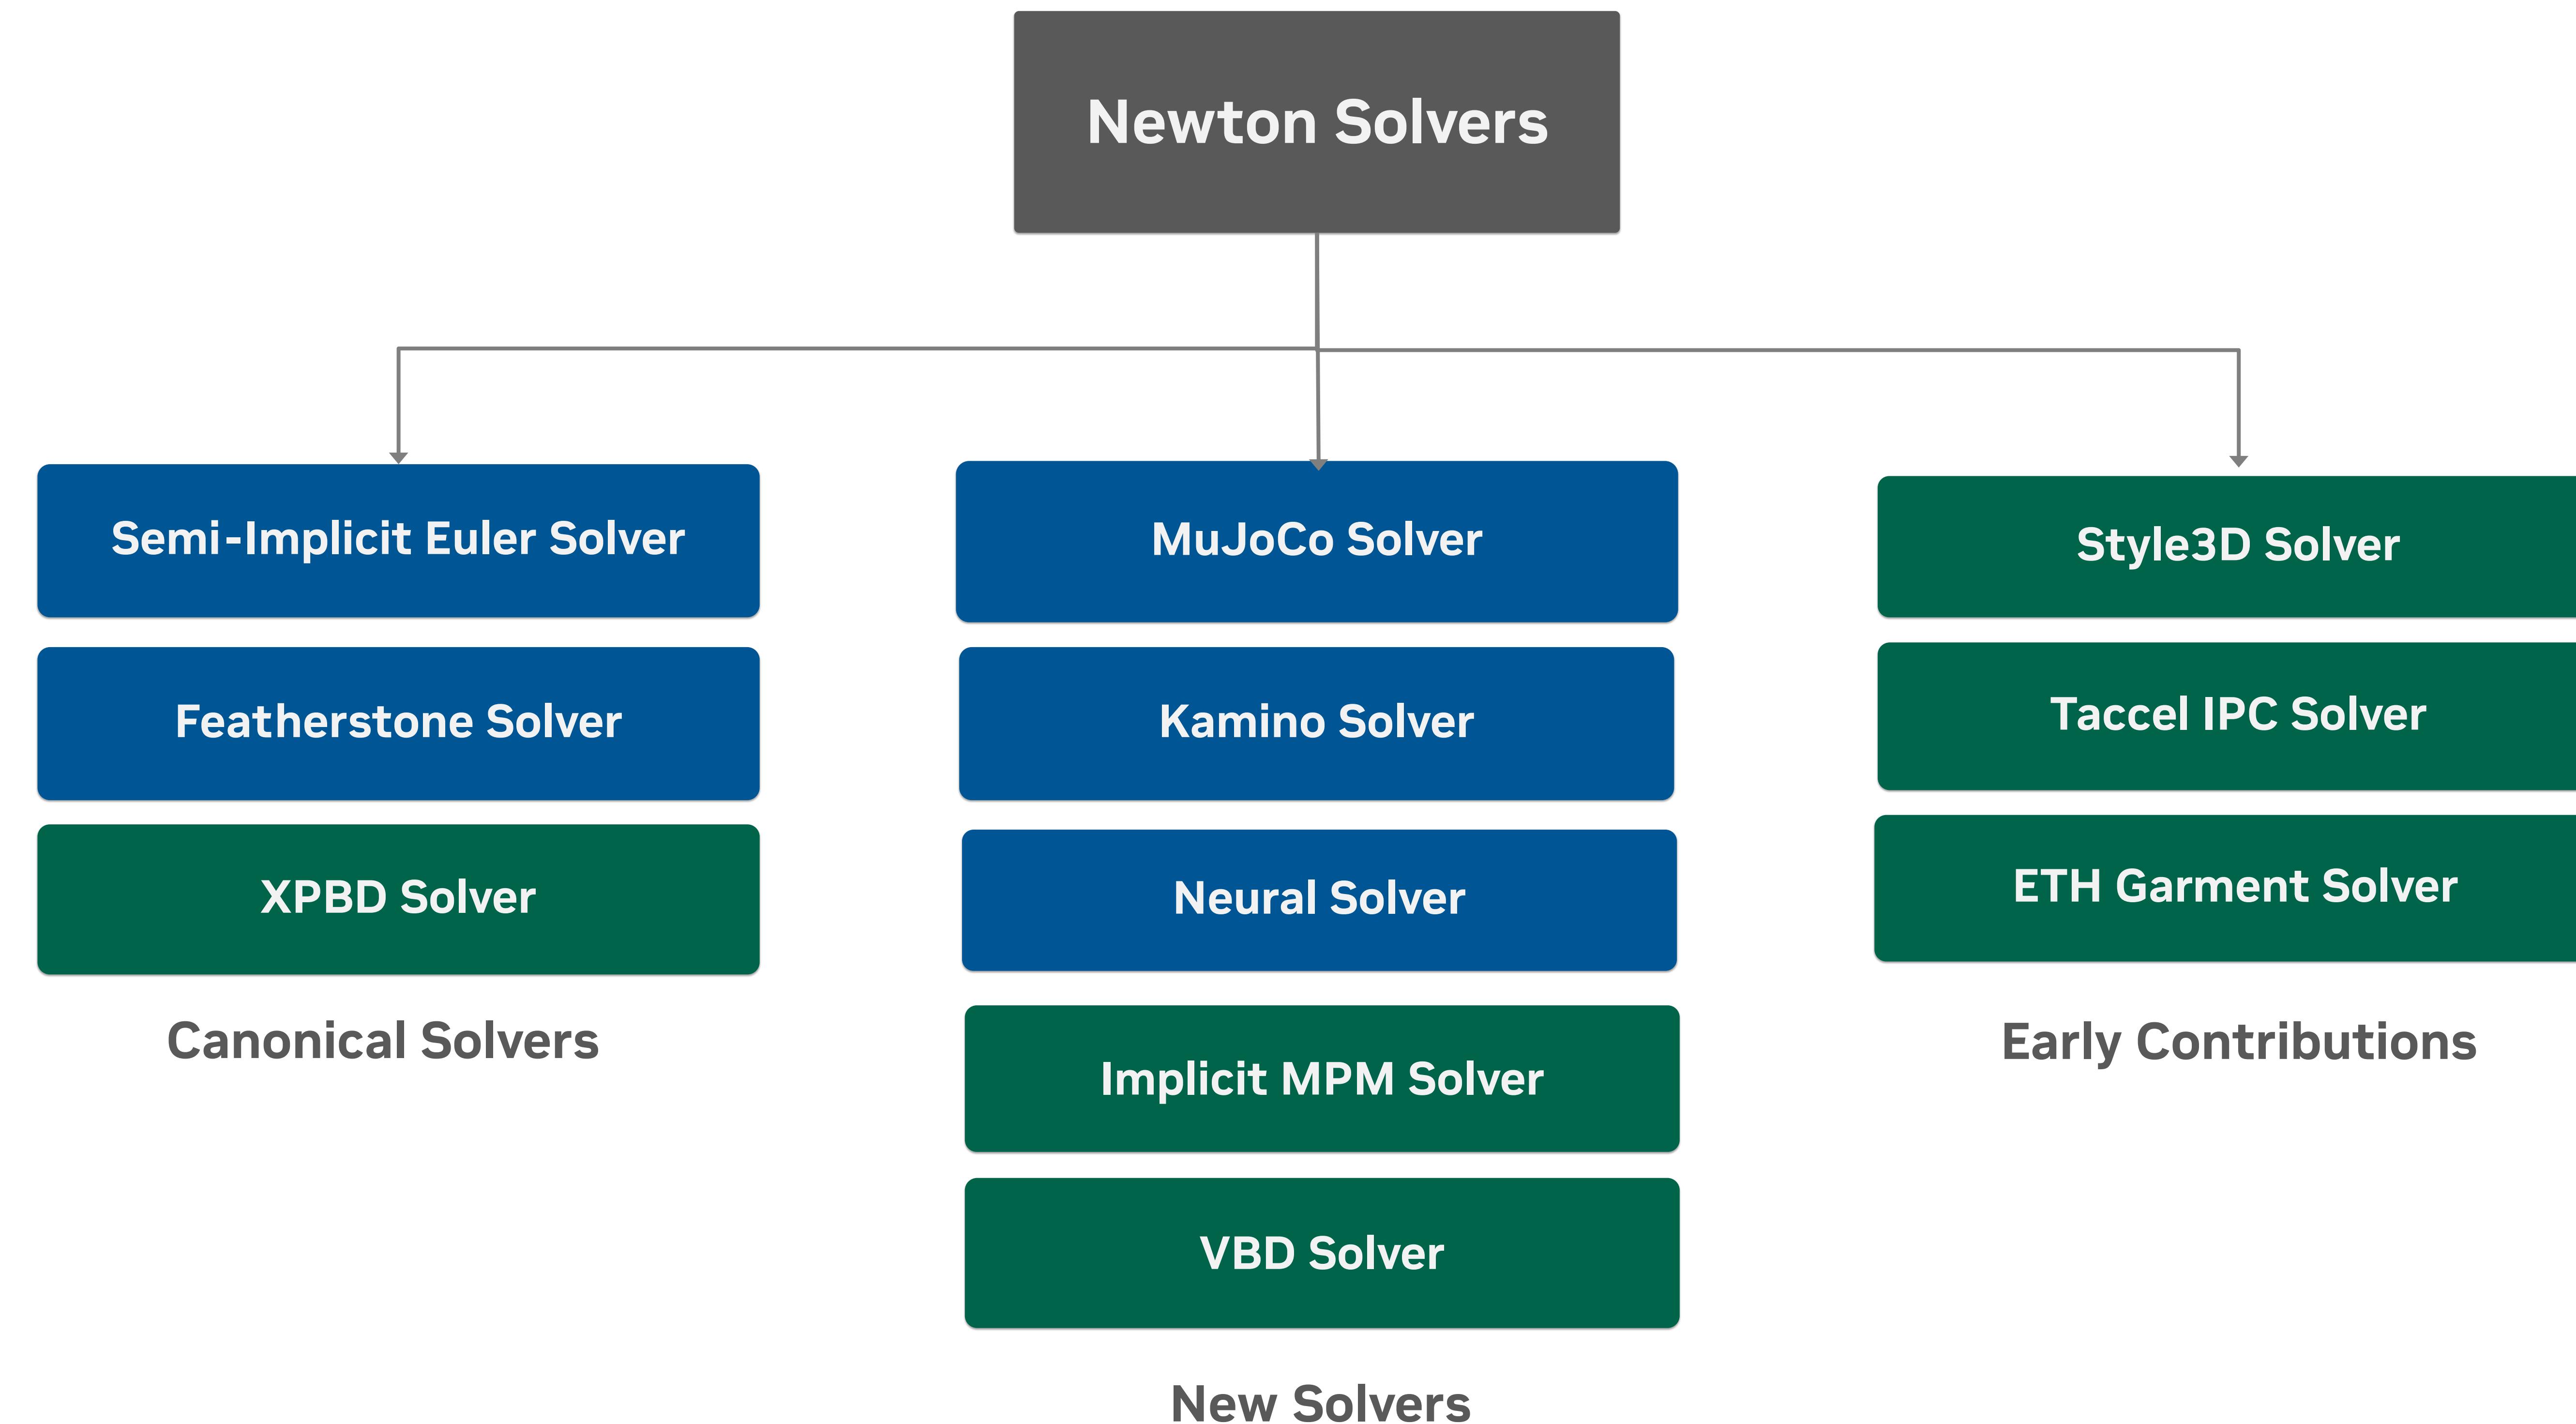
Advanced solvers and Multi-physics coupling support

Rigid Body

Deformable / Hybrid

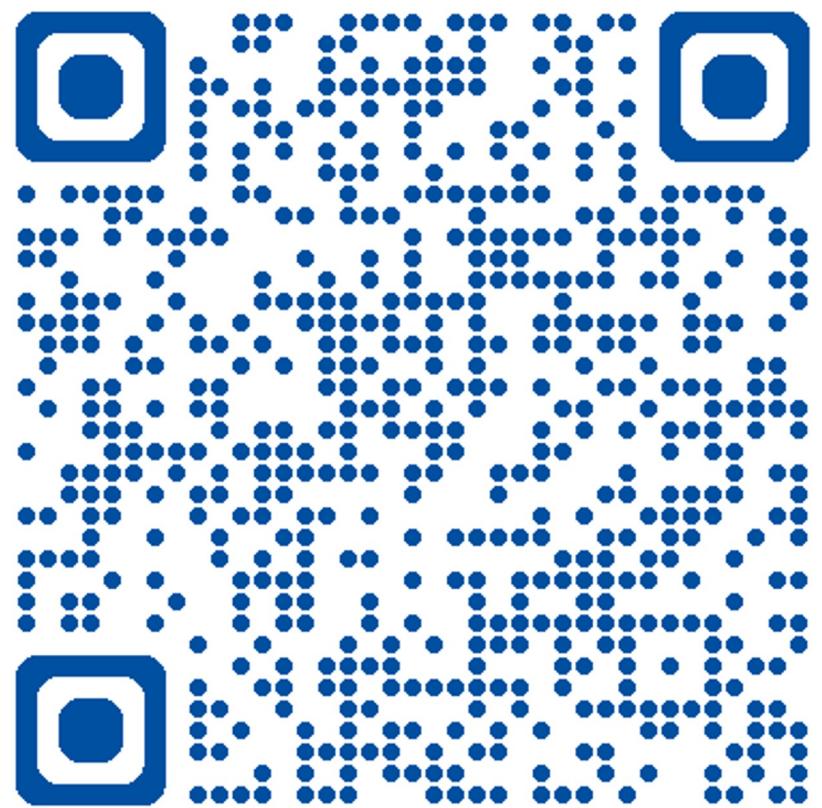
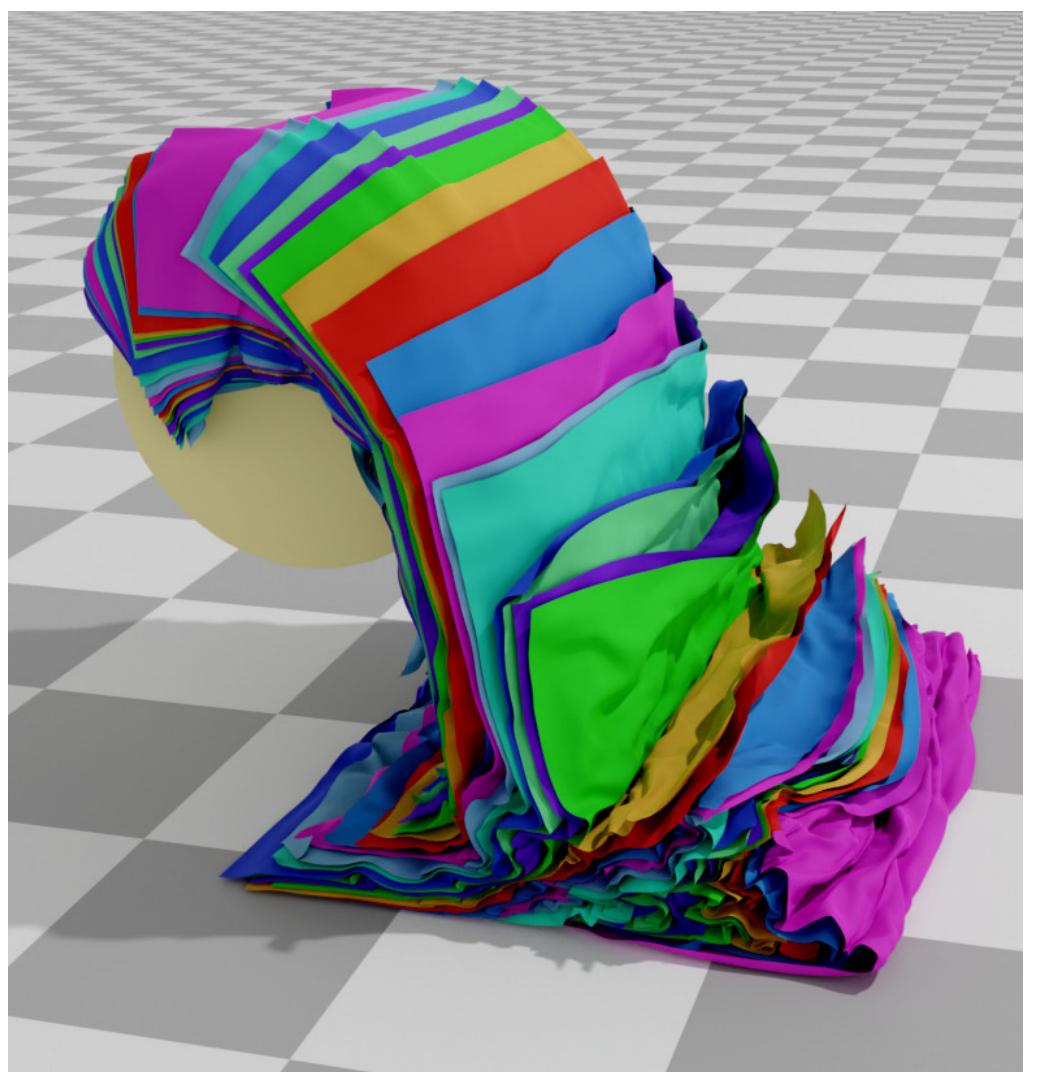
Newton Solvers

Available as Plugins

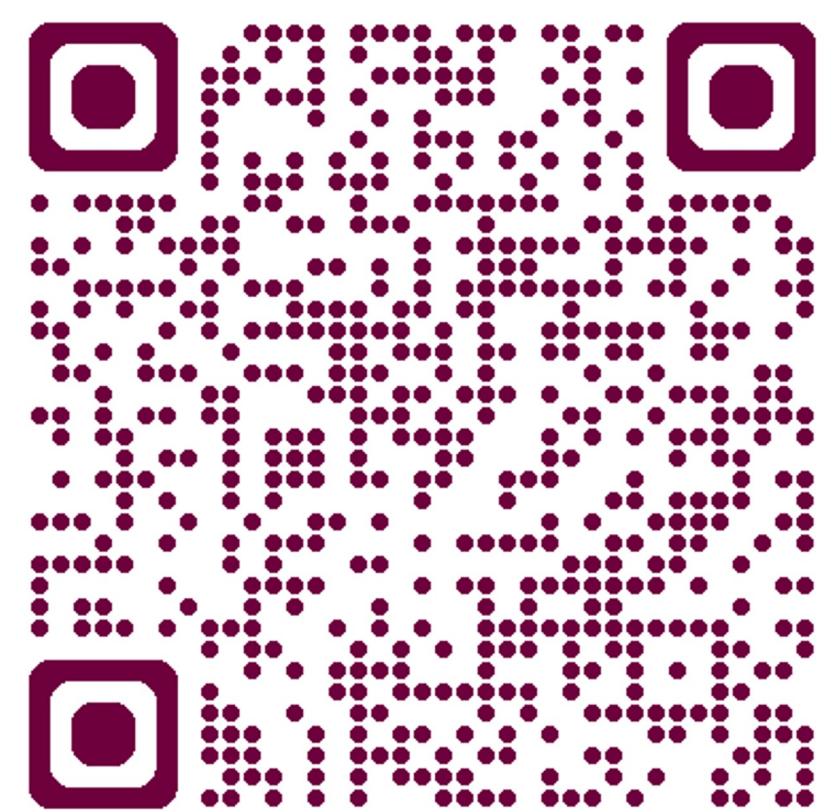


VBD Solver

Vertex Block Descent



**Offset
Geometric
Contact**

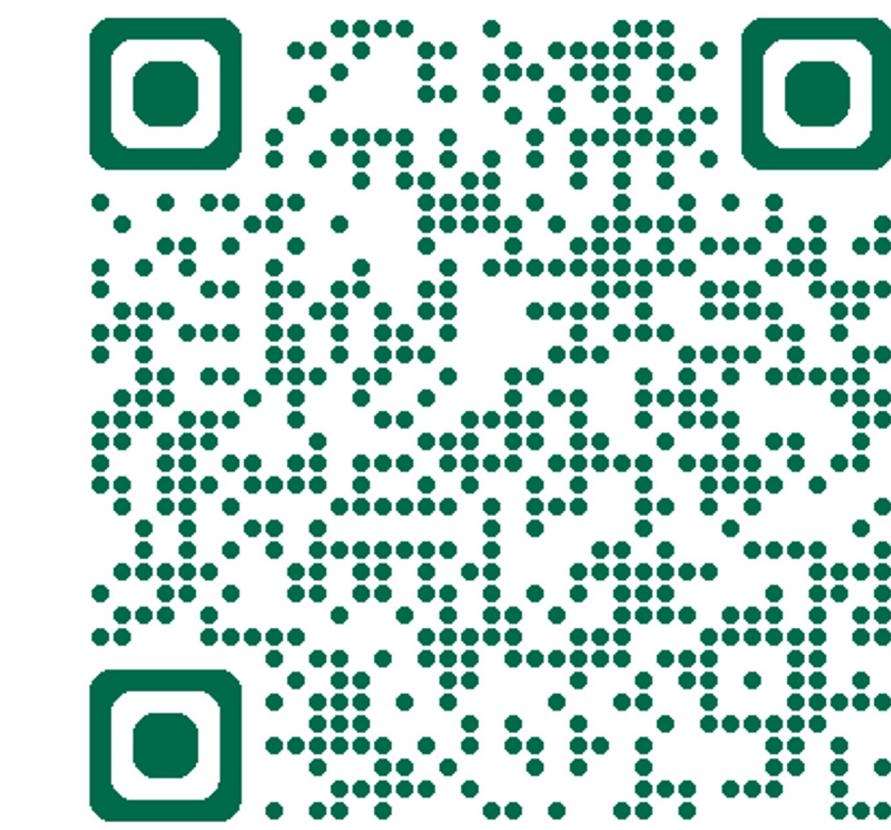


**Vertex
Block
Descent**



Implicit MPM Solver

Material Point Method



**A Semi-Implicit
Material Point
Method for the
Continuum
Simulation of
Granular Materials**



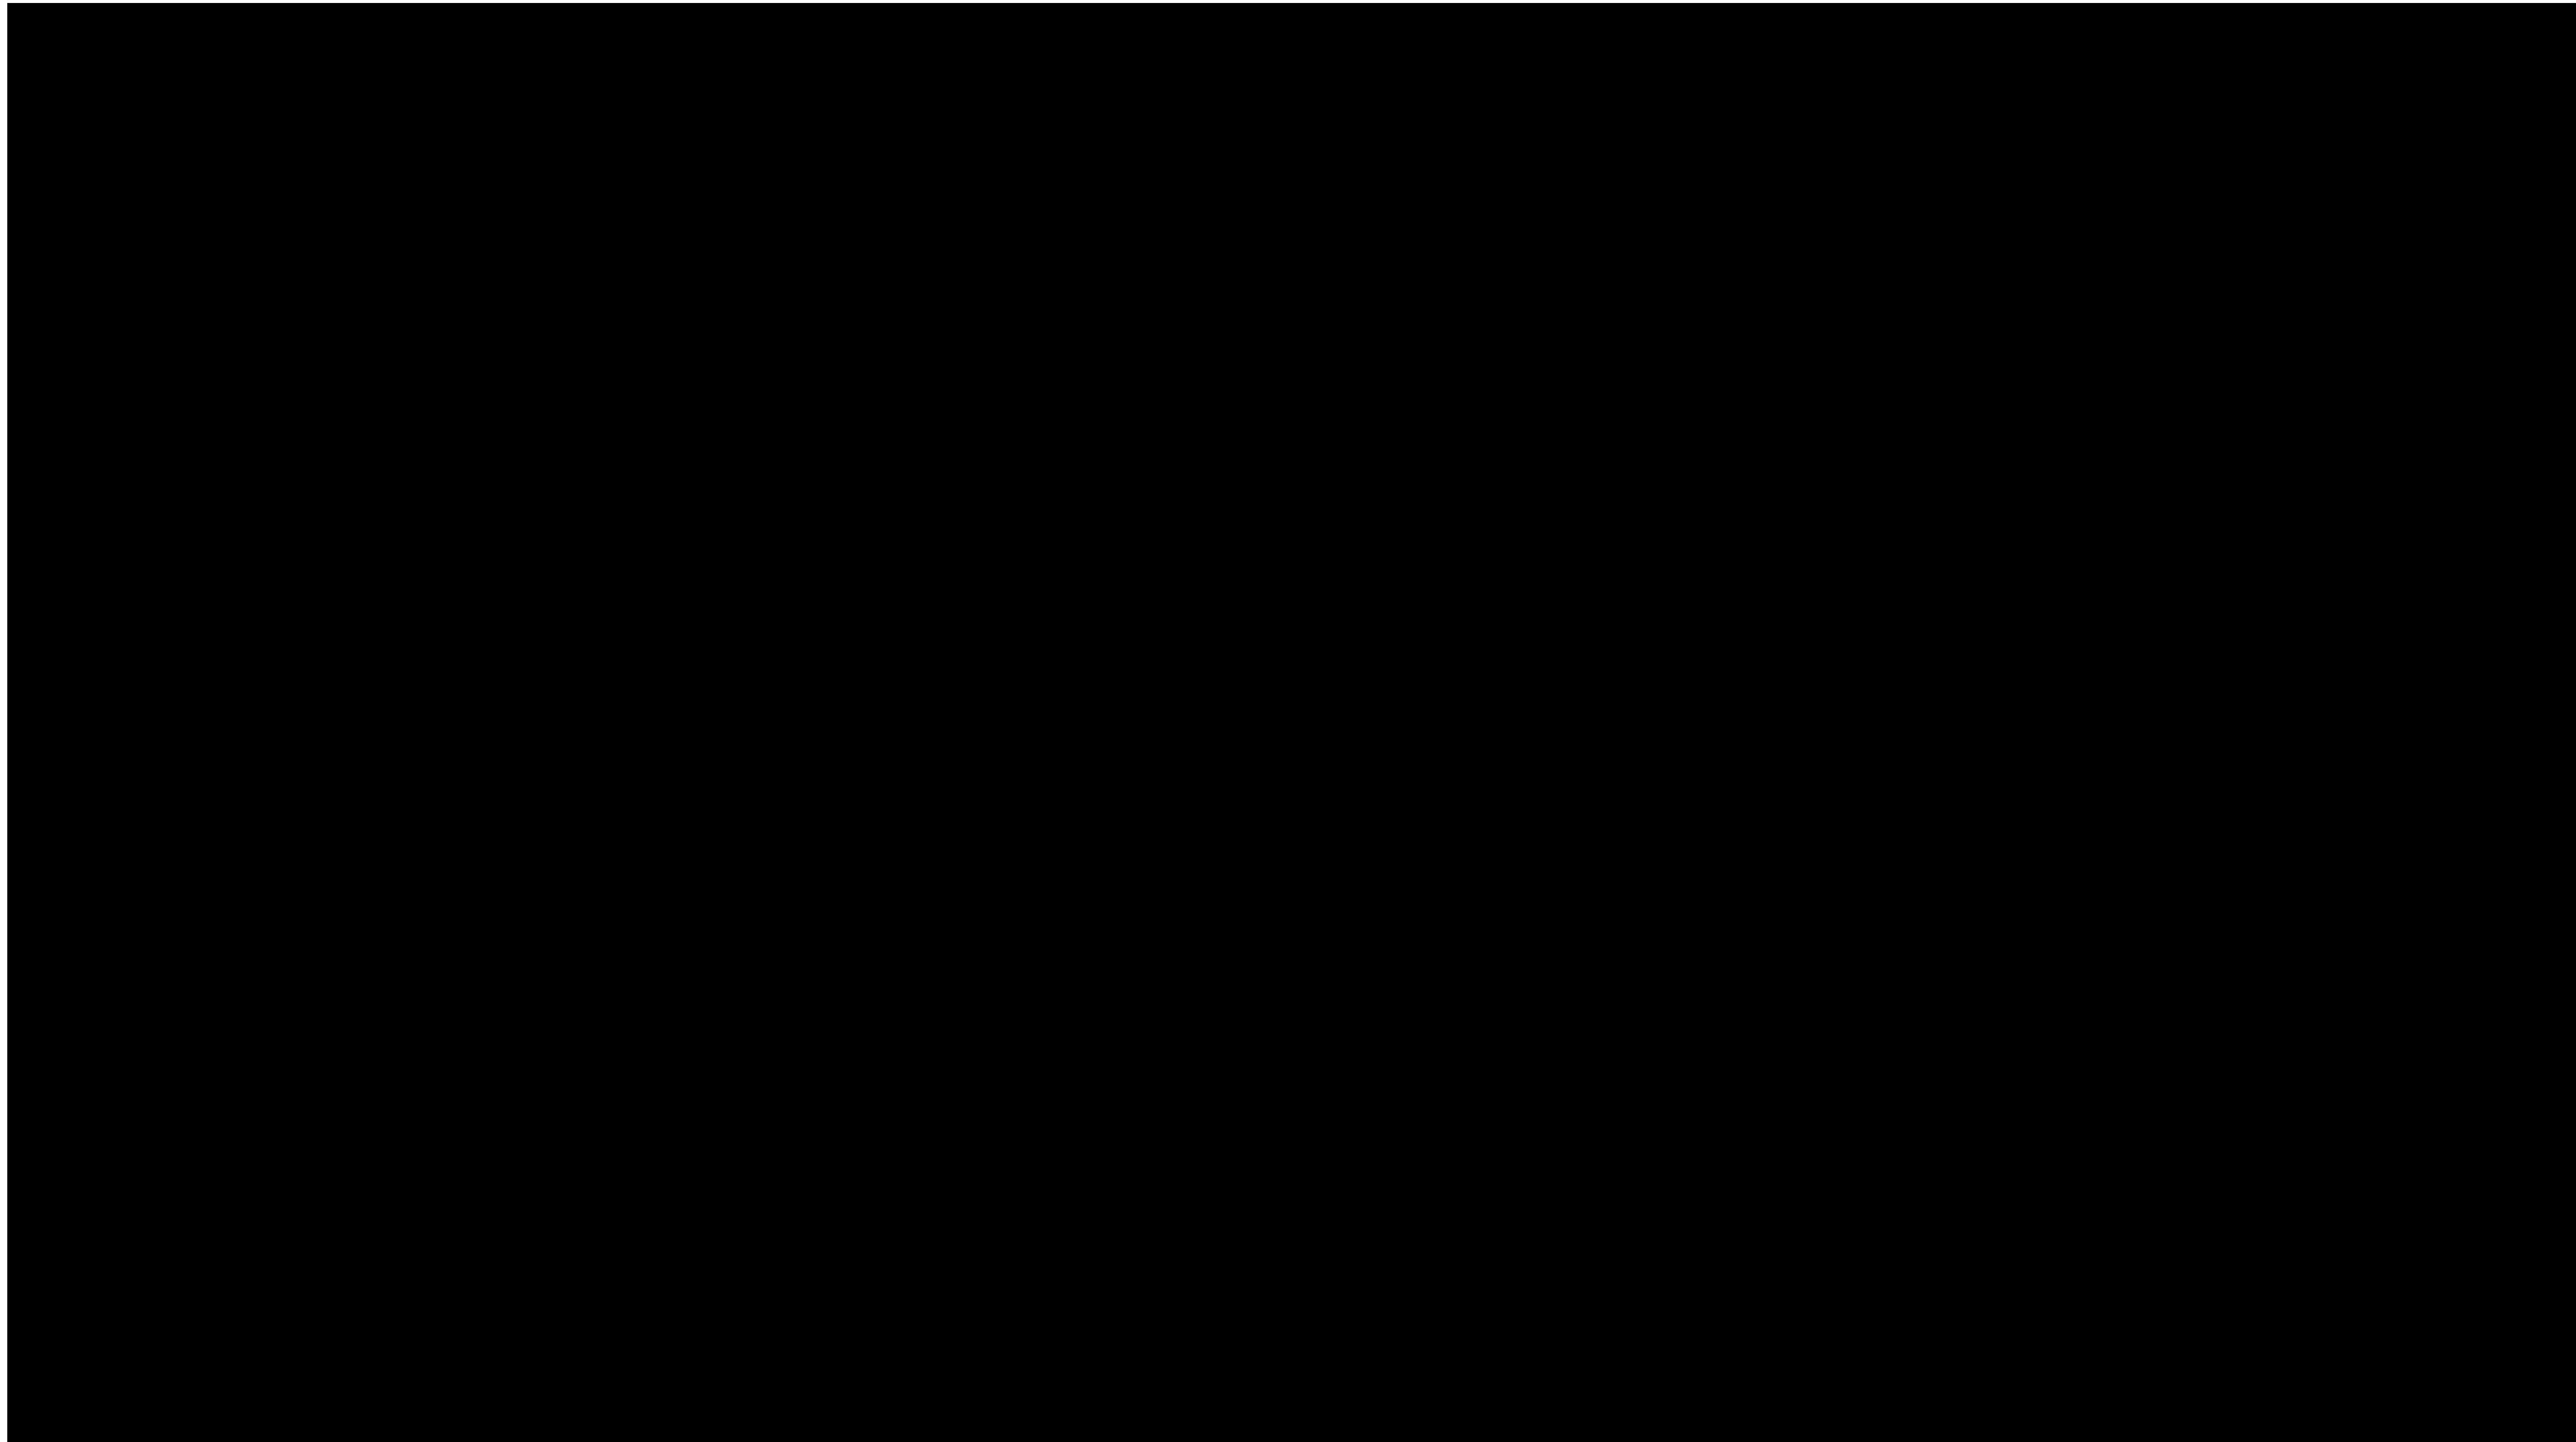
Style3D Solver

cloth simulation



PKU : Taccel IPC Solver

vision-based tactile sensing



Solver Coupling

- Model supports mixed systems: rigid, cloth, soil
- Multiple solvers can run independently
- Currently one-way coupling

```
import newton

builder = newton.sim.ModelBuilder()

# import franka
newton.utils.parse_urdf("franka.urdf", builder)

# add table + cloth
builder.add_body(xform=(0.0, 1.0, 0.0))
builder.add_cloth_grid(pos=(0.0, 2.0, 0.0), dim_x=100,
dim_y=100)

mjc_solver->step(...)
vbd_solver->step(...)
```



Isaac Lab Integration

Experimental Feature Branch

- Under active development, only support limited features and workflows
- Purpose is to let the users understand how Newton integration works in Isaac Lab
- Future updates: Performance improvements and more solver integration

Newton Physics Integration

- Installation
 - Training Environments
 - Newton Visualizer
 - Limitations
 - Solver Transitioning
 - Sim-to-Sim Policy Transfer
 - Sim-to-Real Policy Transfer
- ## Resources

Newton Physics Integration

Newton is a GPU-accelerated, extensible, and differentiable physics simulation engine designed for robotics, research, and advanced simulation workflows. Built on top of **NVIDIA Warp** and integrating MuJoCo Warp, Newton provides high-performance simulation, modern Python APIs, and a flexible architecture for both users and developers.

Newton is an Open Source community-driven project with contributions from NVIDIA, Google Deep Mind, and Disney Research, managed through the Linux Foundation.

This **experimental feature branch** of Isaac Lab provides an initial integration with the Newton Physics Engine, and is under active development. Many features are not yet supported, and only a limited set of classic RL and flat terrain locomotion reinforcement learning examples are included at the moment.

Selection API

Flexible views onto simulation data

- For RL we often need to select a subset of DOFs batched by environment
- Query based on name instead of index to avoid relying on joint orderings
- The Selection API allows creating batched "views" onto the Model and State properties
- Use regex & paths to create [ArticulationViews](#)
- Views isolate data & handle strides automatically

```
from newton.selection import ArticulationView

# create view onto data for one articulation in the Model
ants = ArticulationView(model, "/World/envs/*/Robot/torso")

# get strided joint arrays
ant_q = ants.get_attribute("joint_q", state_0)
ant_qd = ants.get_attribute("joint_qd", state_0)

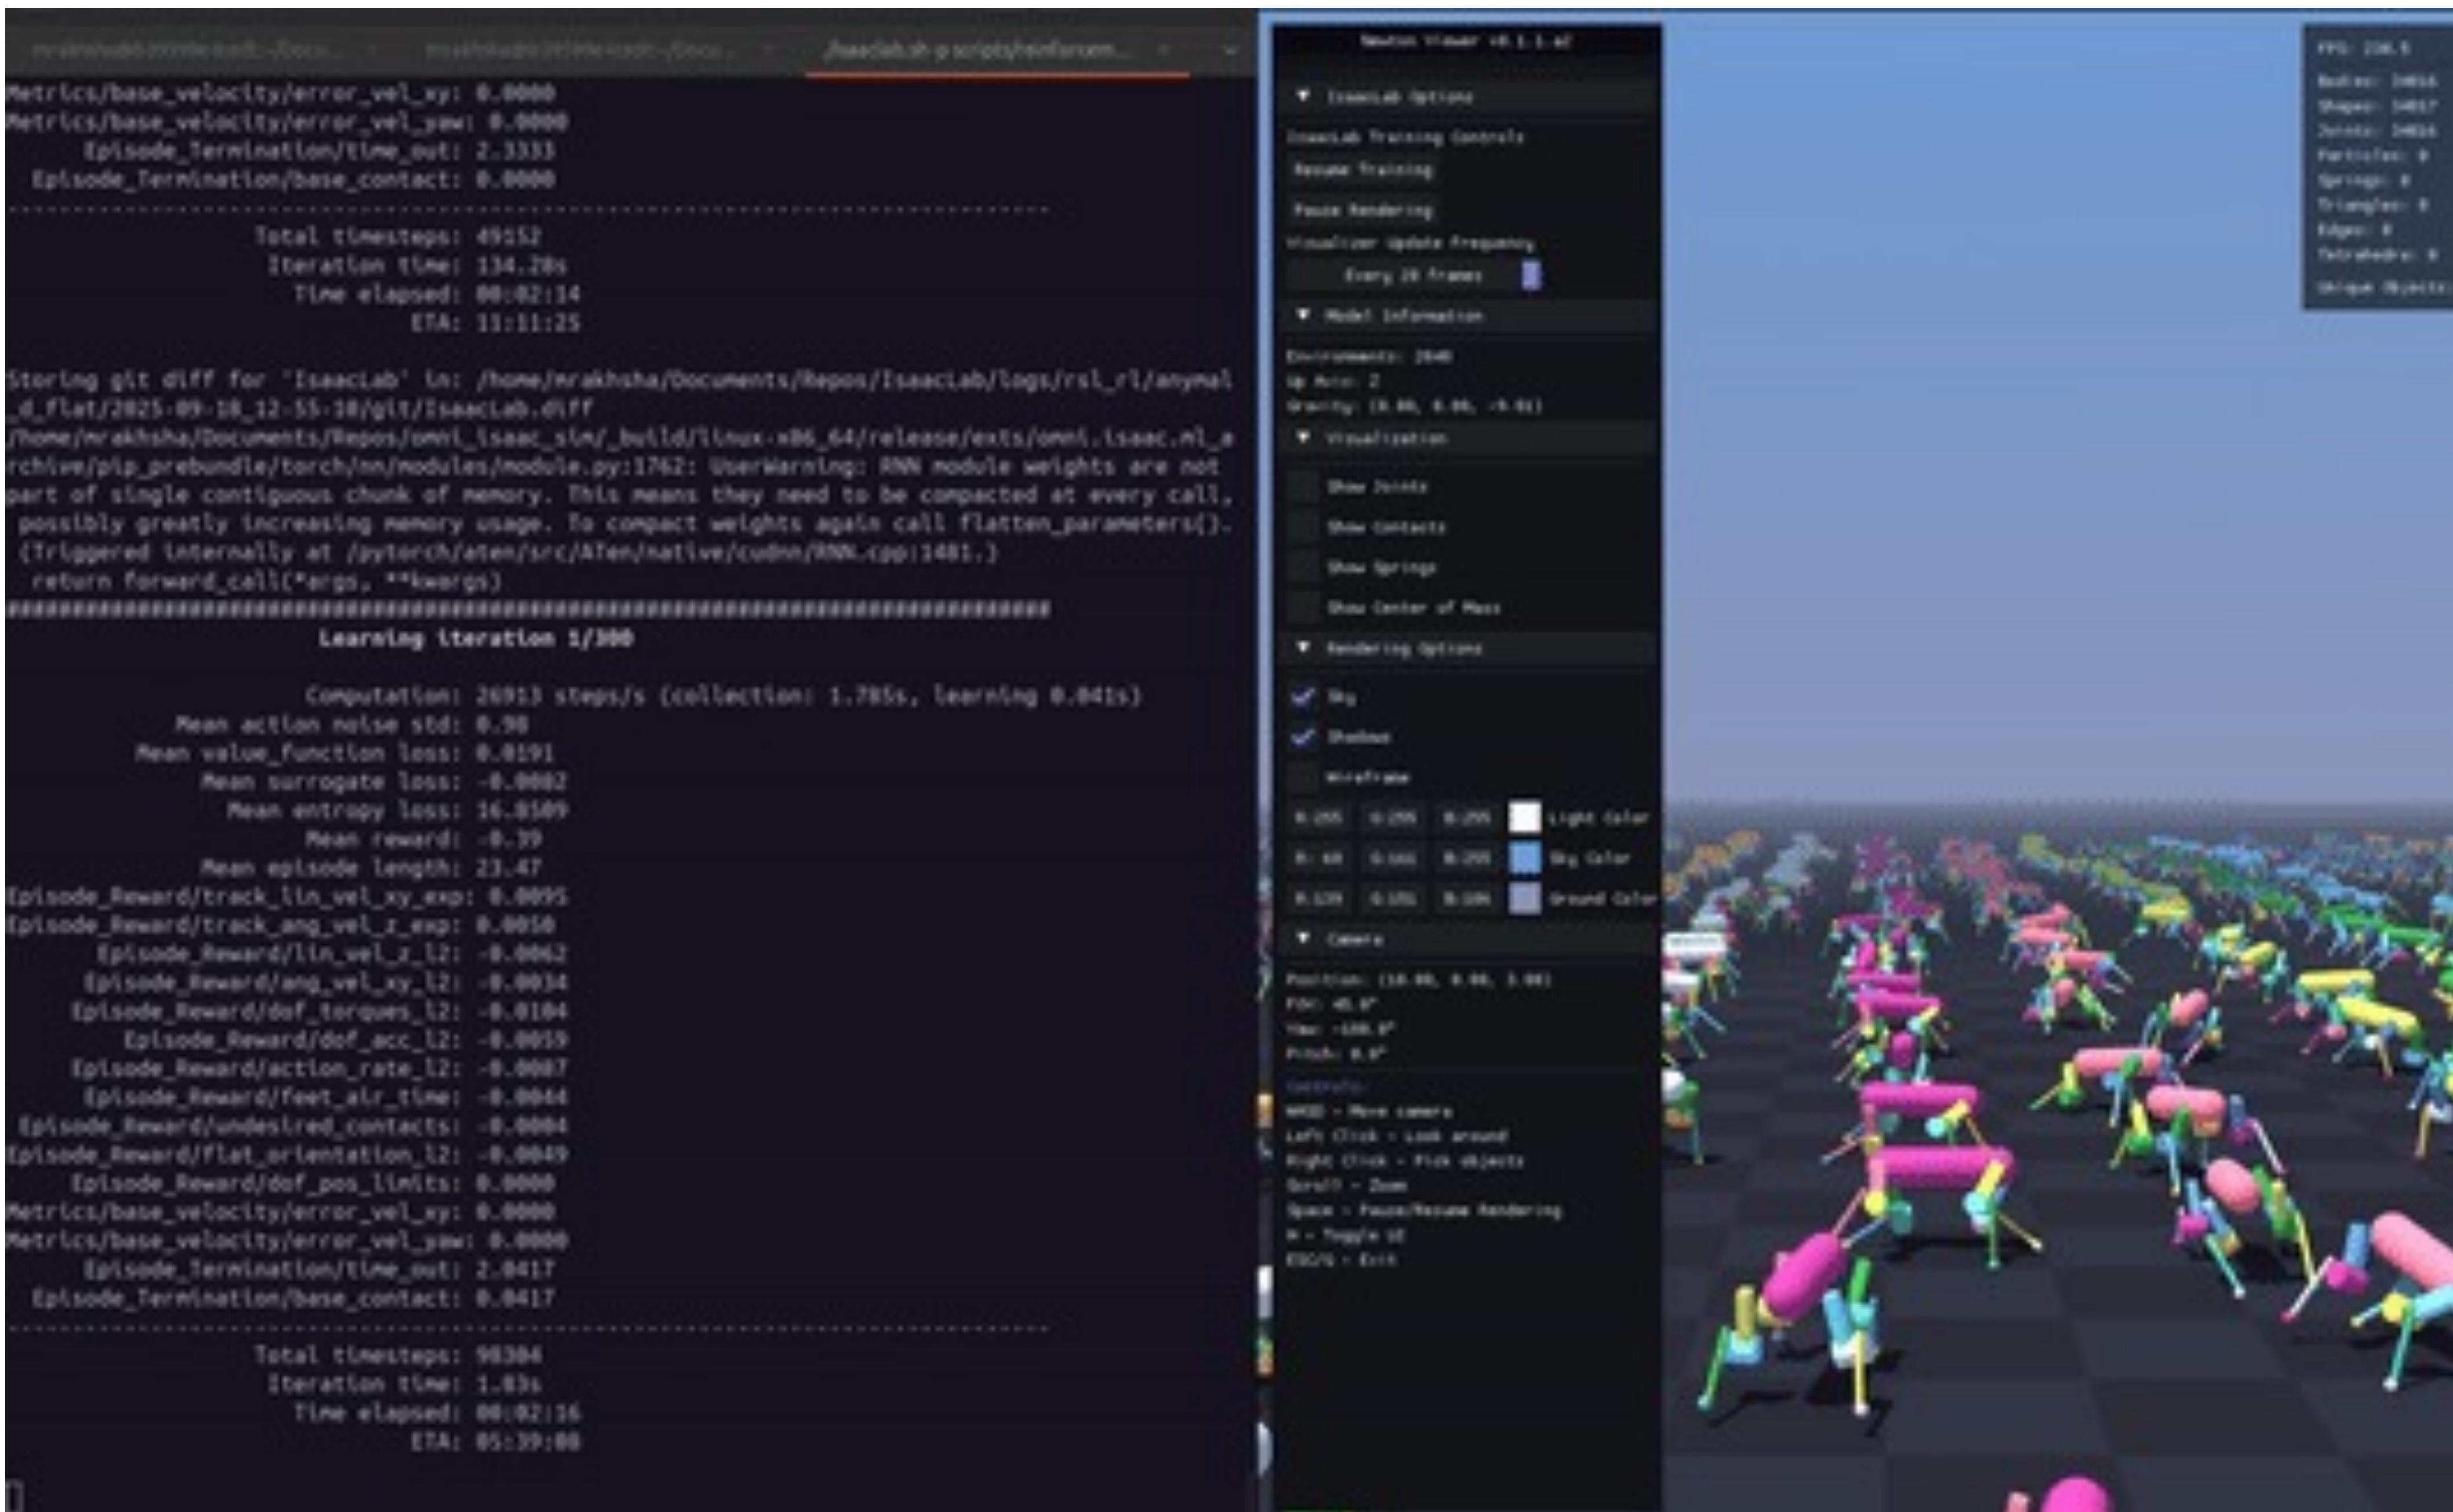
# set dof states (masked)
ants.set_attribute("joint_q", state_0, xform, mask=mask)
ants.set_attribute("joint_qd", state_0, velocity, mask=mask)

# helpers for manipulating root state
ants.set_root_transforms(state_0, xform, mask=mask)
ants.set_root_velocities(state_0, velocity, mask=mask)
```

Isaac Lab Integration Highlights

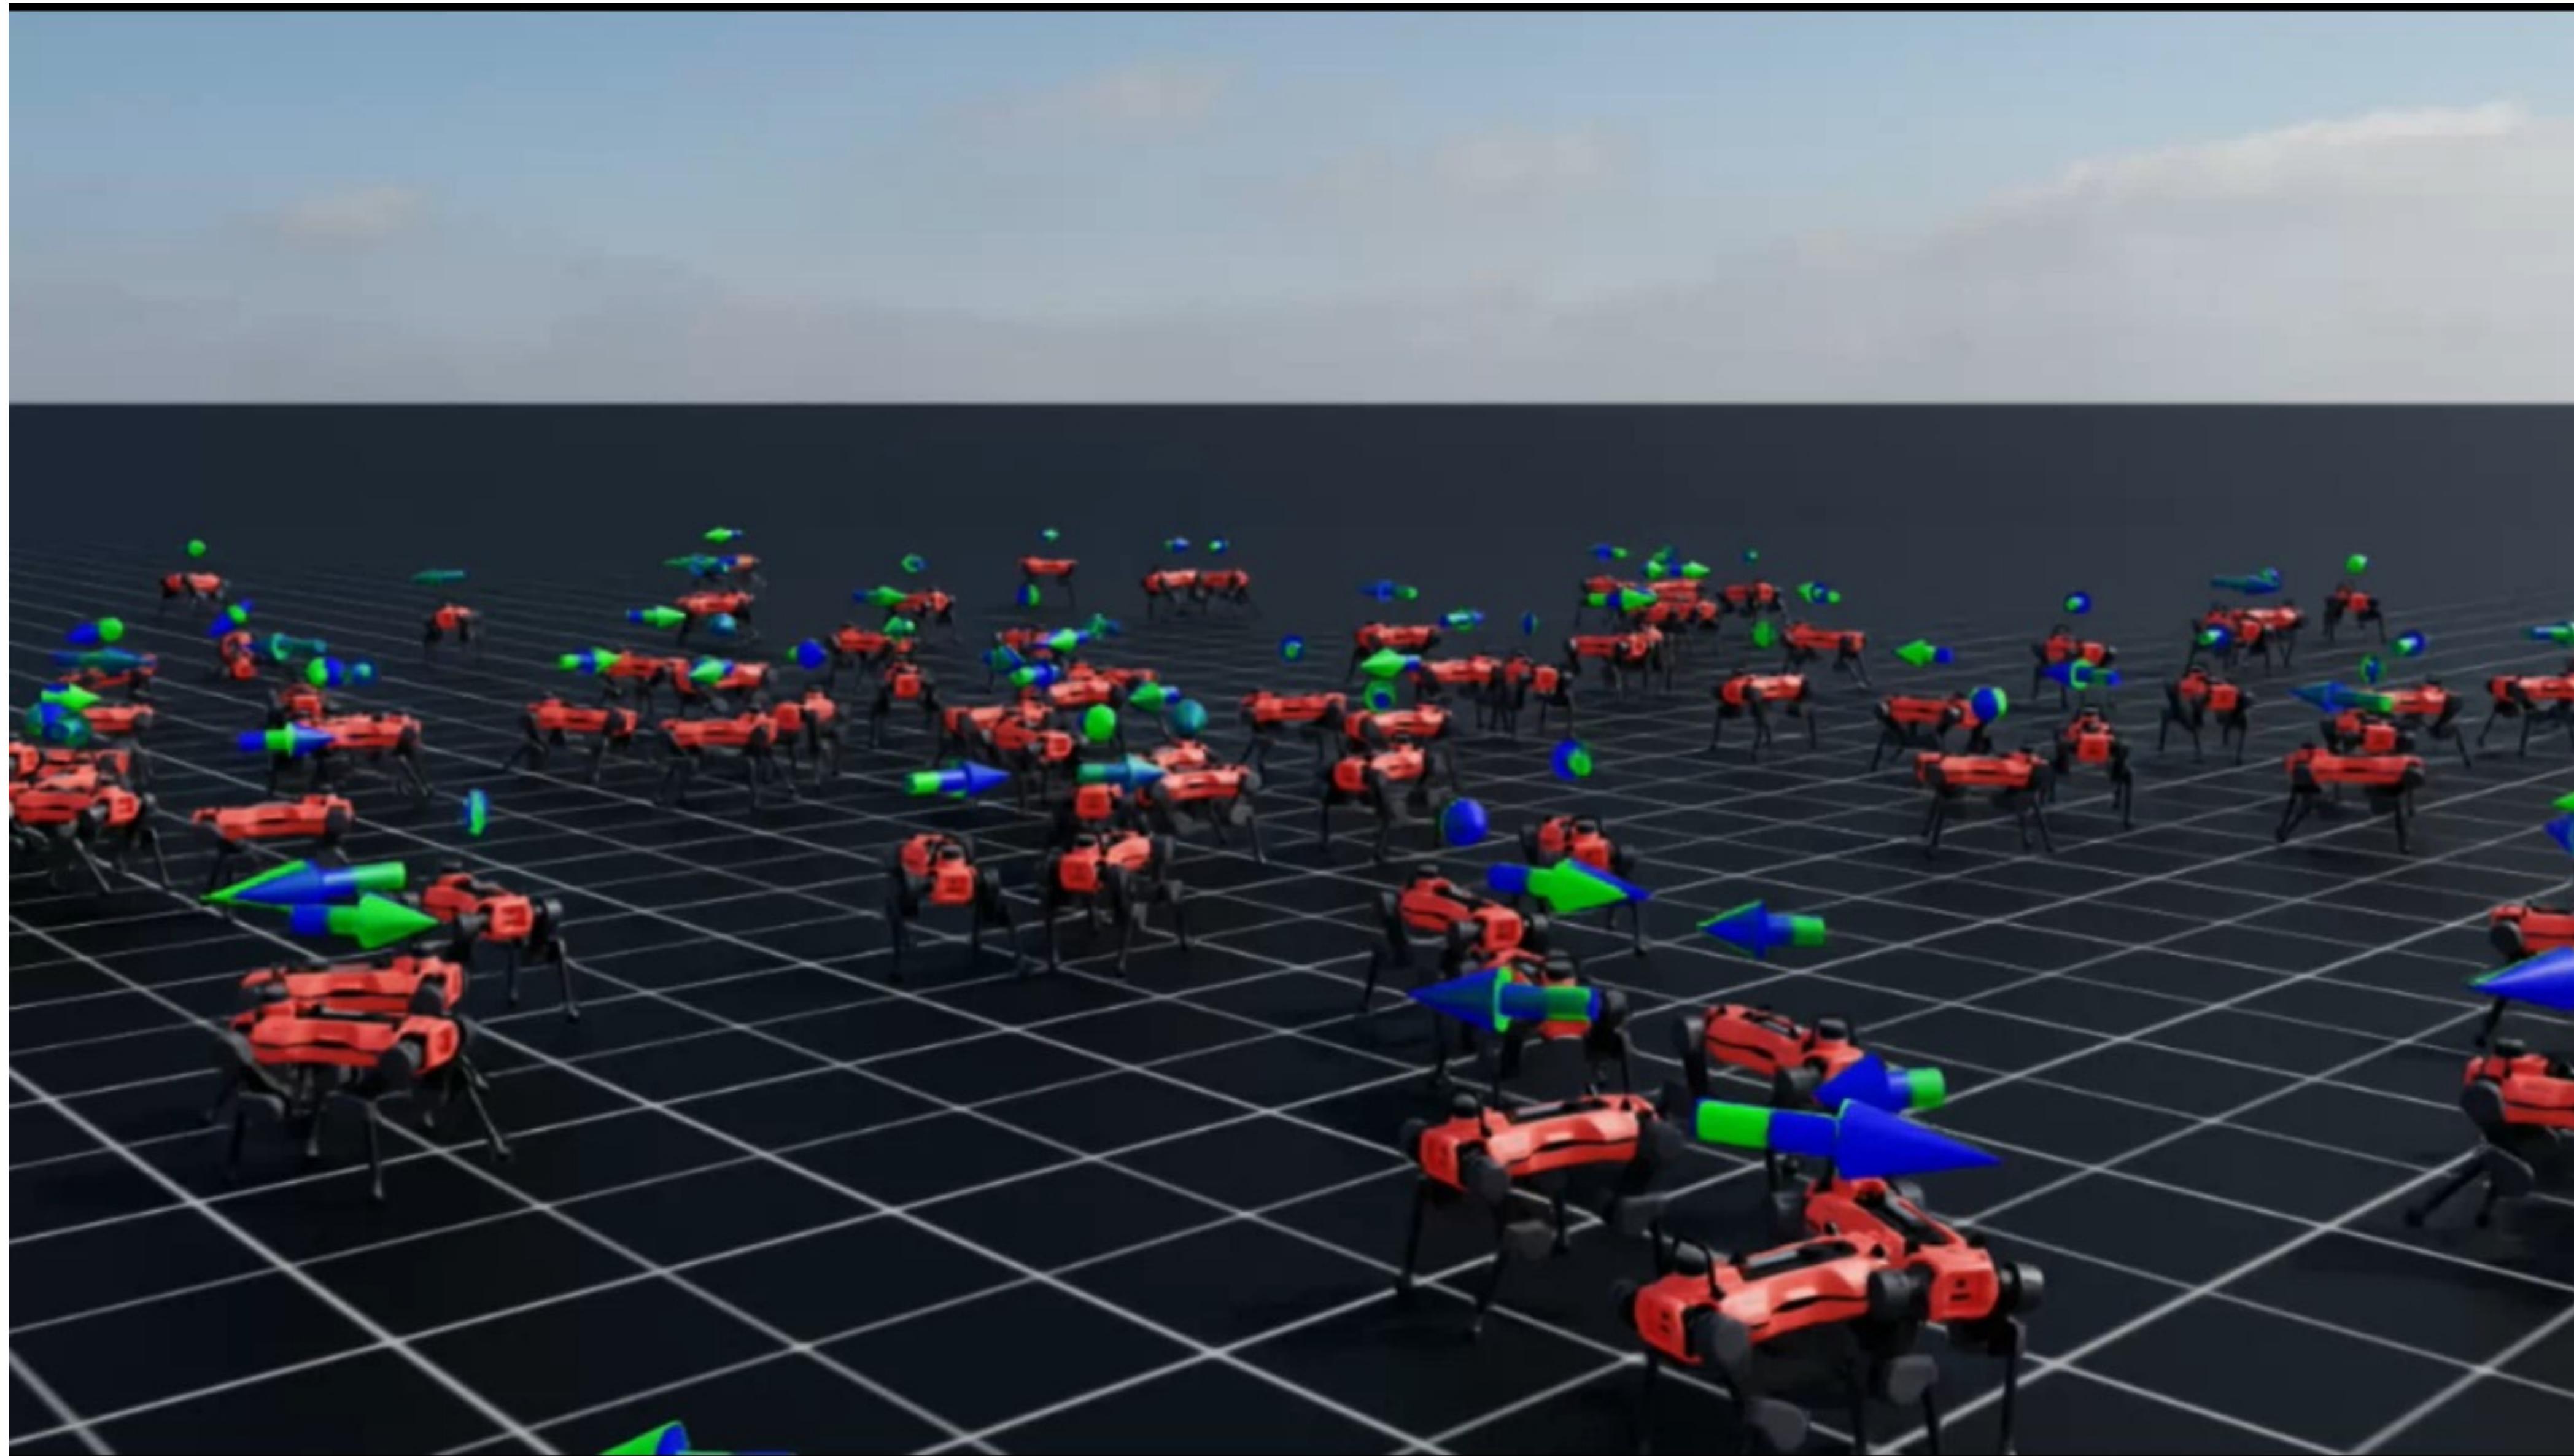
Training

```
./isaaclab.sh -p scripts/reinforcement_learning/rsl_rl/train.py  
--task Isaac-Velocity-Flat-Anymal-D-v0 --num_envs 4096 --headless  
--newton_visualizer
```



Isaac Lab Integration Highlights

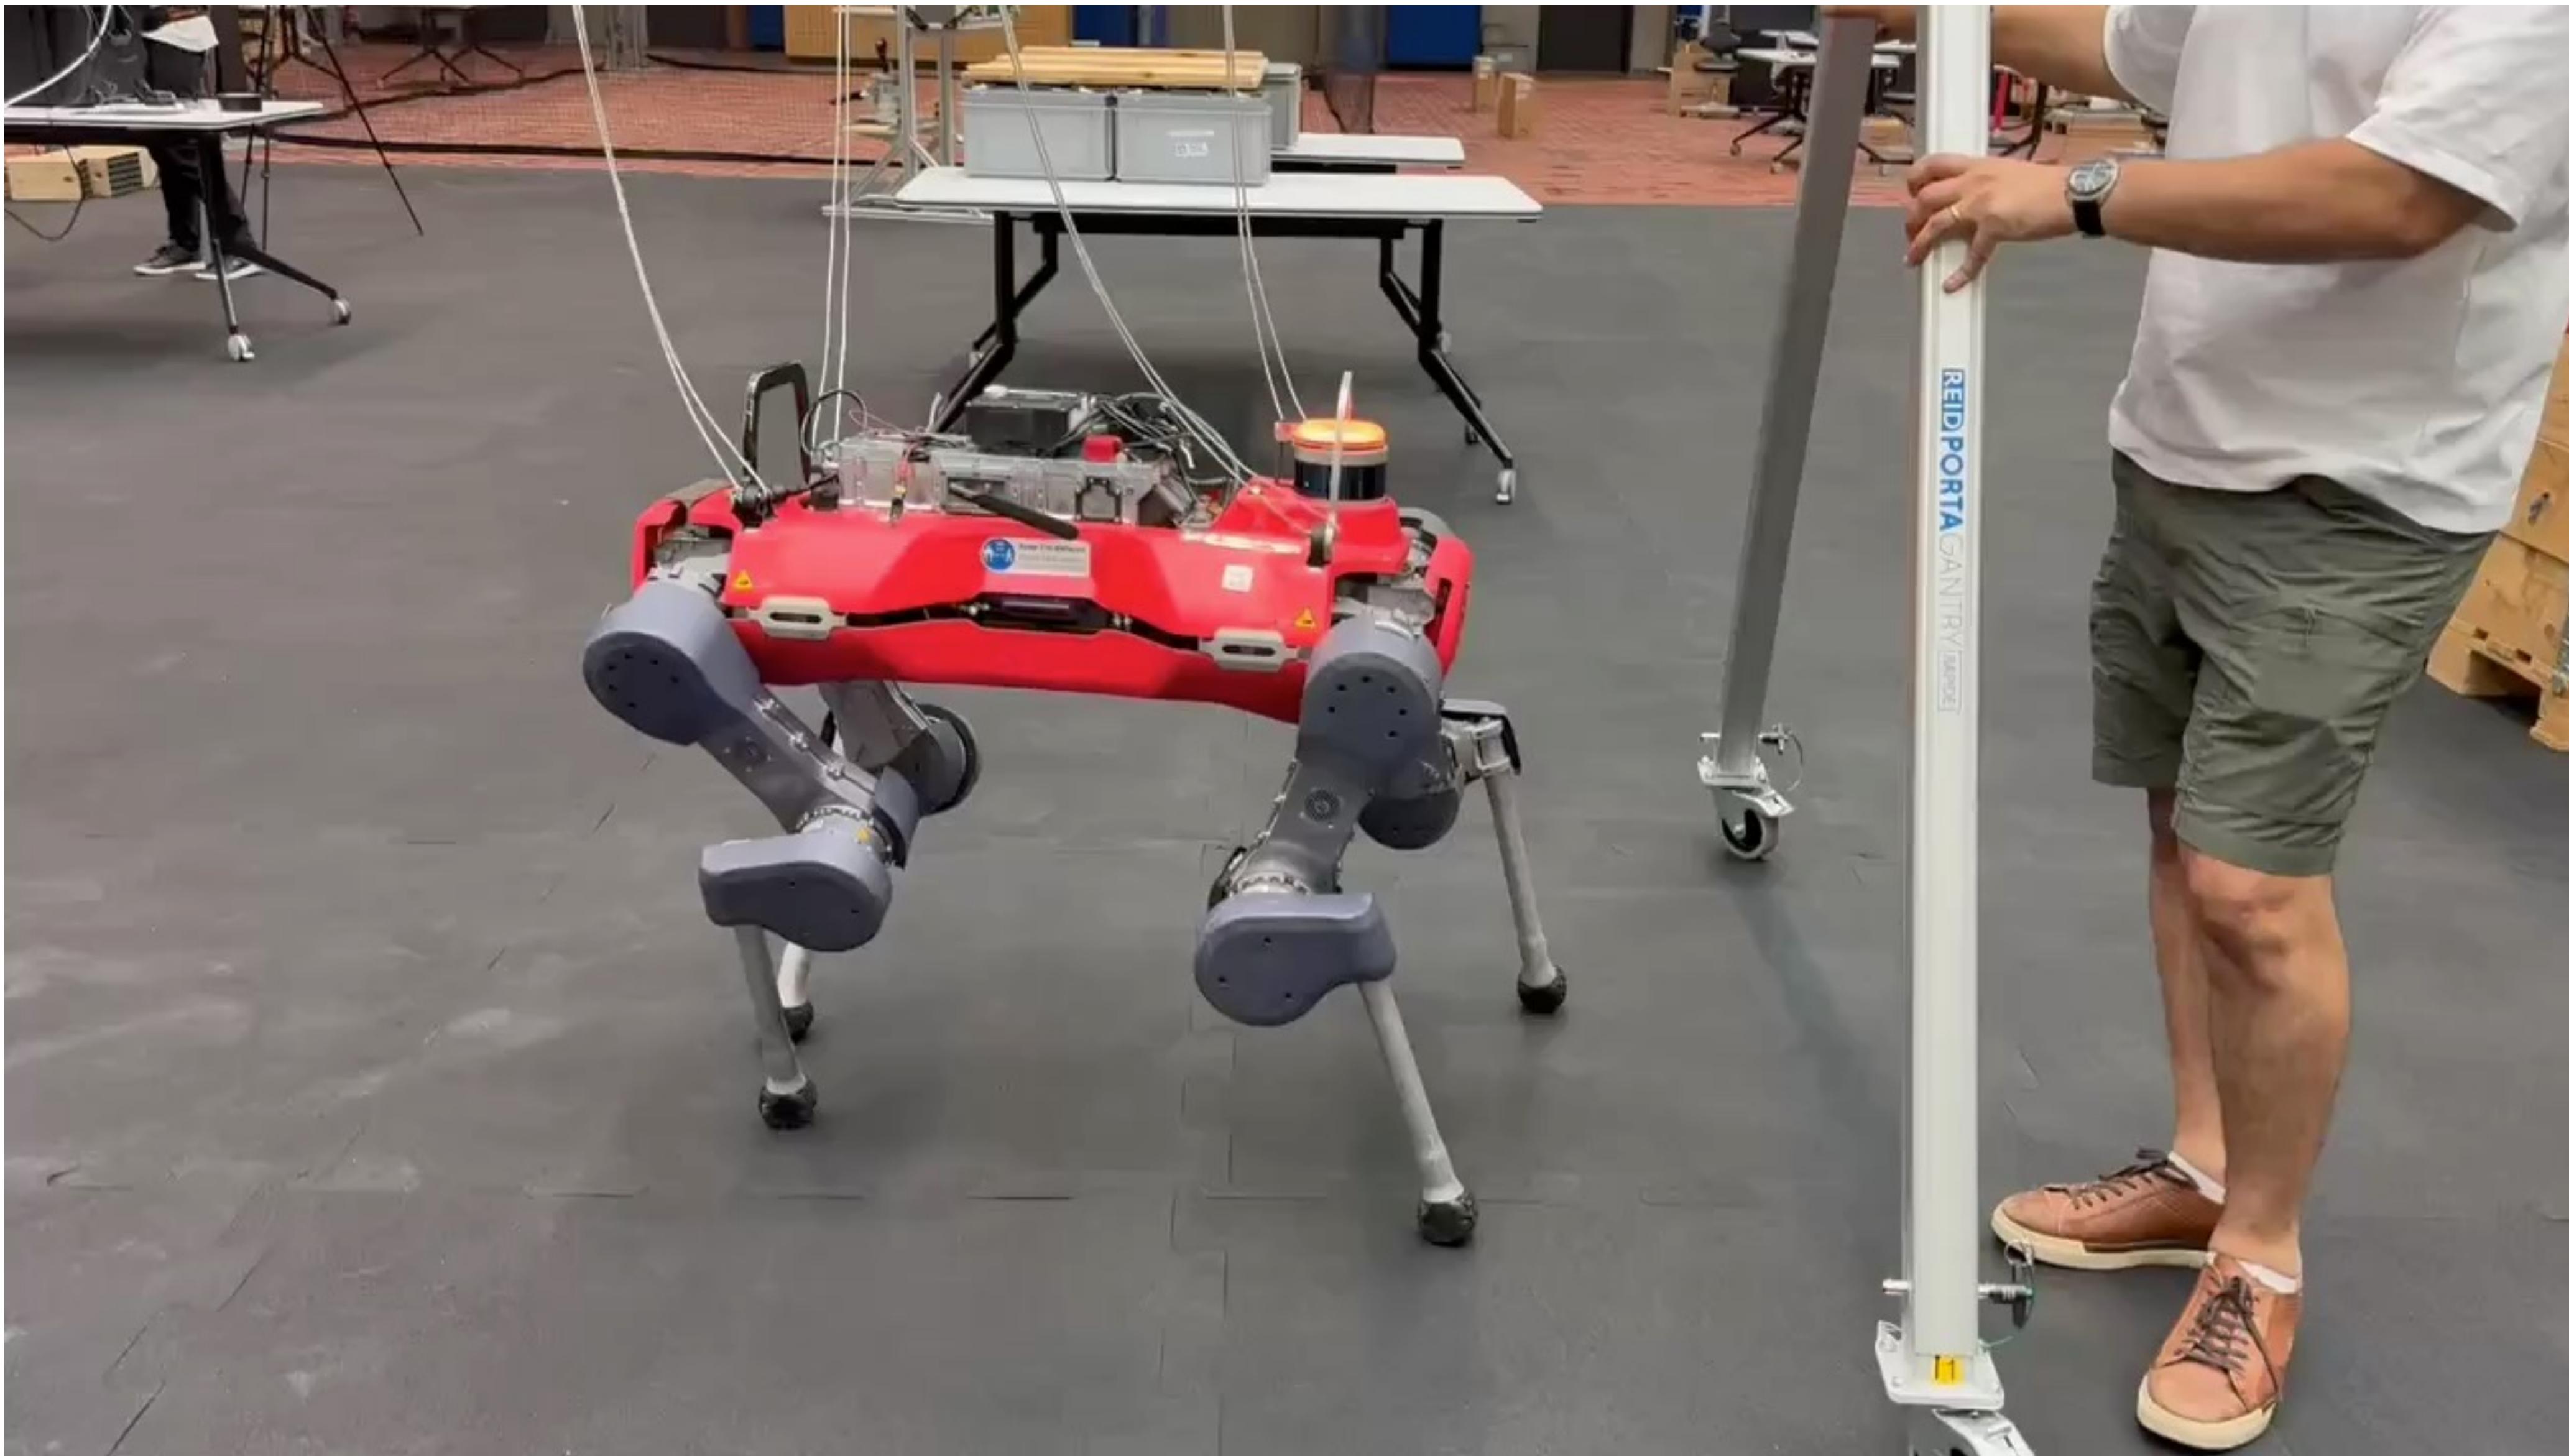
Sim2Sim



Newton to PhysX

Isaac Lab Integration Highlights

Sim2Real

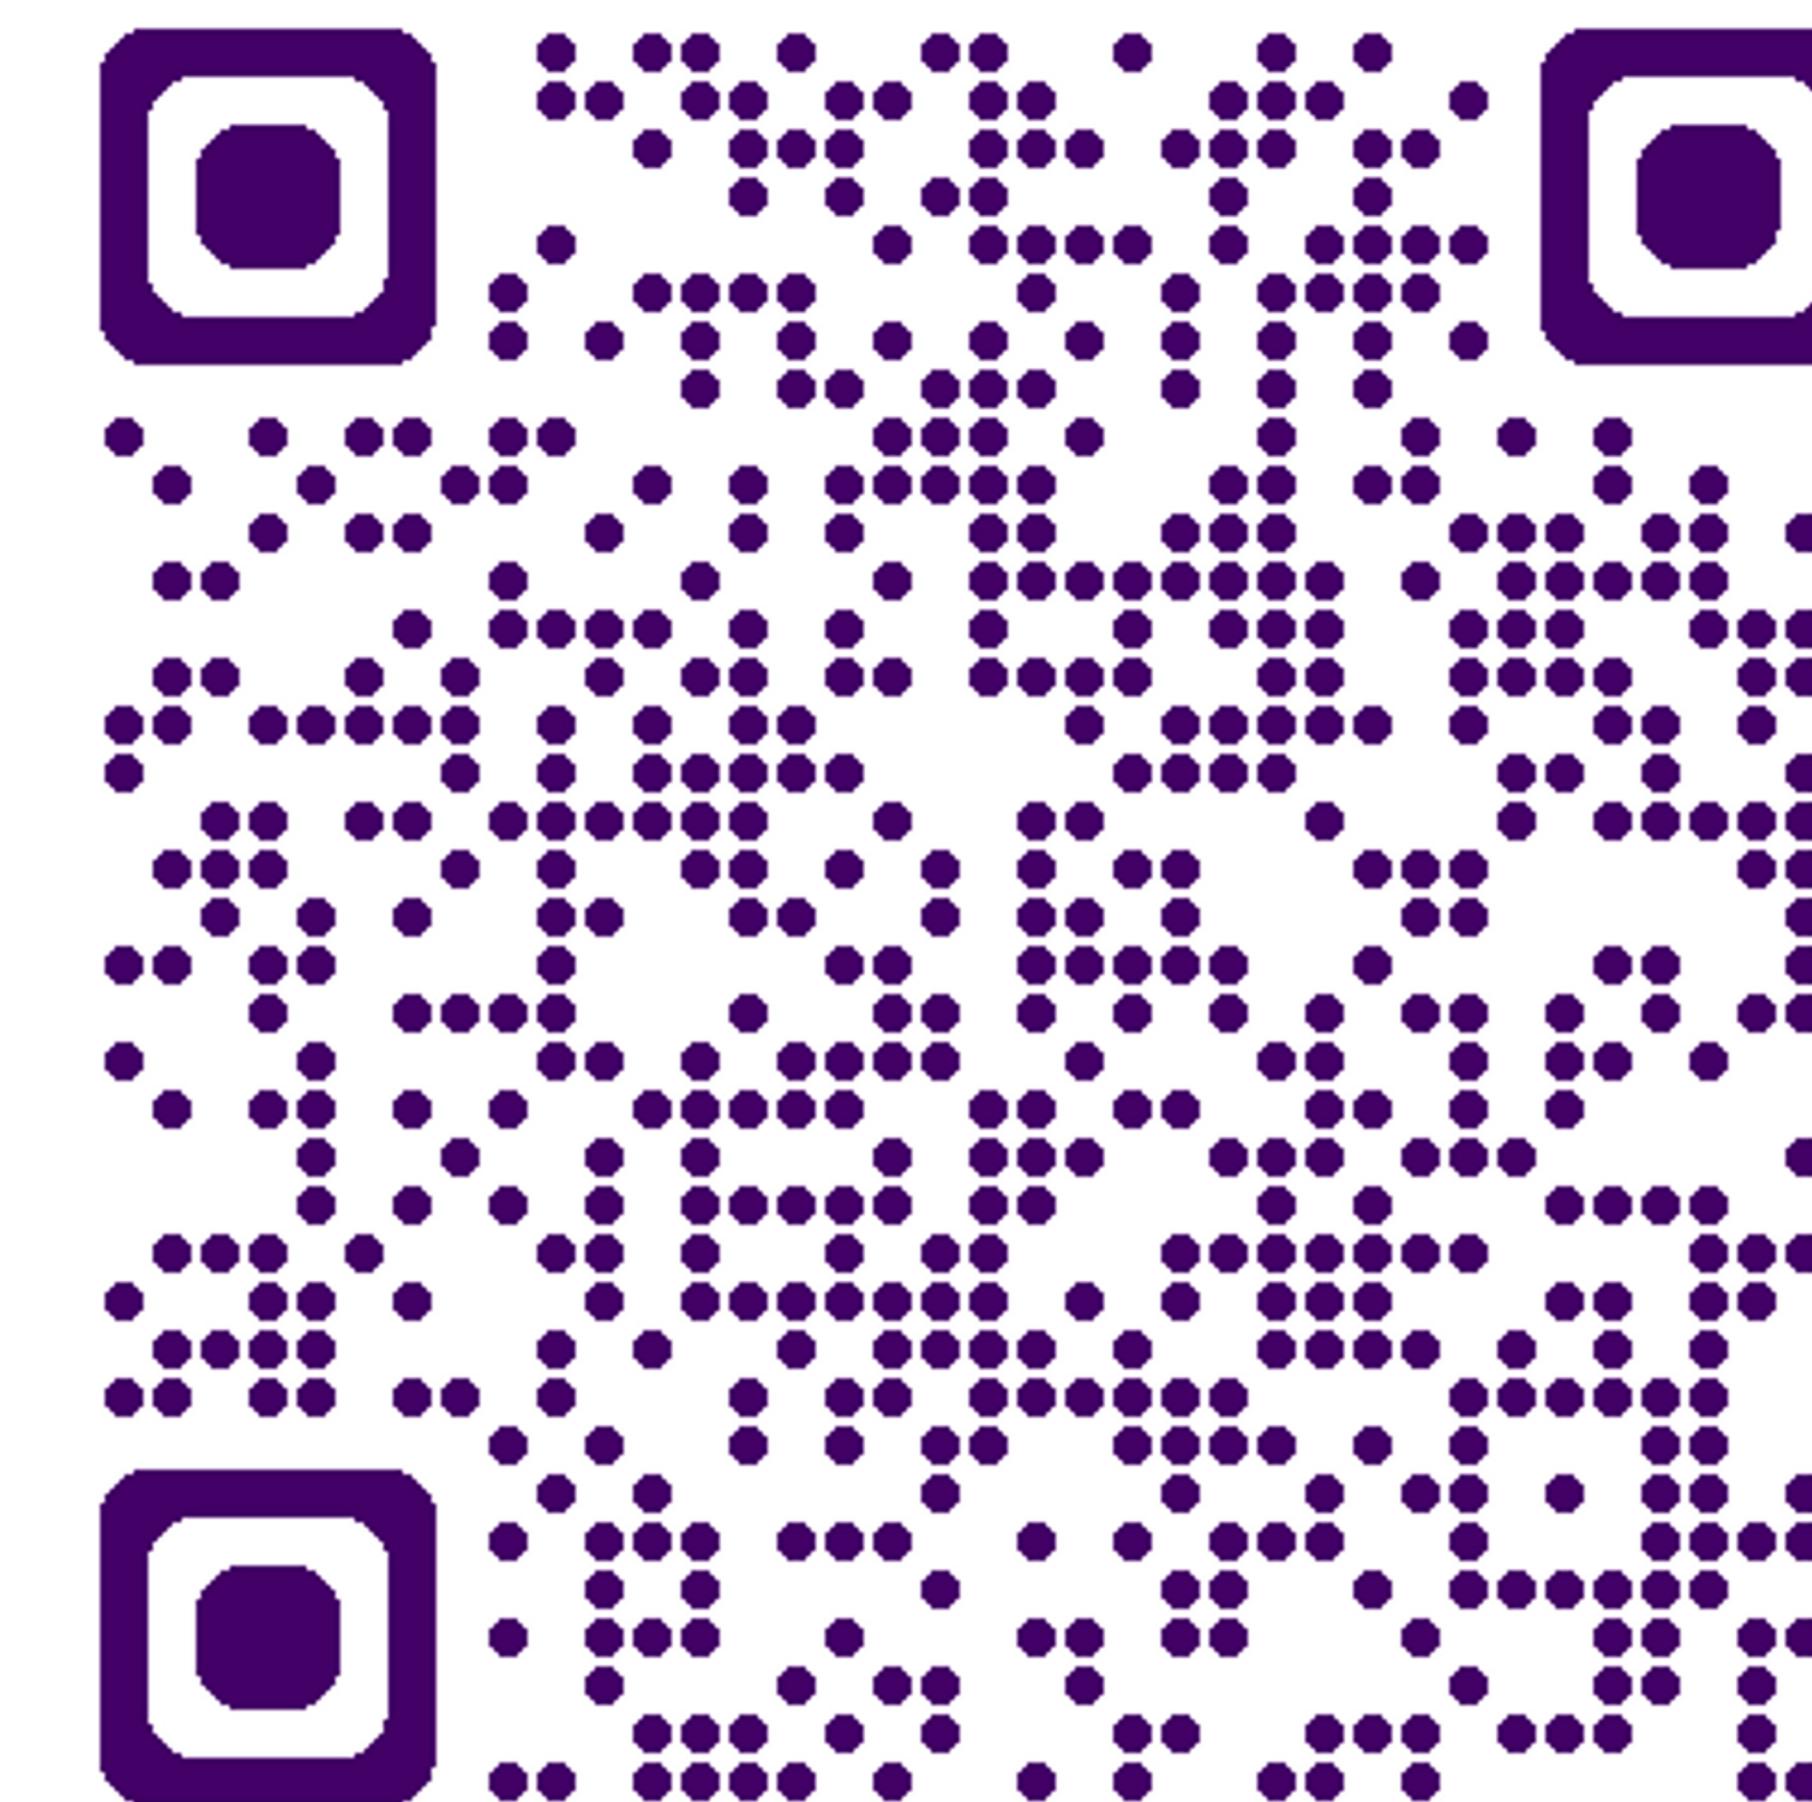


Newton Timeline

MVP, Alpha, Beta, GA

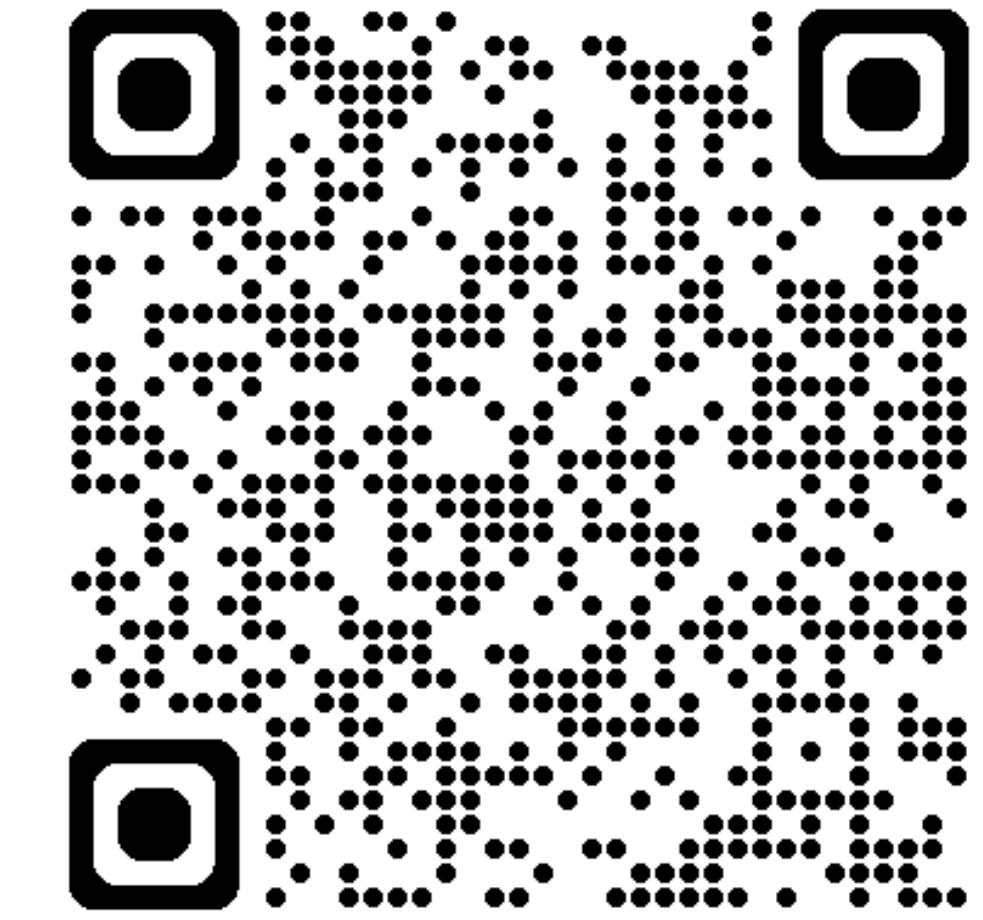


Newton is Open Source under Apache 2.0



<https://github.com/newton-physics>

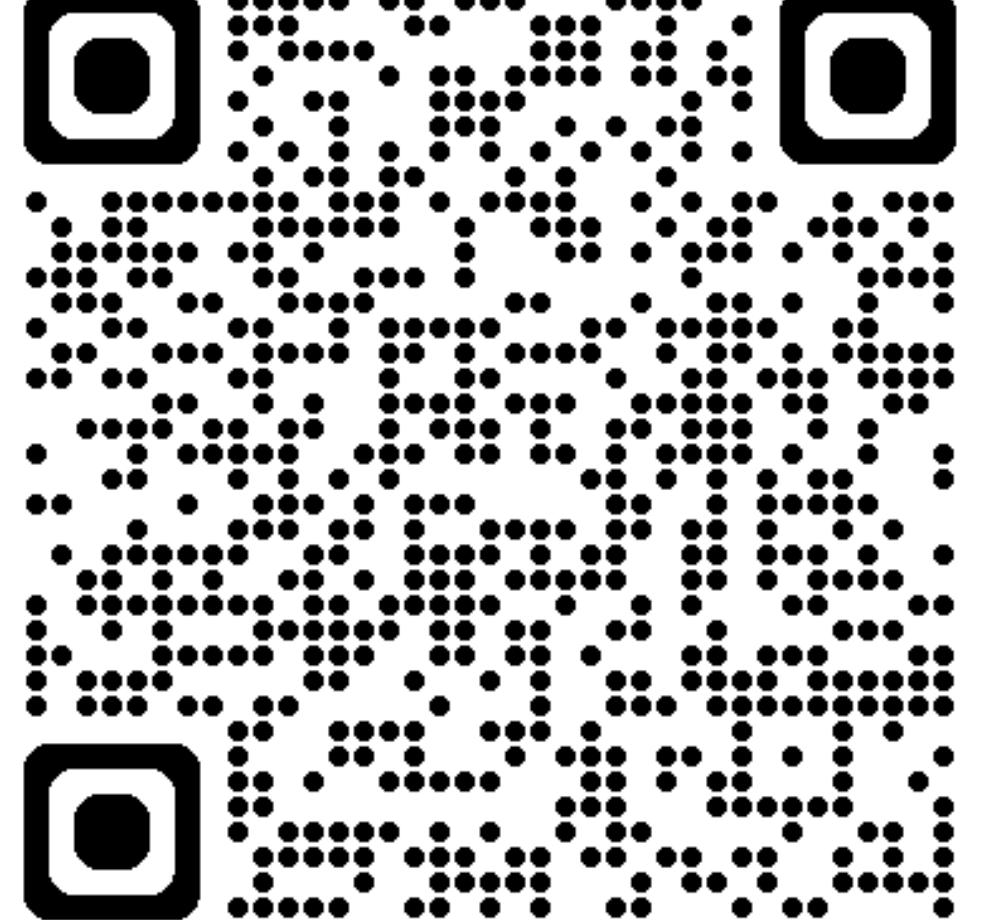
Newton Talks



Announcing MuJoCo-Warp and Newton: How Google DeepMind and NVIDIA are Supercharging Robotics Development

Erik Frey, Google DeepMind

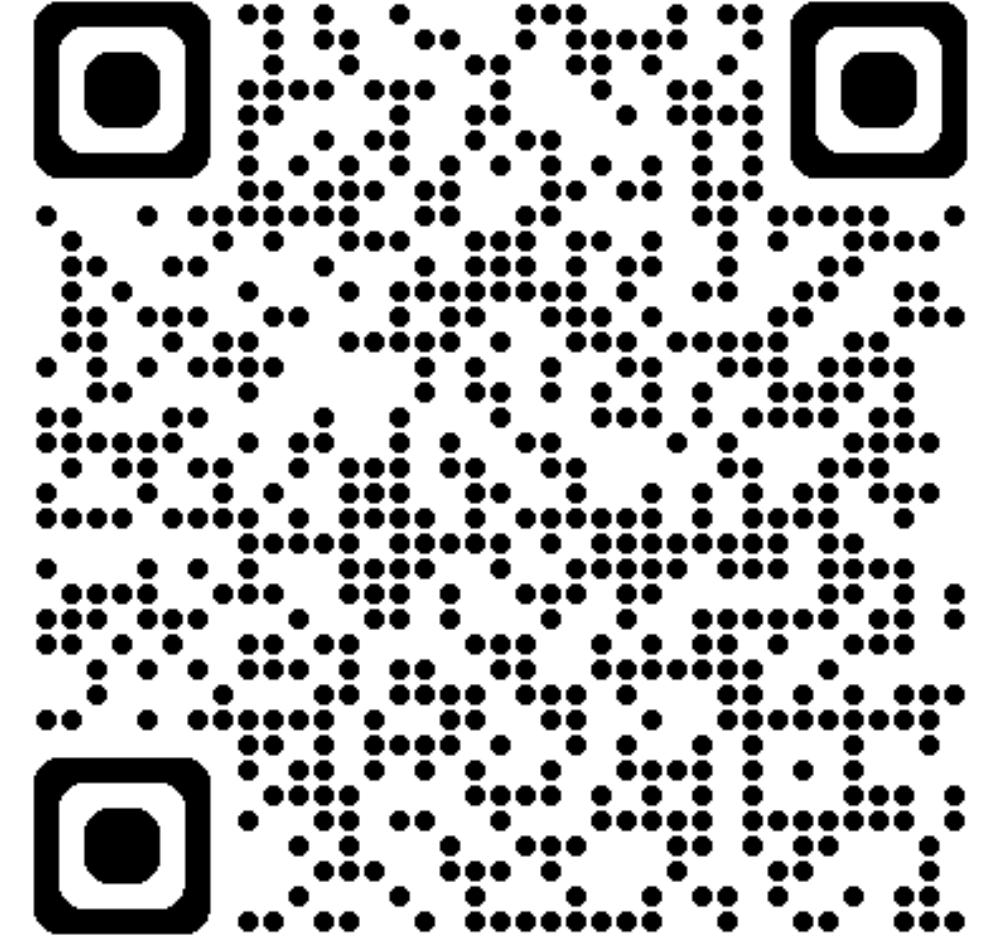
GTC, San Jose, March 2025



How Robots Learn To Be Robots in Simulation With the Newton Physics Engine

Yuval Tassa, Google DeepMind
Miles Macklin, NVIDIA

GTC/VivaTech, Paris, June 2025



How Disney Droids Come to Life With Physics Simulation

Moritz Bächer, Disney Research
Mohammad Mohajerani, NVIDIA

SIGGRAPH, Vancouver, August 2025



Thanks

More Resources

- Newton: <https://github.com/newton-physics/newton>
- Warp: <https://nvidia.github.io/warp/> , <https://nvidia.github.io/warp/publications.html>
- Mujoco Warp: https://github.com/google-deepmind/mujoco_warp
- IsaacLab Integration: <https://isaac-sim.github.io/IsaacLab/main/source/experimental-features/newton-physics-integration/index.html>
- Neural Robot Dynamics : <https://neural-robot-dynamics.github.io/>
- Taccel: <https://taccel-simulator.github.io/supercharging>